

FATTO

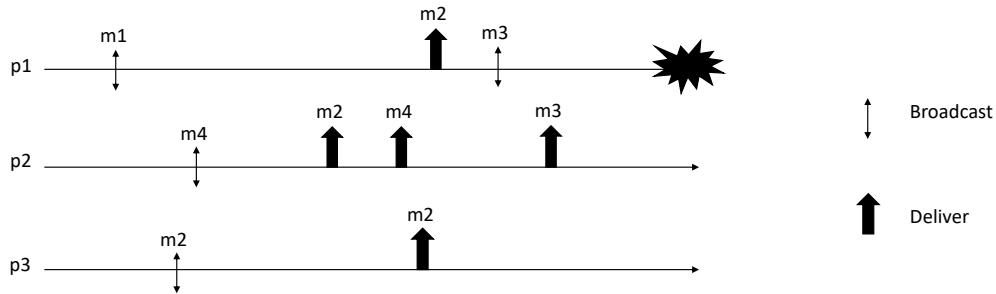
Distributed Systems (9 CFU)

05/07/2022

Family Name _____ Name _____ Student ID _____

Ex 1: Provide the specification of the (1, N) Regular Register and describe the majority voting algorithm discussed during the lectures.

X Ex 2: Consider the message pattern shown in the Figure

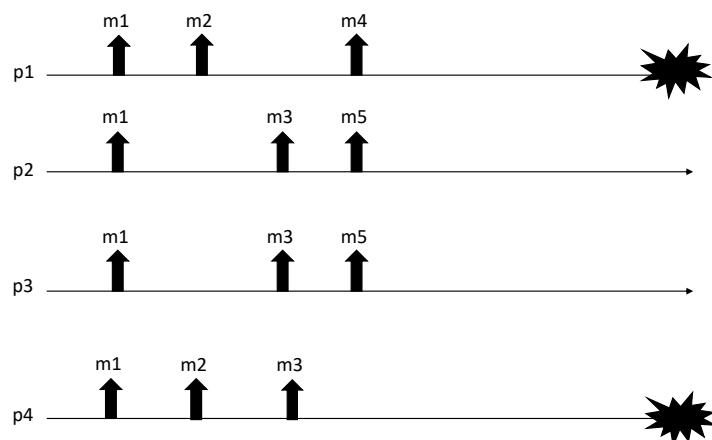


Answer to the following questions:

1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

NOTE: To solve the exercise you can just add deliveries of messages and new broadcast (if needed)

X Ex 3: Consider the execution depicted in the Figure



Answer to the following questions:

1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO (UA, WUTO) but not TO (UA, SUTO)

3. Modify the run in order to obtain an execution satisfying TO (NUA, WNUTO) but not TO(NUA, WUTO).

NOTE: To solve the exercise you can just add deliveries of messages.

X Ex 4: Let us consider a Regular Reliable Broadcast primitive satisfying the following properties:

- *Validity*: If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement*: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Let us consider a distributed system composed of N processes executing the Eager algorithm (reported in figure)

Algorithm 3.3: Eager Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb .

```

upon event (  $rb$ , Init ) do
   $delivered := \emptyset$ ;

upon event (  $rb$ , Broadcast |  $m$  ) do
  trigger (  $beb$ , Broadcast | [DATA, self,  $m$ ] );

upon event (  $beb$ , Deliver |  $p$ , [DATA,  $s$ ,  $m$ ] ) do
  if  $m \notin delivered$  then
     $delivered := delivered \cup \{m\}$ ;
    trigger (  $rb$ , Deliver |  $s$ ,  $m$  );
    trigger (  $beb$ , Broadcast | [DATA,  $s$ ,  $m$ ] );
  
```

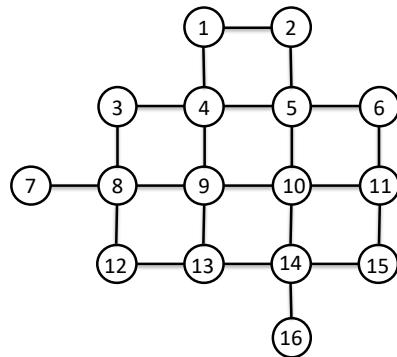
Answer to the following questions:

- X** 1. assuming that up to f processes may commit omission failures and no other failures may happen, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).
- NOT DONE** 1. assuming that up to f processes may be **Byzantine** faulty but constrained to have a symmetric behaviour¹, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).

X Ex 5: Consider a distributed system composed by n processes each one having a unique identifier. Processes communicate by exchanging messages through perfect point-to-point links and are connected through a grid (i.e., each process p_i can exchange messages only with processes located at *nord*, *sud*, *east* and *west* when they exist).

An example of such network is provided in the following figure:

¹ A Byzantine process has a symmetric behaviour if it can change the content of every message it is going to send, but it cannot send different values to different processes when invoking the $bebBroadcast$. Summarizing, it can cheat but it will do it in a consistent way.



Processes are not going to fail, and they initially know only the number of processes in the system N and the identifiers of their neighbors.

Processes in the system must agree on a color assignment satisfying the following specification:

Module

Name: k-Color assignment

Events:

Request: $\langle ca, \text{Propose} \mid c \rangle$: Proposes a color to be adopted.

Indication: $\langle ca, \text{Decide} \mid c \rangle$: Outputs a decided color to be adopted by the process

Properties:

Termination: Every process eventually decides a color.

Validity: If a process decides a color c , then c was proposed by some process or $c = \text{default}$.

Integrity: No process decides twice.

Weak Agreement: If two processes decide c_i and c_j then either $c_i = c_j$ or one of the two is default

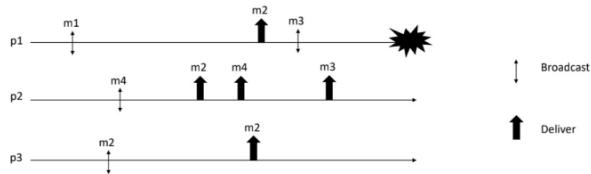
Color Selection: Let C be the set of proposed colors, if $|C| > 0$ then there exists at least a color $c_i \in C$ that is decided by k processes.

Assuming that $1 < k < N$ and that k is known by every process, write the pseudo-code of an algorithm implementing the k-Color assignment primitive.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

Signature: _____

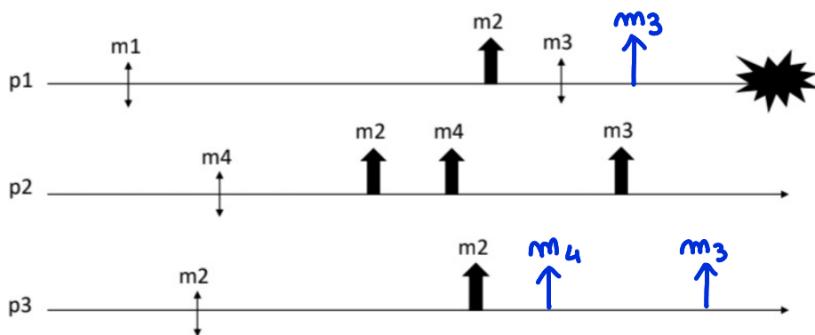
Ex 2: Consider the message pattern shown in the Figure



Answer to the following questions:

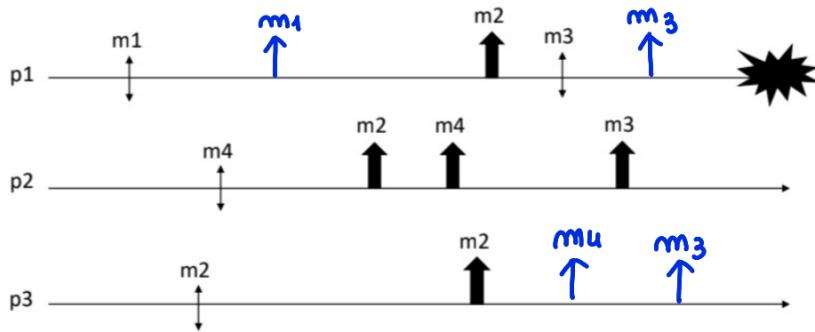
1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

1) URB : faulty subset of correct



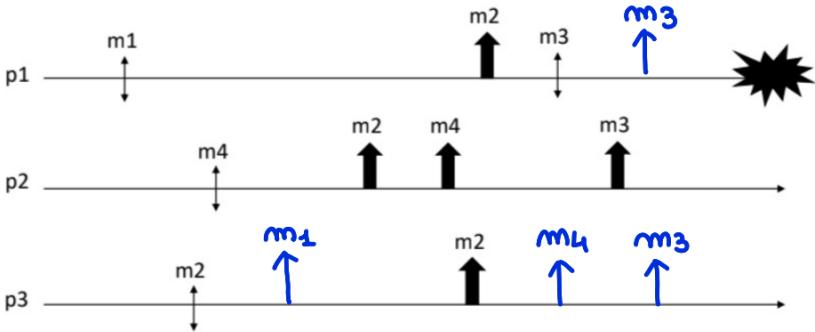
C: m2 m4 m3
F: m2 m3

2) RB, mot URB : correct same set

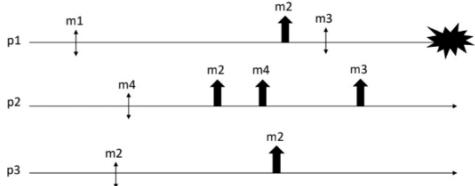


C: m2 m3 m4
F: m2 m3 m1

3) BEB, mot RB : validity
m4
m2



4) Casual and Total



FIFO
+ LOCAL : $m_1 \rightarrow m_3$
 $m_2 \rightarrow m_3$

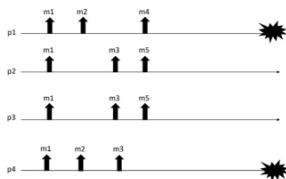
TOTAL: $m_2 \rightarrow m_4 \rightarrow m_3$

$m_1 \ m_2 \ m_4 \ m_3$

$m_2 \ m_1 \ m_4 \ m_3$

$m_2 \ m_4 \ m_1 \ m_3$

Ex 3: Consider the execution depicted in the Figure



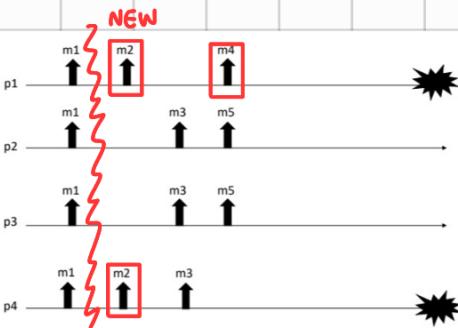
Answer to the following questions:

1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO (UA, WUTO) but not TO (UA, SUTO)

3. Modify the run in order to obtain an execution satisfying TO (NUA, WNUTO) but not TO(NUA, WUTO).

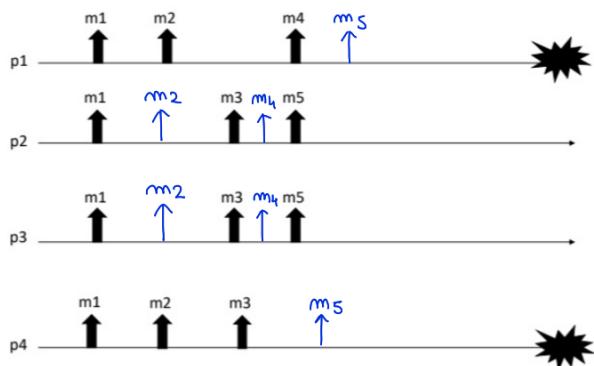
NOTE: To solve the exercise you can just add deliveries of messages.

1) NUA: no agreement (\nvdash no subset of C).
WUTO \rightarrow TO(NUA, WUTO)



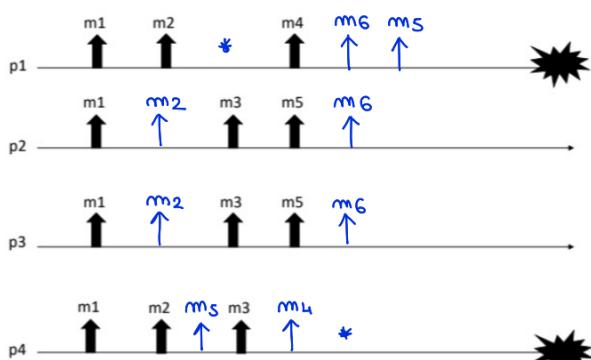
p1 $m_1 (m_2 \ m_4)$
p4 $m_1 (m_2) \ m_3$

2) TO(UA, WUTO), not TO(UA, SUTO)



$m_1 \ m_2 \ m_4 \rightarrow$ SUBSET
 $m_1 \ m_2 \ m_3 \ m_4 \ m_5$
 $m_1 \ m_2 \ m_3 \ m_4 \ m_5$
 $m_1 \ m_2 \ m_3 \rightarrow$ SUBSET

3) To(NUA, WNUTO), mat To(NUA, WUTO)



*: omission

WUTO: 1 2 3 5
1 3 4 5 *

WNUTO: 1 2 3 4 6
1 4 3 5 6 *

1 2 * 4 6 5

1 2 3 5 6

1 2 3 5 6

1 2 3 4

- Ex 4:** Let us consider a Regular Reliable Broadcast primitive satisfying the following properties:
- **Validity:** If a correct process p broadcasts a message m , then p eventually delivers m .
 - **No duplication:** No message is delivered more than once.
 - **No creation:** If a process delivers a message m with sender s , then m was previously broadcast by process s .
 - **Agreement:** If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Let us consider a distributed system composed of N processes executing the Eager algorithm (reported in figure)

```

Algorithm 3.3: Eager Reliable Broadcast
Implements:
    ReliableBroadcast, instance rb.

Uses:
    BestEffortBroadcast, instance beb.

upon event (rb, Init) do
    delivered := ∅;

upon event (rb, Broadcast | m) do
    trigger (beb, Broadcast | [DATA, self, m]);

upon event (beb, Deliver | p, [DATA, s, m]) do
    if m ∉ delivered then
        delivered := delivered ∪ {m};
        trigger (rb, Deliver | s, m);
        trigger (beb, Broadcast | [DATA, s, m]);

```

Answer to the following questions:

1. assuming that up to f processes may commit omission failures and no other failures may happen, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).
1. assuming that up to f processes may be Byzantine faulty but constrained to have a symmetric behaviour¹, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).

f : omit failures

The eager reliable broadcast is the asynchronous RB protocol that cannot provide the list / set of correct processes and cannot know when the retransmission is needed or not. The protocol (to stay safe) will retransmit every message that receive. This means that the FD is not perfect.

①

Validity: yes it is still valid. Because if a correct process broadcast a message $[DATA, s, m]$, p received the m and if $m \notin \text{delivered}$, p delivered $[DATA, s, m]$ through perfect link.

No duplication: yes it is valid because the protocol check if the message is delivered ($m \in \text{delivered}$) in order to avoid duplication of messages.

No creation: yes it is valid because in the broadcast of a message the protocol specify the source s , that

is the sender. When a message is delivered we are sure that the msg is broadcasted by the sender s. The faulty processes can't generate fake messages.

Agreement: yes it is valid, because can happen that faulty processes omit a msg but the correct still broadcast and deliver msgs from correct processes.

② NOT DONE

5) n processes, unique ID

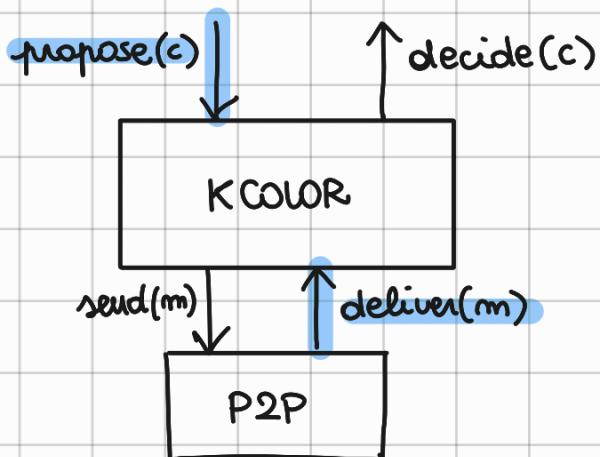
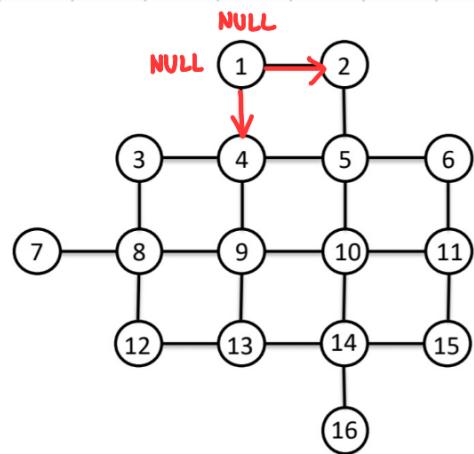
P2P links

grid (N, S, E, W)

N all correct

get - neighbors

no failures



We have a grid of processes that are connected through perfect link. A process is connected to its neighbors and can communicate only with π_N , π_S , π_E and π_W . Assuming that π are not going to fail, each process start with a value c and have to exchange this value c with the other π .

At first we start a round, the π_i starts to propose the value c to its neighbors (if they exist) and c is added to proposal set. When the set is full of value c and all the correct process have proposed, start the decide part. The final value c have to be proposed by K processes accordingly to $1 \leq K \leq N$.

upon event $\langle \text{Kcolor}, \text{INIT} \rangle$ do

$N = \pi$

$c = \text{default}$

$\text{proposal-set} = \emptyset$

$\text{decision} = \perp$

$\text{neighbours} = \{N, S, E, W\}$

$K = *$ it is known

upon event $\langle \text{Kcolor}, \text{PROPOSE} | c \rangle$ do

$\text{proposal} = c$

$\text{proposal-set} = \text{proposal-set} \cup \{c\}$

for each $p_j \in \text{neighbours}$ do

trigger $\langle p_2p, p_2p - \text{SEND} | (\text{PROPOSAL},$
 $\text{proposal-set},$
 $p_i) \rangle$ to p_j

upon event $\langle p_2p, p_2p - \text{DELIVER} | (\text{PROPOSAL},$

$\text{proposal-set},$

$p_t) \rangle$ from p_k

if $N \subseteq \text{proposal-set}$

if $\text{decision} = \perp$ do

$\text{decision} = \text{FUNCTION FOR DECISION}$

else

$\text{proposal-set} = p_s \cup p_s$ update of proposal set

for each $p_j \in \text{neighbours}$ do

trigger $\langle p_2p, p_2p - \text{SEND} | (\text{PROPOSAL},$
 $\text{proposal-set},$
 $p_t) \rangle$ to p_x

upon event FUNCTION FOR DECISION

color = lighter-color (proposal-set)

if self $\leq k$ do

decision = color

trigger < kcolor, DECIDE | decision >

else

decision = ⊥

trigger < kcolor, DECIDE | decision >

OTHERWISE you can insert the function in the p2p deliver:

upon event < p2p, p2p-DELIVER | (PROPOSAL,
proposal-set,
 p_t) > from p_k

if $N \subseteq \text{proposal-set}$

if decision = ⊥ do

color = lighter-color (proposal-set)

if self $\leq k$ do

decision = color

trigger < kcolor, DECIDE | decision >

else

decision = ⊥

trigger < kcolor, DECIDE | decision >

else

proposal-set = $p_s \cup p_s$ update of proposal set

for each $p_j \in$ neighbours do

trigger < p2p, p2p-SEND | (PROPOSAL,
proposal-set,
 p_t) > to p_x