

FATTO

Distributed Systems

10/06/2020

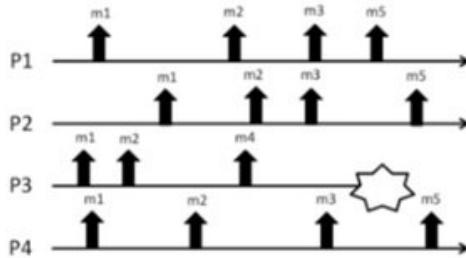
Family Name _____ Name _____ Student ID _____

Please, tick the appropriate option:

- | | | | |
|--------------------------|---|--------------------------|---------|
| <input type="checkbox"/> | Master of Science in Engineering in Computer Science | <input type="checkbox"/> | Erasmus |
| <input type="checkbox"/> | Master of Science in Artificial Intelligence and Robotics | <input type="checkbox"/> | Other |

X Ex 1: Concerning software replication techniques, describe the primary-backup and the active replication approach with particular emphasis on how each technique handles failures.

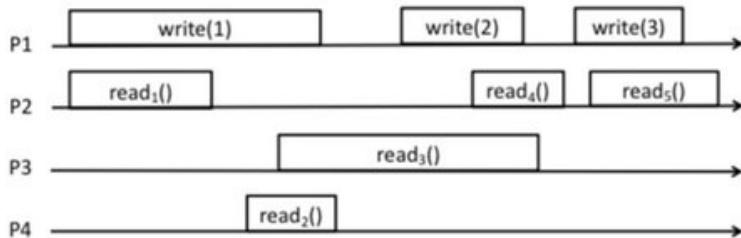
X Ex 2: Consider the run depicted in the figure:



1. Which type of total ordering is satisfied by the run? Specify both the agreement and the ordering properties.
2. Modify the run in order to satisfy TO(UA, WUTO) but not TO (UA SUTO)
3. Modify the run in order to satisfy TO(NUA, WNUTO) but not TO(NUA, WUTO)

NOTE: In order to solve the exercise, you can just ADD broadcast, deliveries and failures.

X Ex 3: Consider the execution depicted in the following figure and answer the questions



1. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.
2. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.
3. Let us assume that values retuned by read operations are as follow: read₁() → 1, read₂() → 0, read₃() → 1, read₄() → 2, read₅() → 3. Is the run depicted in the Figure linearizable?

X Ex 4: Let us consider the following algorithm implementing a (1, N) atomic register in synchronous system.

<pre> 1. upon event { onar, Init } do 2. (ts, val) := (0, ⊥); 3. correct := Π; 4. writeset := ∅; 5. readval := ⊥; 6. reading := FALSE; 7. uponevent(P, Crash p) do 8. correct := correct \ {p}; 9. upon event { onar, Read } do 10. reading := TRUE; 11. readval := val; 12. trigger { beb, Broadcast [WRITE, ts, val] }; 13. upon event { onar, Write v } do 14. trigger { beb, Broadcast [WRITE, ts + 1, v] }; </pre>	<pre> 15. upon event { beb, Deliver p, [WRITE, ts', v'] } do 16. if ts' > ts then 17. (ts, val) := (ts', v'); 18. trigger { pl, Send p, [ACK] }; 19. upon event { pl, Deliver p, [ACK] } then 20. writeset := writeset ∪ {p}; 21. upon correct ⊆ writeset do 22. writeset := ∅; 23. if reading = TRUE then 24. reading := FALSE; 25. trigger { onar, ReadReturn readval }; 26. else 27. trigger { onar, WriteReturn }; </pre>
--	--

Assuming that messages are sent by using perfect point-to-point links and that the broadcast is best effort answer the following questions:

- MANCA*
1. Discuss what happen to every atomic register property (i.e., termination, validity and ordering) if messages can be lost.
 2. Discuss what happen to every atomic register property (i.e., termination, validity and ordering) if we change line 12 (i.e., **trigger** { beb, Broadcast | [WRITE, ts, val] }; in the read handler) with **trigger** { beb, Broadcast | [WRITE, ts+1, val] };

X Ex 5: Consider a distributed system composed of n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through fair loss point-to-point links. Processes are connected through a directed ring (i.e., each process p_i can exchange messages only with processes $p_{i+1 \text{ mod } n}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $new_next(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Regular Reliable Broadcast.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

Signature: _____

1) The primary-backup and the active replication are two techniques that implement liveability.

The first can be explained as: primary means that receives invocation from clients and sends back the answers; also given an object x , $\text{prim}(x)$ returns the primary of x . While backup means that it interacts with $\text{prim}(x)$ and is used to guarantee fault tolerance by replacing a primary when there're crashes.

In the case of NO CRASH scenario, before sending back the response to the client, the primary replica sends an update to all the other correct backups and, only after it gets an ACK from them, it will send a response to the client.

In the case of CRASH, we have different scenarios:

① the primary fails after the client receives the answer and there are two cases:
the client doesn't receive the response due to PP2P link and the client needs to retransmit the requests;
the client receives the answer and it's fine.

② the primary fails before sending update msgs:
client doesn't get answer and needs to resend the requests that is handled as new.

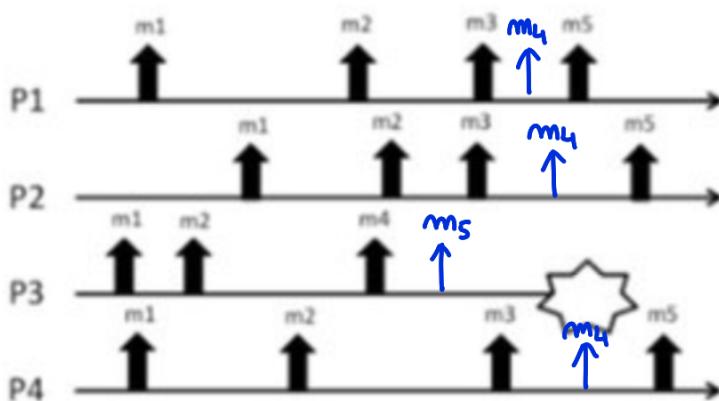
③ the primary fails after sending update msgs but before receiving all ACK.

The update is received either by all or by no one, so the primary have to be elected among all replicas.

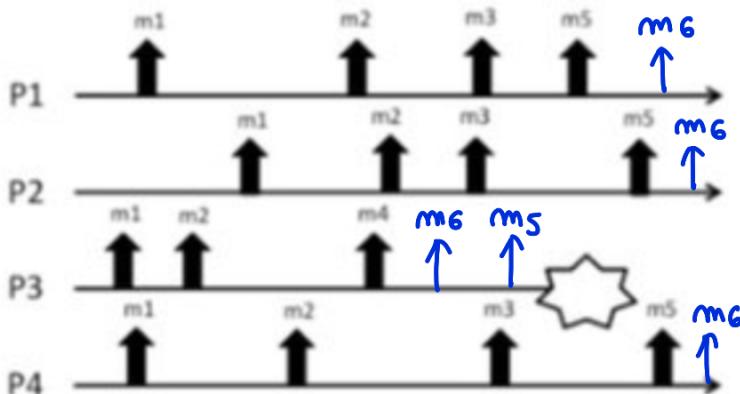
In the **active replication** there isn't a coordinator, all replicas have the same role. Each replica is deterministic, so the client receives the same response for all replicas. It have to ensure linearizability and needs to preserve **atomicity**: if a replica executes an invocation, all correct replicas execute the same invocation, and **ordering**: no two correct replicas have to execute two invocations in different order. It doesn't need recovery action upon the failure of a replica.

2) 2.1 TO(NUA, SUTO)

2.2 (UA, WUTO), mot (UA, SUTO)



2.3 (NUA, WNUTO), mot (NUA, WUTO)



3) 3.1 Regular register

R₁: 0,1

R₂: 0,1

R₃: 0,1,2

R₄: 1,2

R₅: 2,3

3.2 Atomic register

R₁: 0,1

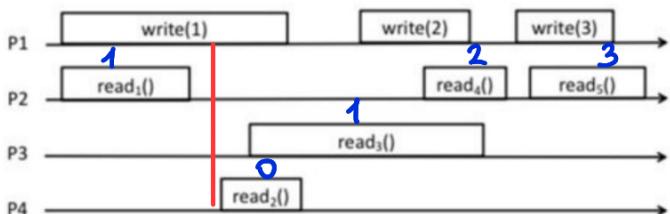
R₂: 0,1

R₃: 0,1,2 & 1,2 → depends on R₁ because if R₁ has (0), R₃ can have (0,1,2)
otherwise (1,2)

R₄: 1,2

R₅: 2,3

$$3.3 \quad R_1=1 \quad R_2=0 \quad R_3=1 \quad R_4=2 \quad R_5=3$$

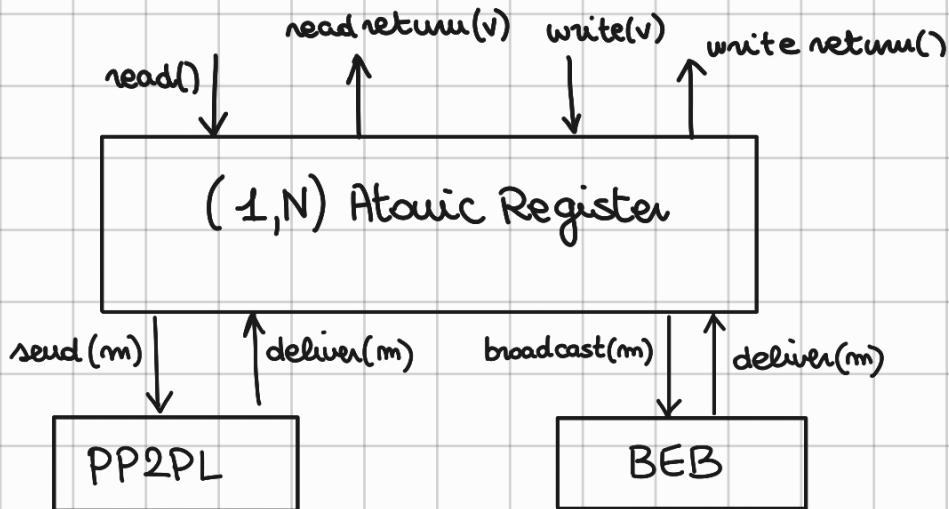


No, because R₁ and R₂ are not concurrent and R₂ returns 0 after R₁=1.
Violation!

4) PP2PL

BEB

(1,N) atomic register



4.1 Messages can be lost

Termination \rightarrow if a correct process invokes all operations, then the operation eventually completes

Validity \rightarrow A read that is not concurrent with a write, returns the last value written;

A read that is concurrent with a write, returns the last value written o the value concurrently written.

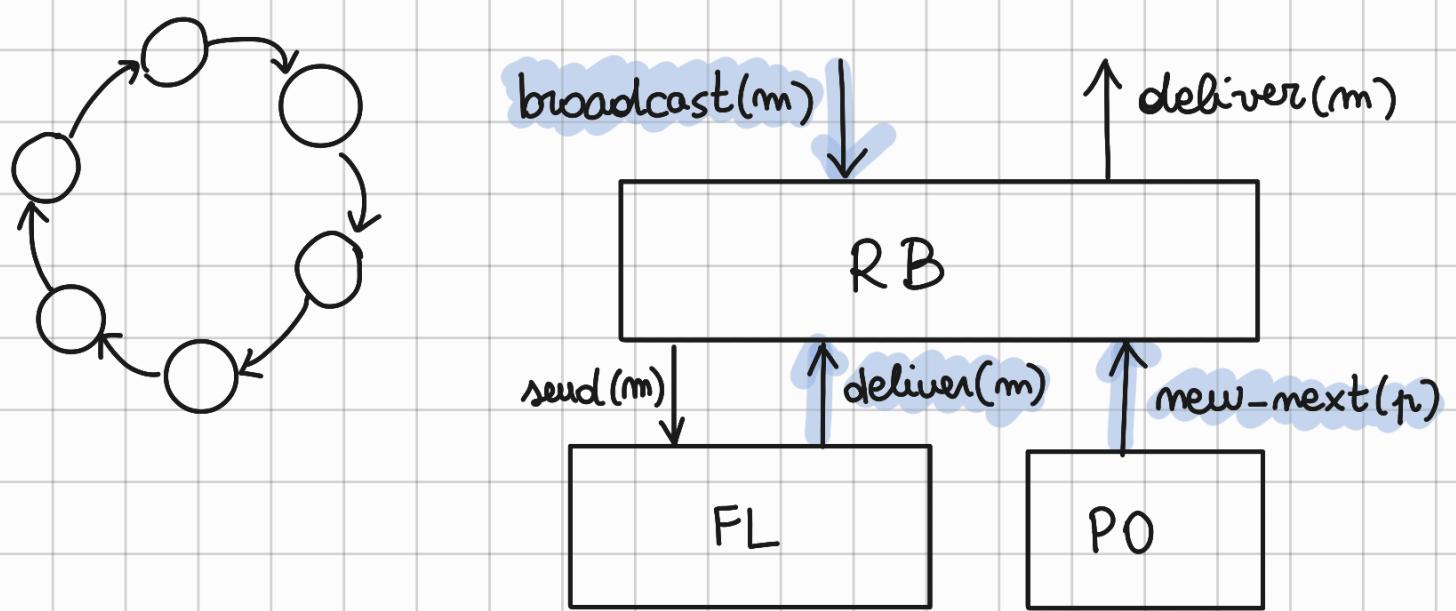
Ordering \rightarrow If a read returns a value v and a subsequent read returns a value w , then the write of w does not precede the write of v .

$$R(v) \prec R(w)$$

$$W(v) \prec W(w)$$

Ex 5: Consider a distributed system composed of n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through fair loss point-to-point links. Processes are connected through a directed ring (i.e., each process p_i can exchange messages only with processes $p_{i+1 \text{ mod } n}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $\text{new_next}(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Regular Reliable Broadcast.



The protocol knows the set of processes and when receive the broadcast event, spread the msg through the fair loss link. We know that using this type of communication channel, msgs can be lost (properties: fair-loss, finite duplication, no creation).

upon event $\langle \text{rb}, \text{INIT} \rangle$ do

correct = Π

mext = $p_{(i+1) \text{ mod } n}$

$\text{from}[p] = [\emptyset]^n$

upon event $\langle \text{rb}, \text{BROADCAST } | m \rangle$ do

