

FATTO

Distributed Systems

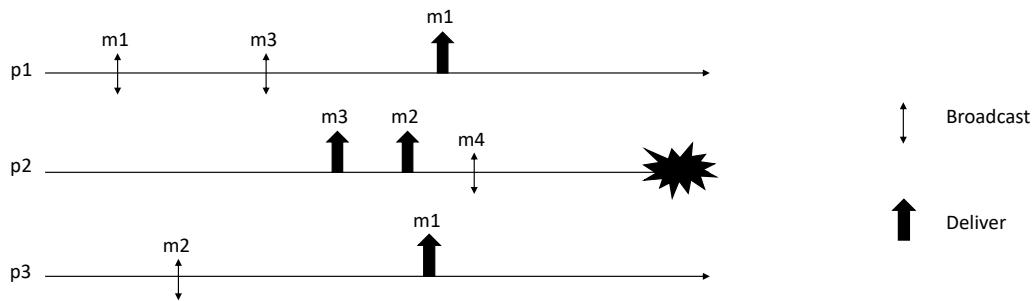
29/01/2021

Exam A

Family Name _____ Name _____ Student ID _____

Ex 1: Provide the specification of the (1, N) Regular Register and describe the majority voting algorithm discussed during the lectures.

X **Ex 2:** Consider the message pattern shown in the Figure

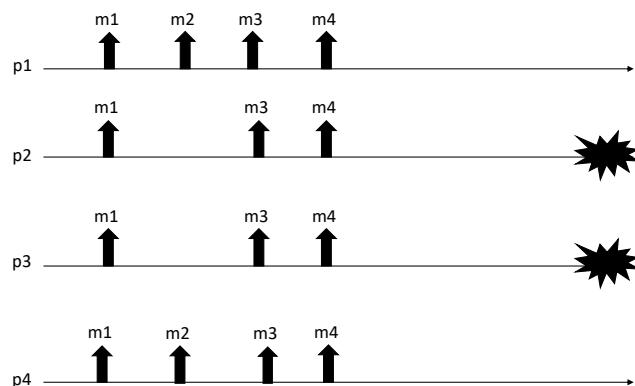


Answer to the following questions:

1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

NOTE: To solve the exercise you can just add deliveries of messages and new broadcast (if needed)

X **Ex 3:** Consider the execution depicted in the Figure



Answer to the following questions:

1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO(UA, WNUTO) but not TO(UA, WUTO)

3. Modify the run in order to obtain an execution satisfying TO(NUA, SUTO) but not TO(UA, SUTO).

NOTE: To solve the exercise you can just add deliveries of messages.

Ex 4: Let us consider a distributed system composed of N processes executing the algorithm reported in figure

Algorithm 3.12: Broadcast with Sequence Number

Implements:
FIFOReliableBroadcast, **instance** *frb*.

Uses:
ReliableBroadcast, **instance** *rb*.

```

upon event <frb, Init > do
    lsn := 0;
    pending :=  $\emptyset$ ;
    next :=  $[1]^N$ ;

upon event <frb, Broadcast | m > do
    lsn := lsn + 1;
    trigger <rb, Broadcast | [DATA, self, m, lsn] >;

upon event <rb, Deliver | p, [DATA, s, m, sn] > do
    pending := pending  $\cup \{(s, m, sn)\}$ ;
    while exists  $(s, m', sn') \in pending$  such that  $sn' = next[s]$  do
        next[s] := next[s] + 1;
        pending := pending  $\setminus \{(s, m', sn')\}$ ;
        trigger <frb, Deliver | s, m' >;

```

Let us assume that

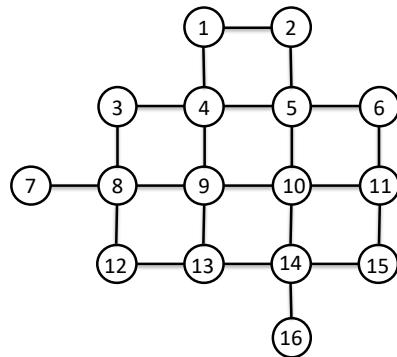
1. up to f processes may be **Byzantine** faulty and
2. A Byzantine process is not able to compromise the underline Reliable Broadcast primitive (i.e., when a **Byzantine** process sends a message through the *rbBroadcast* interface, the message will be reliably delivered to every correct process).

For each of the following properties, discuss if it can be guaranteed when $f=1$ and motivate your answer (also by using examples)

- *Validity*: If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement*: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- *FIFO delivery*: If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

X Ex 5: Consider a distributed system composed by n processes each one having a unique identifier. Processes communicate by exchanging messages through perfect point-to-point links and are connected through a grid (i.e., each process p_i can exchange messages only with processes located at *nord*, *sud*, *east* and *west* when they exist).

An example of such network is provided in the following figure:



Processes are not going to fail, and they initially know only the number of processes in the system N and the identifiers of their neighbors.

Processes in the system must agree on a color assignment satisfying the following specification:

Module

Name: Color assignment

Events:

Request: $\langle ca, \text{Propose} \mid c \rangle$: Proposes a color to be adopted.

Indication: $\langle ca, \text{Decide} \mid c \rangle$: Outputs a decided color to be adopted by the process

Properties:

Termination: Every process eventually decides a color.

Validity: If a process decides a color c , then c was proposed by some process.

Integrity: No process decides twice.

Fairness: If a color c is proposed by a process p_i , c will be eventually decided by some process p_j

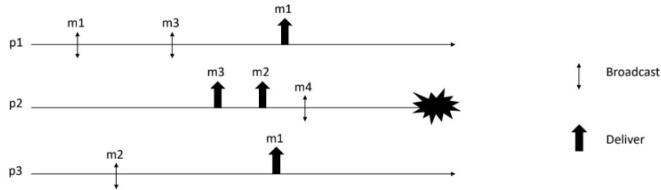
Color Balance: Let C be the set of proposed colors, every $c_i \in C$ is decided by at most $k = N/|C|$ processes.

Write the pseudo-code of an algorithm implementing the Color assignment primitive.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

Signature: _____

Ex 2: Consider the message pattern shown in the Figure

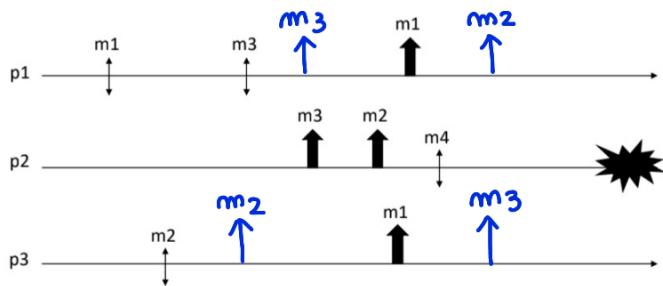


Answer to the following questions:

1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

NOTE: To solve the exercise you can just add deliveries of messages and new broadcast (if needed)

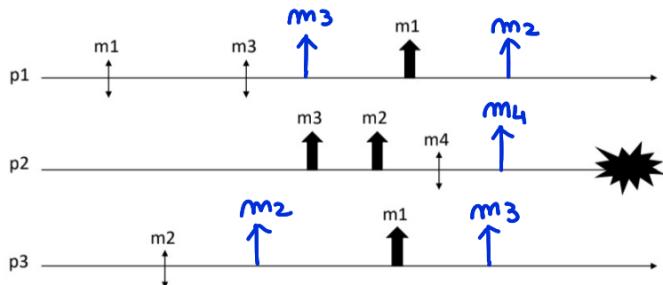
1) URB



correct: m₁ m₂ m₃

faulty subset : m₃ m₂

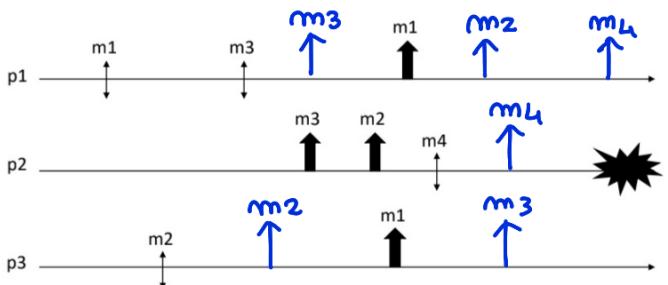
2) RB, not URB



correct: m₃ m₁ m₂
m₂ m₁ m₃

faulty: m₃ m₂ m₄

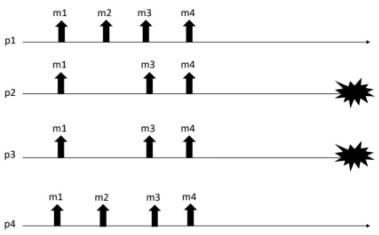
3) BEB, not RB



correct: m₃ m₁ m₂ m₄
m₂ m₁ m₃

faulty: m₃ m₂ m₄

Ex 3: Consider the execution depicted in the Figure



Answer to the following questions:

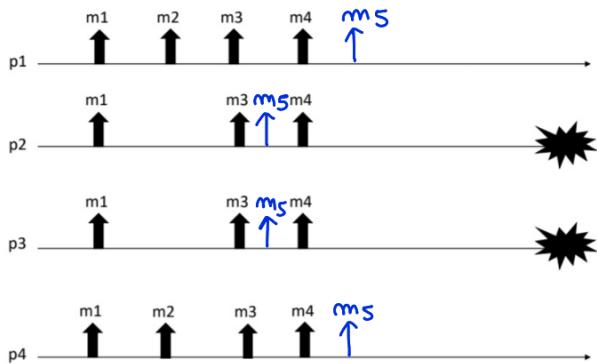
1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO(UA, WNUTO) but not TO(UA, WUTO)
3. Modify the run in order to obtain an execution satisfying TO(NUA, SUTO) but not TO(UA, SUTO).

NOTE: To solve the exercise you can just add deliveries of messages.

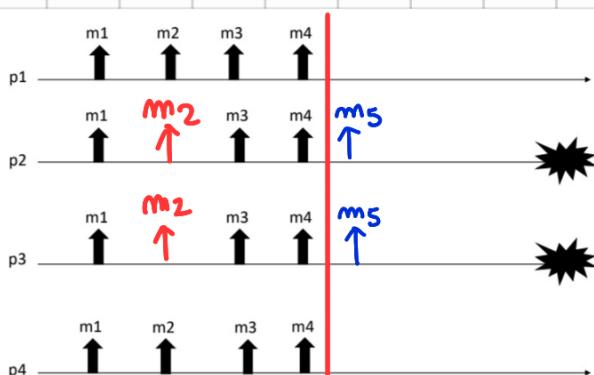
1) UA: the faulty are a subset of correct
 WUTO : no restrictions on the set of delivered message

2) TO(UA, WNUTO), not TO(UA, WUTO)

WNUTO
 1 2 3 4
 1 4 3



3) TO(NUA, SUTO), not TO(UA, SUTO)

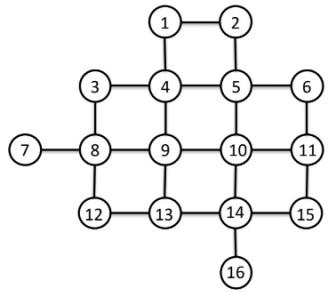


Ex 5: Consider a distributed system composed by n processes each one having a unique identifier. Processes communicate by exchanging messages through perfect point-to-point links and are connected through a grid (i.e., each process p_i can exchange messages only with processes located at *nord*, *sud*, *east* and *west* when they exist).

An example of such network is provided in the following figure:

Processes are not going to fail, and they initially know only the number of processes in the system N and the identifiers of their neighbors.

Processes in the system must agree on a color assignment satisfying the following specification:



Module

Name: Color assignment

Events:

Request: { ca, Propose | c }: Proposes a color to be adopted.

Indication: { ca, Decide | c }: Outputs a decided color to be adopted by the process

Properties:

Termination: Every process eventually decides a color.

Validity: If a process decides a color c , then c was proposed by some process.

Integrity: No process decides twice.

Fairness: If a color c is proposed by a process p_i , c will be eventually decided by some process p_j .

Color Balance: Let C be the set of proposed colors, every $c_i \in C$ is decided by at most $k = N/|C|$ processes.

Write the pseudo-code of an algorithm implementing the Color assignment primitive.

5) n processes, unique ID

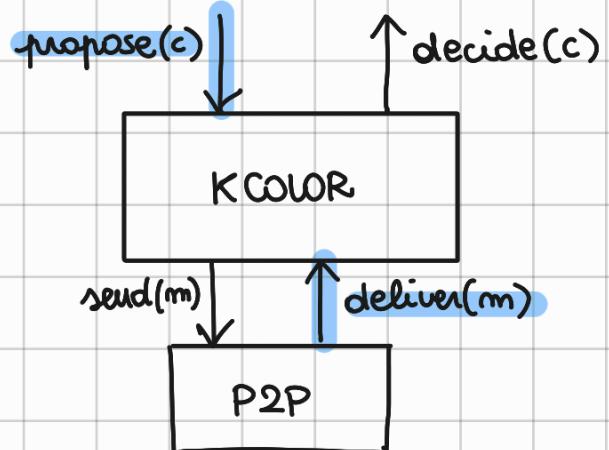
P2P links

grid (N, S, E, W)

N all correct

get - neighbors

no failures



We have a grid of processes that are connected through perfect link. A process is connected to its neighbors and can communicate only with p_N , p_S , p_E and p_W . Assuming that p are not going to fail, each process start with a value c and have to exchange this value c with the other p .

At first we start a round, the p_i starts to propose the value c to its neighbours (if they exist) and c is added to proposal set. When the set is full of value c and all the correct process have proposed, start the decide

part. The final value c have to be proposed by K processes accordingly to $\frac{N}{|C|}$.

upon event $\langle K \text{color}, \text{INIT} \rangle$ do

$$N = \pi$$

$$\text{proposal-set} = \emptyset$$

$$\text{decision} = \perp$$

$$\text{neighbours} = \{ N, S, E, W \}$$

$$K = \frac{N}{\text{proposal-set}}$$

upon event $\langle K \text{color}, \text{PROPOSE} | c \rangle$ do

$$\text{proposal} = c$$

$$\text{proposal-set} = \text{proposal-set} \cup \{ c \}$$

for each $p_j \in \text{neighbours}$ do

trigger $\langle p_2p, p_2p\text{-SEND} | (\text{PROPOSAL},$
 $\text{proposal-set},$
 $p_i) \rangle$ to p_j

upon event $\langle p_2p, p_2p\text{-DELIVER} | (\text{PROPOSAL},$
 $\text{proposal-set},$
 $p_t) \rangle$ from p_k

if $N \subseteq \text{proposal-set}$

if $\text{decision} = \perp$ do

$\text{decision} = \text{FUNCTION FOR DECISION}$

else

$\text{proposal-set} = p_s \cup p_s$ *update of proposal set*

for each $p_j \in \text{neighbours}$ do

trigger $\langle p_2p, p_2p\text{-SEND} | (\text{PROPOSAL},$
 $\text{proposal-set},$
 $p_t) \rangle$ to p_x

upon event FUNCTION FOR DECISION

color = lighter_color (proposal_set)

if self $\leq k$ do

decision = color

trigger < kcolor, DECIDE | decision >

else

decision = ⊥

trigger < kcolor, DECIDE | decision >