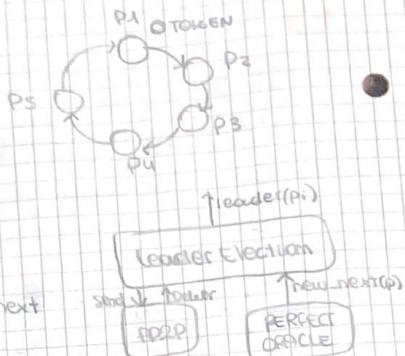


EX.2

```

INIT
ID ← identifier
next = Pi+1(mod n)
leader = i
correct = 0
IF ID == i
    PP2Psend([AUX, self, correct]) to next
UPON EVENT pp2p Deliver([AUX, s])
    IF s < correct
        correct = correct ∨ {s}
        TRIGGER PP2Psend([AUX], self, correct) to next
UPON EVENT new-next(p)
    correct = correct \ {ID(p-1)}
    TRIGGER PP2Psend([CRASH], p-1) to next
UPON EVENT pp2p Deliver([CRASH], q)
    IF q < correct
        correct = correct \ {q}
        TRIGGER PP2Psend([CRASH], q) to next
UPON leader ≠ max(node(correct))
    leader = max(node(correct))
    TRIGGER Eject(leader)
  
```



EX.2

IDEA:
when the write has round in the ring, it's

```

INIT
value = null
writing = false
next = 18(pn)
  
```

```

UPON EVENT < i
    IF writing == FF
        writing = 0
        value = i
        TRIGGER
  
```

```

UPON EVENT < i
    value = v
    IF (!writing == FF)
        writing = 1
        value = i
        TRIGGER
    ELSE
        TRIGGER
  
```

UPON EVENT TRIGGER

WEEK 9

EX.1

1. Regular Register

$$\begin{aligned} r_2() &\rightarrow 0,1,2 \\ r_1() &\rightarrow 1,2 \\ r_4() &\rightarrow 1,2 \\ r_3() &\rightarrow 1,2,3 \\ r_5() &\rightarrow 2,3 \end{aligned}$$

2. Atomic Register

$$\begin{aligned} r_2() &\rightarrow 0,1,2 \\ r_1() &\rightarrow 1,2 \\ r_4() &\rightarrow 1,2 \\ r_3() &\rightarrow \text{IF } r_2() \geq 2 \text{ THEN } r_3() \geq 2,3 \\ &\quad \text{OTHERWISE } r_3() \geq 1,2,3 \\ r_5() &\rightarrow 2,3 \end{aligned}$$

3.

$$S = w(1), \underbrace{r_2(2)}, r_4() \geq 1, w(2), r_1() \geq 2, w(3), r_3() \geq 3, r_5() \geq 3$$

not legal

NOT LINEARIZABLE

OTOKEN

P₂

P₃

leader(p₁)

election

PERFECT
ORACLE

EX.2

IDEA:

when the write has completed round in the ring, it can return

INIT

value = NULL
writing = FALSE
next = 1B(p₁)

UPON EVENT < wr, writeInv > DO value = v

IF writing == FALSE

writing = TRUE

TRIGGER < p2p | send | [WRITE], value, ... > to next

UPON EVENT < p2p | Deliver | [WRITE], v, ... >

value = v

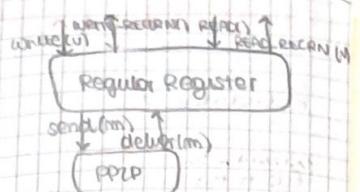
IF (!writing)

Trigger < p2p | send | [WRITE], value, no next >

ELSE

TRIGGER < wr, writeReturn >

UPON EVENT < rr | Read >
TRIGGER < ReadReturn | value >



is operating
the instant

at failure,
accuracy of

let them
time 6
550,480,

(14.1000)

the first time

failure

EXERCISE 1

PER FARE rb-deliver UN QUALSIASI PROCESSO DEVE RICEVERE UNA COPIA DEL MESSAGGIO, QUINDI IL TEMPO ATTESA DI DELIVERY SI PUÒ STIMARE COME IL TEMPO CHE IMPIEGA UN QUALSIASI PROCESSO A FAR DELIVERY DI UN MESSAGGIO.

Quanti messaggi riceve ogni processo $\rightarrow 1 \text{ msg ogni } 20 \text{ s, quindi } \frac{1}{20} = 0.05 \text{ msg/s. I processi sono } 10 \text{ quindi il rate di esecuzione è } 10 \cdot 0.05. \text{ Nell protocollo eager, ogni messaggio viene ripropagato da ogni processo. Quindi l'annual rate è a cui ogni processo è soggetto e } \lambda = 0.05 \cdot 10 \cdot 10 = 5 \text{ req/sec.}$

Ogni processo è in grado di gestire $\mu = 10 \text{ msg/sec.}$

Quindi ~~stabilità condition on~~ $\lambda \leq \mu$

$$R = \frac{1}{\mu - \lambda} = \frac{1}{10 - 5} = 0.2 \text{ s} \quad (\text{stabilità condition on } \lambda \leq \mu)$$

EXERCISE 2

2) Entrambi i protocolli richiedono che un messaggio faccia il giro dell'intero ring. Nel caso di Uniring chi esegue la get genera un solo messaggio, nel caso di Binring sono 2.

$$\lambda_{uni} = 1 \cdot 10 = 10 \text{ req/sec}$$

$$\lambda_{bi} = 2 \cdot 10 = 20 \text{ req/sec}$$

$$\mu = 10 \text{ msg/sec}$$

Siccome un messaggio deve attraversare 10 nodi, risultato finale è $R_i = 10$

$$1) Aser = \prod A_i = \left(\frac{HTBF}{HTBF + HTIR} \right)^{10} = 0,1$$

EXERCISE 3

Entro il tempo Δ ogni processo deve essere in grado di contattare tutti i processi e ricevere risposta. Siccome tutti i processi stanno eseguendo if failure detector, ogni processo deve essere in grado in $\Delta/2$ di ricevere tanti messaggi quanti sono i processi. Assumiamo una distru. esponenziale.

$$\frac{\Delta}{2} = 1 \left(10 \cdot 30 \cdot \frac{\Delta}{2} \right)$$

WEEK 10

EX.4

1. T The performance of a system is mainly measured by the time taken to perform the service, the rate at which the service is performed and the resource consumed.
2. F The response time depends on arrival rate and service rate

$$R = \frac{1}{\lambda - \mu}$$
3. T
4. F In M/M/1/k we consider the probability of an arrival
5. F Since we assume that the ~~asynchronous~~ faults occurrences and repairs are independent among them other.

B42AM

B42ANT

- devi
- Crea
- delo
- auti
- act
- rea

Theis
of th

How

usin

pe'

cr

cr

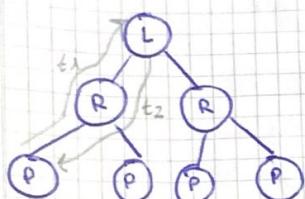
re

A1

A

A

EX.2



$$\lambda_0 = 10 \text{ req/min} \quad 2) A_{tot} = Ar \times As \times Ap = 0.9 \cdot 0.9 \cdot 1 = 0.81$$

$$R_1 = 0.5 \text{ s}$$

$$60 \text{ msg/min}$$

- 1) To evaluate the expected time that elapses between a request generated by a process P and the receipt of the response, we consider:

- Generation time $\frac{1}{\lambda_0} = 0.1 \text{ s}$
- Transmission time (Req \rightarrow R \rightarrow L \rightarrow R \rightarrow P)
 since each process can send/receive 60 msg/min, the average time to transmit a message is $1/60 = 0.016 \text{ min}$
- Computation Time: $0.5/60 = 0.008 \text{ min}$
- Relock Time: 0.016 min $T = \frac{1}{\lambda_0} + 3 \cdot \frac{1}{60} + 0.5 \cdot \frac{1}{60} + 2 \cdot \frac{1}{60} = 0.18 \text{ min}$

EX.3

$$\lambda = 10 \text{ req/s}$$

$$\mu = 5 \text{ req/s}$$



$$R = \frac{1}{\lambda \mu + \lambda} = \frac{1}{20 \cdot 5} = 0.1 \text{ s} \quad \text{expected response time}$$

BYZANTINE CONSISTENCY WEEK 14

EX.1

~~Validity = GUARANTEED.~~ Since the byzantine can not compromise the reliable delivery

~~No duplication = NOT GUARANTEED.~~ The byzantine can send a message already delivered, and the other correct never check if the message was already delivered, so they can deliver it again.

~~No creation = NOT GUARANTEED.~~ The byzantine can create whatever message

EX.1

~~Validity. Not Satisfied.~~ BYZANTINE p BROADCAST [DATA, q, m', -1] IN such a way when q receives it, q delivers m' AND next[q] = 2. Now a broadcast m and q will receive [DATA, q, m, 1] and q will never deliver m because next[q] != 1

~~No Duplication. Not Satisfied.~~ Byzantine can deliver more than once the same message.

~~No creation. Not satisfied.~~ See validity, q delivers m' with sender q but m' was not broadcast by q.

~~Agreement. Not Satisfied.~~ Byzantine p broadcasts through REB [DATA, p, m1, sn] and [DATA, p, m2, sn] (same sn). Suppose two corrects q1, q2 have next[p] = sn. It can happen that q1 receives [DATA, p, m1, sn] and delivers m1 and next[p] = sn + 1 so q1 never deliver m2; q2 receives [DATA, p, m2, sn] and delivers m2 and so q2 will never be able to deliver m1.

~~FIFO Delivery. Not Satisfied.~~ Byzantine broadcasts m1 and then broadcasts m2 both with inverted sn. In such way processes will deliver them in inverted order.

EX.2

~~Validity. Satisfied.~~ Guaranteed by the validity of REB, since the byzantine can not send different values to different processes.

~~No duplication. Not Satisfied.~~ Byzantine can deliver a message more than once.

~~No creation. Not Satisfied.~~ Byzantine can broadcast [DATA, q, m]. When a process deliver it with sender q, but q never broadcast m.

~~Agreement. Satisfied.~~ If the sender is byzantine then we know that for the symmetric behaviour of the byzantine all corrects will receive the same message.

use the
a message
or check
few con

whatever

IN
Now
never

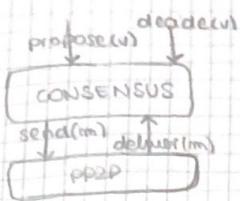
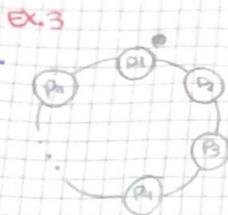
se the

ter q

[DATA, p, m1, sn]
next[p] = sn.
m1
d. delivers

decis
ted

e



IDEA:
who have the token can
propose. When all the
ring has been traversed
then the processes can decide

INIT

next = $p_{(i+1) \bmod n}$

n < n.of processes

decision = 0

proposalset = {}

IF next == p1

TRIGGER <pp2p, Send1> [TOKEN]

UPON EVENT <pp2p, Deliver1> [TOKEN]

TRIGGER <con, Propose1>

UPON EVENT <con, Propose1> n

proposalset = proposalset $\cup \{n\}$

TRIGGER <pp2p, Send1> [PROPOSE, proposalset] to next

TRIGGER <pp2p, Send1> [TOKEN] to next

UPON EVENT <pp2p, Deliver1> [PROPOSE, proposal] do

proposalset = proposalset

IF #(proposalset) == n

decision = min (proposalset)

TRIGGER <con, Decide1> [decision]

next

2. If the processes may crash, the problem is that when the token is passed to next it remains blocked since the next process never pass the token.

In order to solve this issue we need an oracle that propose by new-next (p) to each process in order to be sure that the token is passed to the next correctly.

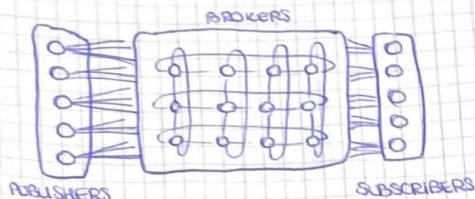
3. Termination. Not Satisfied. Byzantine can never send the token and the ring is broken.

Validity. Not satisfied. Byzantine can change the proposalset, in such a way processes may decide false value.

Integrity. Not Satisfied. Byzantine may decide twice

Agreement. Not Satisfied. Byzantine can change the message in such a way previous correct decide a value and next other value

EX.4



k-connected

1. Subscription-flooding dissemination

↳ each subscription is copied on every broker, in order to build complete subscription tables.

Code for each broker

INIT

```
publishers = getPublishers()
subscribers = getSubscribers()
brokers = getBrokers()
subscriptions[] = {}
messages = {}
```

UPON EVENT pp2pDeliver([SUB, s, p]) from q
IF (s not in subscriptions[p])
subscriptions[p] = subscription[p] U {s}
FOR b in Brokers \ {q}
TRIGGER pp2pSend([SUB, s, p]) to b

UPON EVENT pp2pDeliver([MES, m, p]) from q
IF (m not in messages)
messages = messages U {m}
FOR s in self.subscription
IF (m match subscription[s])
TRIGGER pp2pSend([MES, m, p]) to s
FOR b in Brokers \ {q}
TRIGGER pp2pSend([MES, m, p]) to b

Code for publisher

INIT

```
brokers = getBrokers()
```

UPON EVENT pp2pSend(m)
FOR b in Brokers
TRIGGER pp2pSend([CKES, m, self]) to b

Code for each subscriber

INIT

```
messages = {}
brokers = getBrokers()
```

UPON EVENT Subscribe(s)
FOR b in Brokers
TRIGGER pp2pSend([SUB, s, self]) to b

UPON EVENT pp2pDeliver([MES, m, p]) from q
IF (m not in messages)
messages = messages U {m}
TRIGGER Deliver(m) FROM p.

2. The network
the algorithm

3. The idea
the mode
subscription
neighbour
disjoint
The algo

EX.5

Specification
broadcast = 1
Safety
Safetystep = 1

FOR Glob
Community
must exist
for k
MIMO w



2. The network is k -connected, so the number of crashes that the algorithm can tolerate is $k-1$

3. The idea can be that of attaching to each message / subscription the modes that have relayed it. A neighbor of the source of the subscription / message can consider it as trusted; while a non neighbour waits for $f+1$ copies of the subscription / message with disjoint paths. The algorithm works if and only if $k \geq 2f+1$.

EX. 5

Specification of Reliable communication are:

~~live ness~~ = IF A CORRECT PROCESS P DELIVERS A CONTENT C FROM Q,
Safety THEN C WAS PREVIOUSLY SENT BY Q. (no dup. and no creation)

~~safe ness~~ live ness = IF P IS A CORRECT PROCESS AND SENDS A
CONTENT C TO A CORRECT PROCESS Q, THEN Q EVENTUALLY
DELIVERS C.

FOR Globally Bounded Model the correctness condition of reliable communication is that at least $f+1$ node-disjoint path must exists. In case of authenticated link we need node-connectivity $> 2f$.
For locally bounded instead the sufficient condition is that $K \geq 2f+1$, necessary condition $K=f+1$.