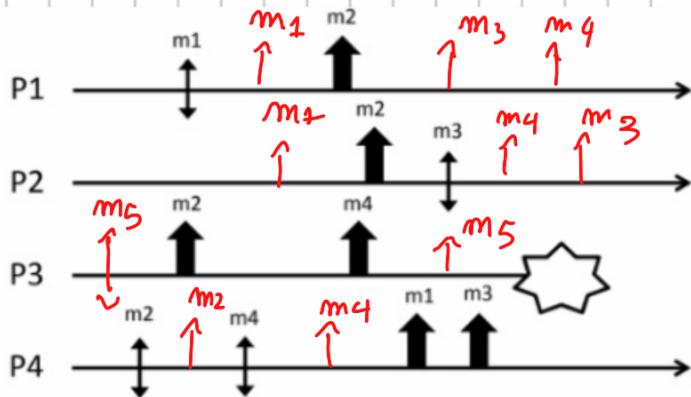


Exercise 4

1. Reliable broadcast but not URB



2. Causal order and total order

$m_2 \rightarrow m_3$ and $m_2 \rightarrow m_4$

- m_1, m_2, m_3, m_4
- m_1, m_2, m_4, m_3
- m_2, m_1, m_5, m_4
- m_2, m_1, m_4, m_3
- m_2, m_4, m_1, m_3

3. Violating causal order, TO (UA, WNUTO) but not satisfying TO (UA, SUTO)

- m_1, m_3, m_2, m_4

- m_1, m_3, m_4, m_2

- m_4, m_1, m_2, m_3

- m_4, m_2, m_1, m_3

- m_4, m_1, m_3, m_2

↳ only the correct processes in P_3 we don't add any deliver and any combination violate SUTO

Exercise 5

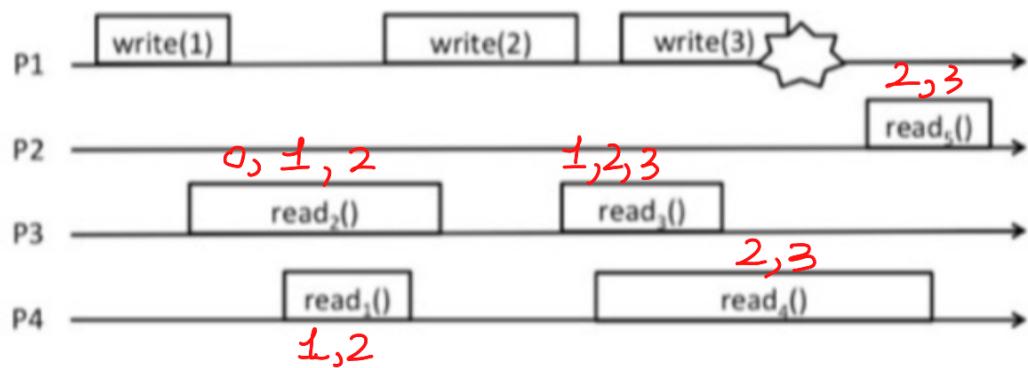
- a. F, we have NVA all correct processes have the same set of delivery, but p_2 faulty have m_4
- b. F, all the correct processes have the same set of deliveries, NVA is satisfied
- c. F, because p_2 deliver m_3 but not m_2 can't violate SUTO
- d. V, because p_1 and p_2 have the same order and p_2 respect the order of m_2 and m_3 , is WUTO
- e. F, this run satisfy WUTO that implies SNUTU
- f. V, because is missing m_2 in p_2 that not satisfies SUTO
- g. F, because like f p_2 deliver m_3 must deliver m_2 also for satisfy SUTO
- h. T, this run:



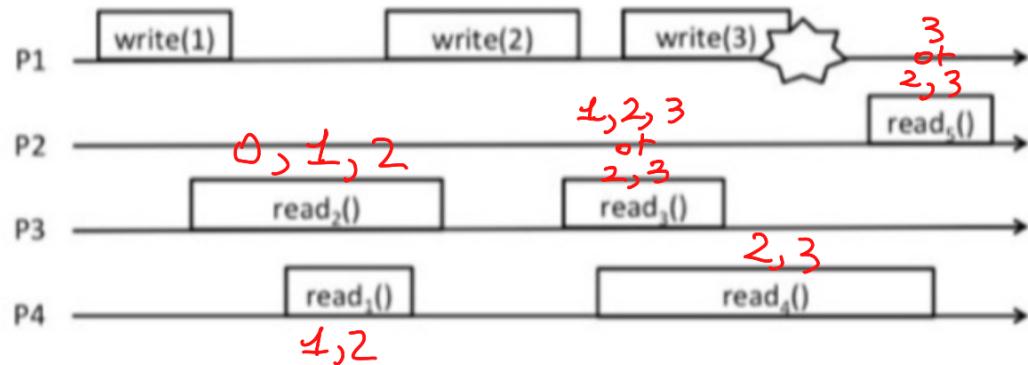
- i. F, because p_2 is correct and deliver m_4 but other correct don't deliver it, NVA is violated
- j. F, p_2 is missing m_2 and there is m_4 before m_3 .

Exercise 6

1. Rx possible value for Regular Register



2. Rx possible value for Atomic Register

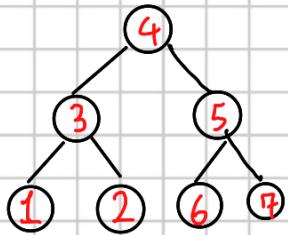


↳ because R₂ and R₃ are concurrent can have causal value

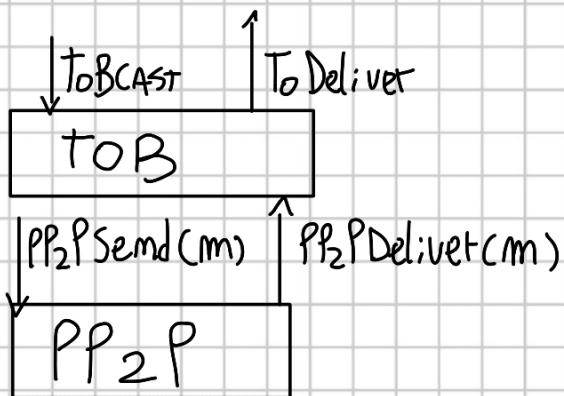
• R₃ depend on value chosen by R₁, if R₁ choose 2 then R₃ must return 2,3 otherwise can return 1,2,3

• R₅ depend on value chosen by R₃ if R₃ choose 3 then R₅ must return 3 otherwise can return 2,3

Exercise 7 - Solution



- processes have unique names.
- each process can talk only with near processes.
- links are P2P link.



$$P_1: \text{FATHER}_1 = P_3$$

$$\text{L_CHILD}_1 = \perp$$

$$\text{R_CHILD}_1 = \perp$$

One single process have FATHER; = \perp , here P_4 .

init:

FATHER; = get_father()

L_CHILD; = get_left()

R_CHILD; = get_right()

Sequence_number; = SM; = 0

LAST_Delivered; = 0

pending; = \emptyset

Upon event ToBcast(m)

if FATHER; $\neq \perp$

→ type of message

trigger PP2PSend('UNORDERED', m) to FATHER;

else

SM; = SM; + 1

trigger PP2PSend('ORDERED', m, SM;) to LEFT;

trigger PP2PSend('ORDERED', m, SM;) to Right;

Upon event PP₂PDeliver ('UNORDERED', m) from P_j

If FATHER_i ≠ L

trigger PP₂PSend ('UNORDERED', m) to FATHER_i

else

SM_i = SM_i + 1

trigger PP₂PSend ('ORDERED', m, SM_i) to LEFT_i and RIGHT_i

Upon event PP₂PDeliver ('ORDERED', m, s)

pending_j = pending_j ∪ {<m, s>}

trigger PP₂PSend ('ORDERED', m, s) to LEFT_i and RIGHT_i

Upon $\exists \langle m, s_m \rangle \in \text{pending}_i$ s.t. $s_m = \text{LAST-DELIV.} + 1$

pending_j = pending_j / {<m, s>}

trigger ToDeliver(m)

LAST-DELIVERED = SM_i

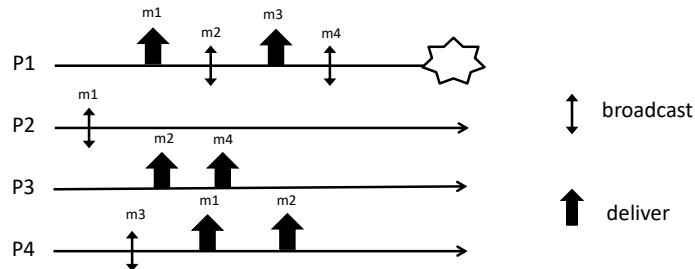
NO ID PROCESS

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 18 – Exercises
November 9th, 2022

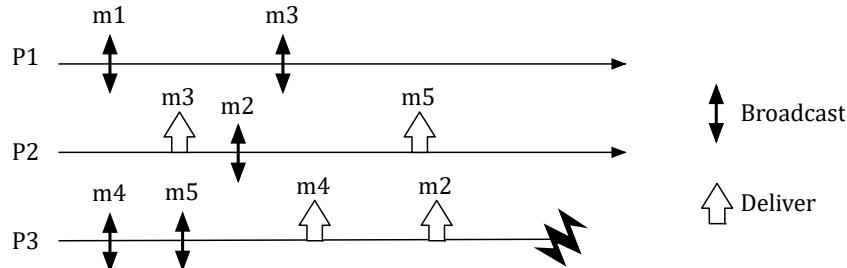
Ex 1: Consider the execution depicted in the Figure



Answer to the following questions:

1. Provide all the delivery sequences that satisfy both causal order and total order
2. Complete the execution in order to have a run satisfying TO(UA, WNUTO), FIFO order but not causal order

Ex 2: Consider the partial execution shown in the Figure and answer to the following questions:



1. Complete the execution in order to have a run satisfying the Regular Reliable Broadcast specification but not Uniform Reliable Broadcast one.
2. For each process, provide ALL the delivery sequences satisfying FIFO Reliable Broadcast but not satisfying causal order.
3. For each process, provide ALL the delivery sequences satisfying total order and causal order.

Ex 3: Consider a distributed system composed of N processes p_1, p_2, \dots, p_N , each having a unique identifier myID . Initially, all processes are correct (i.e. $\text{correct} = \{p_1, p_2, \dots, p_N\}$). Consider the following algorithm:

```

upon event  $\text{xbroadcast}(m)$ 
     $\text{mysn} = \text{mysn} + 1;$ 
     $\forall p \in \text{correct}$ 
         $\text{pp2pSend}(\text{"MSG"}, m, \text{mysn}, \text{myId});$ 

upon event  $\text{pp2pReceive}(\text{"MSG"}, m, \text{sn}, i)$ 
     $\text{mysn} = \text{mysn} + 1;$ 
    if ( $m \notin \text{delivered}$ )
        trigger  $\text{XDeliver}(m);$ 
         $\text{delivered} = \text{delivered} \cup \{m\};$ 

upon event  $\text{crash}(p_i)$ 
     $\text{correct} = \text{correct} / \{p_i\}$ 

```

Let us assume that: (i) links are perfect, (ii) the failure detector is perfect and (iii) initially local variables are initialized as follows $\text{mysn}=0$ e $\text{delivered}=\emptyset$.

Answer to the following questions:

1. Does the $\text{xbroadcast}()$ primitive implement a Reliable Broadcast, a Best Effort Broadcast or none of the two?
2. Considering only the ordering property of broadcast communication primitives discussed during the lectures (FIFO, Causal, Total), explain which ones can be satisfied by the $\text{xbroadcast}()$ implementation.

Provide examples to justify your answers.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through FIFO perfect point-to-point links and are structured through a line (i.e., each process p_i can exchange messages only with processes p_{i-1} and p_{i+1} when they exist). Processes may crash and each process is equipped with a perfect oracle (having the interface $\text{new_right}(p)$ and $\text{new_left}(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Perfect failure detector primitive.

Ex 5: Consider a distributed system composed of n processes p_1, p_2, \dots, p_n connected through a ring topology. Initially, each process knows the list of correct processes and maintains locally a next variable where it stores the id of the following process in the ring.

Each process can communicate only with its next through FIFO perfect point-to-point channels (i.e. the process whose id is stored in the next variable).

Processes may fail by crash and each process has access to a perfect failure detector.

Write the pseudo-code of a distributed algorithm implementing a $(1, N)$ atomic register.

Ex 6: A transient failure is a failure that affects a process temporarily and that alter randomly the state of the process (i.e., when the process is affected by a transient failure, its local variables assume a random value).

Let us consider a distributed system composed by N processes where f_c processes can fail by crash and f_t processes can suffer transient failures between time t_0 and t_{stab} .

Let us consider the following algorithm implementing the Regular Reliable Broadcast specification

Algorithm 3.2: Lazy Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb ;

PerfectFailureDetector, **instance** \mathcal{P} .

```

upon event ⟨  $rb$ , Init ⟩ do
     $correct := \Pi$ ;
     $from[p] := [\emptyset]^N$ ;

upon event ⟨  $rb$ , Broadcast |  $m$  ⟩ do
    trigger ⟨  $beb$ , Broadcast | [DATA, self,  $m$ ] ⟩;

upon event ⟨  $beb$ , Deliver |  $p$ , [DATA,  $s$ ,  $m$ ] ⟩ do
    if  $m \notin from[s]$  then
        trigger ⟨  $rb$ , Deliver |  $s$ ,  $m$  ⟩;
         $from[s] := from[s] \cup \{m\}$ ;
        if  $s \notin correct$  then
            trigger ⟨  $beb$ , Broadcast | [DATA,  $s$ ,  $m$ ] ⟩;

upon event ⟨  $\mathcal{P}$ , Crash |  $p$  ⟩ do
     $correct := correct \setminus \{p\}$ ;
    forall  $m \in from[p]$  do
        trigger ⟨  $beb$ , Broadcast | [DATA,  $p$ ,  $m$ ] ⟩;

```

Answer to the following questions:

1. For every property of the Regular Reliable Broadcast specification, discuss if it guaranteed between time t_0 and t_{stab} and provide a motivation for your answer.
2. For every property of the Regular Reliable Broadcast specification, discuss if it eventually guaranteed after t_{stab} and provide a motivation for your answer.
3. Assuming that the system is synchronous, explain how you can modify the algorithm (no pseudo-code required) to guarantee that No Duplication, Validity and Agreement properties will be eventually guaranteed after t_{stab} .

Ex 7: A service is delivered through 3 components that sequentially handles the requests and it receives around 75 requests per second. The components, A, B and C, serve respectively 100, 80 and 90 requests per second. The clients are complaining about the response time of the system. There is the possibility either 1) to double one component employing a perfect load balancer or 2) to substitute one component with an improved one serving 40% more requests per second.

Which component upgrade and option reduce the response time the most?
Compute the expected response time.

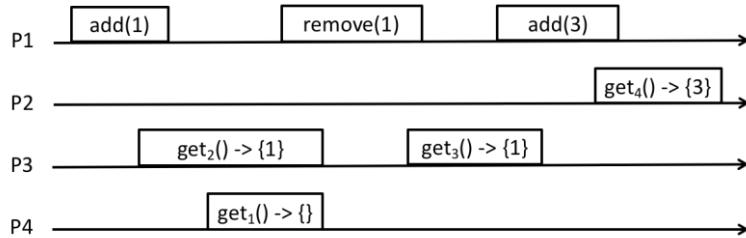
Ex 8: Consider a set object that can be accessed by a set of processes. Processes may invoke the following operations on the object:

- `add(v)`: it adds the value v in to the set
- `remove(v)` it removes the value v from the set
- `get()`: it returns the content of the set.

Informally, every `get()` operation returns all the values that have been added before its invocation and that have not been removed by any `remove()`.

For the sake of simplicity, assume that a value can be added/removed just once in the execution.

Consider the distributed execution depicted in the Figure



Answer to the following questions:

1. Is the proposed execution linearizable? Motivate your answer with examples.
2. Consider now the following property: “every `get()` operation returns all the values that have been added before its invocation and that have not been removed by any `remove()`. If an `add(v)/remove(v)` operation is concurrent with the `get`, the value v may or may be not returned by the `get()`”. Provide an execution that satisfy `get` validity and that is not linerizable.

Exercise 1

1. From the figure we have this CAUSAL ORDER:

$$m_1 \rightarrow m_2 \rightarrow m_4, m_3 \rightarrow m_4$$

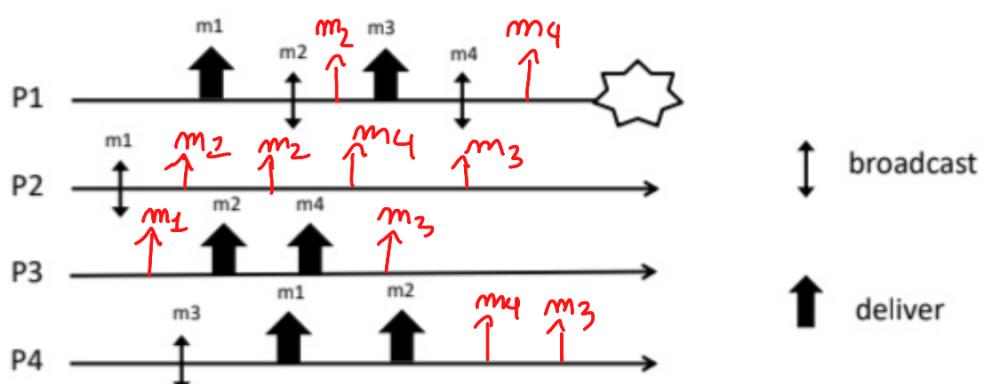
All possible deliveries that satisfy causal order and Total order:
for total order we must notice that m_3 must be delivered after m_1 (P_1), also m_2 must be delivered before m_4 (P_2)

$$m_1, m_2, m_3, m_4$$

$$m_1, m_3, m_2, m_4$$

(P_1 is faulty connect differently, but satisfies causal order)

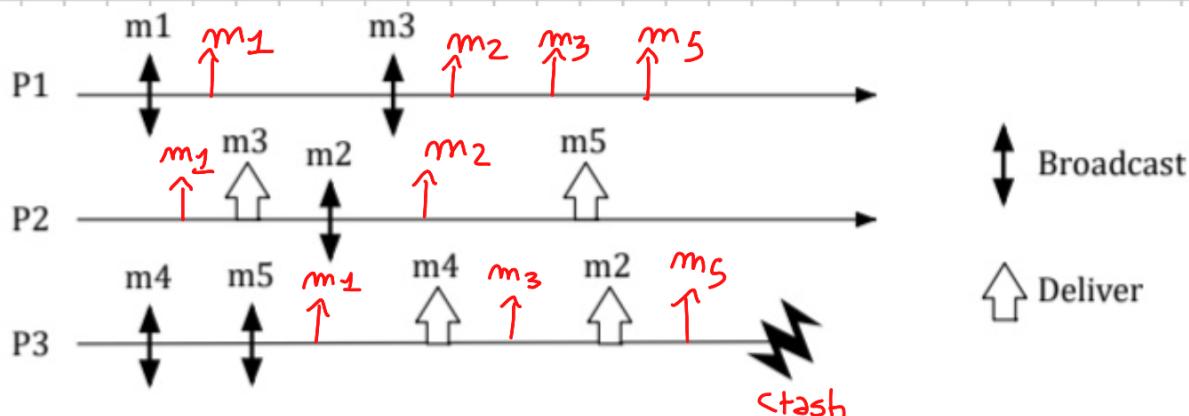
2. TO (UA, UNUTO), FIFO ORDER but not CAUSAL ORDER



FIFO ORDER: $m_2 \rightarrow m_4$ $m_4 \rightarrow m_3$

Exercise 2

1. RRB but not URB



P_3 faulty, m_4 is not delivered by P_3 it's on P_2 , not URB

2. FIFO: $m_4 \rightarrow m_5$, $m_1 \rightarrow m_3$

CAUSAL: $m_3 \rightarrow m_2$

All delivery satisfying FIFO Reliable Broadcast but not causal order for each process:

We must deliver also m_4 because we deliver m_5 for have FIFO order, we satisfy URB that implies RRB!

P_1 : $(m_1, m_2, m_3, m_4, m_5)$, $(m_1, m_2, m_4, m_3, m_5)$
 $(m_1, m_2, m_4, m_5, m_3)$, $(m_2, m_1, m_3, m_4, m_5)$
 $(m_2, m_1, m_4, m_3, m_5)$, $(m_2, m_1, m_4, m_5, m_3)$
 $(m_4, m_1, m_2, m_3, m_5)$, $(m_4, m_1, m_2, m_3, m_5)$
 $(m_4, m_1, m_2, m_5, m_3)$, $(m_4, m_2, m_1, m_3, m_5)$
 $(m_4, m_2, m_1, m_3, m_5)$, $(m_4, m_2, m_1, m_5, m_3)$
 $(m_1, m_4, m_2, m_3, m_5)$, $(m_1, m_4, m_2, m_5, m_3)$
 $(m_1, m_4, m_5, m_2, m_3)$, $(m_2, m_4, m_1, m_5, m_3)$
 $(m_2, m_4, m_5, m_1, m_3)$

P_2 : if we deliver m_4 before m_2 we violate the causal order in P_3 , for the other delivery we can have only combination but satisfy $m_4 \rightarrow m_5$ and $m_1 \rightarrow m_3$
EX.: $(m_4, m_2, m_3, m_1, m_5)$

P_3 : if in P_2 deliver m_4 before m_2 , we can have any combination, only respecting FIFO ORDER. EX. $(m_1, m_2, m_4, m_5, m_3)$

3. $m_4 \rightarrow m_5$ and $m_1 \rightarrow m_3 \rightarrow m_2$, in P_3 m_4 before m_2
 and in P_2 m_3 before m_5
- (m_4, m_1, m_3, m_2, m_5)
 - (m_1, m_4, m_3, m_2, m_5)
 - (m_1, m_3, m_4, m_2, m_5)
 - (m_4, m_1, m_3, m_5, m_2)
 - (m_1, m_4, m_5, m_3, m_2)

For each process, there are possible delivery.

Exercise 3

N processes : $\{P_1, \dots, P_N\}$

Upon event $X_{broadcast}(m)$

$$mySM = mySM + 1$$

$\forall p \in \text{correct}$

$P_{P_2} P_{send} ("MSG", m, mySM, myID)$

Upon event $P_{P_2} P_{Receive} ("MSG", m, sm, i)$

$$mySM = mySM + 1$$

If ($m \notin \text{delivered}$)

$\text{trigger } X_{Deliver}(m);$

$$\text{delivered} = \text{delivered} \cup \{m\}$$

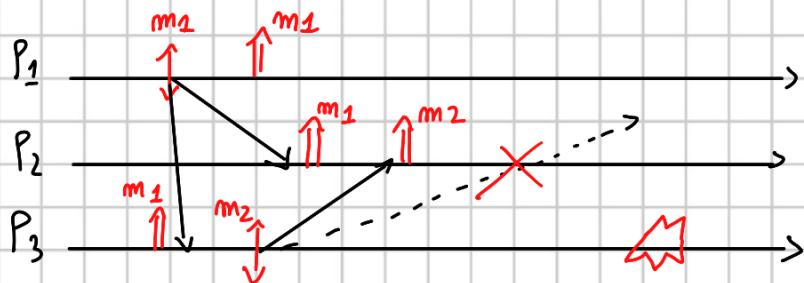
Upon event $\text{crash}(p_i)$

$$\text{correct} = \text{correct} / \{p_i\}$$

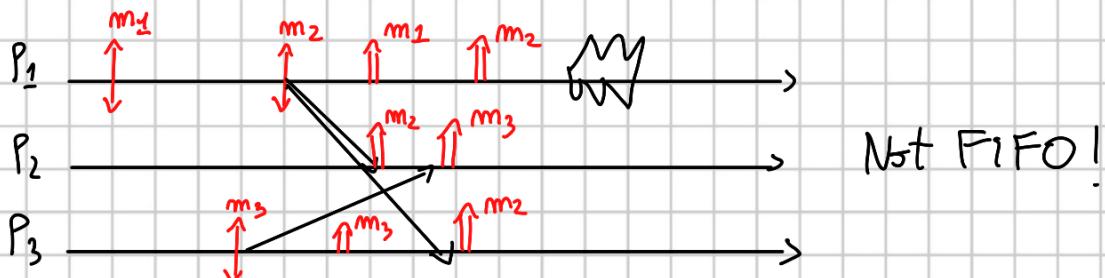
We have Perfect link, & perfect failure detector and
 in init: $mySM = 0$ and $\text{delivered} = \emptyset$

1. XBroadcast implements a Best Effort Broadcast because we satisfy No duplication, Validity, No creation given that we have P2P link, but don't satisfy RB because a message sent by a faulty process could be delivered only by a part of correct processes, we don't retransmitt.

Example: \uparrow X-broadcast \uparrow X-Deliver



2. Considering only the ordering, we have no guarantee on the ordered because we do a simple Broadcast, the sequence number is not used, given that don't satisfy nothing. This example is only BEB:



Exercise 4

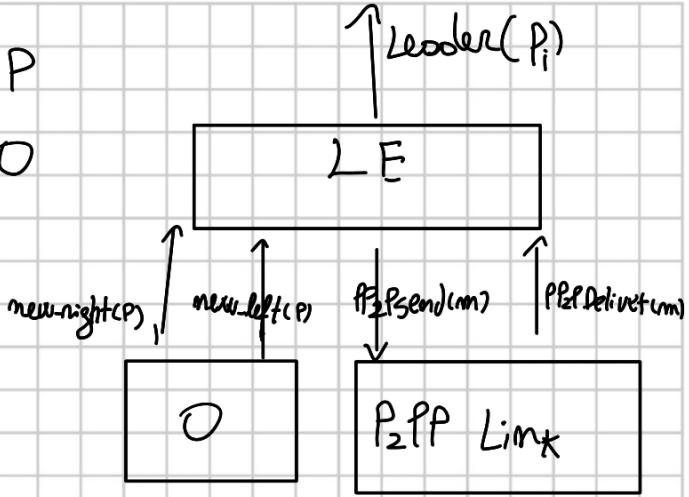
$\Pi = \{P_1, \dots, P_m\}$ m processes, FIFO P2P link, we have a line:



P_i can communicate only with P_{i-1} and P_{i+1} . We have on each with primitive next-right (P) and next-left (P). Implement a Perfect failure detector.

implements: Perfect Failure Detector, P

Uses: Oracle O, P2P link PL FIFO



Upon event $\langle P, \text{init} \rangle$ do

LEFT = get-left()

RIGHT = get-right()

alive = \top

detected = \emptyset

Upon event $\langle O, \text{new-right } | p \rangle$ do

Ex-right = RIGHT

RIGHT = P

detected = detected $\cup \{Ex\text{-RIGHT}\}$

trigger crash(Ex-right)

IF RIGHT \neq null do

trigger PP2PSend(|CRASH|, ExRight) to RIGHT

IF LEFT \neq null

trigger PP2PSend(|CRASH|, ExRight) to LEFT

Upon event $\langle O, \text{new-left } | p \rangle$ do

Ex-left = LEFT

LEFT = P

detected = detected $\cup \{Ex\text{-LEFT}\}$

trigger crash(Ex-left)

IF RIGHT \neq null do

trigger PP2PSend(|CRASH|, detected) to RIGHT

IF LEFT \neq null

trigger PP2PSend(|CRASH|, detected) to LEFT

Upon event $PP_2P\text{Deliver}(|CRASH|, P_i)$ from P_j

$\text{detected} = \text{detected} \cup P_i$

~~trigger $crash(P_i)$~~

IF $P_j = \text{LEFT}$ AND $P_j \neq \text{null}$

~~trigger $PP_2P\text{Send}(|CRASH|, \text{detected})$ to RIGHT~~

else IF $P_j = \text{RIGHT}$ AND $P_j \neq \text{null}$

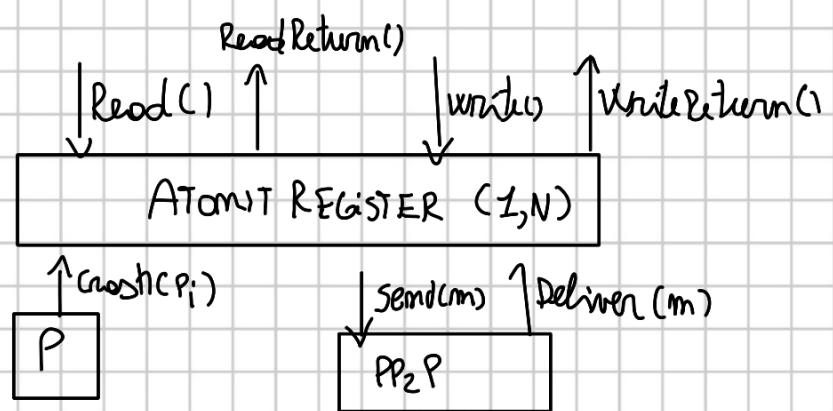
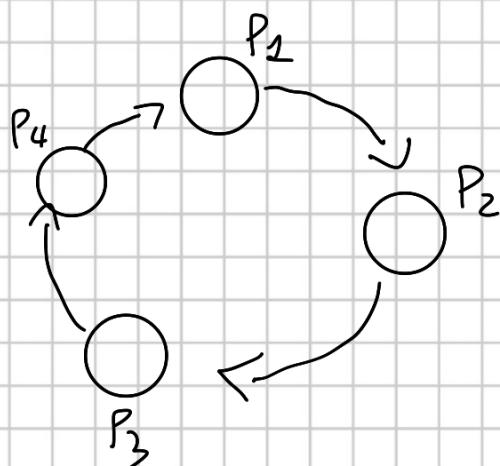
~~trigger $PP_2P\text{Send}(|CRASH|, \text{detected})$ to LEFT~~

X NO MULTIPLE
CRASH ADD

In this way we satisfy Strong Completeness and Accuracy given that we have PP_2P link and the oracle.

Exercise 5

Ring topology with N processes. We have PP_2P link and a perfect failure detector P :



Idea is to write the value on a token and send it in the ring, invoke $writeReturn()$ only when turn lock, and on read a process only when have the token. When a process crash and is our next, change next & recreate a new token.

Use sequence number for discard past token

implements: (1-N)-Atomic Register , AR

Users: PP2P link , PL

Perfect Failure Detector , P

int :

alive = π

val = \perp , sm = 0

sequence-number of token sem

Reading = False

Write = False

next = $P_{(i+1) \bmod N}$

If self = P_0

sm = sm + 1

TOKEN = val

trigger PP2PSend (<|READING|, sm>, TOKEN) to next

Upon event <AR, Write | v> do

write = TRUE

~ move when recieve token can WRITE

val = v

and wait for return

Upon event <AR, Read> do

Reading = TRUE

Upon event PP2P Deliver (<|READING|, s>, TOKEN)

If $s = sm$

do nothing , have received same TOKEN

else if WRITE = TRUE

TOKEN = val , sm = sm + 1

trigger PP2PSend (<|WRITING|, sm>, TOKEN) to next

else if READING = TRUE

val = TOKEN , READING = FALSE , sm = sm + 1

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

trigger Read Return (val)

else

val = TOKEN, sm = sm + 1

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

Upon event PP2P Deliver (<|WRITING1, ss>, TOKEN)

if sm = s

do nothing, have received same TOKEN

else if write = TRUE

sm = sm + 1

trigger Write Return ()

trigger PP2P Send (<|READING1, sm>, TOKEN) to next

else if READING = TRUE

val = TOKEN, READING = FALSE, sm = sm + 1

trigger PP2P Send (<|WRITING1, sm>, TOKEN) to next

trigger Read Return (val)

else

val = TOKEN, sm = sm + 1

trigger PP2P Send (<|WRITING1, sm>, TOKEN) to next

Upon event Crash (Pi)

alive = alive / {Pi}

if Pi = next

return the alive next

next = next_alive ()

trigger PP2P Send (<|READING1, sm>, TOKEN) to next



if next have same sm have received
the precedent token

Exercise 6

1.

No duplication: is not guaranteed because if a process that have a set of delivery, it is affected by a transient failure and his variable From[EP] is changed with other values, there is the possibility that in `ebroadcast` the condition $m \notin \text{From}[S]$ is satisfied multiple times and have multiple delivery of some message.

No creation: is not guaranteed because if a process is affected by a transient failure, its From[EP] is changed randomly, also with values not given by a broadcast. If the same process in the same interval time crash, then the other correct processes may deliver some wrong messages that are in From[EP], not previously broadcast by a process.

Validity: is guaranteed, because we consider a transient failure processes as actually failed, if a correct process p broadcast a message m, thanks to "`if $m \notin \text{From}[S]$ "`, p will deliver it at some time.

Agreement: is guaranteed because if we consider a transient failure processes as actually failed, all the correct processes will deliver the same set of messages, but some messages could be not satisfy no creation. This is guarantee by the From[S] variable that

check if the message is already delivered and by the fact that in every crash we broadcast all messages received from that processes, in this way all correct processes receive missing messages.

2. After t_{stab} the processes will not be affected by transient failure, but some processes are been effected, the algorithm is the same, now the precedent processes that was affected by transient failure return correct, for these also validity and agreement is not eventually guarantee.

Validity: if in the transient failure of p , $f_{from[p]}$ is charged with randomly messages and after t_{stab} that process want to broadcast a message and is already in $f_{from[p]}$, he will not deliver it, violating agreement.

Agreement: we have the same problem of validity, for example like we say for validity p could not deliver message m , but other correct processes could deliver it, and we have that correct processes have different set of messages, violating agreement.

3. For guarantee the No Duplication, validity and agreement after t_{stab} we can use the assumption that by N processes F_t can suffer transient failure. We use a Quorum

with $N - F_C$ (faulty not transient) messages, we send with every message the $f_{\text{from}}[p]$ variable, when we receive $N - F_C$ from $[p]$ we do an interception between all $f_{\text{from}}[p]$, and eventually we achieve all the property, because the missing messages are retransmitted and the randomizing is correctly.

$$E[R] \stackrel{\text{MM1}}{=} \frac{1}{100-75} + \frac{1}{160-95} + \frac{1}{90-75}$$

Exercise 7

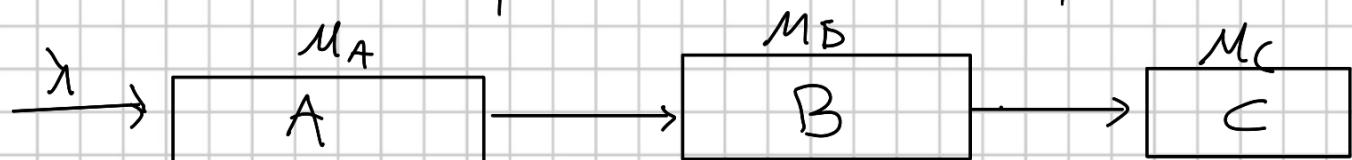
3 components handles sequentially the service.

$\lambda = 75 \text{ req/sec}$ is the arrival rate

$$\begin{aligned} M_A &= 100 \text{ req/sec} \\ M_B &= 80 \text{ req/sec} \\ M_C &= 90 \text{ req/sec} \end{aligned} \quad \left. \begin{array}{l} \\ \\ \end{array} \right\} \text{service rate}$$

1) double one component employing a perfect load balancer.

2) substitute one component with one 40% faster.



1) in the first case we have a perfect load balancer, we double the slowest component $M_B = 160 \text{ req/sec}$.

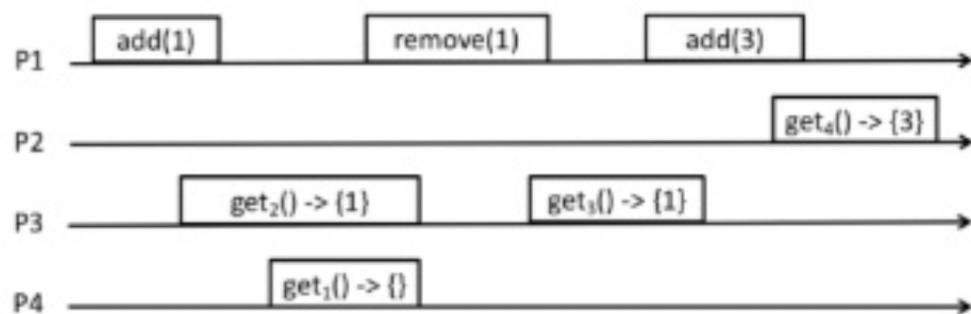
$$\begin{aligned} E[R] \stackrel{\text{MM1}}{=} & \frac{1}{100-75} + \frac{1}{160-95} + \frac{1}{90-75} = \\ & = 0,045 + 0,011 + 0,066 = 0,117 \end{aligned}$$

2) With second option upgrade always the slowest component and obtain $M_B = 112 \text{ reg/sec}$

$$E[R]^{M/M/2} = \frac{1}{25} + \frac{1}{37} + \frac{1}{15} = \\ = 0,045 + 0,0275 + 0,065 = 0,1375$$

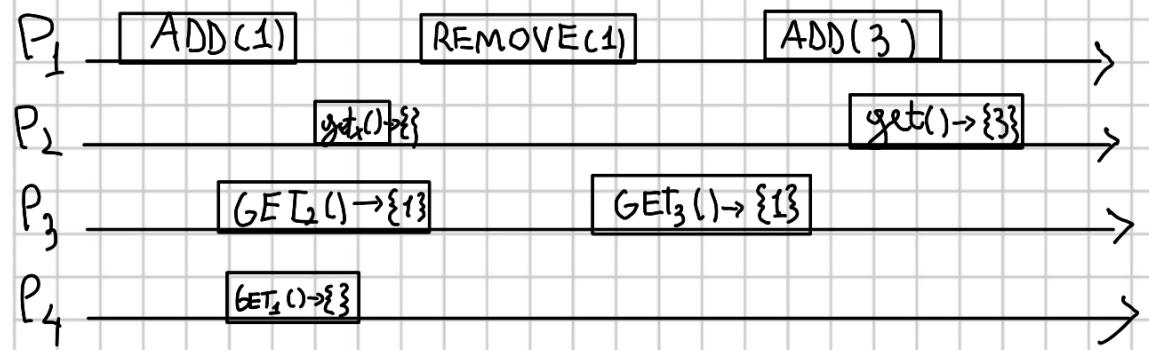
We chose option 1) and upgrade component B!

Exercise 8



1. No is not linearizable because of $\text{get}_1() \rightarrow \{\}$ and $\text{get}_3() \rightarrow \{3\}$ are in contradiction, we have no problem with $\text{get}_1() \rightarrow \{\}$, $\text{get}_2() \rightarrow \{1\}$ and Remove(1) because are concurrent and we can rearrange, but $\text{get}_3()$ we can drop only with remove(1). we can not provide a sequence is legal.

2. The execution with the get-validity provided in the text, it is linearizable, we provide a max execution that is not linearizable and satisfy get validity.



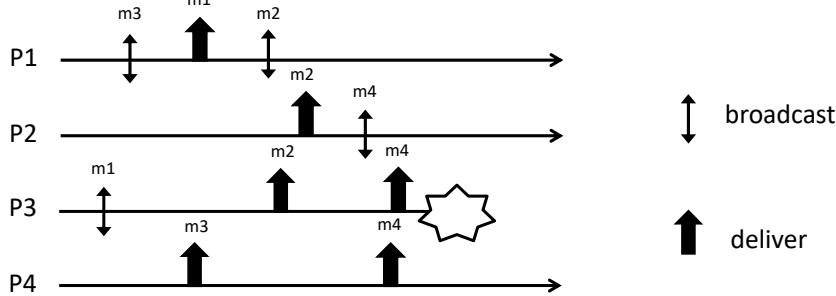
We change the execution, now we have three concurrent events get_x , get_z and get_y but get_x violate validity because return 1 but no after $ADD(1)$.

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 19 – Exercises
November 16th, 2022

Ex 1: Consider the partial execution depicted in the Figure



$$m_2 \rightarrow m_2$$

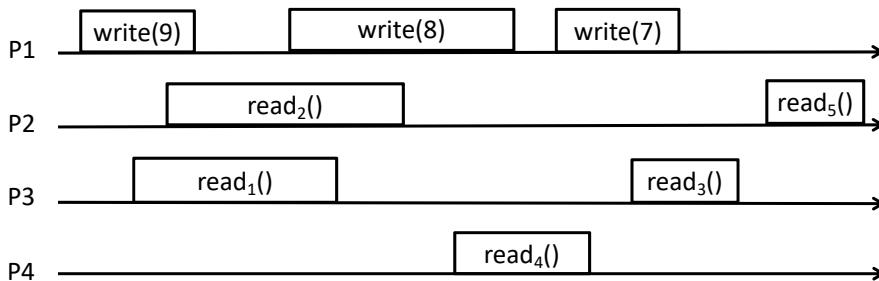
$$m_2 \rightarrow m_4$$

Answer to the following questions:

1. Provide ALL the possible delivery sequences that satisfies causal order and TO (UA, SUTO)
2. Complete the execution in order to have a run satisfying TO (UA WNUTO), FIFO order Broadcast but not Causal Order Broadcast
3. Complete the execution in order to have a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast and not satisfying Total Order.

NOTE: In order to solve the exercise, you can only add broadcast, deliveries and failures.

Ex 2: Consider the partial execution depicted in the Figure



Answer to the following questions:

1. Define ALL the values that can be returned by read operations (Rx) assuming that the run refers to a regular register.
2. Define ALL the values that can be returned by read operations (Rx) assuming that the run refers to an atomic register.

3. Assign to each read operations (R_x) a return value that makes the execution linearizable.

Ex 3: Let us consider the following algorithm

```

upon event < $frb$ , Init > do
   $lsn := 0;$ 
   $pending := \emptyset;$ 
   $next := [1]^N;$ 

upon event < $frb$ , Broadcast |  $m$  > do
  for each  $p \in \Pi$  do
    trigger < $Send$  | [DATA, self,  $m$ ,  $lsn$ ] > to  $p$ ;
   $lsn := lsn + 1;$ 

upon event < $Deliver$  |  $p$ , [DATA,  $s$ ,  $m$ ,  $sn$ ] > do
   $pending := pending \cup \{(s, m, sn)\};$ 

  while exists  $(s, m', sn') \in pending$  such that  $sn' = next[s]$  do
     $next[s] := next[s] + 1;$ 
     $pending := pending \setminus \{(s, m', sn')\};$ 
    trigger < $frb$ , Deliver |  $s, m'$  >;
  
```

Let us consider the following properties:

- **Validity:** If a correct process p broadcasts a message m , then p eventually delivers m .
- **No duplication:** No message is delivered more than once.
- **No creation:** If a process delivers a message m with sender s , then m was previously broadcast by process s .
- **Agreement:** If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- **FIFO delivery:** If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

Assuming that every process may fail by crash, address the following points:

1. Considering that messages are sent by using *perfect point to point links*, for each property mentioned, discuss if it satisfied or not and provide a motivation for your answer:
2. Considering that messages are sent by using *fair loss links*, for each property mentioned, discuss if it satisfied or not and provide a motivation for your answer.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links and are structured through a ring (i.e., each process p_i can exchange messages only with processes $p_{i+1 \text{ mod } n}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $new_next(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Uniform Reliable Broadcast communication primitive.

Ex 5: Let us consider the following algorithm implementing a (1, N) atomic register in synchronous system.

<pre> 1 upon event { onar, Init } do 2 (ts, val) := (0, ⊥); 3 correct := Π; 4 writeset := ∅; 5 readval := ⊥; 6 reading := FALSE; 7 uponevent(P,Crash p) do 8 correct := correct \ {p}; 9 upon event { onar, Read } do 10 reading := TRUE; 11 readval := val; 12 trigger { beb, Broadcast [WRITE, ts, val] }; 13 upon event { onar, Write v } do trigger { beb, Broadcast [WRITE, ts + 1, v] }; </pre>	<pre> 14 upon event { beb, Deliver p, [WRITE, ts', v'] } do 15 if ts' > ts then 16 (ts, val) := (ts', v'); 17 trigger { pl, Send p, [ACK] }; 18 upon event { pl, Deliver p, [ACK] } then 19 writeset := writeset ∪ {p}; 20 upon correct ⊆ writeset do 21 writeset := ∅; 22 if reading = TRUE then 23 reading := FALSE; 24 trigger { onar, ReadReturn readval }; 25 else 26 trigger { onar, WriteReturn }; </pre>
---	---

Assuming that messages are sent by using perfect point-to-point links and that the broadcast is best effort answer the following questions:

1. Discuss what does it happen to every atomic register property (i.e., termination, validity and ordering) if the failure detector is eventually perfect and not perfect
2. Discuss what does it happen to every atomic register property (i.e., termination, validity and ordering) if we change line 12 with **trigger** { beb, Broadcast | [WRITE, ts+1, val] };

Ex 6: Consider a distributed system composed of n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links. Processes are connected through a directed ring (i.e., each process p_i can exchange messages only with processes $p_{i+1(\text{mod } n)}$). Processes may crash and each process is equipped with a perfect oracle (having the interface $\text{new_next}(p)$) reporting a new neighbor when the previous one is failing.

Write the pseudo-code of an algorithm implementing a Leader Election primitive at every process p_i .

Exercise 1

1. $m_3 \rightarrow m_2 \rightarrow m_4$, $m_1 \rightarrow m_2$

$m_3 \rightarrow m_4$ for TO

m_3, m_1, m_2, m_4

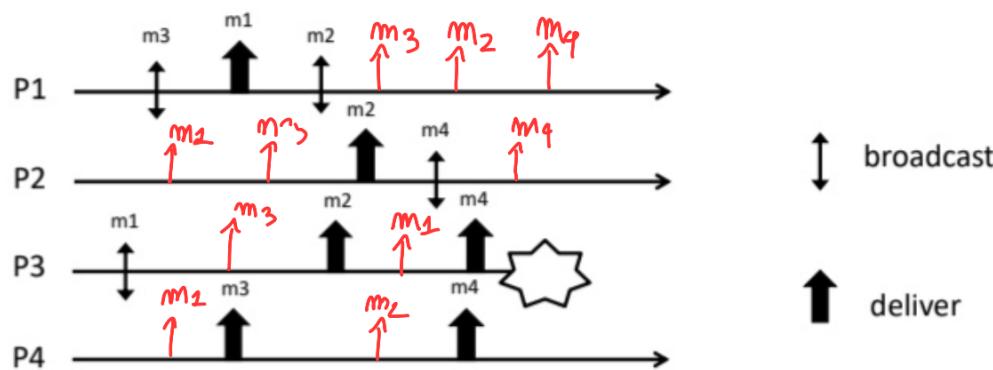
m_1, m_3, m_2, m_4

must be the possible sequences for all processes for satisfy TO(UA, SUT) and causal order.

2. $m_3 \rightarrow m_2$ FIFO

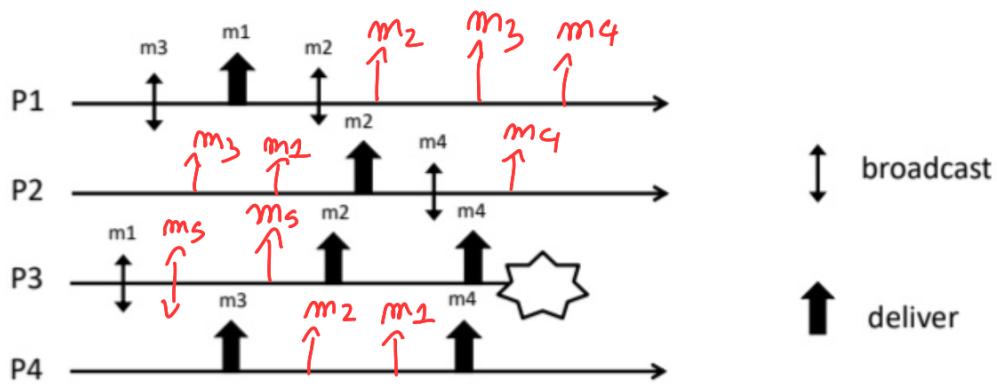
$m_1 \rightarrow m_2 \rightarrow m_4$ CAUSAL ORDER

$m_3 \rightarrow m_4$



Satisfy TO(UA, UNTO) and FIFO order broadcast but not causal, because all correct processes have some set of delivery and also faulty P3 this satisfy UA, UNTO because correct processes have some order of delivery but the faulty have an inversion ($m_3 \rightarrow m_1$). Not causal because in P3 we have $m_1 \rightarrow m_2$.

3. We show on execution that satisfy regular reliable broadcast but not uniform and not total order.



Without adding a broadcast in P3 and the delivery of m_5 we can not have RB but not have UA, because we must have that the faulty process have a different set of messages respect to corrects one.

Exercise 2

1. Rx values with regular register

$R_2 : 0, 9, 8$, $R_1 : 0, 9, 8$, $R_3 : 8, 7$, $R_4 : 9, 8, 7$, $R_5 : 7$

2. Rx values with an atomic register

$R_1 : 0, 9, 8$, $R_2 : 0, 9, 8$, $R_5 : 7$

$R_4 : 9, 8, 7$ or $8, 7$ because if R_1 or R_2 return 8 then R_4 cannot return 9.

$R_3 : 8, 7$ or 7 because if R_4 return 7 then R_3 cannot return 8.

3. With these assignment this execution is linearizable.

$R_1 : 0$, $R_2 : 9$, $R_4 : 8$, $R_3 : 7$, $R_5 : 7$

We respect causal order and also validity.

Exercise 3

1. Considering we have P2P link:

Validity: is not satisfied because all processes set $lm = 0$ and $next[1]^N$, when they do the first broadcast send the message with $sm = 0$ and increment lm only after the send. The message will be added in pending but never delivered because we never satisfy $0 = next[5]$.

No duplication: is satisfied because we use sequence number for messages, send the messages only one time in the broadcast and we continuously check if we have something in pending with $sm = next[5]$ and increment sm when deliver the message.

No creation: is satisfied because we have P2P link that guarantee this property and also by the algorithm we add in pending only messages sent by a previous broadcast.

Agreement: is not satisfied because if a faulty process broadcast a message m and is added in pending by some correct processes, they will deliver it, but they don't retransmit it or neither when a process crash, and for this if a correct process don't receive the same broadcast it never deliver m .

FIFO delivery: is not satisfied because the second message sent by any process with $sm = 1$ will be always delivered before the first because it will remain forever in pending with $sm = 0$, for the other

messages after the second the property is satisfied, but the second violate FIFO delivery.

2. Using fair loss links:

Validity: Some as before and also we don't have the guarantee that when we send the message it will arrive at destination, and the algorithm doesn't have any retransmission.

No duplication: is guarantee because we have the sequence number and the check in the while if we succeed in send the message in FLL. Otherwise we have no problem because we don't send the message and can not be delivered.

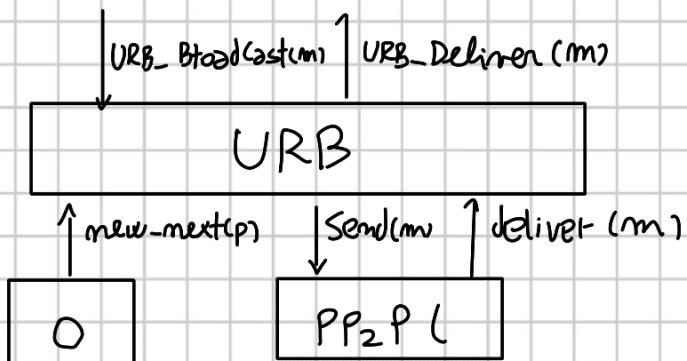
Agreement: is not satisfied like before for the problem with faulty processes and also we can lost the messages on the channel for these some faulty processes can deliver a message m but other correct processes could not deliver it.

No creation: is guaranteed by the FFL and by the algorithm itself

FIFO Delivery: is not satisfied because for the same problem that we show in 1., the second message is deliver before the first.

Exercise 4

m processes : $\Pi = \{P_1, \dots, P_m\}$, unique identifiers, we use PP₂P links and we have a ring structure. We can have crashes. we have an oracle with $\text{next_next}(p)$ when neighbor is failing. Implementing leader election for each P_i .



We use a token, every process that want to do a broadcast must have the token, and can deliver only when the token has complete one loop. We manage the crash retransmitting and check if the token is a duplication using the last delivery tag and S.M.

Implementation: Uniform Reliable Broadcast, URB

Used: Perfect Point-to-point link, PP₂P link

Oracle, O

init:

Pending = ⊥

next = $P_{i+1} \pmod m$

delivered = \emptyset

wait = FALSE

S.M. = 1 \longrightarrow sequence number

last_delivered = ⊥

If self = P_1 do

TOKEN = ⊥

{Trigger PP₂P Send (FREE | TOKEN, sm) to next}

Upon event URB-Broadcast(m)

WAIT = TRUE \rightarrow wait for token for Broadcast m

Pending = m

Upon event P2P Deliver (FREE1, <TOKEN, S>)

Sm = S + 1

IF WAIT = TRUE do

TOKEN = Pending,

trigger PP2PSend (BROADCAST1, <TOKEN, Sm>) to next

Pending = 1

else

trigger PP2PSend (FREE1, <TOKEN, S>) to next

Upon event P2P Deliver (BROADCAST1, <TOKEN, S>)

IF TOKEN = last-delivered do

IF Sm > S

< DISCARD> TOKEN is a DUPLICATIONS

else

Sm = S + 1

trigger PP2PSend (FREE1, <TOKEN, S>) to next

else IF WAIT = TRUE and pending = 1

WAIT = FALSE

delivered = delivered \cup TOKEN

last-delivered = TOKEN

trigger URB-Deliver (TOKEN)

TOKEN = 1, Sm = S + 1

trigger PP2PSend (FREE1, <TOKEN, Sm>) to next

else

$$sm = s + 1$$

$$\text{delivered} = \text{delivered} \cup \text{TOKEN}$$

$$\text{last-delivered} = \text{TOKEN}$$

trigger URB-DELIVR(TOKEN)

trigger PP2PSend(|BROADCAST|, <TOKEN, sm>) to next

Upon event next-next(p)

$$next = p$$

$$\text{TOKEN} = \text{last-delivered}$$

trigger PP2PSend(|BROADCAST|, <TOKEN, sm>) to next

We guarantee no duplication of TOKEN
checking CAST-DELIVER and the sequence number.
If the next that received the new TOKEN finds
that his sm is greater than the sender, discard
the TOKEN.

Exercise 5

(1-N) atomic register, P2P links and BEB.

1. we have eventually failure detector:

Termination: is satisfied because if a correct process is suspected for error then the condition $\text{correct} \subseteq \text{writerset}$ is satisfied with less ACK, but the write or read will reach the termination.

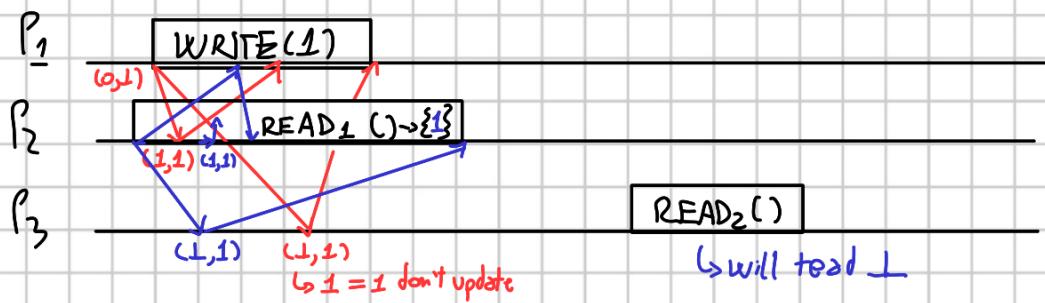
Validity: is not satisfied because if a process p that want to read is suspected for error and a process obs a `WriteReturn` because believe that p is crashed and also p doesn't receive the new value because writer has used BEB, it is possible that p that is not concurrent with last write, it read the precedent value violating validity.

Ordering: is not satisfied for the same motivation of validity. A correct process p that is been suspected would not receive the other messages and if he read a value could return a value that violating the correct order.

2. in line 12: trigger <Beb, Broadcast [WRITE, $t_0 + 1, w]$)

Termination: is satisfied because we change only the timestamp, the process will receive the ACK like before the change and satisfied " $\text{correct} \subseteq \text{writerset}$ "

Validity: is not satisfied because if we have this situation:



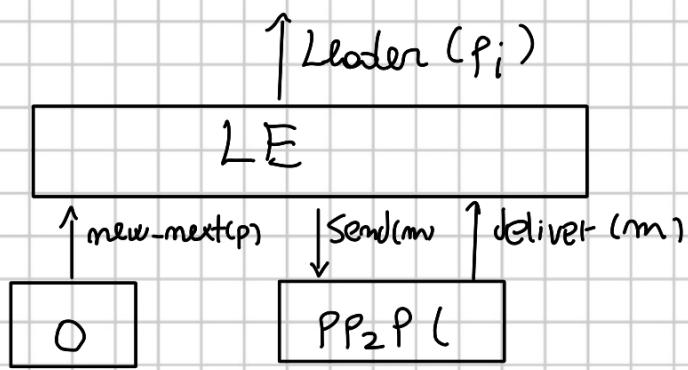
We can have 2 process after a write that return a value written before it, violate validity.

Agreement: is not satisfied like we see in the example
we can have a read_2 that is after read_1 that return a value that violate the ordering.

Exercise 6

n processes : $\Pi = \{P_1, \dots, P_n\}$, unique identifiers, we use P2P links and we have a ring structure. we can have crashes. we have an oracle with $\text{next_meet}(p)$ when neighbor is failing.

Implementing Uniform Reliable Broadcast:



We initialize P_m as the leader for all processes. When a process receives $\text{new-meet}(p)$, check if it is the leader, the next that is crashed, then if it is true become himself the leader and send to other the communication. otherwise send the actual leader to next.

then if it is true become himself the leader and send to other the communication. otherwise send the actual leader to next.

Implementation: Leader Election , LE

Used: Perfect Point to Point link, P2P link
Oracle , O

init:

$$\text{next} = P_{(i+1) \bmod n}$$

$$\text{alive} = \Pi$$

$$\text{leader} = P_m$$

$$\text{WAIT} = \text{FALSE}$$