

Upon event new_next (p)

IF leader > p AND self > p

leader = self

trigger leader (leader)

next = p

Wait = TRUE \rightarrow Wait all confirm new leader

trigger PP2PSend (NEW-LEADER, leader) to next

else

next = p

trigger PP2PSend (NEW-LEADER, leader)

Upon event PP2PDeliver (NEW-LEADER,)

IF leader = ()

IF wait = TRUE

wait = FALSE

< Stop resending m, i'm new leader >

else

< Discard message, it is already received the messages >

else

leader = ()

trigger leader (leader)

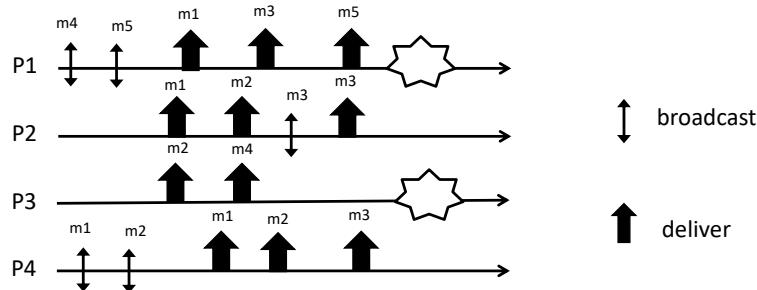
trigger PP2PSend (NEW-LEADER, leader) to next

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 23 – Exercises
November 23th, 2022
(Estimated time to complete all exercises: 3 hours)

Ex 1: Consider the execution depicted in the Figure

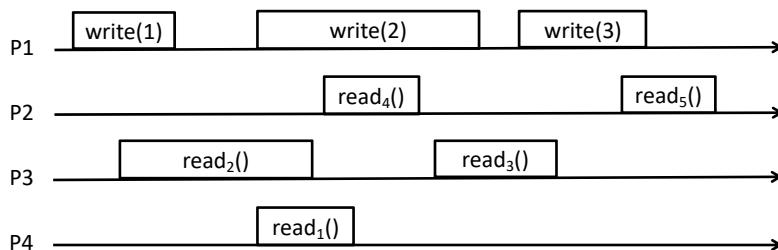


Answer to the following questions:

1. Which is the strongest TO specification satisfied by the proposed run? Motivate your answer.
2. Does the proposed execution satisfy Causal order Broadcast, FIFO Order Broadcast or none of them?
3. Modify the execution in order to satisfy TO(UA, WUTO) but not TO(UA, SUTO).
4. Modify the execution in order to satisfy TO(NUA, WNUTO) but not TO(UA, WNUTO).

NOTE: In order to solve point 3 and point 4 you can only add messages and/or failures.

Ex 2: Consider the execution depicted in the following figure and answer the questions



1. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to a regular register.

2. Define ALL the values that can be returned by read operations (Rx) assuming the run refers to an atomic register.
3. Let us assume that values retuned by read operations are as follow: $\text{read}_1() \rightarrow 2$, $\text{read}_2() \rightarrow 2$, $\text{read}_3() \rightarrow 3$, $\text{read}_4() \rightarrow 1$, $\text{read}_5() \rightarrow 3$. Is the run depicted in the Figure linearizable?

Ex 3: Consider the algorithm shown in the Figure

<pre> upon event { Init } do delivered := Ø; pending := Ø; correct := Π; forall m do ack[m] := Ø; upon event { urb, Broadcast m } do pending := pending ∪ {(self, m)}; trigger { beb, Broadcast [DATA, self, m] }; upon event { beb, Deliver p, [DATA, s, m] } do ack[m] := ack[m] ∪ {p}; if (s, m) ∈ pending then pending := pending ∪ {(s, m)}; trigger { beb, Broadcast [DATA, s, m] }; </pre>	<pre> upon event {◊P,Suspect p} do correct := correct \ {p}; upon event {◊P,Restore p} do correct := correct ∪ {p}; function candeliver(m) returns Boolean is return (correct ⊆ ack[m]); upon exists (s, m) ∈ pending such that candeliver(m) do delivered := delivered ∪ {m}; trigger { urb, Deliver s, m }; </pre>
---	--

Assuming that the algorithm is using a Best Effort Broadcast primitive and an Eventually Perfect Failure Detector $\diamond P$ discuss if the following properties are satisfied or not and motivate your answer

- *Validity*: If a correct process p broadcasts a message m, then p eventually delivers m.
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s, then m was previously broadcast by process s.
- *Uniform agreement*: If a message m is delivered by some process (whether correct or faulty), then m is eventually delivered by every correct process.

Ex 4: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_N\}$ with unique identifiers that exchange messages through perfect point-to-point links and are structured in a ring topology (i.e., each process p_i can exchange messages only with processes $p_{(i+1) \bmod N}$ and stores its identifier in a local variable `next`).

Each process p_i knows the initial number of processes in the system (i.e., every process p_i knows the value of N).

1. Assuming that processes are not going to fail, write the pseudo-code of an algorithm that implements a $(1, N)$ regular register.

Ex.5: Answer true or false to the following claims providing a motivation:

1. The performance of a system can be analyzed independently from its load.
2. Let us assume a single service with a single class of requests (i.e. a single workload component). If we are under the stability condition ($\lambda < \mu$) then the expected response time of the system is independent from the arrival pattern.
3. The workload parameters that mostly influence the performance of system are the arrival pattern and the service demands.
4. The time needed to run a simulation is upper-bounded by the simulated time
5. The reliability function $R(t)$ associated to a component may increase after the restoration of a component

Ex.6: A service is provided through 3 types of components (e.g. database, webserver, and an application), referred with A, B and C. At least one instance of each component must be available to provide the service. One only instance is available for components A and C, whereas two instances of component B are available. Assuming that all failures and restorations are independent among them, that the failures rates of the three components A, B and C are respectively 0.2 , 0.5 and 0.3 faults per day, that the mean time to repair the components A, B, and C are respectively 2h, 3h, and 1.5h, is the service available at least 98% of the time? If not, is the availability target met deploying an additional instance of one of the components?

Ex.7: Let us assume a probabilistic version of the Eager Reliable Broadcast protocol in which every process retransmits a rb-delivered message only once with probability 0.5. The rb-broadcast operation is periodically triggered, and the calls follows an exponential distribution with average time between consecutive calls of 10 seconds. Assuming a distributed system of 30 processes and no processing delay, what is the minimum in-channel capacity per process (in terms of message per seconds) to ensure that the expected time for a rb-delivery is below 0.5 seconds?

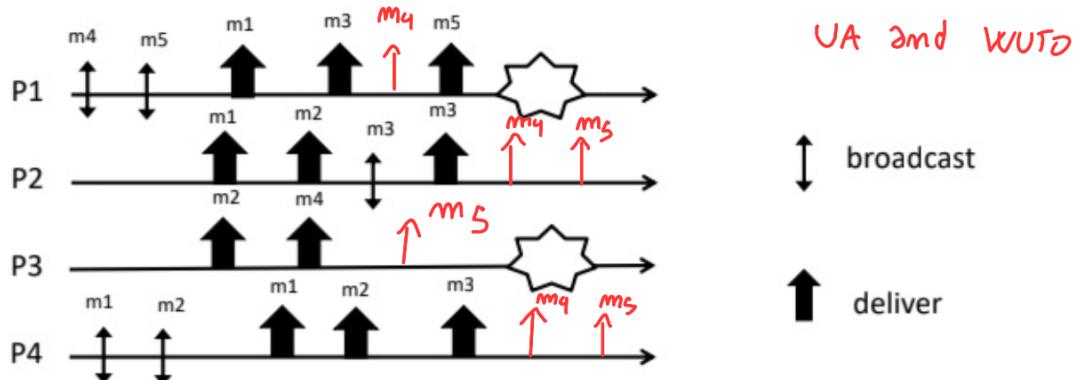
Exercise:

1. We don't have UA because P_1 is faulty process have a set of delivery that is not a subset of the correct one, for m_5 . We have WUTO because the order of pair is respected but not SUTO because we have holes in delivery:

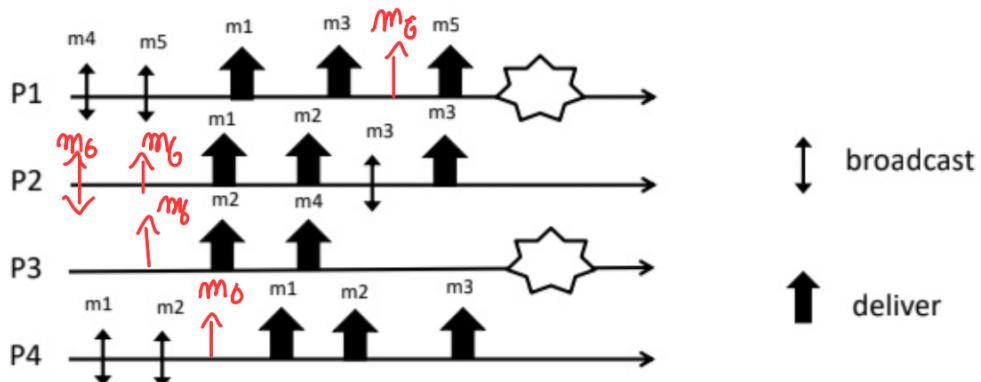
TO(NUA, WUTO)

2. For Fifo we have $m_4 \rightarrow m_5$, $m_1 \rightarrow m_2$
For causal order also $m_1 \rightarrow m_3$ and $m_2 \rightarrow m_3$
Causal order and FIFO broadcast are satisfied
because the order in delivery is satisfied in all
correct processes.

3.



4.



Exercise 2

1. $R_1 = \{1, 2\}$, $R_2 = \{0, 1, 2\}$, $R_3 = \{1, 2, 3\}$, $R_4 = \{1, 2\}$, $R_5 = \{2, 3\}$
2. $R_2 = \{0, 1, 2\}$, R_4 depend on R_2 if $R_2 = 2$ then $R_4 = 2$ else $R_4 = \{1, 2\}$, R_1 is concurrent with R_2 and R_4 then can return 1, 2 independently from the other 2, R_3 depend on the value before if R_1 or R_4 or R_2 return 2 then $R_3 = \{2, 3\}$ else $\{1, 2, 3\}$, R_5 depend on R_3 if $R_3 = \{3\}$ then $R_5 = \{3\}$ else $\{2, 3\}$.
3. $R_1 = 2$, $R_2 = 2$, $R_3 = 3$, $R_4 = 1$, $R_5 = 3$ is not linearizable because R_2 and R_4 are not concurrent R_4 return 1 after $R_2 = 2$, violating linearizability.

Exercise 3

Validity: is satisfied because when we broadcast a message p we add it in pending, and wait that the ACK associated to that message is equal to the number of correct processes. We have two case we wrongly suspect a wrong process and deliver the message or we receive all ACKS and deliver, but in any case we deliver p .

No duplication: is not guaranteed because we use pending and don't delete any messages from it, if we wrongly suspect a process, we would deliver the message without an ACK, but when restore the process and arrive his ACK the undeliver(p) check is true, and given that the message is in pending, we will deliver it again.

No Cetision: also it's guarantee we deliver only messages in pending that by the algorithm have only messages previously broadcasted.

Uniform agreement: is violated because if for example we wrongly suspect all other processes and in that period we broadcast a message p , we deliver after receive the delivery of B , but when we restore the other correct processes we don't retransmit again p then we have a delivery that other correct processes don't have.

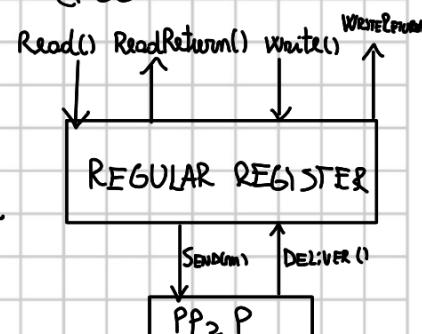
Exercise 4

n processes $\Pi = \{p_1, p_2, \dots, p_n\}$, P2P link, each p_i have $\text{next} = p_{(i+1) \bmod n}$, No failure.

We have only a writer for example p_1 that send the value and wait for return for do the `writeReturn`, the other can read their local value at any time.

Implementation: Regular Register ($1, n$), RR

Used: Perfect Point to Point link, P2P Link



init:

$val = 1$

$\text{next} = p_{(i+1) \bmod n}$

Upon event $\langle \text{RR}, \text{Read} \rangle$ **do**

trigger $\langle \text{RR}, \text{ReadReturn} | val \rangle$

upon event <RR, write, v> do

val = v

trigger PP2PSend (IWRITEI, val) to next

upon event PP2PDeliver (IWRITEI, v) from next

IF val >= v do

trigger writeReturn <val>

else

val = v

trigger PP2PSend (IWRITEI, val) to next

When we want read, we read local value and do readReturn. The writer update value, and send the value in the ring, when return the value do the write return, each reader simply update value and send to next.

Exercise 5

1. False, the performance of a system depend heavily on the characteristics of its load, because the expected performances depend on the characteristic of the inputs, we must define very well the workload.
2. λ is not enough to characterize the workload of a system, the response time depend on the arrival pattern \Rightarrow False
3. True, they are fundamental, because they characterize the arrival pattern, given the information about queuing. The service demands describe the demand on the system.

The service time per completion is important but it is part of service demands, and the response time is not part of workload

4. False, generally there is no relationship between simulated time and the time needed to run a simulation on the computer, the simulation clock is different from the clock of computer.

5. False, Reliability doesn't depend on recovery.

Exercise 6

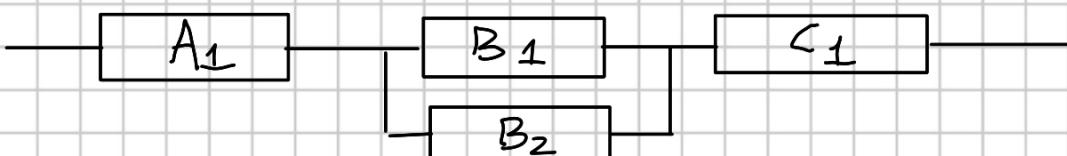
instance for each

3 components A, B, C. At least one available for provide the services. A, C one instance, B two instances.

$\lambda_A = 0.2$, $\lambda_B = 0.5$, $\lambda_C = 0.3$. $MTTR_A = 2 \text{ h}$, $MTTR_B = 3 \text{ h}$, $MTTR_C = 1.5 \text{ h}$

Service is available at least 98% of the time? we achieve it adding one instance to A or B or C?

Because service is guaranteed also if only one instance is available we have this situation:



$$A_{\text{serial}} = \prod A_i = A_A \cdot A_B \cdot A_C$$

$$A_{\text{parallel}} = 1 - \prod (1 - A_i) \Rightarrow A_B = 1 - (1 - A_{B_1})(1 - A_{B_2})$$

$$\lambda = \frac{1}{MTBF} \Rightarrow MTBF = \frac{1}{\lambda}, A_i = \frac{MTBF}{MTBF + MTTR}$$

$$MTBF_A = \frac{1}{0.2 \text{ ft/day}} = 120 \text{ h} \quad MTBF_B = \frac{1}{0.5 \text{ hours}} = 24 \text{ h}$$

$$MTBF_C = \frac{1}{0.3 \text{ hours}} = 20 \text{ h}$$

$$A_A = \frac{120}{2+120} = 0,983$$

$$A_{B_1} = A_{B_2} = \frac{40}{40+3} = 0,94$$

$$A_B = 1 - (1 - 0,94)^2 = 0,9964$$

$$A_C = \frac{80}{1,5+80} = 0,981$$

$$A_{TOT} = A_A \cdot A_B \cdot A_C = 0,96 < 0,98$$

No with this inputs not satisfied 98%

Try to optimize the bottleneck C with two instances:

$$A_C = 1 - (1 - 0,981)^2 = 0,999$$

$$A_{TOT} = A_A \cdot A_B \cdot A_{C_{new}} = 0,98$$

We achieve objective adding a instance to C.

Exercise 7

Probabilistic Eager Reliable Broadcast, t_b -delivered retrasmited once with probability 0.5. t_b -broadcast is triggered every 10 seconds, 30 processes and no delay. message per second capacity per process what is the minimum capacity for it to ensure Expected time for a r_b -delivery is less than 0.5 seconds?

We use a queue model with:

$$R = \text{EXPECTED TIME TO DELIVER} = \frac{1}{M_i - \frac{30 \cdot \frac{1}{10} \cdot \frac{1}{2}}{\downarrow \text{PROCESSES}}} \hookrightarrow 0.5 \text{ SECONDS}$$

$$\frac{1}{M_i - \frac{3}{2}} < \frac{1}{2} \Rightarrow \frac{2}{2M-3} < \frac{1}{2}$$

$$2M-3 > 4 \Rightarrow M > \frac{7}{2} = 3.5 \text{ messages/sec}$$

Dependable Distributed Systems
Master of Science in Engineering in Computer Science

AA 2022/2023

Lecture 27 – Exercises
November 30th, 2022

Ex 1: Let us consider a distributed system composed of N processes executing the algorithm reported in figure

Algorithm 3.12: Broadcast with Sequence Number

Implements:

FIFOReliableBroadcast, **instance** *frb*.

Uses:

ReliableBroadcast, **instance** *rb*.

```
upon event <frb, Init > do
    lsn := 0;
    pending := {};
    next := [1]N;
    
upon event <frb, Broadcast | m > do
    lsn := lsn + 1;
    trigger <rb, Broadcast | [DATA, self, m, lsn] >;
    
upon event <rb, Deliver | p, [DATA, s, m, sn] > do
    pending := pending ∪ {(s, m, sn)};
    while exists (s, m', sn') ∈ pending such that sn' = next[s] do
        next[s] := next[s] + 1;
        pending := pending \ {(s, m', sn')};
        trigger <frb, Deliver | s, m' >;
```

Let us assume that

1. up to f processes may be Byzantine faulty and
2. A Byzantine process is not able to compromise the underline Reliable Broadcast primitive (i.e., when a Byzantine process sends a message through the *rbBroadcast* interface, the message will be reliably delivered to every correct process).

For each of the following properties, discuss if it can be guaranteed when $f=1$ and motivate your answer (also by using examples)

- *Validity:* If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication:* No message is delivered more than once.
- *No creation:* If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement:* If a message m is delivered by some correct process, then m is eventually delivered by every correct process.
- *FIFO delivery:* If some process broadcasts message m_1 before it broadcasts message m_2 , then no correct process delivers m_2 unless it has already delivered m_1 .

Ex 2: Let us consider a distributed system composed of N processes executing the algorithm reported in figure

Algorithm 3.2: Lazy Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb ;
PerfectFailureDetector, **instance** \mathcal{P} .

```

upon event  $\langle rb, Init \rangle$  do
     $correct := \Pi$ ;
     $from[p] := [\emptyset]^N$ ;

upon event  $\langle rb, Broadcast \mid m \rangle$  do
    trigger  $\langle beb, Broadcast \mid [DATA, self, m] \rangle$ ;

upon event  $\langle beb, Deliver \mid p, [DATA, s, m] \rangle$  do
    if  $m \notin from[s]$  then
        trigger  $\langle rb, Deliver \mid s, m \rangle$ ;
         $from[s] := from[s] \cup \{m\}$ ;
    if  $s \notin correct$  then
        trigger  $\langle beb, Broadcast \mid [DATA, s, m] \rangle$ ;

upon event  $\langle \mathcal{P}, Crash \mid p \rangle$  do
     $correct := correct \setminus \{p\}$ ;
    forall  $m \in from[p]$  do
        trigger  $\langle beb, Broadcast \mid [DATA, p, m] \rangle$ ;

```

Let us assume that up to f processes may be Byzantine faulty with a symmetric behaviour i.e., a Byzantine process can change the content of the message it is going to send, but it cannot send different values to different processes when invoking the $bebBroascast$ (they can lie but in a consistent way).

For each of the following properties, discuss if it can be guaranteed when $f=1$ and motivate your answer (also by using examples)

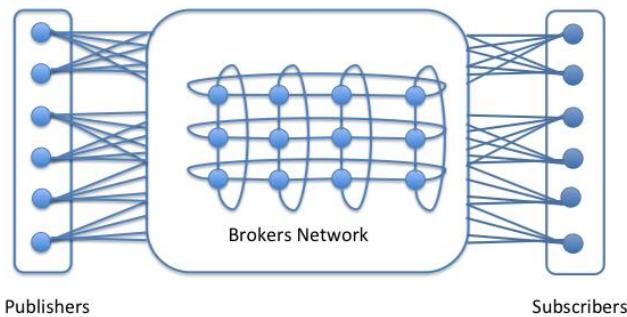
- *Validity:* If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication:* No message is delivered more than once.
- *No creation:* If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement:* If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Ex 3: Consider a distributed system constituted by n processes $\Pi = \{p_1, p_2, \dots, p_n\}$ with unique identifiers that exchange messages through perfect point-to-point links and are arranged in a unidirectional ring (i.e., each process p_i can exchange messages only with process $p_{(i+1) \bmod n}$ and stores its identifier in a local variable `next`).

Each process p_i knows the initial number of processes in the system (i.e., every process p_i knows the value of n).

1. Assuming that processes are not going to fail, write the pseudo-code of an algorithm that implements a consensus primitive
2. Let's now assume that processes may crash and that each correct process has access to a perfect failure detector. Discuss the issues of the implementation provided in point 1 and describe how you should modify the algorithm to make it fault tolerant
3. Considering the algorithm proposed in point 1, discuss which properties may be compromised by the presence of one Byzantine process in the ring.

Ex 4: Let us consider a distributed system composed by publishers, subscribers and brokers. Processes are arranged in a network made as follows and depicted below:



1. Each publisher is connected to k brokers through perfect point-to-point links;
2. Each subscriber is connected to k brokers through perfect point-to-point links;
3. Each broker is connected to k brokers through perfect point-to-point links and the resulting broker network is k -connected (4 -connected in the example);

Answer to the following questions:

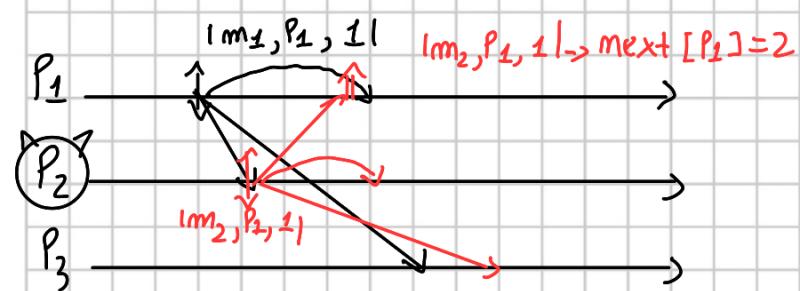
1. Write the pseudo-code of an algorithm (for publisher, subscribers and broker nodes) implementing the subscription-flooding dissemination scheme assuming that processes are not going to fail.
2. Discuss how many crash failures the proposed algorithm can tolerate.
3. Modify the proposed algorithm in order to tolerate f Byzantine processes in the broker network and discuss the relation between f and k .

Exercise 1

N processes, up to F processes may be Byzantine faulty and Byzantine process is not able to compromise the RB primitive. LINK ARE NOT AUTHENTICATED

For $F=1$:

Validity: is not guaranteed, because if a correct process P_1 send a broadcast a message m , $sm=1$, but the Byzantine process send a message m_2 with the same sm , like this:

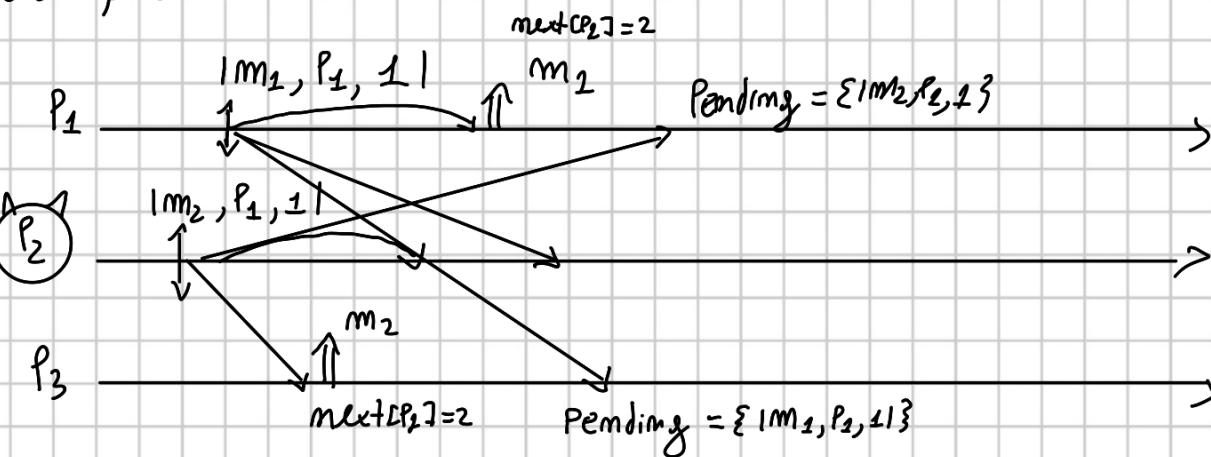


We have that P_2 receive $(m_2, P_2, 1)$, put it in pending and because $sm=1 = next[P_2]$, he will deliver it, but when receive $(m_1, P_1, 1)$ the original message will not be delivered because now $next[P_1] = 1$.

No duplication: is not guaranteed because like before a process P_1 broadcast a message $(m_1, P_1, 1)$, when the Byzantine process receive it, he can broadcast the message $(m_2, P_1, 2)$ that will be delivered by all correct processes because have the next sequence number, also the original message will be delivered with $sm=1$.

No creation: not guaranteed, the motivation is exactly the same for validity, Byzantine process can impersonate another process and achieve the delivery of the message.

Agreement: is not guaranteed because we can share this example:



P_1 deliver m_1 but P_3 deliver m_2 , violating that must deliver the same set of messages. The messages in pending will not be delivered because for the $SM=1$.

FIFO delivery: is not satisfied because byzantine process can broadcast a message m_3 with the same sequence number of m_1 , achieve the delivery and after m_2 will be delivered without m_1 .

Exercise 2

For $f = 1$:

Validity: it is satisfied because we have $from[p]$ variable and when we receive a deliver of a broadcast, we check that m is in the list, the byzantine process can only send the same message impersonating the correct process p or another message, but m is always delivered when we check $m \notin from[s]$.

No duplication: is not guaranteed because if a process i deliver a message m , if the

Byzantine process send the same message but changing the sender, the process i deliver again the message because not from [S] not contain m with the sender.

No creation: not guarantee because we don't use authenticated links, the real sender p is not checked, the byzantine process can create false broadcast with fake sender.

Agreement: is guarantee because is true that byzantine can change the content of the message that he send, but have a symmetric behaviour, we achieve agreement thanks to the retransmission of messages of a crashed process.

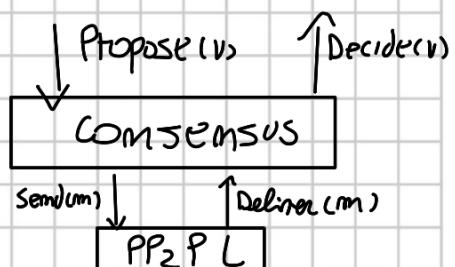
Exercise 3

n processes $\Pi = \{p_1, \dots, p_n\}$, P2P links, ring.

1. no failure, implements consensus primitive, we use a token where each process propose a value and after token is full we take the greater value.

Implementation: Consensus, C

Uses: P2P links, PL



init

state = wait

next = $p_{(i+1) \bmod N}$

$\Pi = n$

IF $\text{self} = p_0$:

$$\text{TOKEN} = \text{TOKEN} \cup \{v\}$$

trigger $\text{PP}_2\text{P}\text{Send}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{P}\text{Deliver}(\text{TOKEN})$

IF $\text{size}(\text{TOKEN}) == n$

$$\text{decision} = \text{max}(\text{TOKEN})$$

trigger $\text{Decide}(\text{decision})$

trigger $\text{PP}_2\text{P}\text{Send}(\text{decision})$ to next

else

$$\text{TOKEN} = \text{TOKEN} \cup \{v\}$$

trigger $\text{PP}_2\text{P}\text{Send}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{P}\text{Deliver}(\text{decision})$

IF $\text{State} == \text{WAIT}$

$\text{State} = \text{done}$

trigger $\text{Decision}(\text{decision})$

trigger $\text{PP}_2\text{P}\text{Send}(\text{decision})$ to next

else

<reach consensus>

2. We have a problem with crashes, messages can be lost and not reach consensus, we must add:

init

$$\text{pending} = \emptyset$$

$$\text{decision} = \text{null}$$

state = wait

next = $P_{(i+1) \bmod N}$

IT = n

IF $\text{self} = P_0$:

$\text{TOKEN} = \text{TOKEN} \cup \{V\} \rightarrow \text{pending} =$

$\text{pending} = \text{TOKEN}$ my

trigger $\text{PP}_2\text{PSend}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{PDeliver}(\text{TOKEN})$

IF $\text{TOKEN} == \text{pending}$:

< discard duplicate token >

object

else IF $\text{size}(\text{TOKEN}) == \overline{n}$

number of c

$\text{decision} = \max(\text{TOKEN})$

processes

trigger $\text{Decide}(\text{decision})$

trigger $\text{PP}_2\text{PSend}(\text{decision})$ to next

else

$\text{TOKEN} = \text{TOKEN} \cup \{V\}$

$\text{pending} = \text{TOKEN}$

trigger $\text{PP}_2\text{PSend}(\text{TOKEN})$ to next

Upon event $\text{PP}_2\text{PDeliver}(d)$

IF $\text{State} == \text{WAIT}$ AND $\text{decision} == \text{null}$

$\text{state} = \text{done}$, $\text{decision} = d$

trigger $\text{Decision}(d)$

trigger $\text{PP}_2\text{PSend}(d)$ to next

else

< reach Consensus > or < duplicate token >

Upon event crash(p)

$$\Pi = \Pi - 1$$

If $p == \text{next}$

$\text{next} = \text{next connect()}$

If $\text{decision} = \text{null}$

$\text{TOKEN} = \text{pending}$

trigger $\text{PP}_2 \text{Psend}(\text{TOKEN})$ to next

else

trigger $\text{PP}_2 \text{Psend}(\text{decision})$ to next

We have odd pending where we store the token value we have and decision where we store the decided value. In case of crash we update alive processes and if the next is crashed, we update next and resend the token or the decision. We check in deliver that the message is a duplicate.

3 Taking algorithm 1., we loose termination because the byzantine process can decide to not resend the token or the decision. We loose validity because byzantine process can change the TOKEN arbitrary and a process p can decide a value V never proposed. Integrity is guaranteed because correct processes have variable wait, we must add that is not uniform because of Byzantine. Agreement is not guaranteed because

if a process decide a value v and the next is the buy-in time process, it can send another decision V_1 .

Exercise 4 publish, subscribe, notify, unsubscribe

1. In subscription-flooding dissemination scheme, we have that each publisher when want to publish a value, he broadcast the value to all the K brokers which he is connected. The subscribers when want to subscribe send the message to each of the K broker which he is connected. When a broker receive a publish event start a notify event flooding and save the value, but before check if has saved a subscription with that value, and in affirmative we send the value to the subscriber. Subscriber get the notify event, and save the value.

implementation: subscription-flooding

Uses: P2P links, PL

init:

delivered = \emptyset Publisher, subscriber and broker.

neighbours = $\{b_1, \dots, b_K\}$ Publisher and subscriber

↳ brokers can have links also with publisher and subscribers

subscription = \emptyset subscriber and broker

(subscription is an array of couples that are composed like (constraint, subscriber) where constraint is a function that take a value and return a boolean)

Upon event publish (v)

if $v \notin$ delivered :

delivered = delivered $\cup \{v\}$

for each $m \in$ neighbours : // send to brokers

trigger PP2PSend (NOTIFY, v) to m

Upon event subscribe (constraint)

subscription = subscription $\cup \{(\text{constraint}, \text{self})\}$

for each $m \in$ neighbours : // send to brokers

trigger PP2PSend (SUBSCRIPTION | <constraint, self>) to m

upon event PP2PDeliver (SUBSCRIPTION | <C, p>) from S

subscription = subscription $\cup \{(\mathcal{C}, p)\}$

↳ broker save the subscription !

Upon event unsubscribe (constraint)

subscription = subscription $\setminus \{(\text{constraint}, \text{self})\}$

for each $m \in$ neighbours :

trigger PP2PSend (UNSUBSCRIPTION | <constraint, self>) to m

upon event PP2PDeliver (UNSUBSCRIPTION | <C, p>) from S

subscription = subscription $\setminus \{(\mathcal{C}, p)\}$

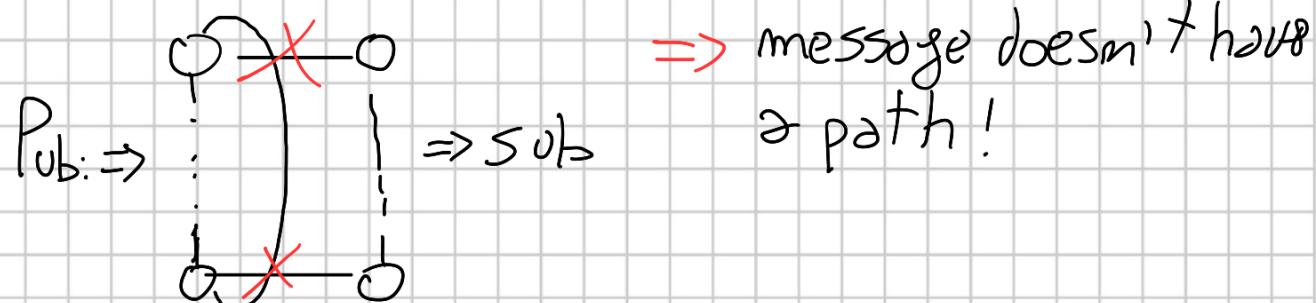
Upon event PP_{2P} Deliver (NOTIFY l, v) from P
if $v \notin \text{delivered}$ do
 for each $s \in \text{subscription s.t. } s.\text{constraint}(v) = \text{True}$
 trigger PP_{2P} Send (PUBLISH l, v) to $s.\text{subscriber}$
 for each $m \in \text{neighbours s.t. } m \neq p$:
 trigger PP_{2P} Send (NOTIFY l, v) to m
 ↳ flooding of the event.
 $\text{delivered} = \text{delivered} \cup \{v\}$

Upon EVENT PP_{2P} Deliver (PUBLISH l, v)
 $\text{delivered} = \text{delivered} \cup \{v\}$
 ↳ subscriber save the value

Publisher with publish send his value to the k broker. The subscriber can subscribe/ unsubscribe sending the message to the k brokers. The brokers send the value received to a subscriber if exist a subscription saved and broadcast the value to the k broker near to him.

2. We can tolerate at most $k-1$ failures, because each time a failure, the neighbours of some brokers decrease, a partition is created, but until $k-1$ failures is impossible to have that a partition is isolated from an other, but with k failures can happen that a

partition is isolated and the NOTIFY events can't arrive to the subscribers:



At the beginning we have K partition of K brokers each, each partition have K connection to other partitions.

3. For tolerate the Byzantine processes we must use authenticated links, in each event we must check that the sender is equal to the sender in the event. Subscribers must check the value received by brokers that is really a value on which is interested and that doesn't already received, also we can accept a value only when we have received $F+1$ events with the same value.

It is necessary that K is at least equal to $2F+1$ because we must guarantee a reliable communication.

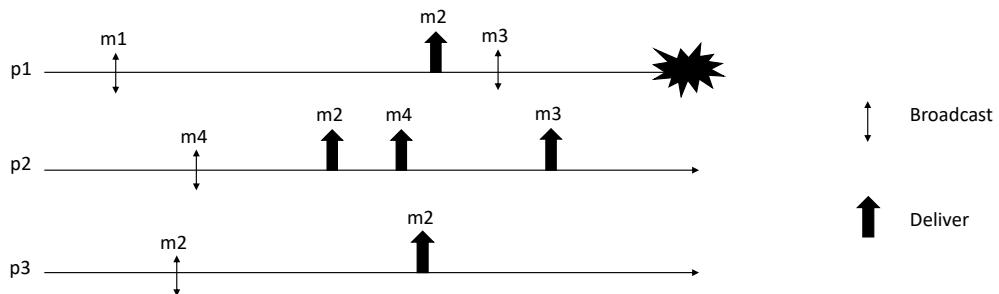
Distributed Systems (9 CFU)

05/07/2022

Family Name _____ Name _____ Student ID _____

Ex 1: Provide the specification of the (1, N) Regular Register and describe the majority voting algorithm discussed during the lectures.

Ex 2: Consider the message pattern shown in the Figure

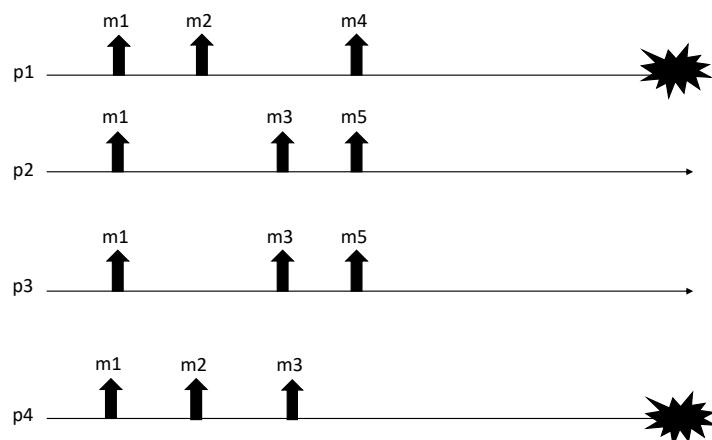


Answer to the following questions:

1. Complete the partial execution in order to obtain a run satisfying Uniform Reliable Broadcast
2. Complete the partial execution in order to obtain a run satisfying Regular Reliable Broadcast but not Uniform Reliable Broadcast
3. Complete the partial execution in order to obtain a run satisfying Best Effort Broadcast but not Regular Reliable Broadcast
4. List ALL the possible sequences satisfying both causal order and total order

NOTE: To solve the exercise you can just add deliveries of messages and new broadcast (if needed)

Ex 3: Consider the execution depicted in the Figure



Answer to the following questions:

1. Which is the strongest Total Order specification satisfied by the proposed run? Provide your answer by specifying both the agreement and the ordering property.
2. Modify the run in order to obtain an execution satisfying TO (UA, WUTO) but not TO (UA, SUTO)

3. Modify the run in order to obtain an execution satisfying TO (NUA, WNUTO) but not TO(NUA, WUTO).

NOTE: To solve the exercise you can just add deliveries of messages.

Ex 4: Let us consider a Regular Reliable Broadcast primitive satisfying the following properties:

- *Validity*: If a correct process p broadcasts a message m , then p eventually delivers m .
- *No duplication*: No message is delivered more than once.
- *No creation*: If a process delivers a message m with sender s , then m was previously broadcast by process s .
- *Agreement*: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

Let us consider a distributed system composed of N processes executing the Eager algorithm (reported in figure)

Algorithm 3.3: Eager Reliable Broadcast

Implements:

ReliableBroadcast, **instance** rb .

Uses:

BestEffortBroadcast, **instance** beb .

```

upon event (  $rb$ , Init ) do
   $delivered := \emptyset$ ;

upon event (  $rb$ , Broadcast |  $m$  ) do
  trigger (  $beb$ , Broadcast | [DATA, self,  $m$ ] );

upon event (  $beb$ , Deliver |  $p$ , [DATA,  $s$ ,  $m$ ] ) do
  if  $m \notin delivered$  then
     $delivered := delivered \cup \{m\}$ ;
    trigger (  $rb$ , Deliver |  $s$ ,  $m$  );
    trigger (  $beb$ , Broadcast | [DATA,  $s$ ,  $m$ ] );
  
```

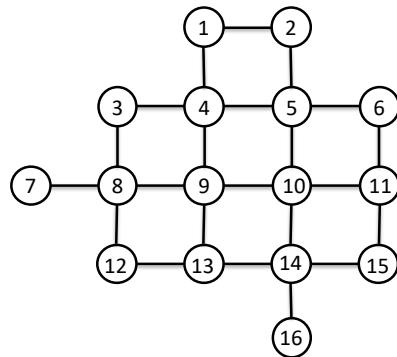
Answer to the following questions:

1. assuming that up to f processes may commit omission failures and no other failures may happen, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).
1. assuming that up to f processes may be Byzantine faulty but constrained to have a symmetric behaviour¹, discuss if the eager algorithm is still able to satisfy the Regular Reliable Broadcast specification (discuss each property individually).

Ex 5: Consider a distributed system composed by n processes each one having a unique identifier. Processes communicate by exchanging messages through perfect point-to-point links and are connected through a grid (i.e., each process p_i can exchange messages only with processes located at *nord*, *sud*, *east* and *west* when they exist).

An example of such network is provided in the following figure:

¹ A Byzantine process has a symmetric behaviour if it can change the content of every message it is going to send, but it cannot send different values to different processes when invoking the $bebBroadcast$. Summarizing, it can cheat but it will do it in a consistent way.



Processes are not going to fail, and they initially know only the number of processes in the system N and the identifiers of their neighbors.

Processes in the system must agree on a color assignment satisfying the following specification:

Module

Name: k-Color assignment

Events:

Request: $\langle ca, \text{Propose} \mid c \rangle$: Proposes a color to be adopted.

Indication: $\langle ca, \text{Decide} \mid c \rangle$: Outputs a decided color to be adopted by the process

Properties:

Termination: Every process eventually decides a color.

Validity: If a process decides a color c , then c was proposed by some process or $c = \text{default}$.

Integrity: No process decides twice.

Weak Agreement: If two processes decide c_i and c_j then either $c_i = c_j$ or one of the two is default

Color Selection: Let C be the set of proposed colors, if $|C| > 0$ then there exists at least a color $c_i \in C$ that is decided by k processes.

Assuming that $1 < k < N$ and that k is known by every process, write the pseudo-code of an algorithm implementing the k-Color assignment primitive.

According to the Italian law 675 of the 31/12/96, I authorize the instructor of the course to publish on the web site of the course results of the exams.

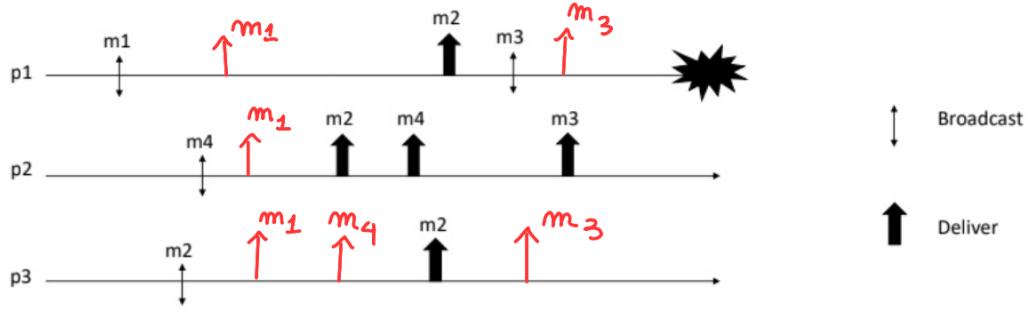
Signature: _____

Exercise 1.

(1, N) Regular Register specification. It has 4 events that are:
• $\langle \text{tt}, \text{Read} \rangle$ that invokes a read operation on the register.
• $\langle \text{tt}, \text{Write } v \rangle$ that invokes a write operation with value v on the register.
• $\langle \text{tt}, \text{ReadReturn } v \rangle$ that completes a read operation on the register with return value v .
• $\langle \text{tt}, \text{WriteReturn} \rangle$ that completes a write operation on the register.
It must satisfy two properties: termination (if a correct process invokes an operation, then the operation eventually completes) and validity (A read that is not concurrent with a write returns the last value written; a read that is concurrent with a write returns the last value written or the value concurrently written). In the majority voting regular registers protocol we assumed a crash failure model with no perfect failure detector. We assumed N processes whose 1 writer and N readers, with a majority of correct processes. It uses P2P link and best-effort broadcast. The main idea of the algorithm is that each process locally stores a copy of the current value of the register. Each written value is univocally associated to a timestamp. The writer and the reader processes use a set of witness processes, to track the last value written. Use a quorum, the intersection of any 2 sets of witness processes is not empty for deliver a write or a read. The read return the highest value witness with the current timestamp when reach a quorum of reply and WriteReturn is done when p reach a quorum of ACKs.

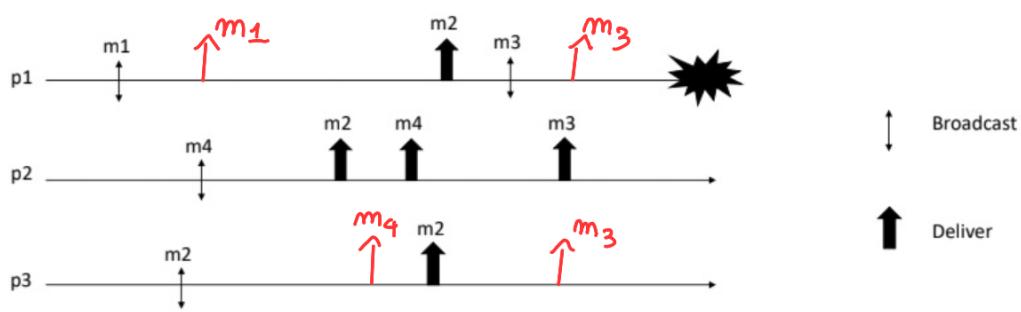
Exercise 2

1.



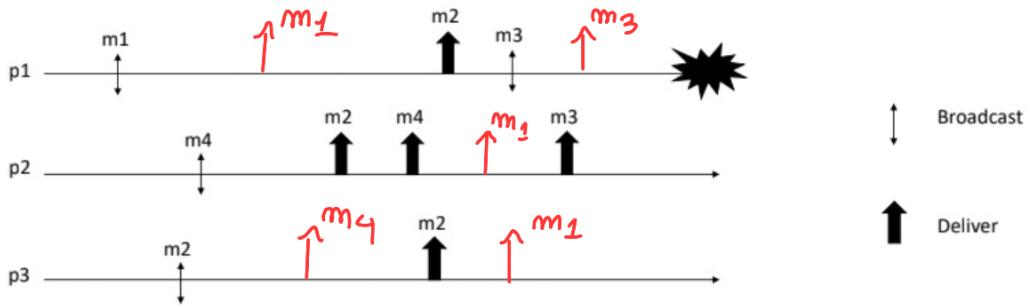
This turn satisfying Uniform Reliable Broadcast, because correct processes have same set of deliveries, the faulty one has a subset.

2.



This turn satisfying Regular Reliable Broadcast, because correct processes have same set of deliveries, but not uniform because faulty one have delivered m1 that other don't have.

3.



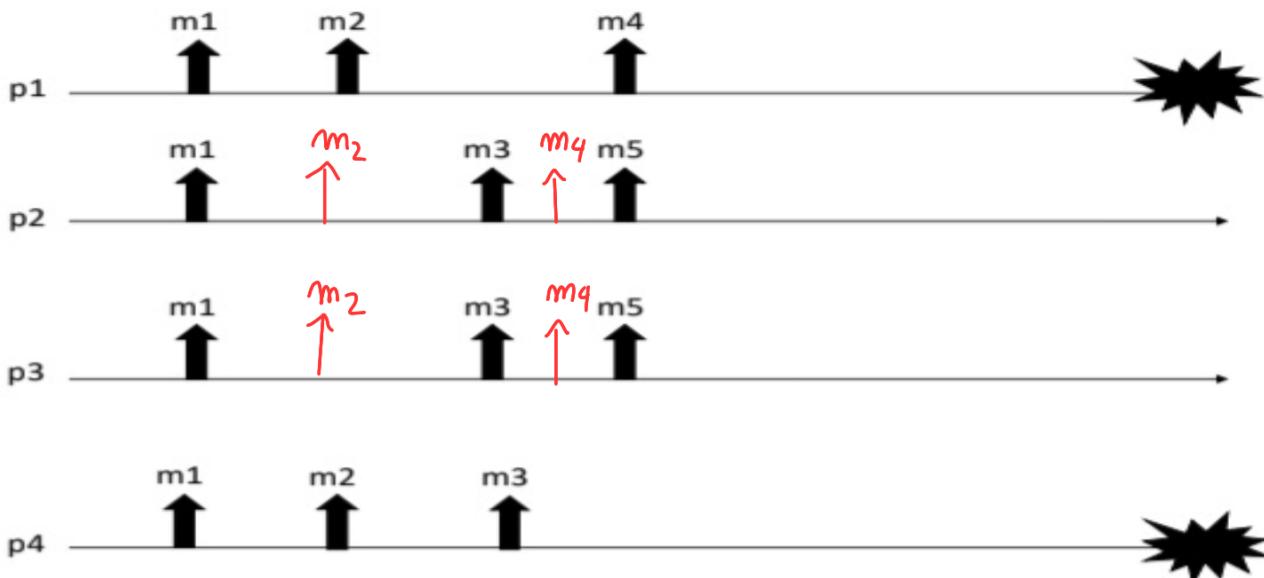
This turn satisfying best effort broadcast because correct processes don't have same set of deliveries.

4. For have causal order we must have $m_1 \rightarrow m_3$ and $m_2 \rightarrow m_3$, For total order because of P₂ we have constraint $m_2 \rightarrow m_4 \rightarrow m_3$, the sequences that satisfying both are: (m_1, m_2, m_4, m_3), (m_2, m_1, m_4, m_3), (m_2, m_1, m_1, m_3), (m_2, m_1, m_1, m_3).

Exercise 3

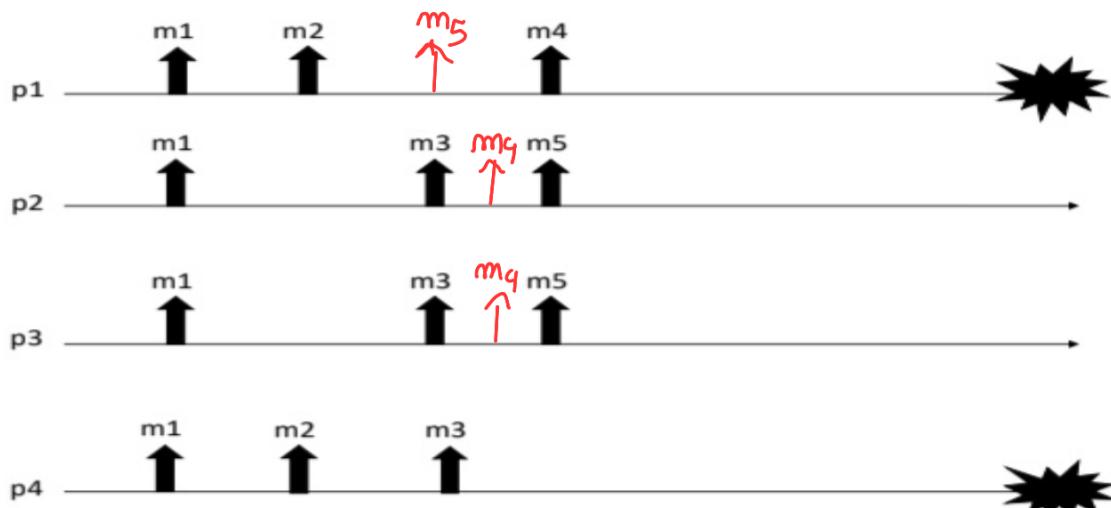
1. This turn doesn't satisfy UA because p_1 and p_4 that are faulty have not a subset of delivery of the correct processes. It satisfies SU TO because every process respect the constraint of WUTO, only pair are respected of other. TO (NUA, WUTO)

2.



it satisfies UA because correct processes have same set of messages, and faulty one is subset. It satisfies WUTO because the order of any pair is respected but not SU TO because p_1 deliver m_4 but not m_3 before

3.



it not satisfies UA because p_1 and p_4 deliver m_2 , it does not satisfy WUTO because p_1 violate order of pair $(m_1, m_2) \rightarrow m_3$

Exercise 4

1. Validity: it is satisfied because if a correct process p broadcast a message m, eventually p receive $(p, l, DATA_S, m)$ and because $m \notin S$ delivered, he will deliver m. All work because b2b use P2P links.

No duplication: it is satisfied because when a process receive a broadcast message m, check that $m \notin S$ delivered, it is impossible to have duplication.

No creation: it is satisfied because if m is delivered, then following the algorithm it must be that s generate a broadcast event with m, the f processes can only do an omission but not generate fake messages.

Agreement: it is satisfied because also if f processes faulty do an omission, when a correct process receive a message do a broadcast, and other correct processes eventually deliver that message

2. Validity: it is satisfied because if a correct process p broadcast a message m, eventually p receive $(p, l, DATA_S, m)$ and because $m \notin S$ delivered, he will deliver m. All work because b2b use P2P links.

No duplication: it is not guarantee because a byzantine process can not insert a message delivered in the set delivered, and to deliver the same message when receive again m.

No creation: it is not guarantee because a byzantine process can forge a message m with a fake sender, and because the algorithm don't use authenticate limit, a process p will deliver this m.

Agreement: it is satisfied thanks to the fact we use P2P links and every correct process when receive a message m , he will generate a new broadcast of the same message, then at some point every correct process reach the same set of deliveries.

Exercise 5

We have m processes each with an unique identifier. We use P2P links, using a topology like in the figure, each process know and can communicate with north, south, east and west, if exist. No failures. Each process know only number of processes N . The algorithm that i propose do: is like a consensus protocol. Each process propose a color c ; and disseminate to his neighbours. When a process receive N different propose (each is a pair of (value, identifier)), it will decide for the darkest color (HP: colors are comparable). After HP choose a value, given that we want that a color is decided by K processes, a process p check if its id is less or equal to K then generate a decide event with color c , otherwise it will decide default ($= \perp$).

Next page pseudo-code:

two events and a passive function.

Init:

$mp = N$ (number of processes)

decision = (\perp , FALSE)

propose = \emptyset , where we collect proposes

neighbours = {nord, sud, east, west} (some not exist)

$K = *$ is given

Upon event $\langle c, \text{Propose} | c \rangle$ do

propose = propose \cup $\langle c, \text{self} \rangle$

for each $p \in \text{neighbours}$ do

trigger PP2PSend (IPROPOSEI, $\langle c, \text{self} \rangle$) to p

Upon event PP2PDeliver (IPROPOSEI, $\langle c, id \rangle$) from s

if $\langle c, id \rangle \notin \text{propose}$ do

propose = propose \cup $\langle c, id \rangle$

for each $p \in \text{neighbours}$ s.t. $p \neq s$ do

trigger PP2PSend (IPROPOSEI, $\langle c, \text{self} \rangle$) to p

else

nothing i have already received

Upon exist $|\text{Propose}| == N$ and $\text{decision}_2 == \text{FALSE}$ do

choose = darkest_color (Propose) \downarrow

If self $\leq K$ do

decision = (choose, TRUE)

return darkest_color proposed

trigger $\langle c, \text{decide} | \text{choose} \rangle$

else

decision = (\perp , TRUE)

trigger $\langle c, \text{decide} | \perp \rangle$

it satisfies termination, because every process disseminate his propose and the received one to neighbours, given that there are no failures and P2P links are used, all processes will derive N different proposal and decide a value. It satisfies validity because when a process decide, it decide the darkest color proposed by a process or the default. it satisfies integrity because when a process decide, set decision = 2 equal to TRUE and the passive check is no longer satisfied. It satisfied weak agreement because no process decide a value different from darkest color of \perp , darkest color is the same for all because all processes have same set propose. It satisfies color selection because we construct algorithm in a way that the darkest color is chosen by the first K processes, exist always a color chosen K times.