

03/10/23

Dependable Distributed Systems
Master of Science in Engineering in Computer
Science

AA 2023/2024

LECTURE 4: LOGICAL CLOCK

Recap

Physical clock synchronization algorithms have the aim to coordinate processes to reach an agreement on a common notion of time

The accuracy of the synchronization is strongly dependent on the estimation of transmission delay

- **ISSUE**: it can be hard to find a good estimation

OBSERVATION

- In several applications it is not important when things happened but in which order they happened

We need to find a reliable way to order events without using clock synchronization!

Happened-Before relation

OBSERVATION

- Two events occurred at some process p_i happened in the same order as p_i observes them
- When p_i sends a message to p_j the send event happens before the deliver event

Lamport introduces the *happened-before relation* to capture causal dependencies between events (causal order relation)


- We note with \rightarrow_i the ordering relation between events in a process p_i
- We note with \rightarrow the happened-before between any pair of events

Happened-Before Relation: Definition

Two events e and e' are related by happened-before relation ($e \rightarrow e'$) if:

- $\exists p_i \mid e \rightarrow_i e'$ local execution history
- $\exists m \mid e = \text{send}(m) \text{ and } e' = \text{deliver}(m)$
- $\exists e, e', e'' \mid (e \rightarrow e'') \wedge (e'' \rightarrow e')$ (happened-before relation is transitive)

= TRUE




RULE

- 2. $\ell_3 \rightarrow \ell_4$
- 1. $\ell_4 \rightarrow \ell_5$

Happened-Before Relation

OBSERVATIONS

- Happened-before relation imposes a partial order over events of the execution history
 - It may exist a pair of events $\langle e_i, e_j \rangle$ such that e_i and e_j are not in happened-before relation
 - If e_i and e_j are not in happened-before relation then they are concurrent ($e_i \parallel e_j$) → **NOT RELATED**
- For any pair of events e_i and e_j in a distributed system only one of the following holds
 - $e_i \rightarrow e_j$,
 - $e_j \rightarrow e_i$
 - $e_i \parallel e_j$



Logical Clock

The Logical Clock, introduced by Lamport, is a software counter that *monotonically* increases its value

A logical clock L_i can be used to *timestamp* events \rightarrow INTEGER

$ts_e = L_i(e)$ is the “logical” timestamp assigned by a process p_i to events e using its current logical clock

PROPERTY

- If $e \rightarrow e'$ then $ts_e < ts_{e'}$ \rightarrow keep track of CAUSALITY

Observation

- The ordering relation obtained through logical timestamps is only a partial order

Scalar Logical Clock: an implementation

Each process p_i initializes its logical clock $L_i=0$ ($\forall i = 1...N$)

p_i increases L_i of 1 when it generates an event (either *send* or *receive*)

- $L_i = L_i + 1$

When p_i **sends** a message m


- creates an event *send*(m)
- increases $L_i = L_i + 1$
- timestamps m with $ts=L_i$

When p_i **receives** a message m with timestamp ts

- Updates its logical clock $L_i = \max(ts, L_i)$
- Produces an event *receive*(m)
- Increases $L_i = L_i + 1$

ensure
MONOTONICITY

Scalar Logical Clock: example



e^j_i is j -th event of process p_i

L_i is the logical clock of p_i

NOTE

- $e^1_1 \rightarrow e^2_1$ and timestamps reflect this property
- $e^1_1 \parallel e^1_3$ and respective timestamps have the same value
- $e^1_2 \parallel e^1_3$ but respective timestamps have different values

e^1_2 clock = 2
 e^1_3 clock = 1
 } CONCURRENT

Limits of Scalar Logical Clock

Scalar logical clock can guarantee the following property

- If $e \rightarrow e'$ then $ts_e < ts_{e'}$

But it is not possible to guarantee

- If $ts_e < ts_{e'}$ then $e \rightarrow e'$

Consequently:

- Using scalar logical clocks, it is not possible to determine if two events are concurrent or related by the happened-before relation

Mattern [1989] and Fridge [1991] proposed an improved version of logical clock where events are timestamped with local logical clock and node identifier

- ***Vector Clock***

Vector Clock : definition

if PROCESSES = 10 \rightarrow vector size = 10

Vector Clock for a set of N processes is composed by an array of N integer counters

Each process p_i maintains a Vector Clock V_i and timestamps events by mean of its Vector Clock

Similarly to scalar clock, Vector Clock is attached to message m

- in this case the timestamp will be an integer vector (i.e., an array of integer)

Vector Clock allows nodes to order events in happens-before just looking at their timestamps

- Scalar clocks: $e \rightarrow e'$ implies $L(e) < L(e')$
- Vector clocks: $e \rightarrow e'$ iff $L(e) < L(e')$

Vector Clock : an implementation

Each process p_i initializes its Vector Clock V_i

- $V_i[j]=0 \quad \forall j = 1 \dots N$

p_i increases $V_i[i]$ of 1 when it generates an event

- $V_i[i] = V_i[i] + 1$


When p_i **sends** a message m

- Creates an event *send*(m)
- Increases V_i
- timestamps m with $ts=V_i$

When p_i **receives** a message m containing timestamp ts

- Updates its logical clock $V_i[j] = \max(ts[j], V_i[j]) \quad \forall j = 1 \dots N$
- Generates an event *receive*(m)
- Increases V_i

Vector Clock: an example



Vector Clock: properties

A Vector Clock V_i

- $V_i[i]$ represents the number of events produced by p_i
- $V_i[j]$ with $i \neq j$ represents the number of events generated by p_j that p_i can know

$V = V'$ if and only if

- $V[j] = V'[j] \quad \forall j = 1 \dots N$


$V \leq V'$ if and only if

- $V[j] \leq V'[j] \quad \forall j = 1 \dots N$


$V < V'$ therefore the event associated to V happened before the event associated to V' if and only if

- $V \leq V' \wedge V \neq V'$
 - $\forall i = 1 \dots N \quad V'[i] \geq V[i]$
 - $\exists i \in \{1 \dots N\} \mid V'[i] > V[i]$

A comparison of Vector Clocks



$V(e) < V'(e')$ then $e \rightarrow e'$



$V(e) \neq V'(e')$ then $e \parallel e'$
concurrent

Differently from Scalar Clock, Vector Clock allows to determine if two events are concurrent or related by a happened-before relation

Logical clock in distributed algorithms

We have seen two mechanisms to represent logical time

- Scalar Clock : *timestamp*
- Vector Clock

Each mechanism can be used to solve different problems, depending on the problem specification

- Scalar Timestamp → Lamport's Mutual Exclusion
- Vector Timestamp → Causal Broadcast


Distributed Mutual Exclusion

The Mutual Exclusion Problem

Let us consider

- a set of processes $\Pi = \{p_1, p_2, \dots, p_n\}$
- a set of resources $R = \{r_1, r_2, \dots, r_m\}$

shared resources



PROBLEM

- Processes need to access resources exclusively and we need to design a distributed abstraction that allows them to coordinate to get access to resources

System Model

Let us consider

- a set of processes $\Pi = \{p_1, p_2, \dots p_n\}$
- a set of resources $R = \{r_1, r_2, \dots r_m\}$
 - For the sake of simplicity let us assume $|R| = 1$

The system is asynchronous *not impo TIME*

Processes are not going to fail *not FAILURE* (they will be always correct)

Processes communicate by exchanging messages on top of perfect point-to-point links

The Mutual Exclusion abstraction

EVENTS

- `request()`: it issues a request to enter into the critical section
- `ok()`: it notifies the process that it can now access the critical section
- `release()`: it is invoked to leave the critical section and to allow someone else to enter

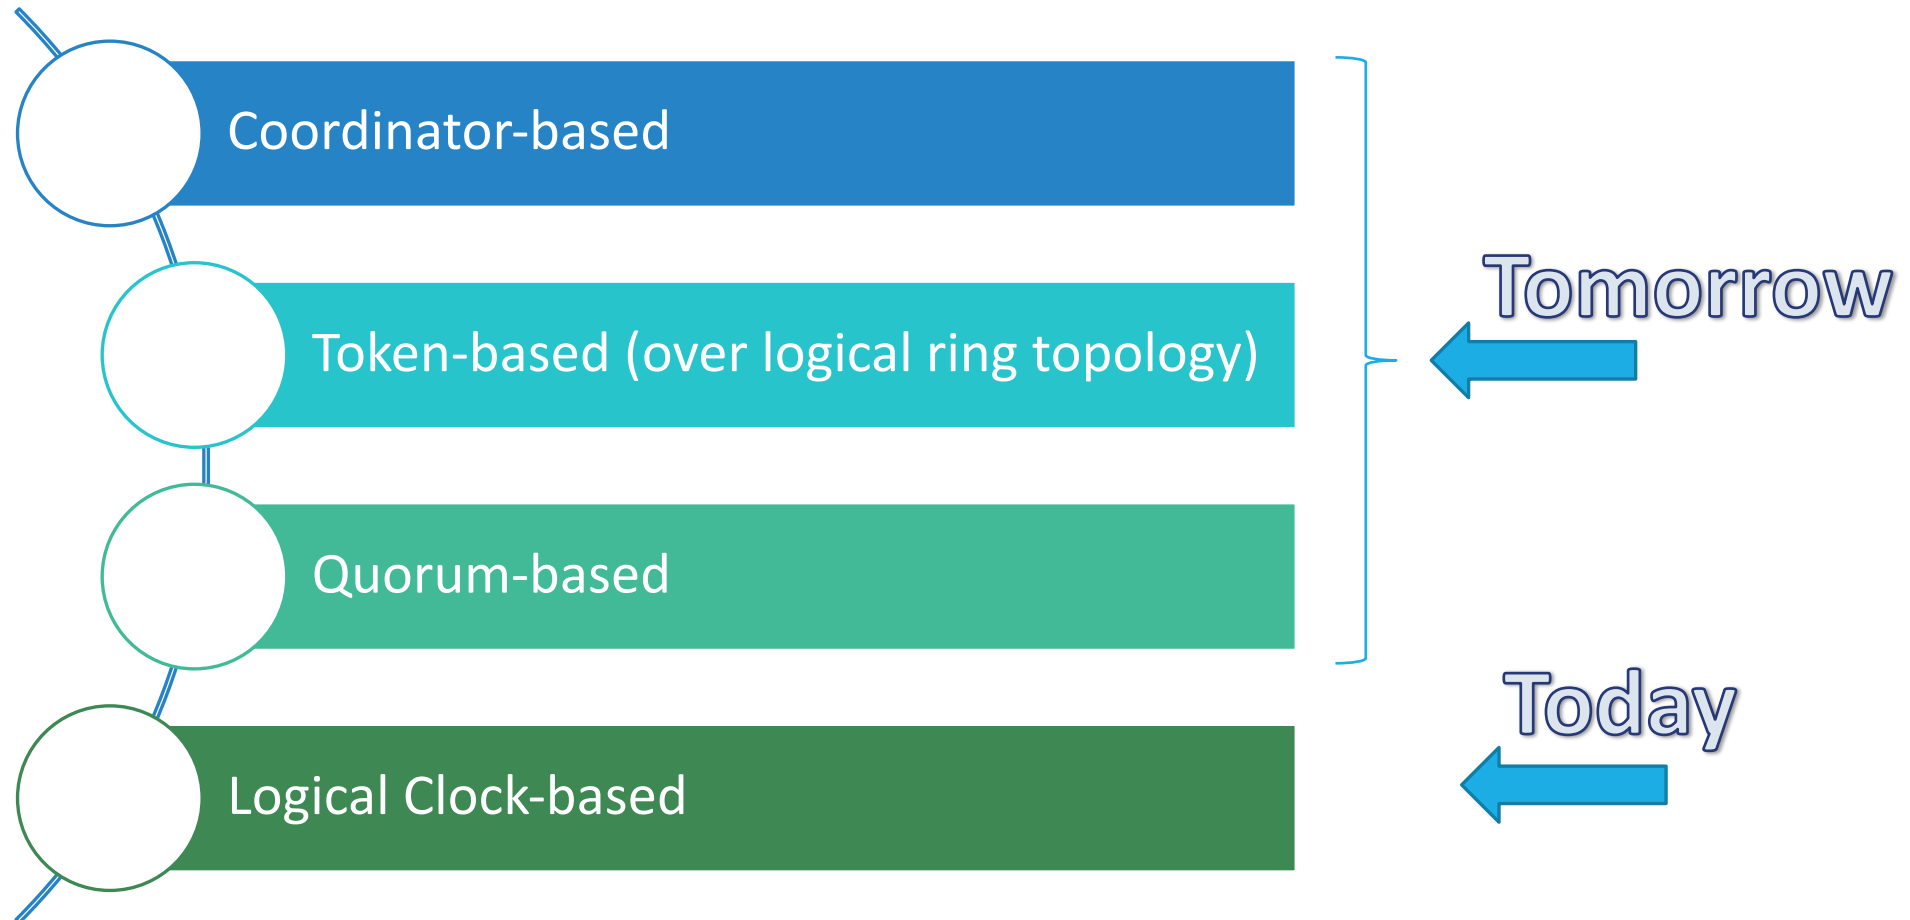
PROPERTIES

- **Mutual Exclusion**: at any time t , at most one process p is running the critical section
- **No-Deadlock**: there always exists a process p able to enter the critical section
- **No-Starvation**: every `request()` and `release()` operation eventually terminate

Critical
section
" Shared
Resources



Different Approaches to Distributed Mutual Exclusion



Timestamp-based algorithm: Lamport's Distributed Mutual Exclusion

Difference from concurrent system

- When a process wants to enter the CS sends a request message to all the other

An history of the operations is maintained by using a counter (timestamp)

Each transmission and reception event is relevant to the computation

- The counter is incremented for each send and receive event
- The counter is incremented also when a message, not directly related to the mutual exclusion computation, is sent or received.

Lamport's algorithm: implementation

Local data structures to each process p_i

- **ck** *timestamp*
 - Is the counter for process p_i
- **Q**
 - Is a queue maintained by p_i where CS access requests are stored

Algorithm rules for a process p_i

- **Request to access the CS**
 - p_i sends a request message, attaching **ck**, to all the other processes
 - p_i adds its request to **Q**
- **Request reception from a process p_j**
 - p_i puts p_j request (including the timestamp) in its queue
 - p_i sends back an ack to p_j

Lamport's algorithm: implementation

Algorithm rules for a process p_i

- **p_i enters the CS iff**

1. p_i has, in its queue, a request with timestamp t
2. t is the small timestamp in the queue
3. p_i has already received an ack with timestamp t' from any other process and $t' > t$

} if I have 3 TRUE
↓
entering
in CS


- **Release of the CS**

- p_i sends a RELEASE message to all the other processes
- p_i deletes its request from the queue

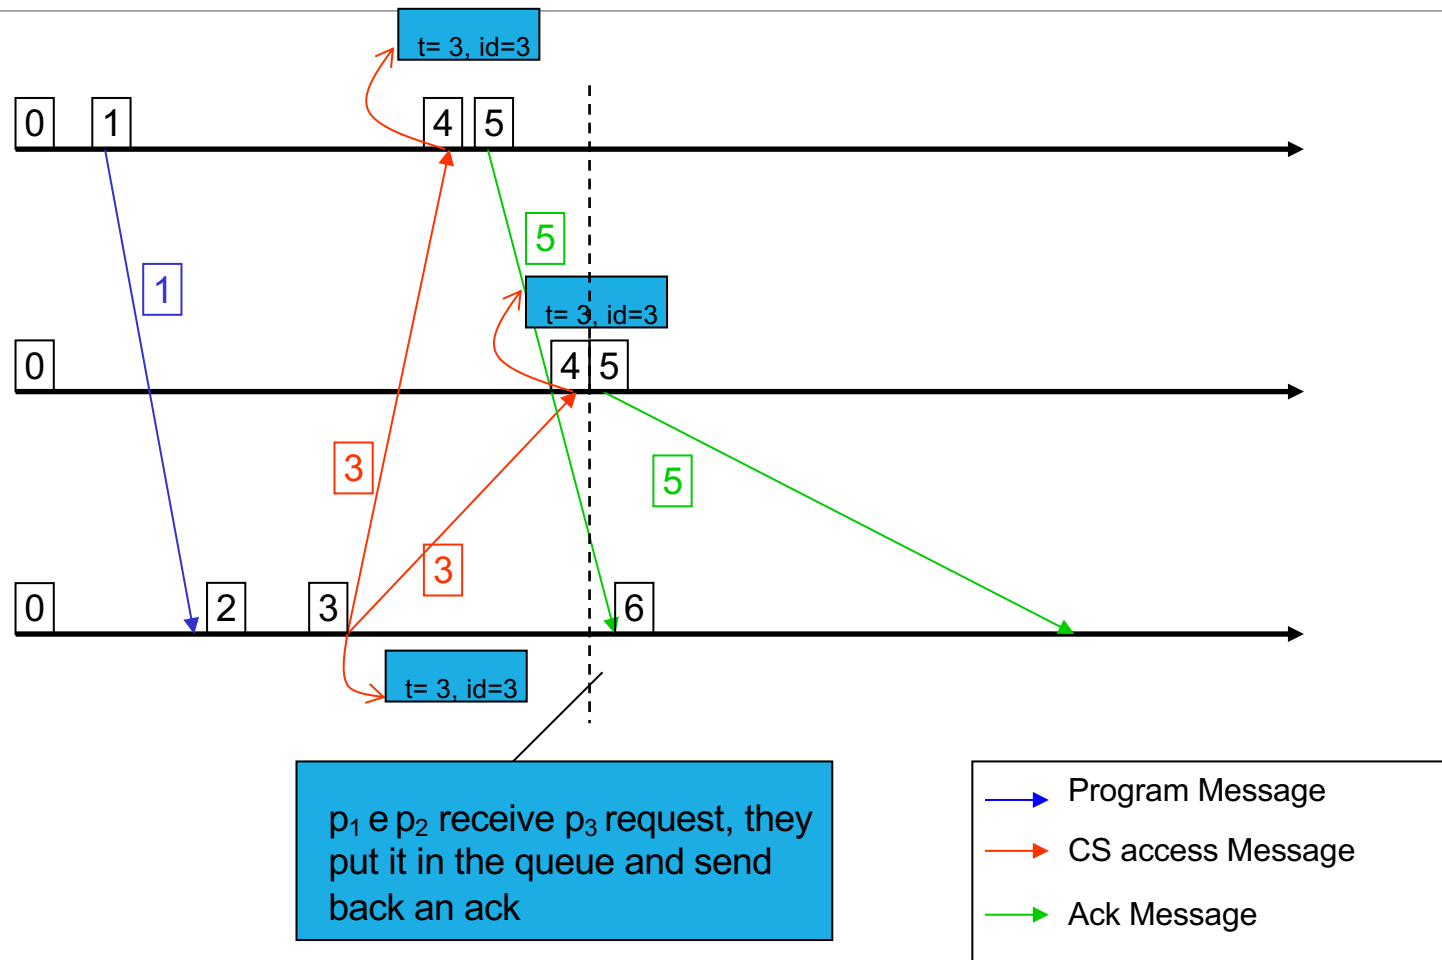
- **Reception of a release message from a process p_j**

- p_i deletes p_j 's request from the queue

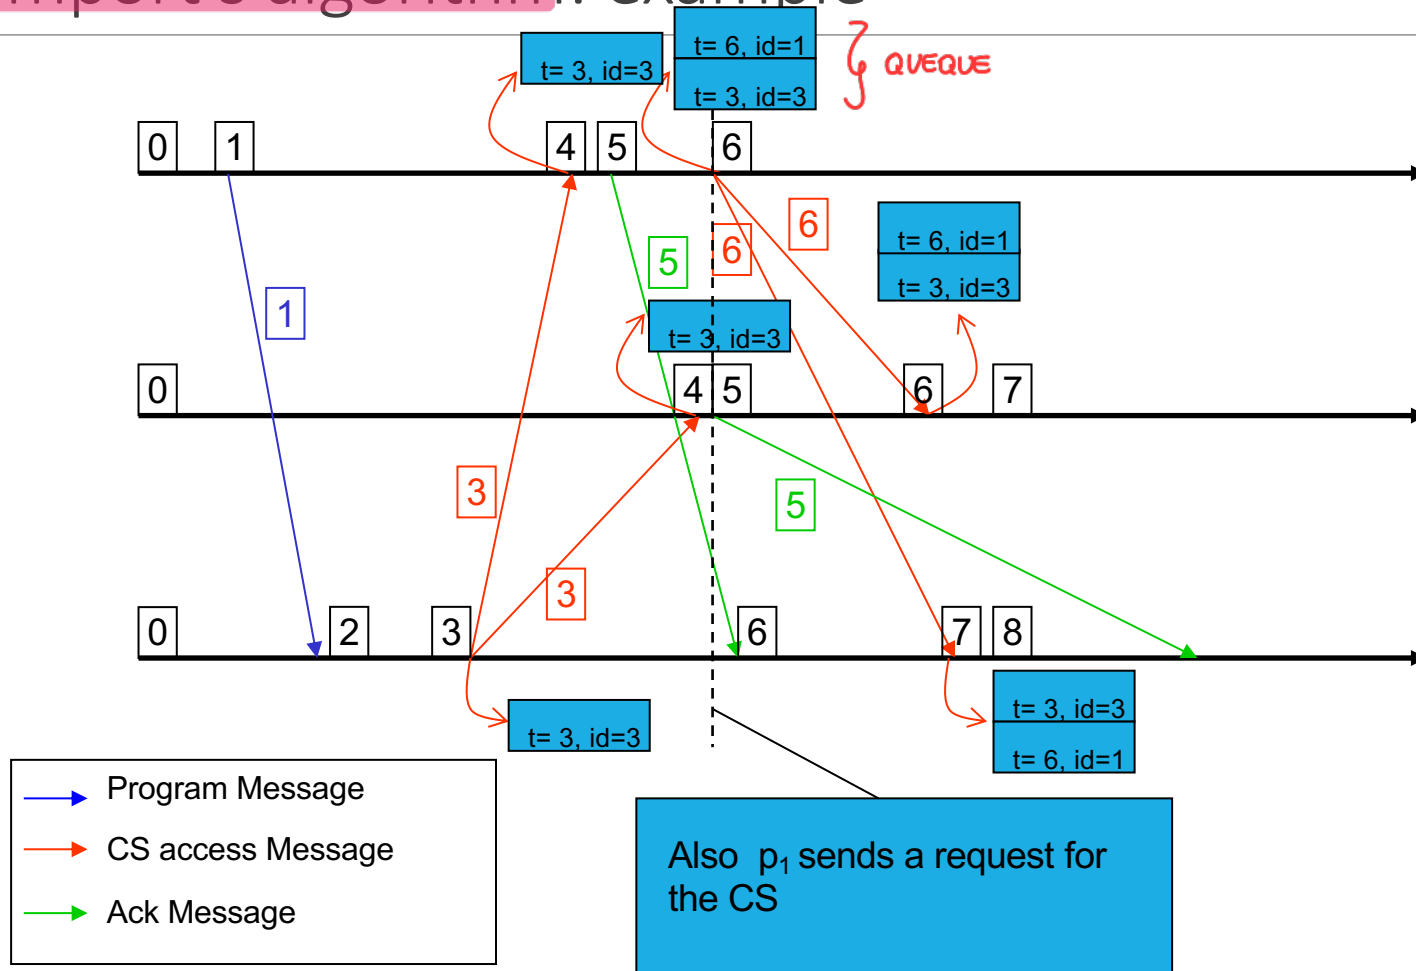
Lamport's algorithm: example



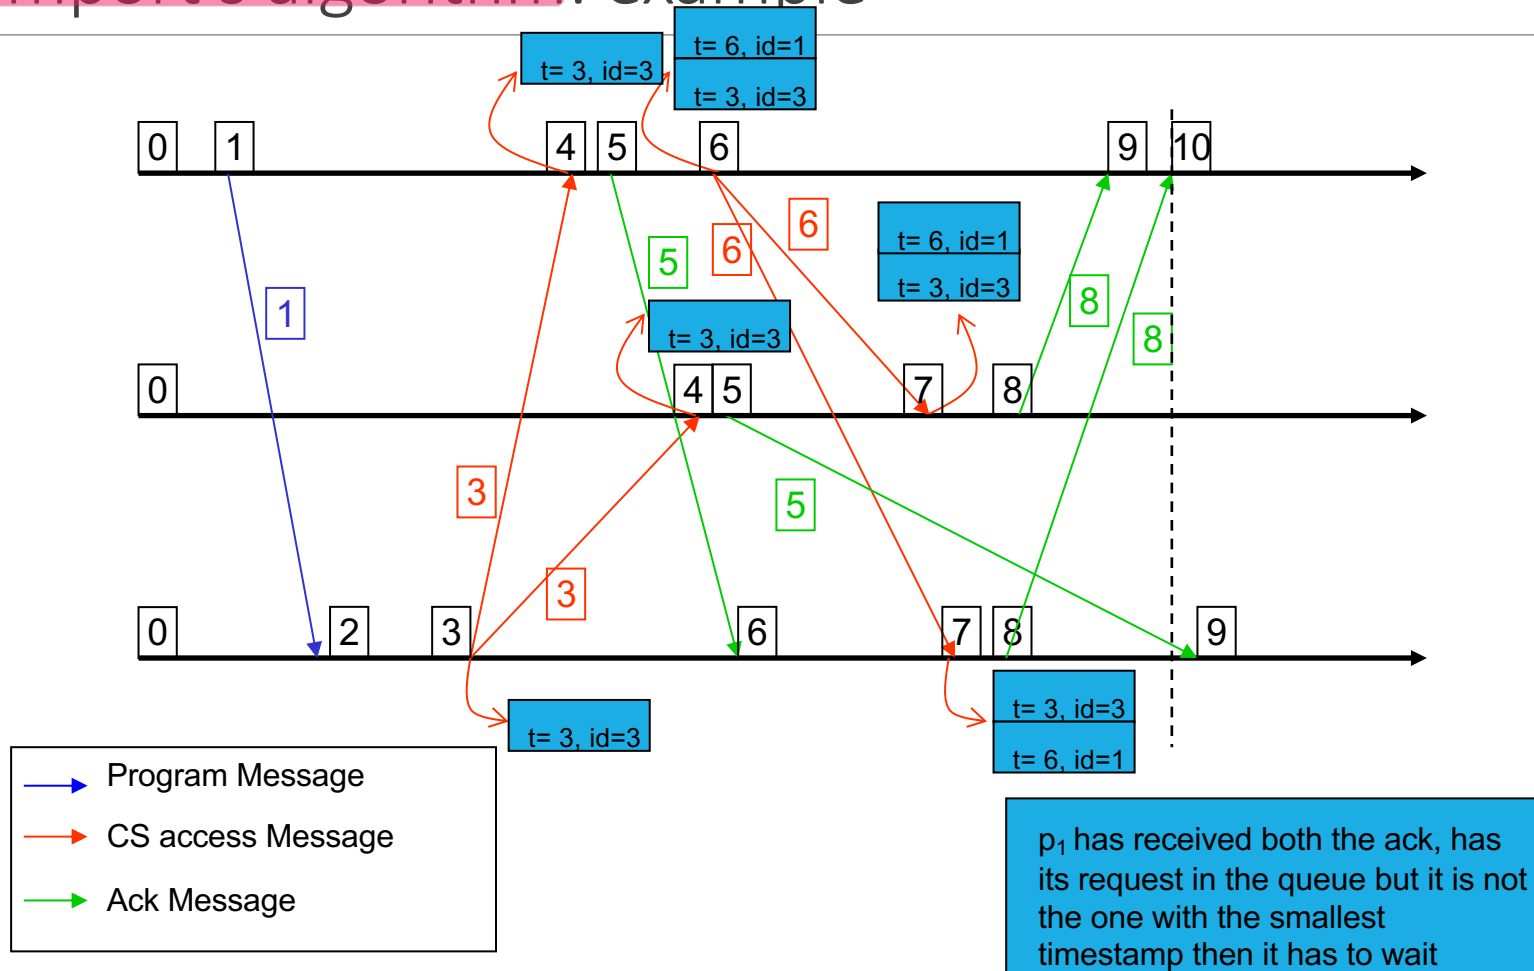
Lamport's algorithm: example



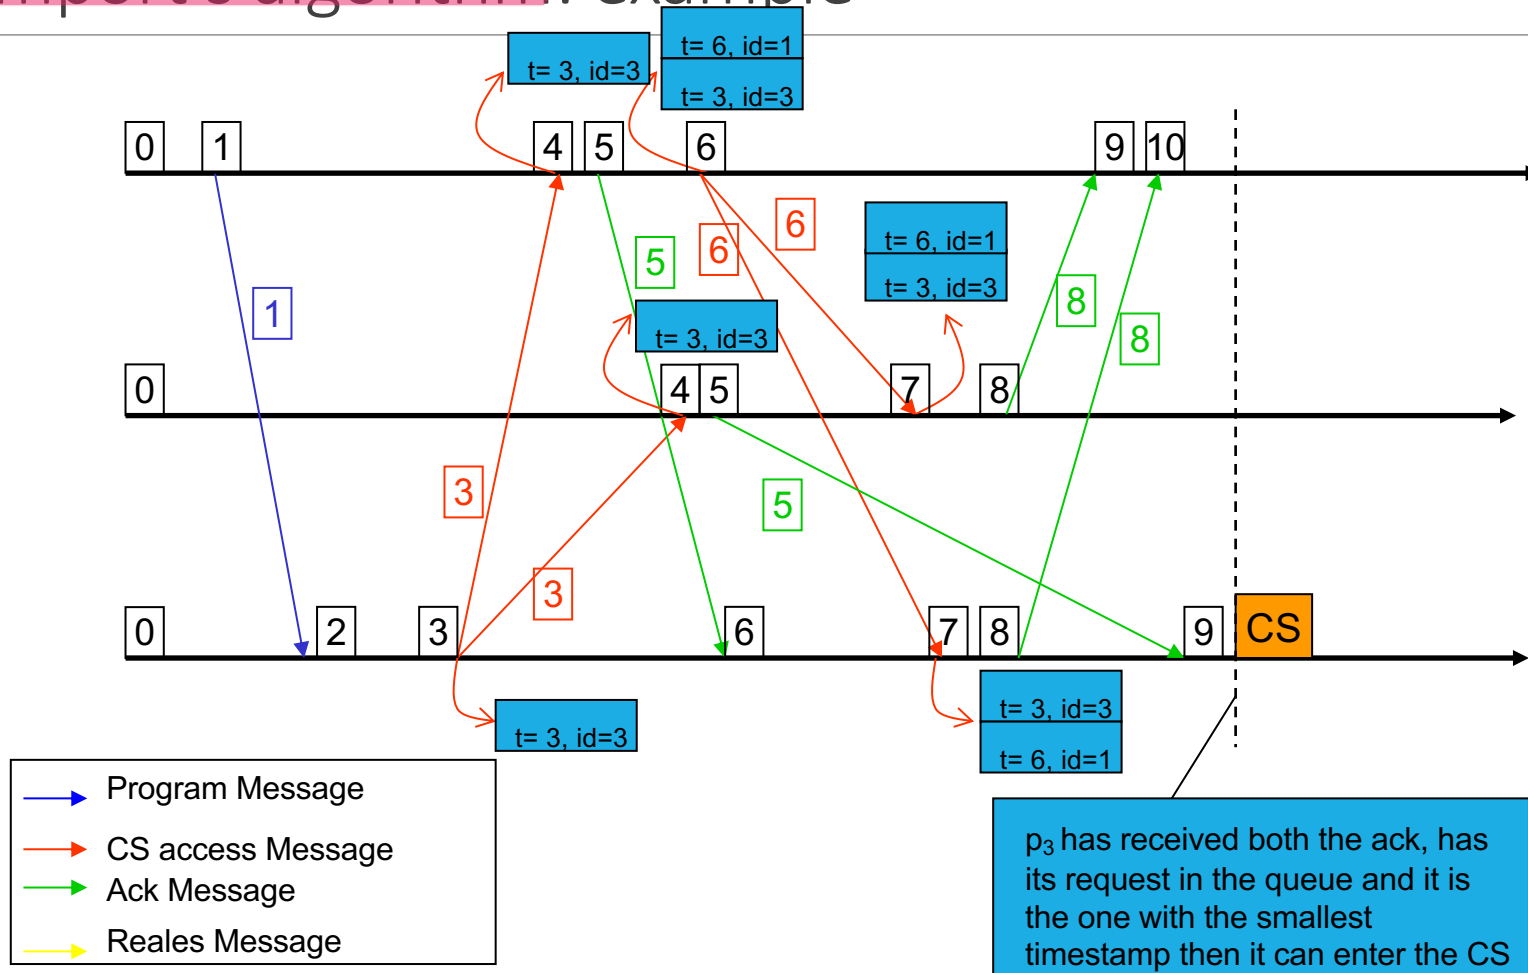
Lamport's algorithm: example



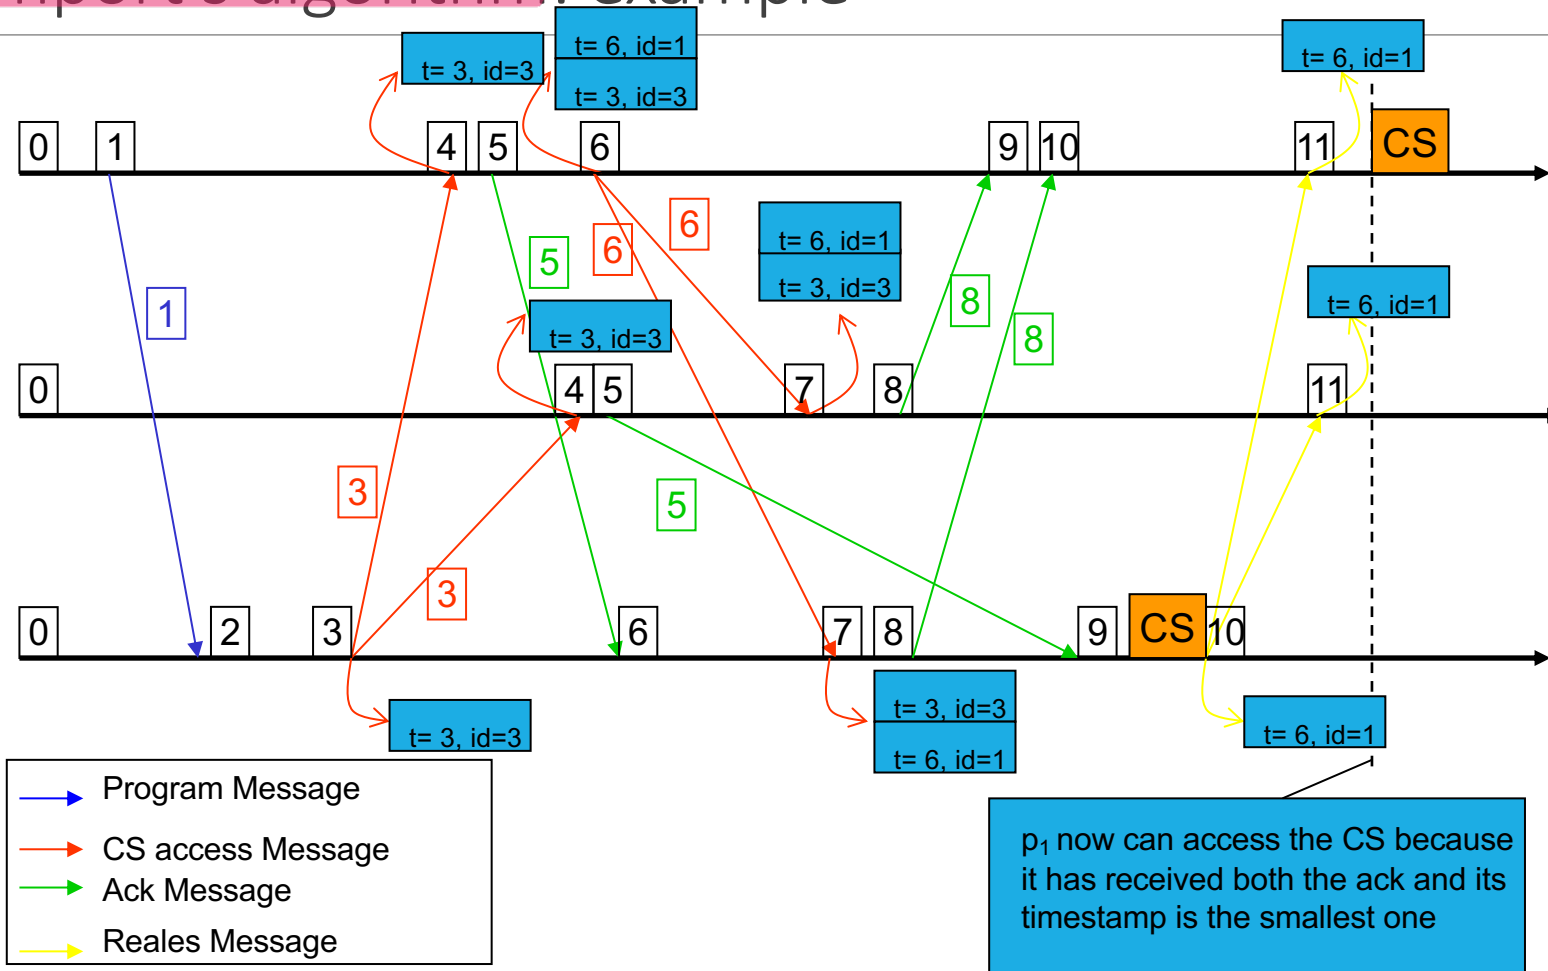
Lamport's algorithm: example



Lamport's algorithm: example




Lamport's algorithm: example



Lamport's algorithm: safety proof

Let us suppose by contradiction that both p_i and p_j enter the CS

- \Rightarrow both the processes have received an ack from any other process and, to enter the CS, the timestamp has to be the smallest in the queue
 - $t_i < t_j < \text{ack}_i.\text{ts}$
 - $t_j < t_i < \text{ack}_j.\text{ts}$



Lamport's algorithm: properties

Fairness is satisfied: different requests are satisfied in the same order as they are generated

- Such order comes from the happened-before relation:
 - If two requests are in happened-before relation then they are satisfied in the same order.
 - If two request are concurrent with respect to the happened before relation then the access can happen in any order

P_1 P_2 concurrently, so $ts(p_1) = ts(p_2)$

id is *unique* and I use it to break it the symmetry } order by *TS*, if $ts(p_1) = ts(p_2)$
↓
order by *ID*

Lamport's algorithm: performances

Lamport's algorithm needs $3(N-1)$ messages for the CS execution

- $N-1$ requests
- $N-1$ acks
- $N-1$ releases

In the best case (none is in the CS and only one process ask for the CS) there is a delay (from the request to the access) of 2 messages

Ricart-Agrawala's algorithm: implementation

Local variables

- #replies (initially 0)
- **State** $\in \{\text{Requesting}, \text{CS}, \text{NCS}\}$ (initially NCS)
processes that are
- Q pending requests queue (initially empty)
- Last_Req
- Num

Algorithm

begin

1. **State=Requesting**
LOGICAL CLOCK
2. Num=num+1; Last_Req=num
TIMESTAMP
3. $\forall i=1 \dots N$ send REQUEST(num) to p_i
4. Wait until #replies=n-1
5. **State=CS**
6. CS
7. $\forall r \in Q$ send REPLY to r *all processes that are waiting in the queue*
8. $Q = \emptyset$; **State=NCS**; #replies=0


Upon receipt REQUEST(t) from p_j

1. **If** State=CS or (State=Requesting and {Last_Req,i}<{t,j})
2. Then insert in Q{t, j}
3. Else send REPLY to p_j
4. Num=max(t,num)


Upon receipt of REPLY from p_j

1. #replies=#replies+1

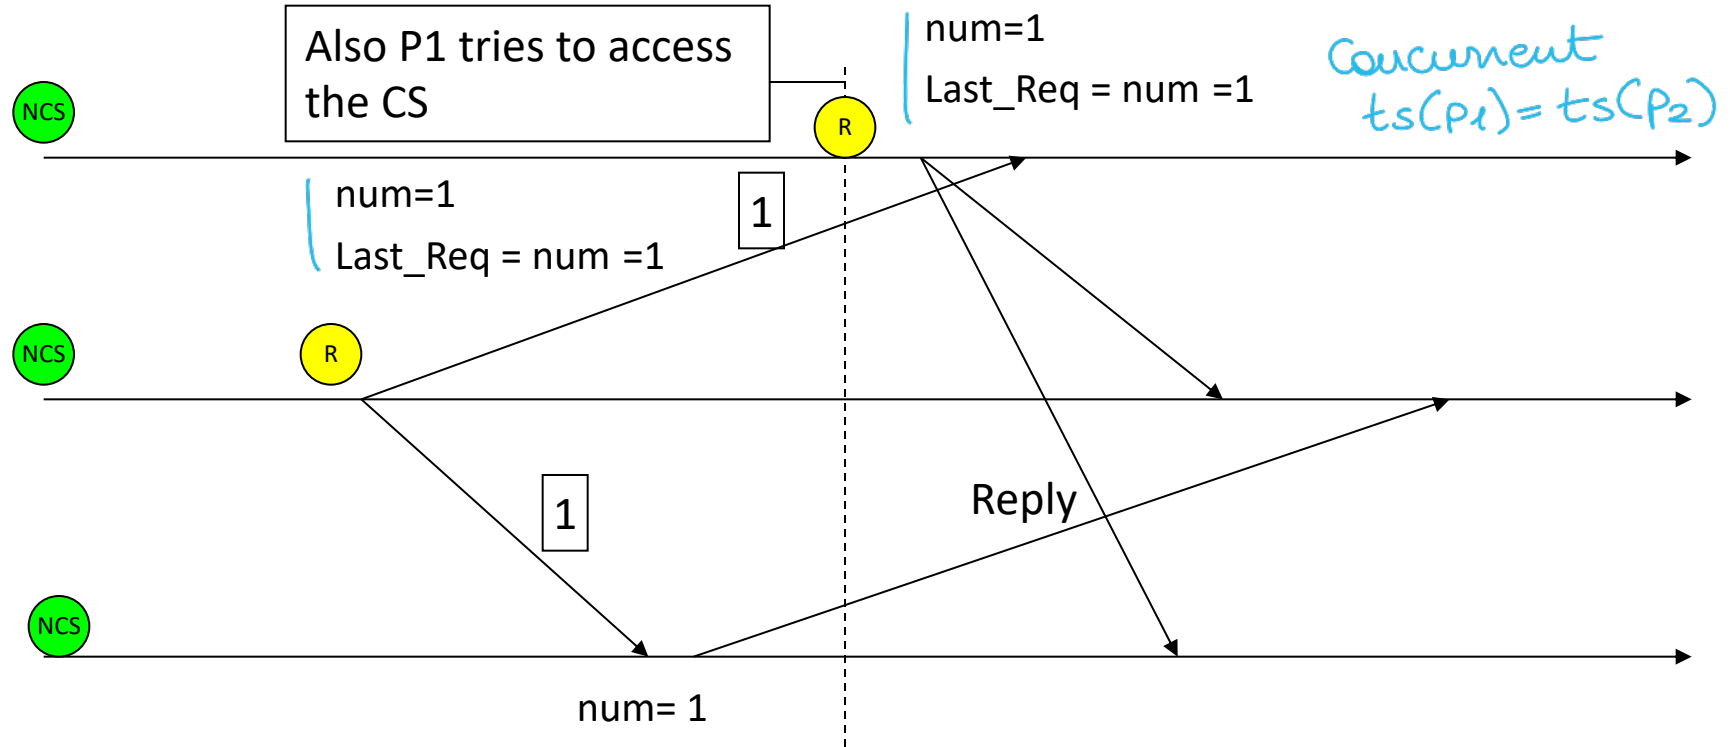
Ricart-Agrawala's algorithm: example



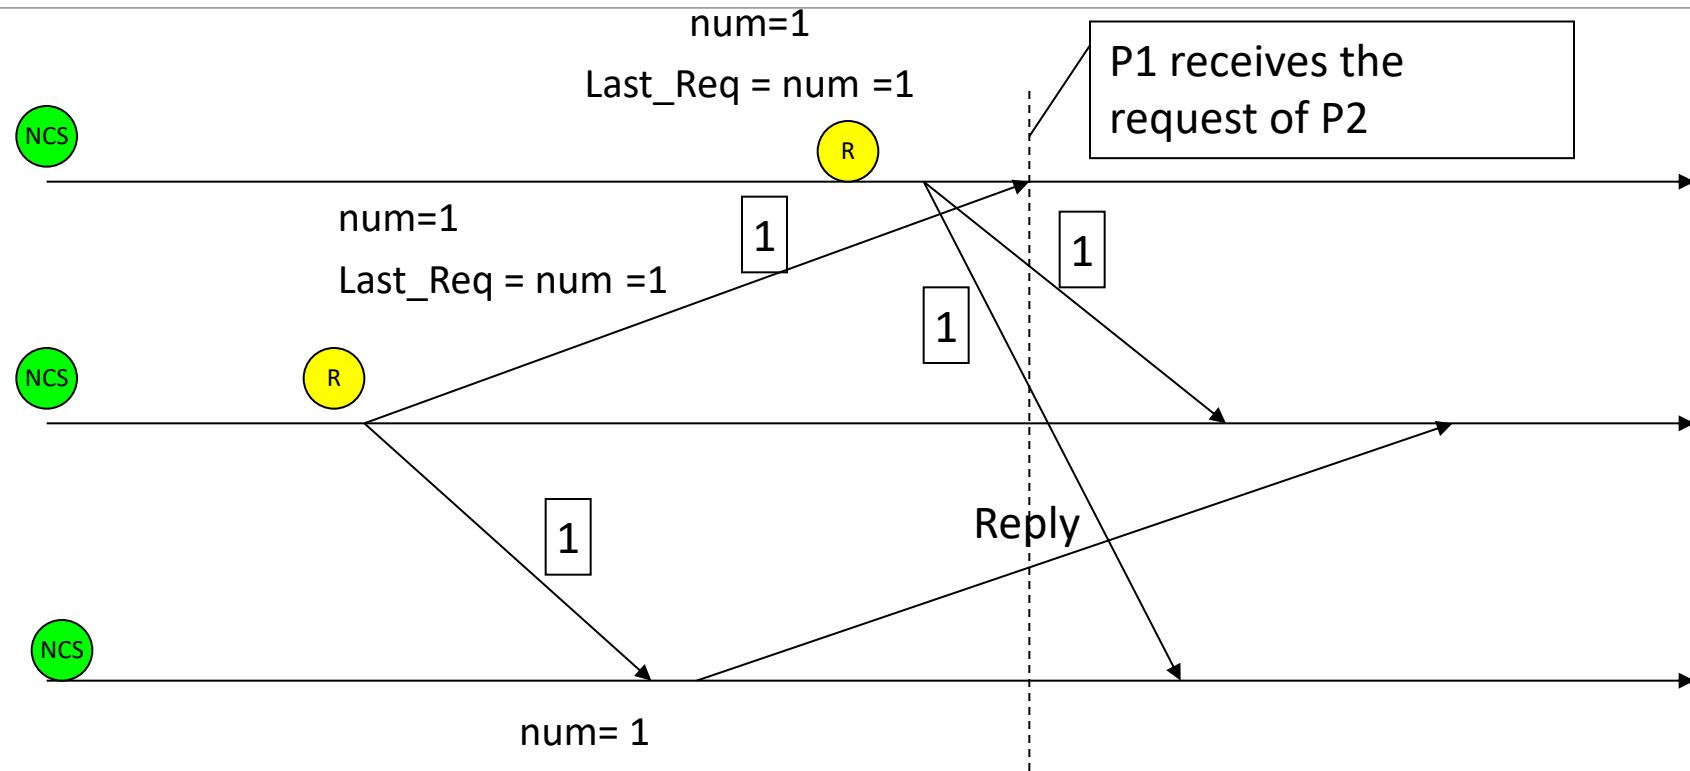
Ricart-Agrawala's algorithm: example



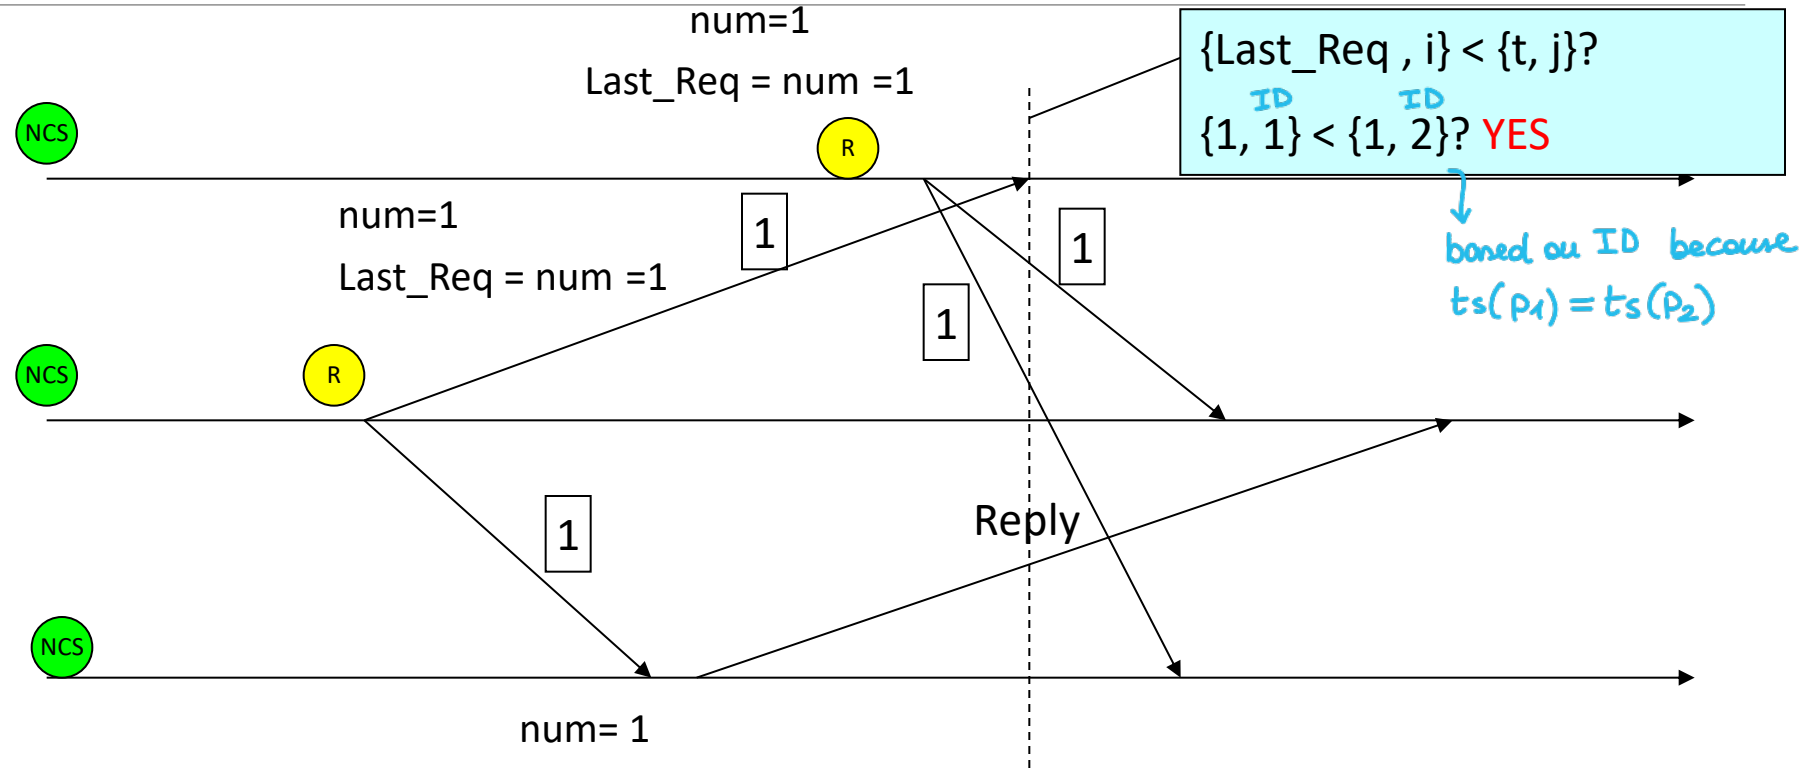
Ricart-Agrawala's algorithm: example




Ricart-Agrawala's algorithm: example

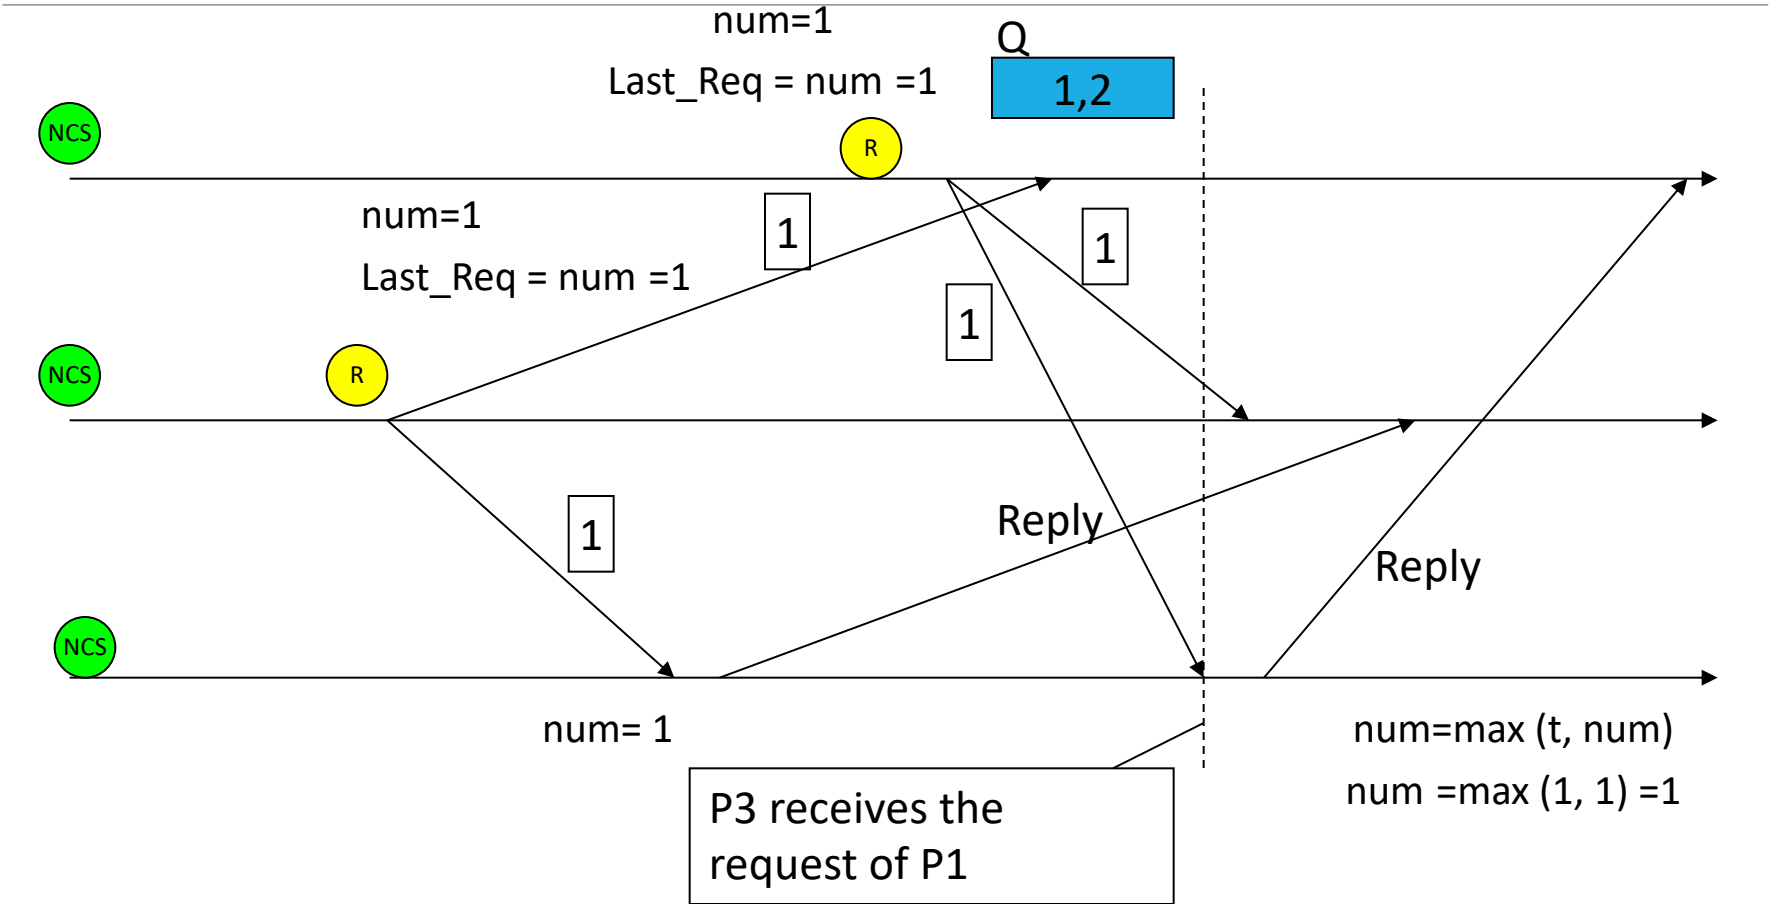


Ricart-Agrawala's algorithm: example




Ricart-Agrawala's algorithm: example

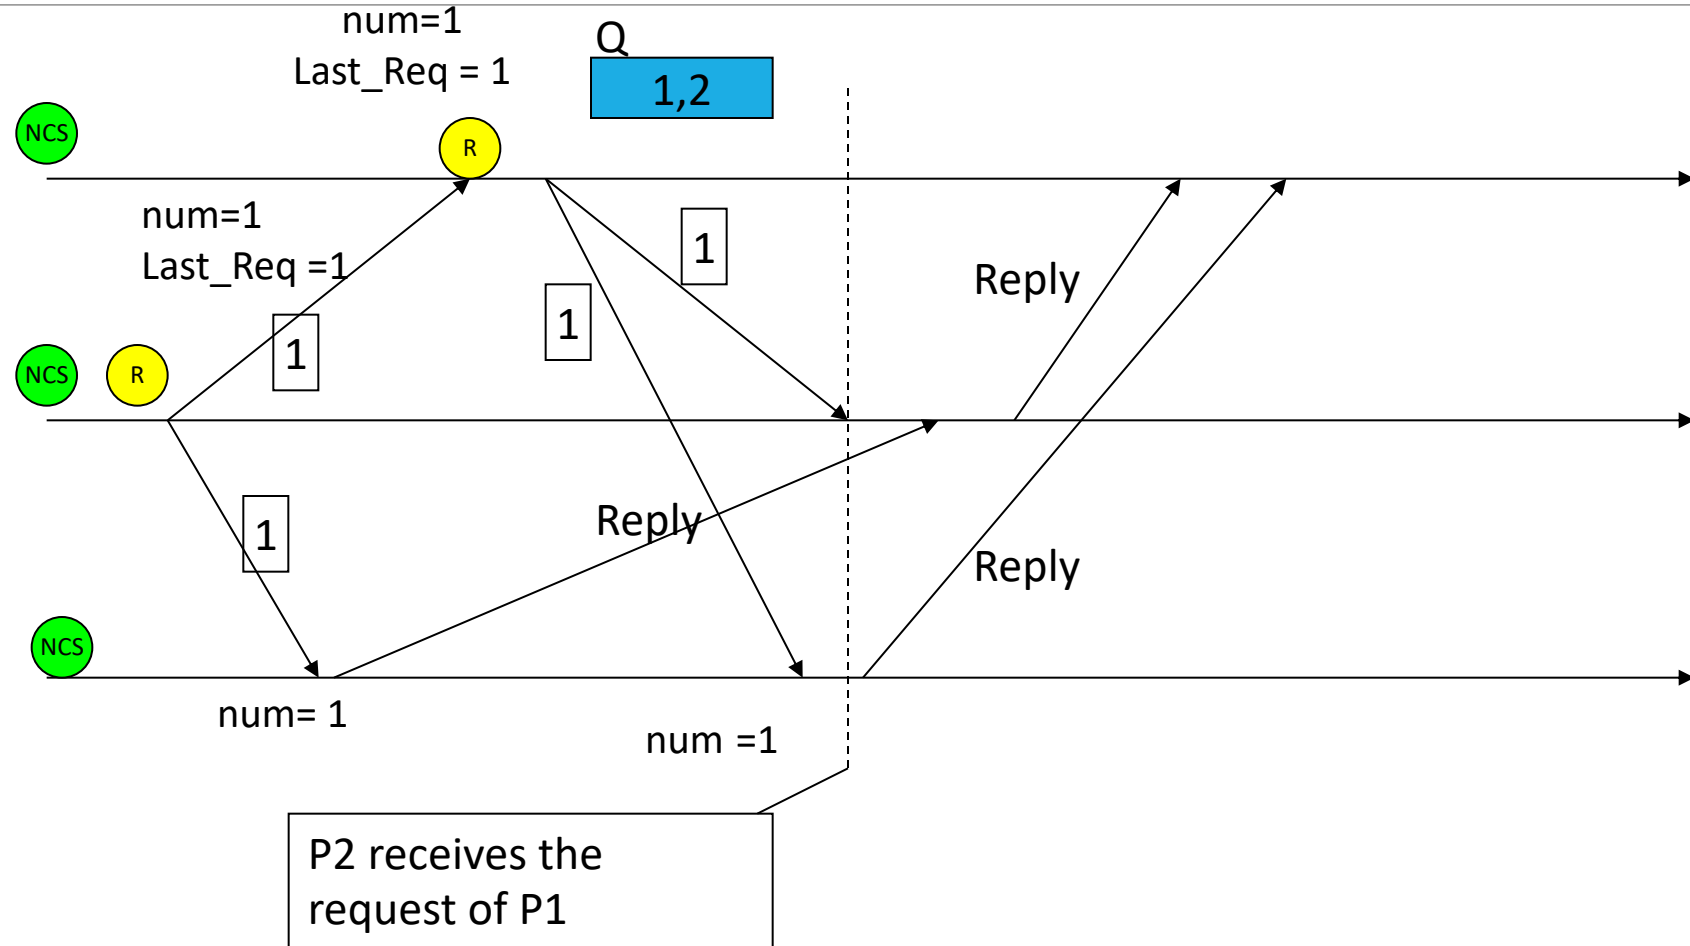





Ricart-Agrawala's algorithm: example




Ricart-Agrawala's algorithm: example



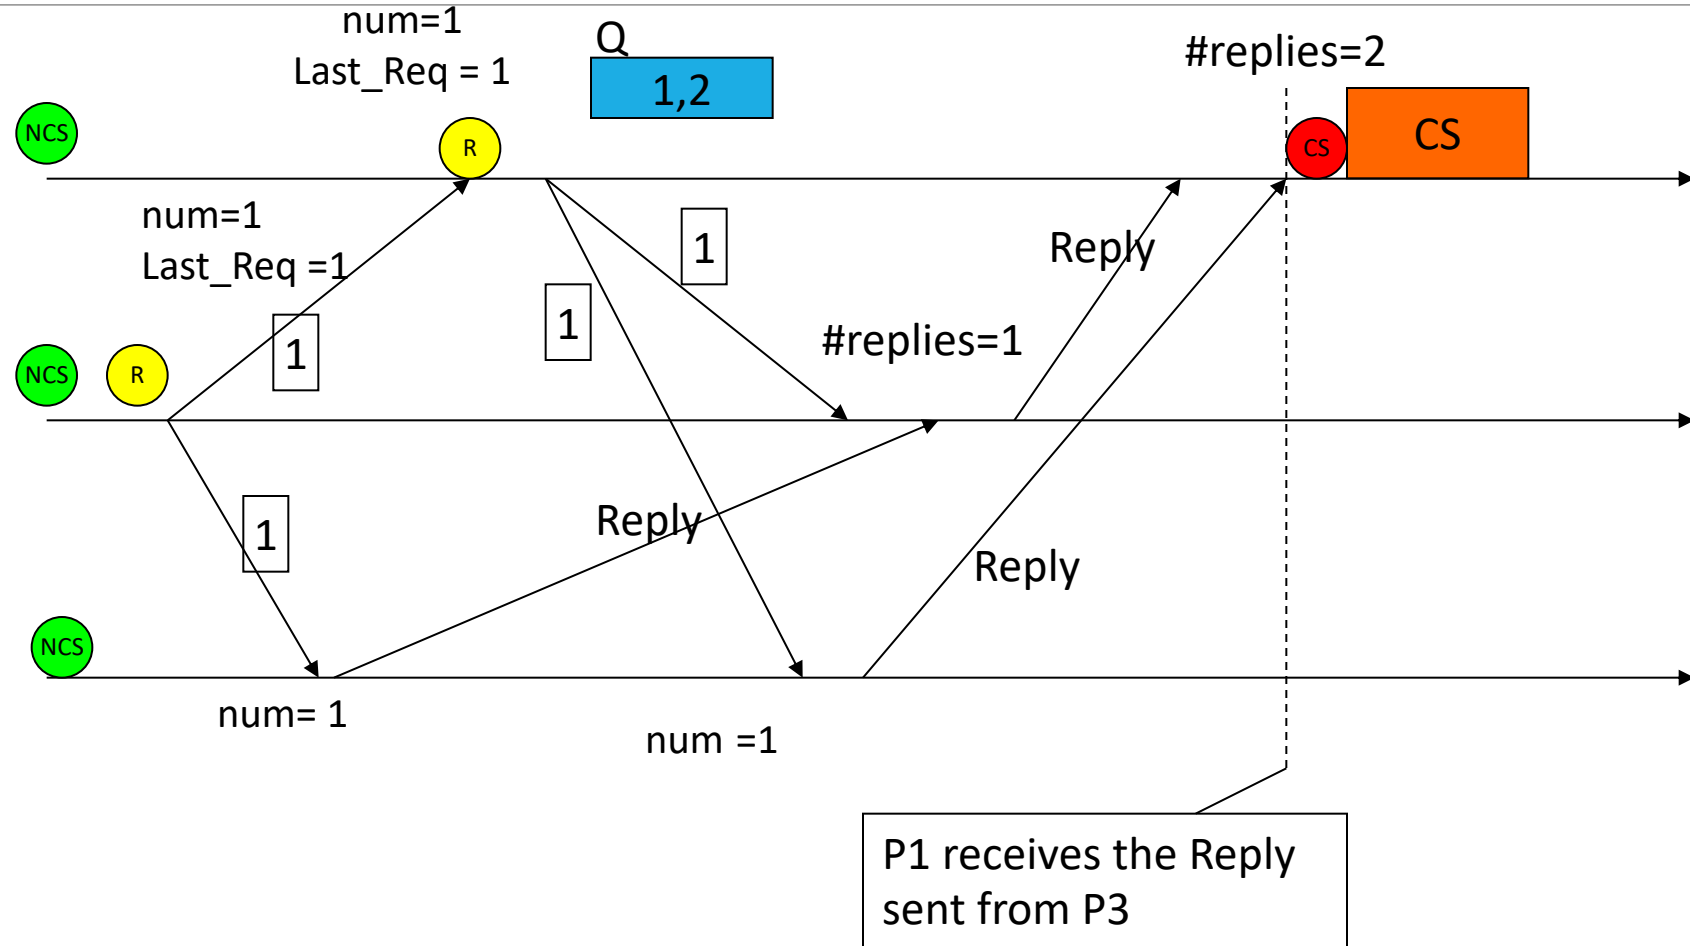
Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example



Ricart-Agrawala's algorithm: example

