

17|10|23

Dependable Distributed Systems

Master of Science in Engineering in Computer Science

AA 2023/2024

LECTURE 9: BROADCAST COMMUNICATIONS (PART 2)

Recap: Best Effort Broadcast (BEB) Specification

Module 3.1: Interface and properties of best-effort broadcast

Module:

Name: BestEffortBroadcast, **instance** *beb*.

Events:

Request: $\langle \text{beb}, \text{Broadcast} \mid m \rangle$: Broadcasts a message *m* to all processes.

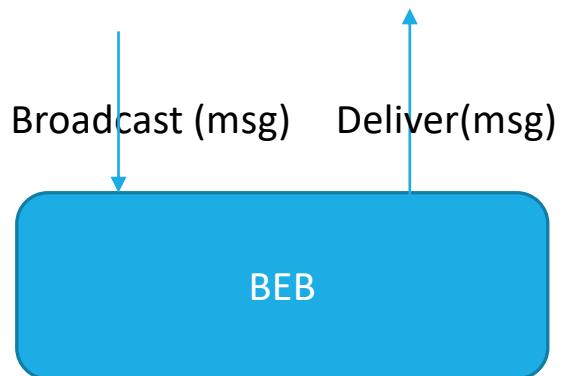
Indication: $\langle \text{beb}, \text{Deliver} \mid p, m \rangle$: Delivers a message *m* broadcast by process *p*.

Properties:

BEB1: *Validity*: If a correct process broadcasts a message *m*, then every correct process eventually delivers *m*.

BEB2: *No duplication*: No message is delivered more than once.

BEB3: *No creation*: If a process delivers a message *m* with sender *s*, then *m* was previously broadcast by process *s*.



Recap: (Regular) Reliable Broadcast (RB)

Module 3.2: Interface and properties of (regular) reliable broadcast

Module:

Name: ReliableBroadcast, **instance** rb .

Events:

Request: $\langle rb, \text{Broadcast} \mid m \rangle$: Broadcasts a message m to all processes.

Indication: $\langle rb, \text{Deliver} \mid p, m \rangle$: Delivers a message m broadcast by process p .

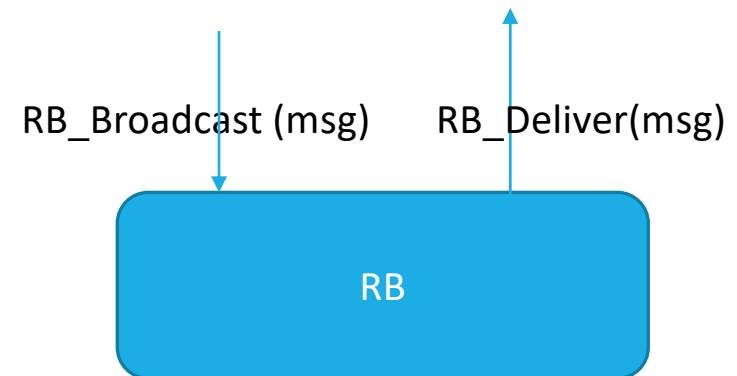
Properties:

RB1: Validity: If a correct process p broadcasts a message m , then p eventually delivers m .

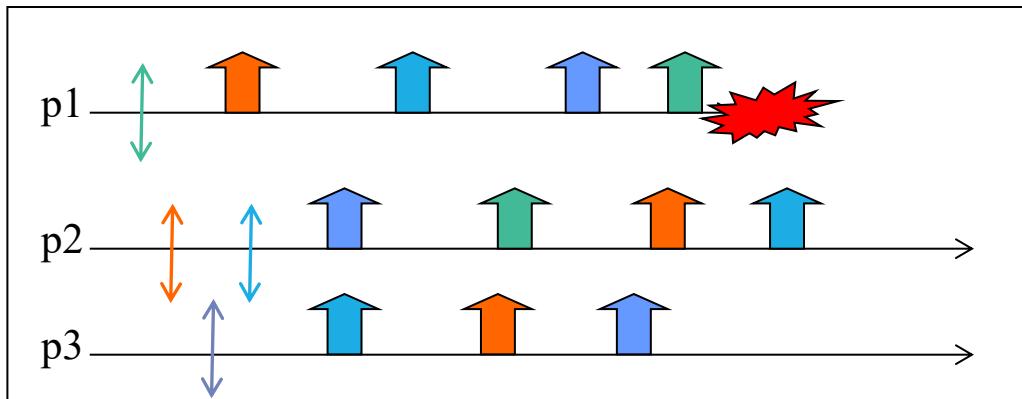
RB2: No duplication: No message is delivered more than once.

RB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

RB4: Agreement: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

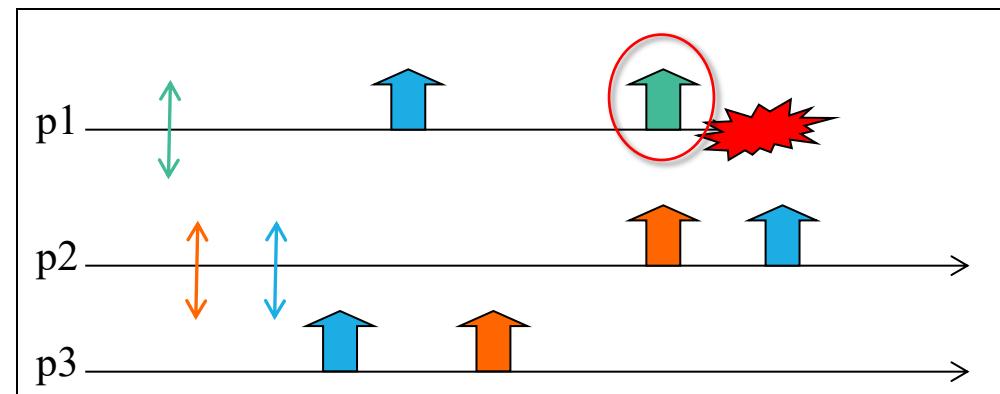


Recap: BEB vs RB



Satisfies BEB but not RB
(violation of the Agreement Property)

Satisfies RB



Uniform Reliable Broadcast (URB) Specification

Module 3.3: Interface and properties of uniform reliable broadcast

Module:

Name: UniformReliableBroadcast, instance *urb*.

Events:

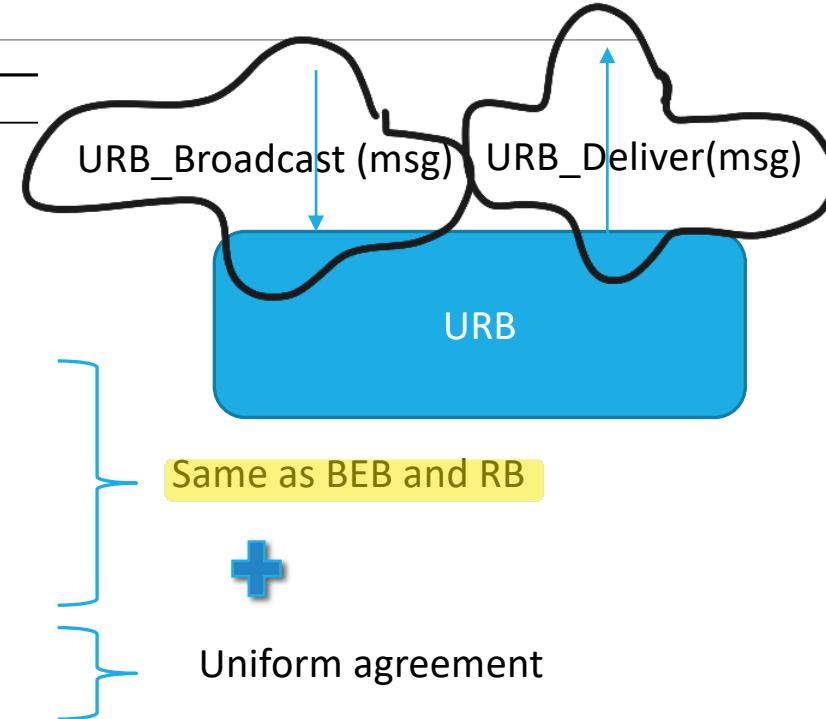
Request: $\langle \text{urb}, \text{Broadcast} \mid m \rangle$: Broadcasts a message *m* to all processes.

Indication: $\langle \text{urb}, \text{Deliver} \mid p, m \rangle$: Delivers a message *m* broadcast by process *p*.

Properties:

URB1–URB3: Same as properties RB1–RB3 in (regular) reliable broadcast (Module 3.2).

URB4: *Uniform agreement*: If a message *m* is delivered by some process (whether correct or faulty), then *m* is eventually delivered by every correct process.

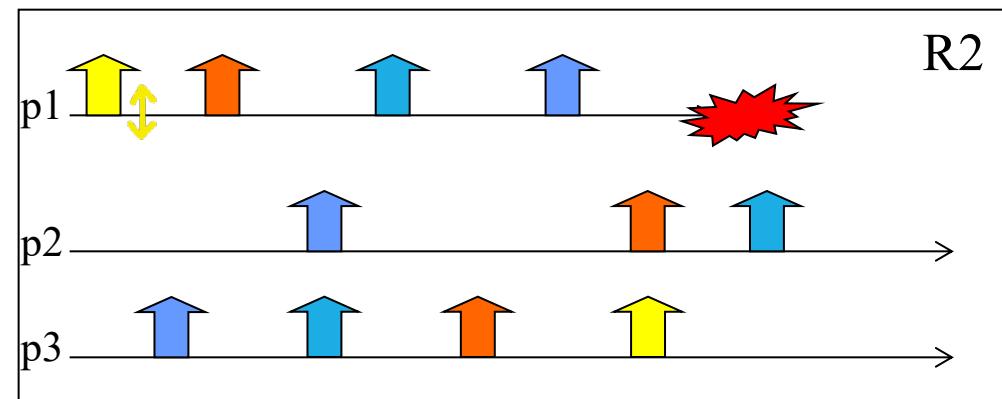
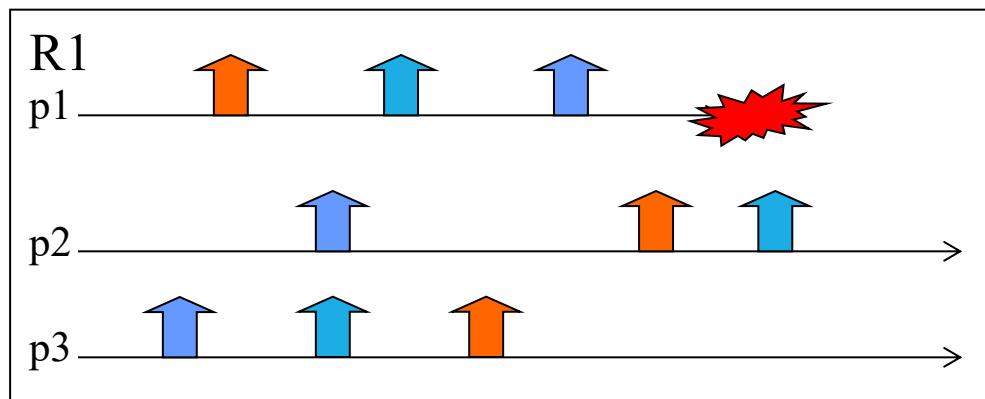


Agreement on a message delivered by any process (crashed or not)!



the set of messages delivered by a correct process is a superset of the ones delivered by a faulty one

BEB vs RB vs URB



RB
URB
BEB

uniform agreement between FAULTY and CORRECT

BEB if yellow message is sent by p1

Non-correct otherwise

RB : NO, set of correct is DIFFERENT

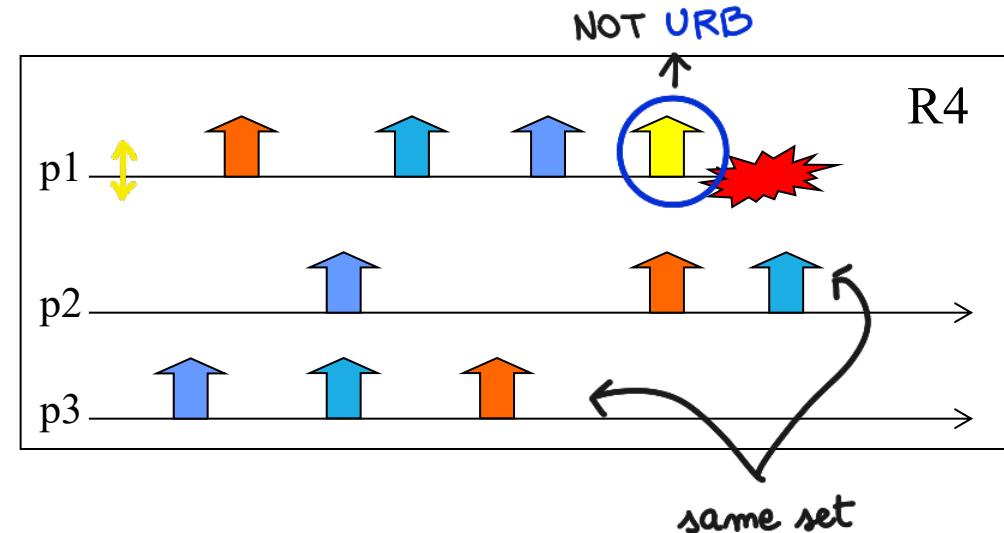
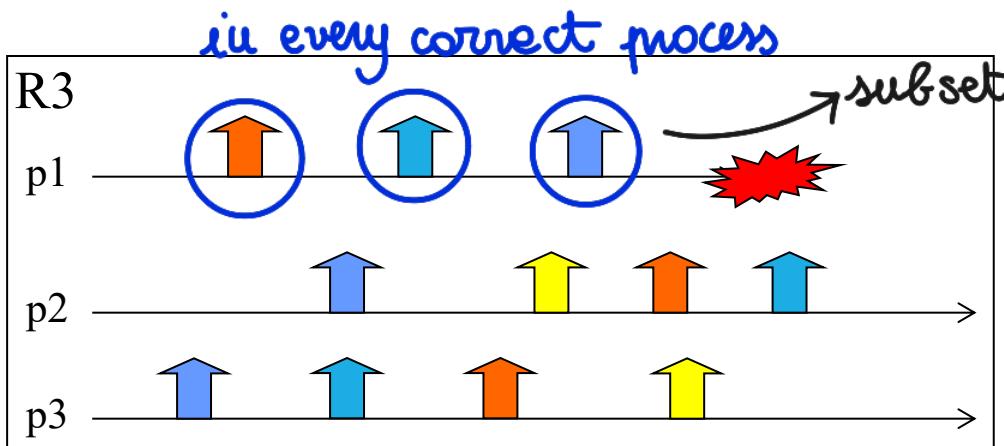
BEB vs RB vs URB

is also BEB, RB
↓
URB

↳ same set

RB if yellow message is sent by p1

Non-correct otherwise



Uniform Reliable Broadcast (URB) Implementation in Synchronous System

Algorithm 3.4: All-Ack Uniform Reliable Broadcast

Implements:

UniformReliableBroadcast, **instance** *urb*.

Uses:

BestEffortBroadcast, **instance** *beb*.

PerfectFailureDetector, **instance** \mathcal{P} .

```

upon event < urb, Init > do
    delivered :=  $\emptyset$ ;
    pending :=  $\emptyset$ ;
    correct :=  $\Pi$ ;
    forall m do ack[m] :=  $\emptyset$ ;

upon event < urb, Broadcast | m > do
    pending := pending  $\cup$  {(self, m)};
    trigger < beb, Broadcast | [DATA, self, m] >

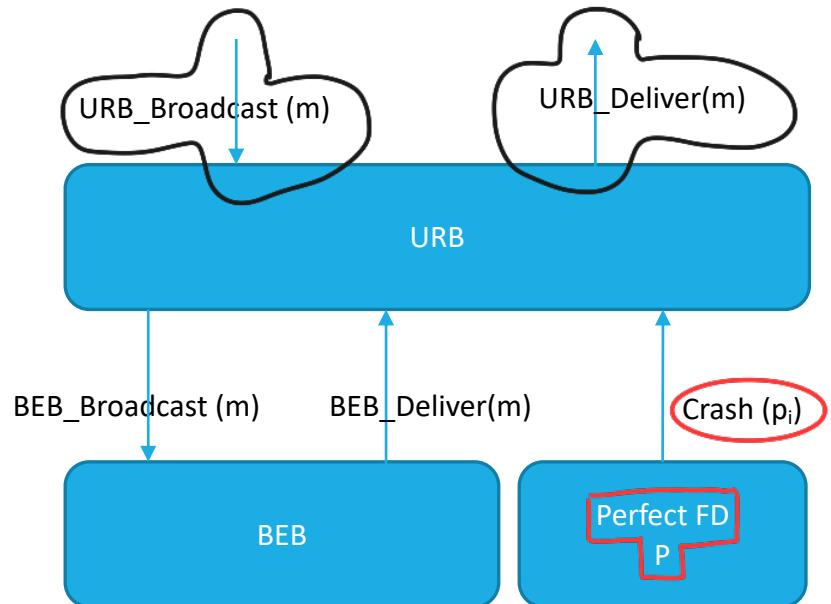
upon event < beb, Deliver | p, [DATA, s, m] > do
    ack[m] := ack[m]  $\cup$  {p};
    if (s, m)  $\notin$  pending then
        pending := pending  $\cup$  {(s, m)};
    trigger < beb, Broadcast | [DATA, s, m] >

upon event <  $\mathcal{P}$ , Crash | p > do
    correct := correct \ {p};

function cadeliver(m) returns Boolean is
    return (correct  $\subseteq$  ack[m]);

upon exists (s, m)  $\in$  pending such that cadeliver(m)  $\wedge$  m  $\notin$  delivered do
    delivered := delivered  $\cup$  {m};
    trigger < urb, Deliver | s, m >

```



- 1) spreading the msg to everybody
- 2) I have to acknowledge the receive of the msg, sending back an ack
- 3) When I have the ack from all the **correct** processes, I know that everybody have seen the msg and we can deliver it

9. URB

Properties:

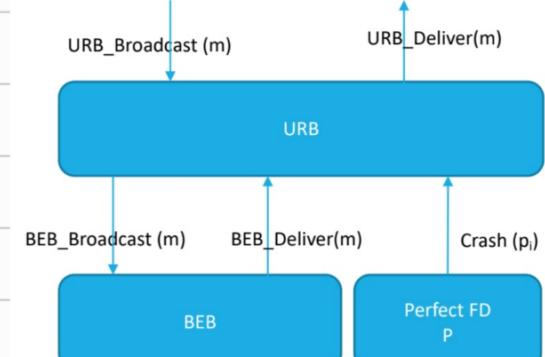
RB1: Validity: If a correct process p broadcasts a message m , then p eventually delivers m .

RB2: No duplication: No message is delivered more than once.

RB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

RB4: Agreement: If a message m is delivered by some correct process, then m is eventually delivered by every correct process.

URB4: Uniform agreement: If a message m is delivered by some process (whether correct or faulty), then m is eventually delivered by every correct process.



Algorithm 3.4: All-Ack Uniform Reliable Broadcast

Implements:

UniformReliableBroadcast, instance urb .

Uses:

BestEffortBroadcast, instance beb .

PerfectFailureDetector, instance \mathcal{P} .

```

upon event < urb, Init > do
  delivered := ∅; set of delivered msg.
  pending := ∅; msg. that I received but waiting the ack
  correct := Π;
  forall m do ack[m] := ∅;

upon event < urb, Broadcast | m > do
  pending := pending ∪ {(self, m)}; processes that send the msg. to all
  trigger < beb, Broadcast | [DATA, self, m] >; processes → store msg. locally
  if (s, m) ∈ pending then msg. is in pending?
    pending := pending ∪ {(s, m)};
    trigger < beb, Broadcast | [DATA, s, m] >; → forward it

upon event < beb, Deliver | p, [DATA, s, m] > do
  ack[m] := ack[m] ∪ {p}; insert name of the sender in ack
  if (s, m) ∈ pending then msg. is in pending?
    pending := pending ∖ {(s, m)};
    trigger < beb, Broadcast | [DATA, s, m] >; → forward it

upon event < P, Crash | p > do
  correct := correct ∖ {p}; Fail ⇒ process removed from correct set

function candeliver(m) returns Boolean is
  return (correct ⊆ ack[m]); ] is TRUE when the ack is from all correct

upon exists (s, m) ∈ pending such that candeliver(m) ∧ m ∉ delivered do for NO DUPLICATION
  delivered := delivered ∪ {m};
  trigger < urb, Deliver | s, m >;
  ↙ can deliver the msg. when the cond. = TRUE
  
```

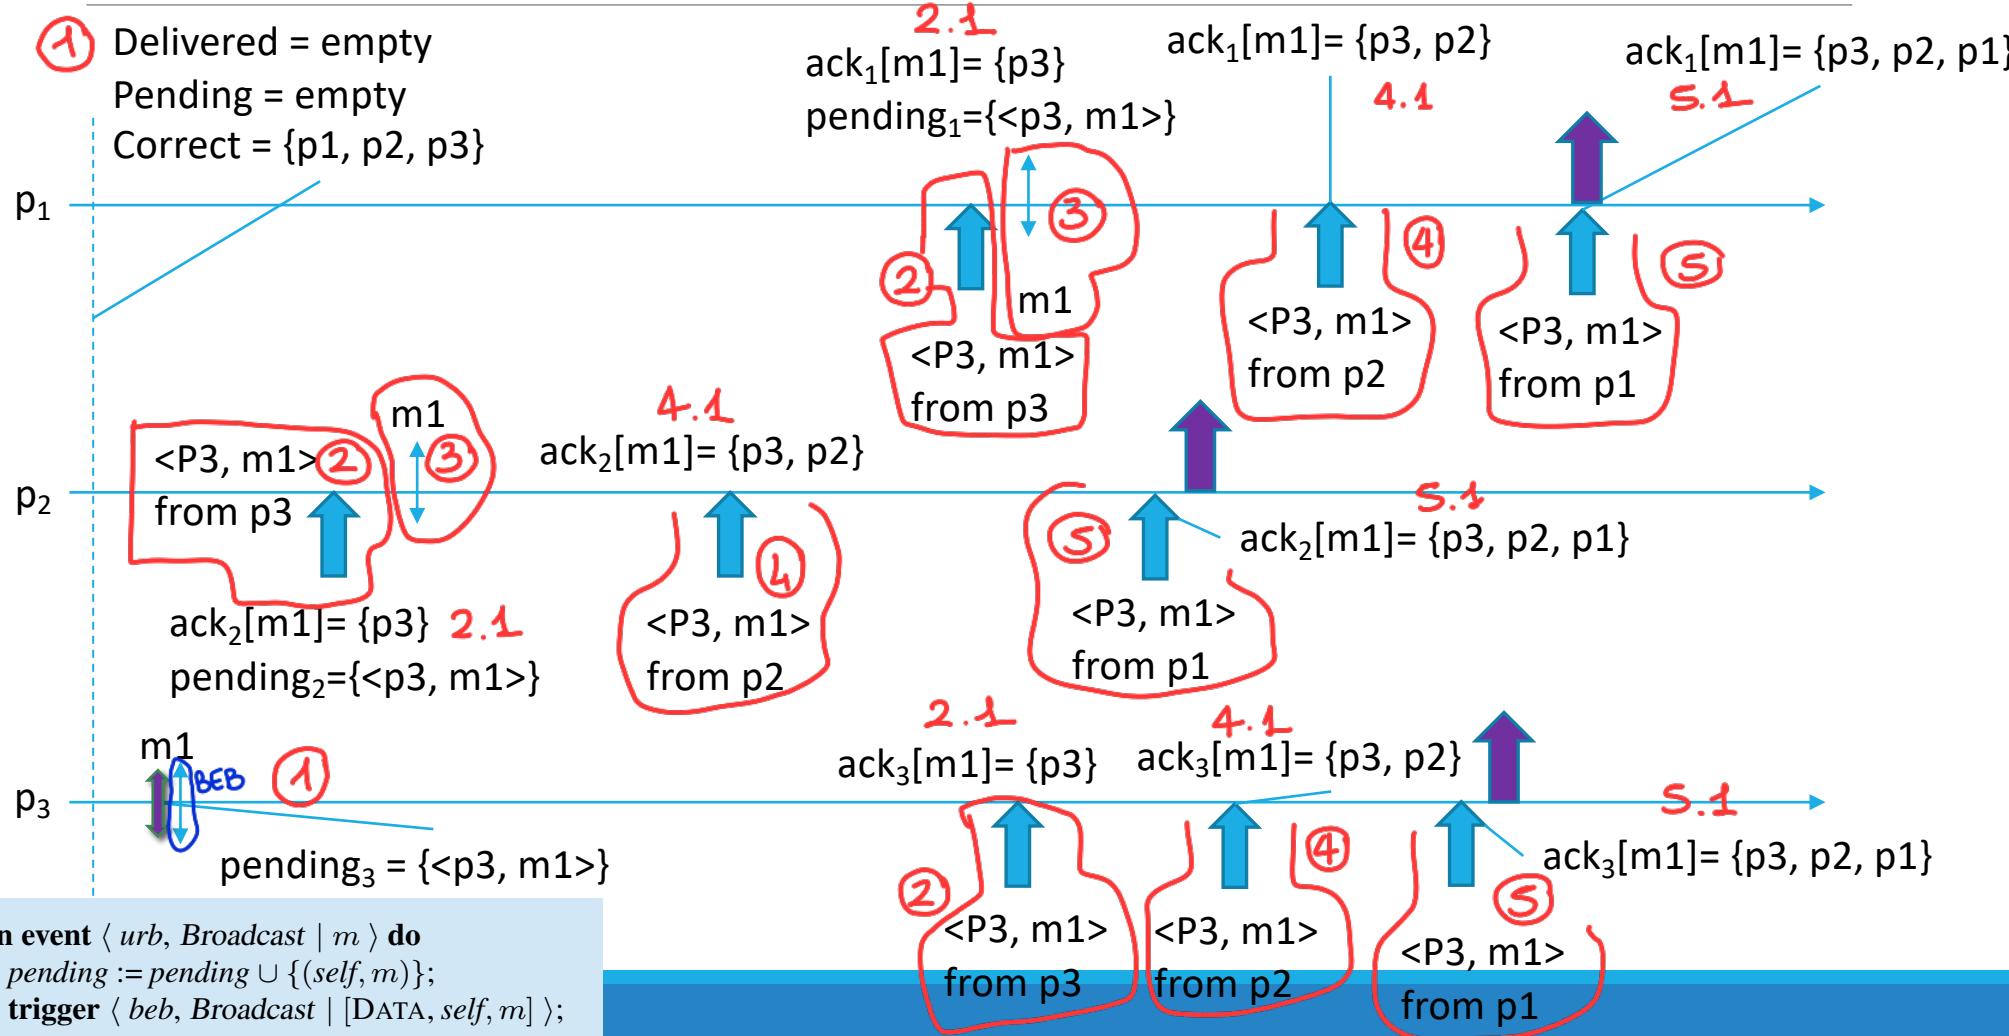
Example

- ① Delivered = empty
- Pending = empty
- Correct = {p1, p2, p3}

② **upon event** $\langle beb, Deliver \mid p, [\text{DATA}, s, m] \rangle$ **do**
 $ack[m] := ack[m] \cup \{p\};$
if $(s, m) \notin pending$ **then**
 $pending := pending \cup \{(s, m)\};$
trigger $\langle beb, Broadcast \mid [\text{DATA}, s, m] \rangle;$

③ **function** $candeliver(m)$ **returns** Boolean **is**
return $(correct \subseteq ack[m]);$

upon exists $(s, m) \in pending$ **such that** $candeliver(m) \wedge m \notin delivered$ **do**
 $delivered := delivered \cup \{m\};$
trigger $\langle urb, Deliver \mid s, m \rangle;$



Example

Delivered = empty
 Pending = empty
 Correct = {p1, p2, p3}

p_1

$ack_1[m1] = \{p3\}$
 $pending_1 = \{<p3, m1>\}$

$m1$
 $<P3, m1>$
 from p3

p_2

$ack_2[m1] = \{p3, p2\}$
 $pending_2 = \{<p3, m1>\}$

$<P3, m1>$
 from p2

p_3

$ack_2[m1] = \{p3\}$
 $pending_2 = \{<p3, m1>\}$

$pending_3 = \{<p3, m1>\}$

```

upon event ( P, Crash | p ) do
  correct := correct \ {p};

function candeliver(m) returns Boolean is
  return (correct ⊆ ack[m]);

upon exists (s, m) ∈ pending such that candeliver(m) ∧ m ∉ delivered do
  delivered := delivered ∪ {m};
  trigger ( urb, Deliver | s, m );
  
```

$<P3, m1>$
 from p3
 $<P3, m1>$
 from p2

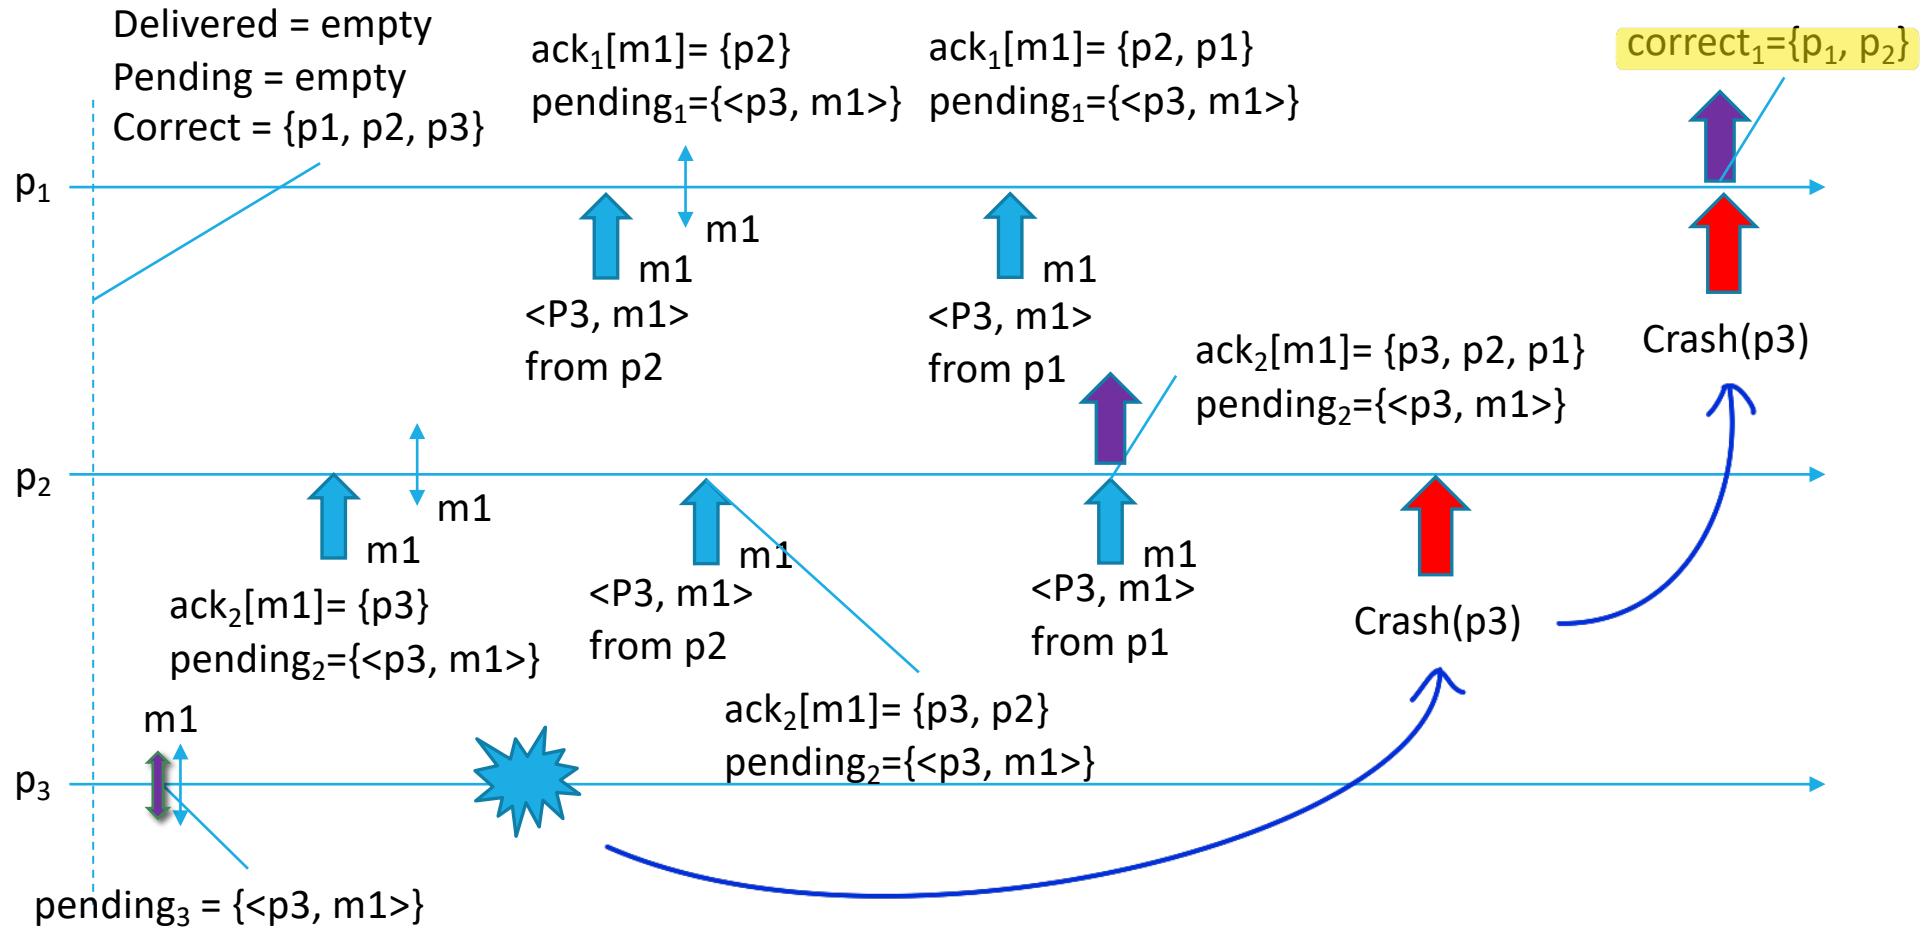
$ack_1[m1] = \{p3, p2, p1\}$
 $pending_1 = \{<p3, m1>\}$

$<P3, m1>$
 from p1
 $correct_2 = \{p_2, p_3\}$

Crash(p1)

Crash(p1)

Example



Uniform Reliable Broadcast (URB) Implementation in Asynchronous System

Algorithm 3.5 Majority-Ack Uniform Reliable Broadcast

Implements:

UniformReliableBroadcast (urb).

Extends:

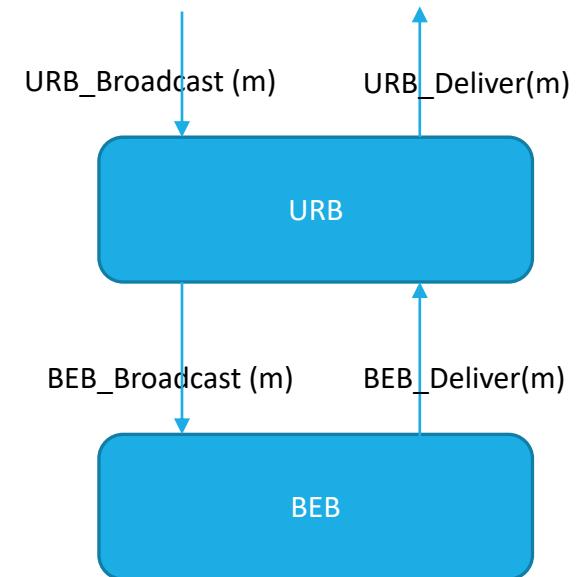
All-Ack Uniform Reliable Broadcast (Algorithm 3.4).

Uses:

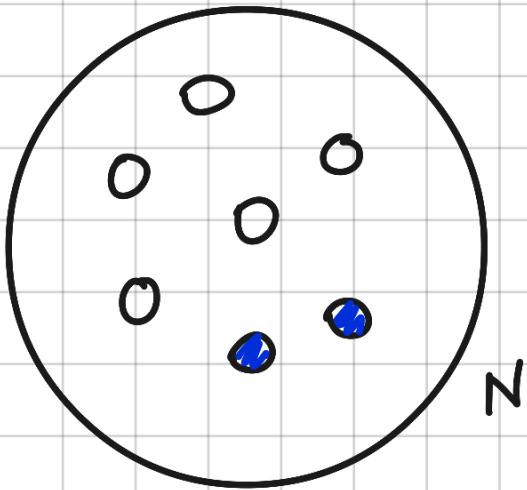
BestEffortBroadcast (beb).

```
function canDeliver(m) returns boolean is
    return (|ackm| > N/2);
```

```
// Except for the function above, and the non-use of the
// perfect failure detector, same as Algorithm 3.4.
```



We need to assume a majority of correct processes

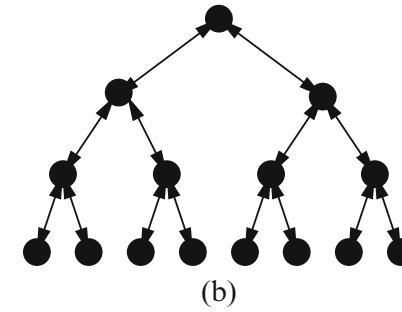
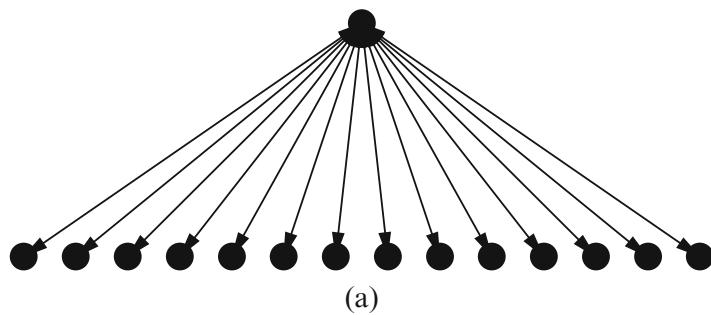


* ACK \geq n - FAILURE

Uniform Reliable Broadcast

- There exists an algorithm for synchronous system using Perfect failure detector
- There exists an algorithm for asynchronous system when assuming a “*majority of correct processes*”
- Can we devise a uniform reliable broadcast algorithm for a partially synchronous system (using an eventually perfect failure detector) but without the assumption of a majority of correct processes?

Ack Implosion and ack tree



Problems

1. Process spends all its time by doing the ack task
2. Maintaining the tree structure

Acks make
reliable broadcast
not scalable!

Probabilistic broadcast

Message delivered 99% of the times

Not fully reliable

Suitable for large & dynamic groups

Probabilistic Broadcast

Module 3.7: Interface and properties of probabilistic broadcast

Module:

Name: ProbabilisticBroadcast, **instance pb**.

Events:

Request: $\langle pb, Broadcast \mid m \rangle$: Broadcasts a message m to all processes.

Pb_Broadcast(msg)

Indication: $\langle pb, Deliver \mid p, m \rangle$: Delivers a message m broadcast by process p .

Pb_Deliver(msg)

PbB

Properties:

PB1: Probabilistic validity: There is a positive value ε such that when a correct process broadcasts a message m , the probability that every correct process eventually delivers m is at least $1 - \varepsilon$.

PB2: No duplication: No message is delivered more than once.

PB3: No creation: If a process delivers a message m with sender s , then m was previously broadcast by process s .

Gossip Dissemination

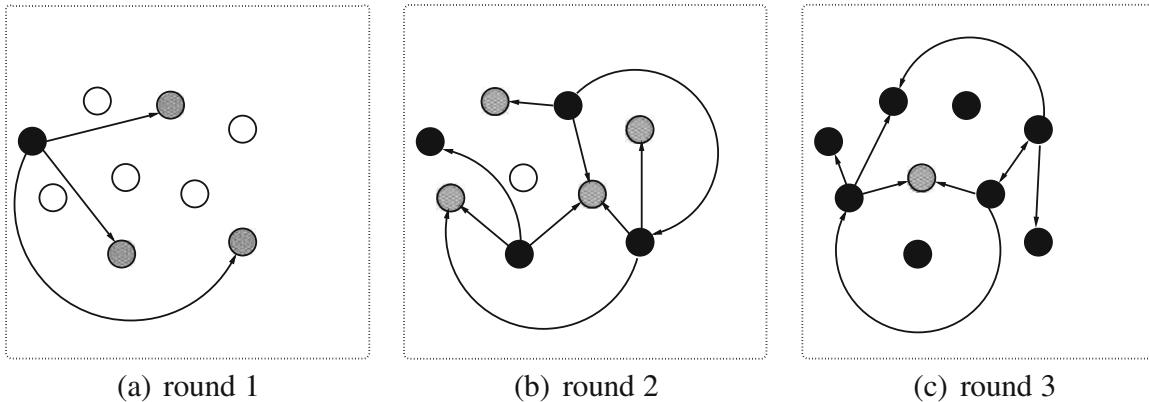


Figure 3.6: Epidemic dissemination or gossip (with fanout 3)

- A process sends a message to a set of randomly chosen k processes
- A process receiving a message for the first time, forwards it to a set of k randomly chosen processes (this operation is also called a round)
- The algorithm performs a maximum number of r rounds

Eager Probabilistic Broadcast

Algorithm 3.9: Eager Probabilistic Broadcast

Implements:

ProbabilisticBroadcast, **instance** *pb*.

Uses:

FairLossPointToPointLinks, **instance** *fl*.

upon event $\langle pb, \text{Init} \rangle$ **do**
 delivered := \emptyset ;

procedure *gossip*(*msg*) **is**
 forall *t* \in *picktargets*(*k*) **do** **trigger** $\langle fl, \text{Send} \mid t, msg \rangle$;

upon event $\langle pb, \text{Broadcast} \mid m \rangle$ **do**
 delivered := *delivered* \cup {*m*};
 trigger $\langle pb, \text{Deliver} \mid self, m \rangle$;
 gossip([GOSSIP, *self*, *m*, *R*]);

upon event $\langle fl, \text{Deliver} \mid p, [\text{GOSSIP}, s, m, r] \rangle$ **do**
 if *m* \notin *delivered* **then**
 delivered := *delivered* \cup {*m*};
 trigger $\langle pb, \text{Deliver} \mid s, m \rangle$;
 if *r* > 1 **then** *gossip*([GOSSIP, *s*, *m*, *r* - 1]);

function *picktargets*(*k*) **returns** set of processes **is**
 targets := \emptyset ;
 while #(targets) $<$ *k* **do**
 candidate := *random*($\Pi \setminus \{\text{self}\}$);
 if *candidate* \notin *targets* **then**
 targets := *targets* \cup {*candidate*};
 return *targets*;

9. PB

PB1: *Probabilistic validity:* There is a positive value ε such that when a correct process broadcasts a message m , the probability that every correct process eventually delivers m is at least $1 - \varepsilon$.

PB2: *No duplication:* No message is delivered more than once.

PB3: *No creation:* If a process delivers a message m with sender s , then m was previously broadcast by process s .

Algorithm 3.9: Eager Probabilistic Broadcast

Implements:

ProbabilisticBroadcast, **instance** pb .

Uses:

FairLossPointToPointLinks, **instance** fll .

upon event $\langle pb, \text{Init} \rangle$ **do**
 $delivered := \emptyset;$

procedure $\text{gossip}(msg)$ **is**
 forall $t \in \text{picktargets}(k)$ **do** **trigger** $\langle fll, \text{Send} \mid t, msg \rangle$;

upon event $\langle pb, \text{Broadcast} \mid m \rangle$ **do**
 $delivered := delivered \cup \{m\};$
 trigger $\langle pb, \text{Deliver} \mid self, m \rangle;$
 $\text{gossip}([\text{GOSSIP}, self, m, R]);$

upon event $\langle fll, \text{Deliver} \mid p, [\text{GOSSIP}, s, m, r] \rangle$ **do**
 if $m \notin delivered$ **then**
 $delivered := delivered \cup \{m\};$
 trigger $\langle pb, \text{Deliver} \mid s, m \rangle;$
 if $r > 1$ **then** $\text{gossip}([\text{GOSSIP}, s, m, r - 1]);$

```
function picktargets(k) returns set of processes is
    targets := ∅;
    while #(targets) < k do
        candidate := random(Π \ {self});
        if candidate ∉ targets then
            targets := targets ∪ {candidate};
    return targets;
```

References

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 3 - from Section 3.9 (except 3.9.6)
- Chapter 6 – Section 6.1