

18|10|23

# Dependable Distributed Systems

## Master of Science in Engineering in Computer Science

AA 2023/2024

---

LECTURE 10: CONSENSUS

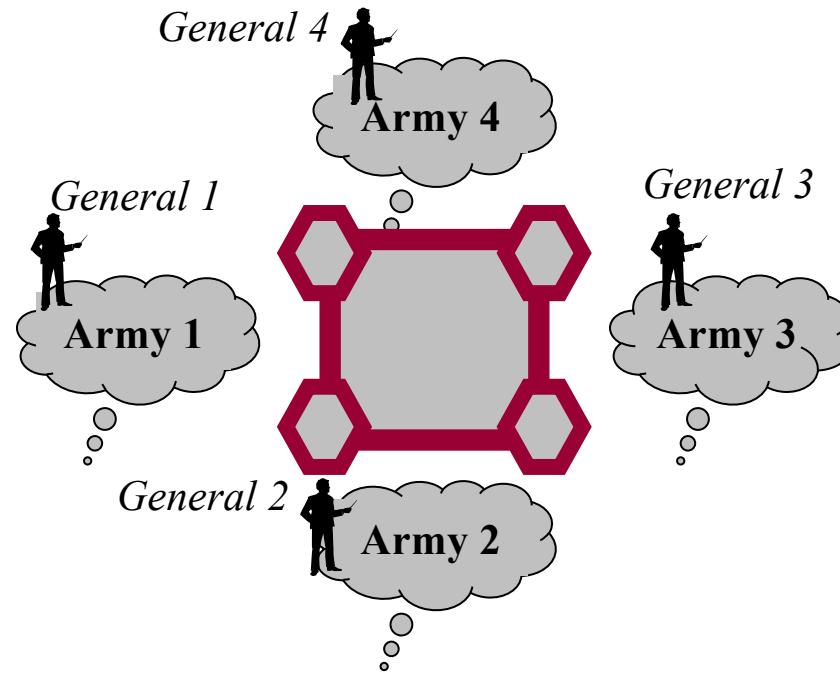
# Consensus Problem

---

- A group of processes must agree on a value that has been proposed by one of them (e.g., commit/abort of a transaction).
- It is an abstraction of a class of problems where processes start with their opinion and then converge on one of them
- It is a fundamental problem
- We study algorithms working on weak models

# Consensus Example

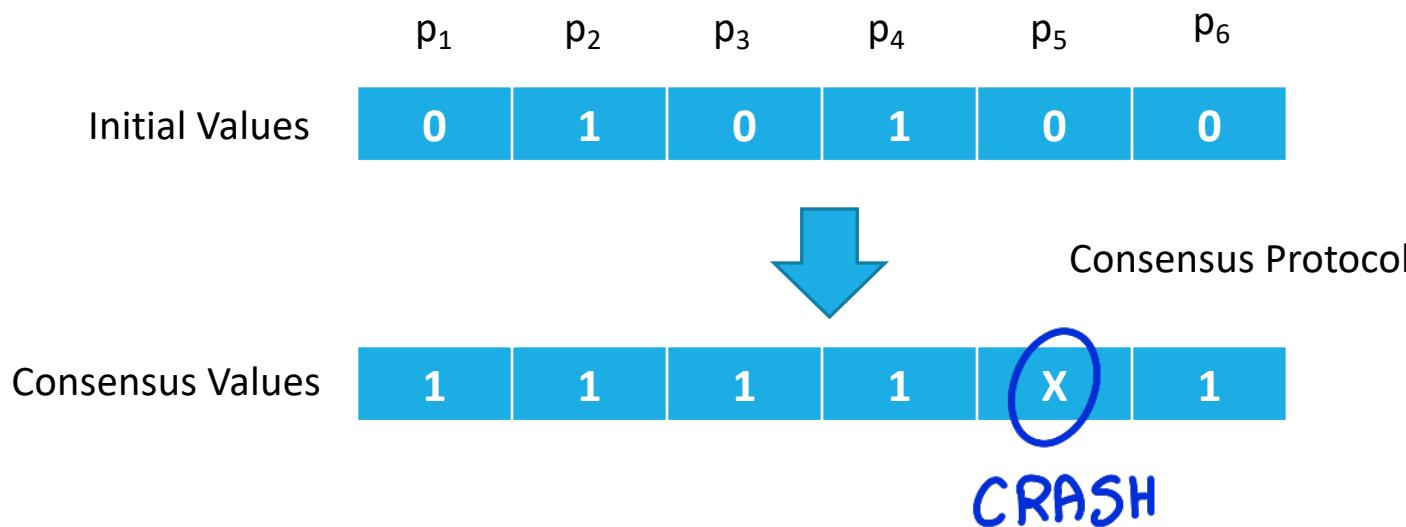
---



Generals have to reach consensus between attack and get back

# Consensus Definition

- Set of initial values  $\in \{0,1\}$ .
- Every process has to decide the same value  $\in \{0,1\}$  based on the initial proposals.



# Consensus Specification

---

**Module 5.1:** Interface and properties of (regular) consensus

**Module:**

**Name:** Consensus, **instance**  $c$ .

**Events:**

**Request:**  $\langle c, \text{Propose} \mid v \rangle$ : Proposes value  $v$  for consensus.

1)

**Indication:**  $\langle c, \text{Decide} \mid v \rangle$ : Outputs a decided value  $v$  of consensus.

2)

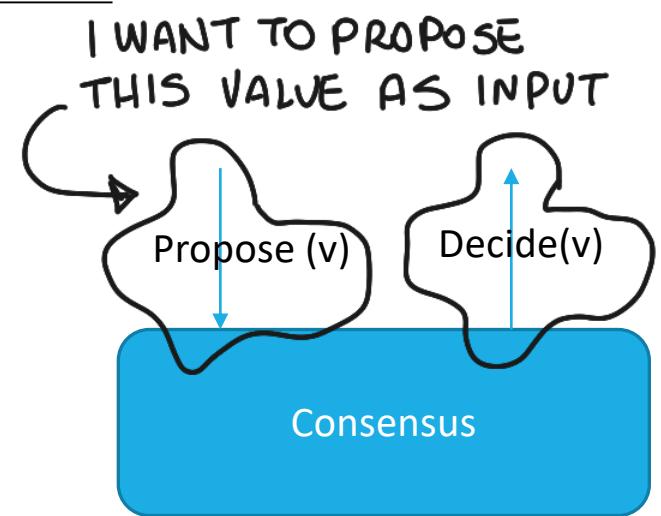
**Properties:**

**C1:** *Termination*: Every correct process eventually decides some value.

**C2:** *Validity*: If a process decides  $v$ , then  $v$  was proposed by some process.

**C3:** *Integrity*: No process decides twice.  $\downarrow$  IMPLIES

**C4:** *Agreement*: No two correct processes decide differently.



FLP

# Impossibility Result



no algorithm can guarantee to reach  
consensus in an asynchronous system, even  
with one process crash failure.

Fisher, Lynch e Patterson (FLP result)

Ref: Journal of the ACM, Vol. 32, No. 2, April 1985.

# A Consensus Implementation in Synchronous Systems: Flooding Consensus

---

Basic Idea:

- Processes exchange their values
- when all the proposals from correct processes are available a one value can be chosen

Problem:

- due to failures, some values can be lost

Solution:

- A value can be selected only when no failures happen during the communication

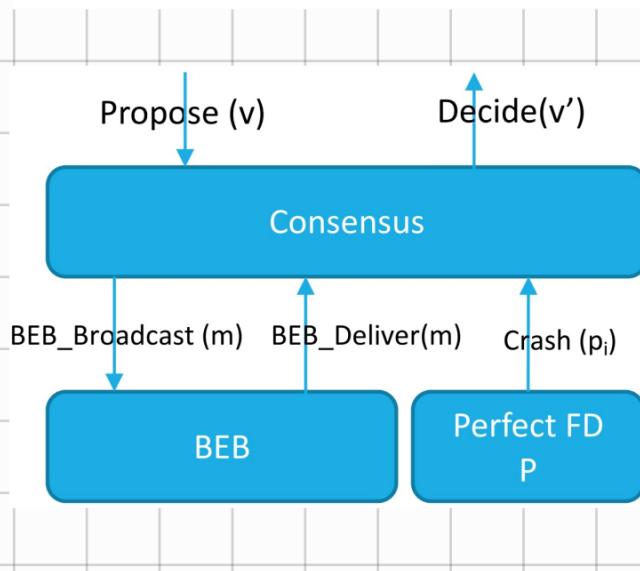
# 10: CONSENSUS

**C1: Termination:** Every correct process eventually decides some value.

**C2: Validity:** If a process decides  $v$ , then  $v$  was proposed by some process.

**C3: Integrity:** No process decides twice.

**C4: Agreement:** No two correct processes decide differently.



**Algorithm 5.1: Flooding Consensus**

**Implements:**

Consensus, **instance**  $c$ .

**Uses:**

BestEffortBroadcast, **instance**  $beb$ ;  
PerfectFailureDetector, **instance**  $\mathcal{P}$ .

```

upon event < c, Init > do
    correct :=  $\Pi$ ; all processes
    round := 1; count no of interactions before decision
    decision :=  $\perp$ ; store value that we are going to decide
    receivedfrom :=  $[\emptyset]^N$ ; list composed by  $n$  entries, one
    proposals :=  $[\emptyset]^N$ ; for each process in the system
    receivedfrom[0] :=  $\Pi$ ; msgs received from a p
    round[0] := 0; in a specific round
    value_proposed :=  $\perp$ ; all processes
    for each process in the system
        proposals[0] :=  $\Pi$ ; msgs received from a p
        round[0] := 0; in a specific round
    endfor
upon event <  $\mathcal{P}$ , Crash |  $p$  > do
    correct := correct \ { $p$ };

```

```

upon event < c, Propose |  $v$  > do
    proposals[1] := proposals[1]  $\cup$  { $v$ };
    trigger < beb, Broadcast | [PROPOSAL, 1, proposals[1]] >;
    ROUND = 1
    set of proposal for round 1

```

```

upon event < beb, Deliver |  $p$ , [PROPOSAL,  $r$ ,  $ps$ ] > do
    receivedfrom[ $r$ ] := receivedfrom[ $r$ ]  $\cup$  { $p$ }; sender of proposal
    proposals[ $r$ ] := proposals[ $r$ ]  $\cup$   $ps$ ;
    round[ $r$ ] :=  $r$ ; proposal set
upon correct  $\subseteq$  receivedfrom[round]  $\wedge$  decision =  $\perp$  do = TRUE
    if receivedfrom[round] = receivedfrom[round - 1] then
        decision := min(proposals[round]); min value in set
        trigger < beb, Broadcast | [DECIDED, decision] >;
        trigger < c, Decide | decision >;
    else = FALSE, some process failed
        round := round + 1;
        trigger < beb, Broadcast | [PROPOSAL, round, proposals[round - 1]] >;
    endif
    not decided yet
    every correct p is in receivedfrom
    check if the p have crashed in the current round, checking the previous round
upon event < beb, Deliver |  $p$ , [DECIDED,  $v$ ] > such that  $p \in$  correct  $\wedge$  decision =  $\perp$  do
    decision :=  $v$ ;
    trigger < beb, Broadcast | [DECIDED, decision] >;
    trigger < c, Decide | decision >;
    IF P1 DECIDE BEFORE P2,
    P1 TAKE THE DECISION
    
```

# A Consensus Implementation in Synchronous Systems: Flooding Consensus

## Algorithm 5.1: Flooding Consensus

### Implements:

Consensus, **instance**  $c$ .

### Uses:

BestEffortBroadcast, **instance**  $beb$ ;  
PerfectFailureDetector, **instance**  $\mathcal{P}$ .

```

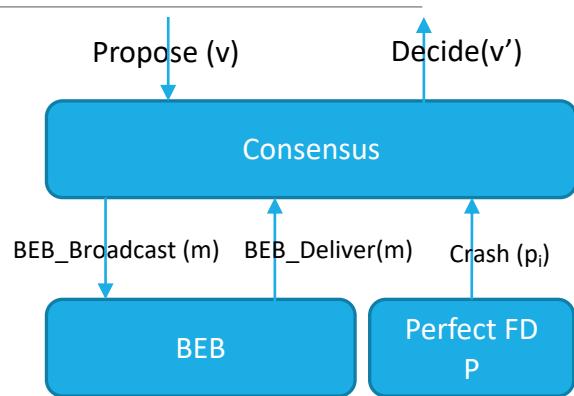
upon event <  $c$ , Init > do
    correct :=  $\Pi$ ; all processes
    round := 1; count no of interaction before decision
    decision :=  $\perp$ ; store value that we are going to decide
    receivedfrom :=  $[\emptyset]^N$ ; list composed by  $n$  entries, one
    proposals :=  $[\emptyset]^N$ ; for each process in the system
    receivedfrom[0] :=  $\Pi$ ; msgs received from a p
    proposals[0] :=  $\emptyset$ ; in a specific round
    round 0
upon event <  $\mathcal{P}$ , Crash |  $p$  > do
    correct := correct \ { $p$ };
upon event <  $c$ , Propose |  $v$  > do
    proposals[1] := proposals[1]  $\cup$  { $v$ };
    trigger <  $beb$ , Broadcast | [PROPOSAL, 1, proposals[1]] >;
    ROUND = 1
    set of proposal
    for round 1
    ROUND

```

```

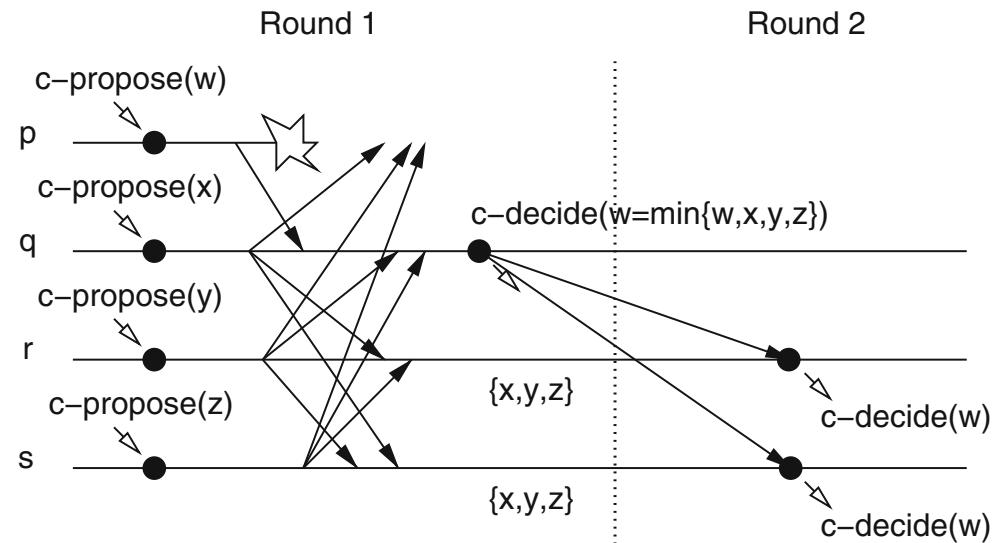
upon event <  $beb$ , Deliver |  $p$ , [PROPOSAL,  $r$ ,  $ps$ ] > do
    receivedfrom[ $r$ ] := receivedfrom[ $r$ ]  $\cup$  { $p$ }; sender of proposal
    proposals[ $r$ ] := proposals[ $r$ ]  $\cup$   $ps$ ;
    every correct p is in receivedfrom
    upon correct  $\subseteq$  receivedfrom[round]  $\wedge$  decision =  $\perp$  do = TRUE
        if receivedfrom[round] = receivedfrom[round - 1] then → check if the p have
            decision := min(proposals[round]); min value in set
            trigger <  $beb$ , Broadcast | [DECIDED, decision] >;
            trigger <  $c$ , Decide | decision >;
        else = FALSE, someone failed
            round := round + 1;
            trigger <  $beb$ , Broadcast | [PROPOSAL, round, proposals[round - 1]] >;
upon event <  $beb$ , Deliver |  $p$ , [DECIDED,  $v$ ] > such that  $p \in$  correct  $\wedge$  decision =  $\perp$  do
    decision :=  $v$ ;
    trigger <  $beb$ , Broadcast | [DECIDED, decision] >;
    trigger <  $c$ , Decide | decision >;
    IF P1 DECIDE BEFORE P2,
    P1 TAKE THE DECISION

```

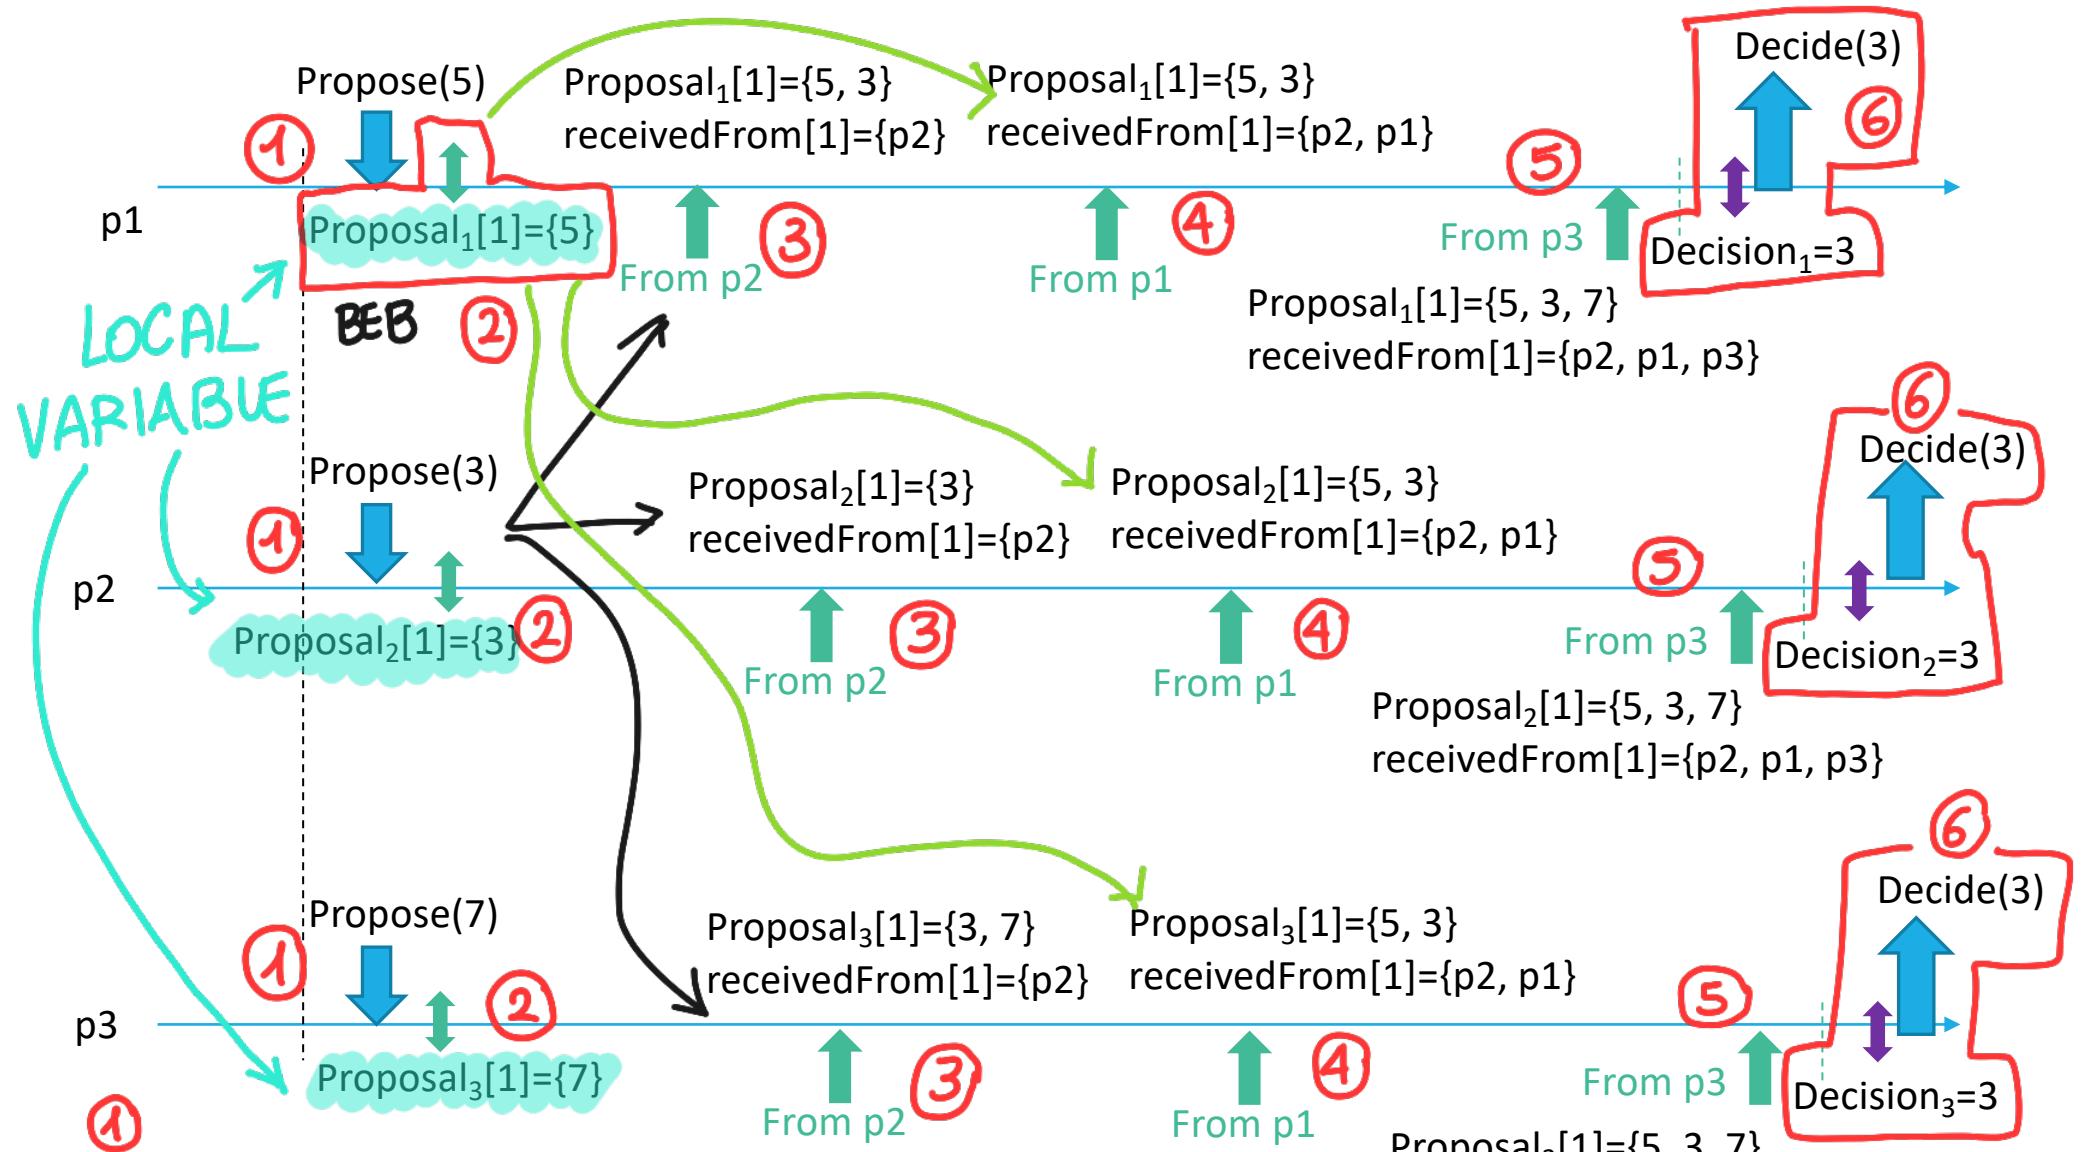


# Example of Flooding Consensus

---



**Figure 5.1:** Sample execution of flooding consensus



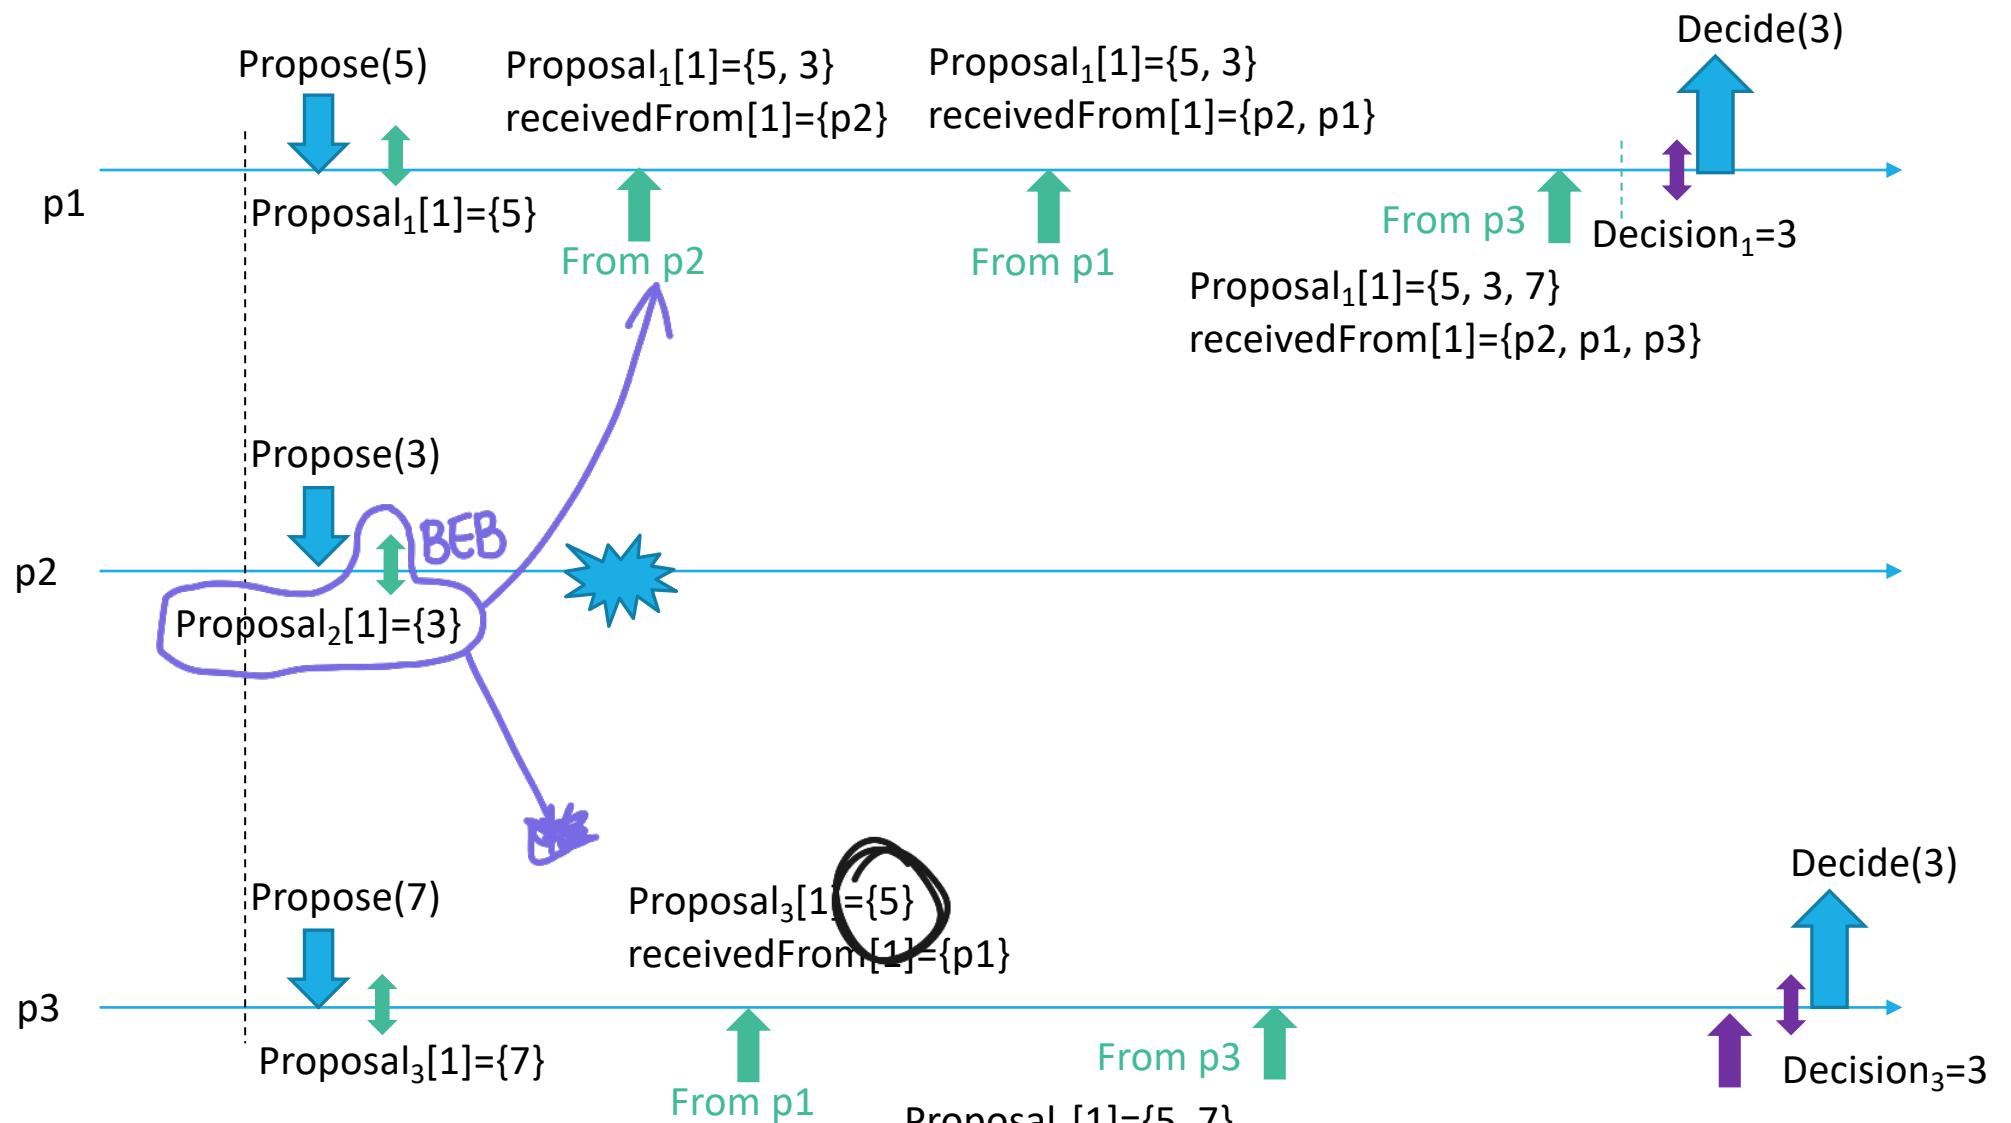
Correct={p1, p2, p3}

Round=1

Decision= $\perp$

receivedFrom[0]={p1, p2, p3} **INIT. STEP**

Proposals[0]={}



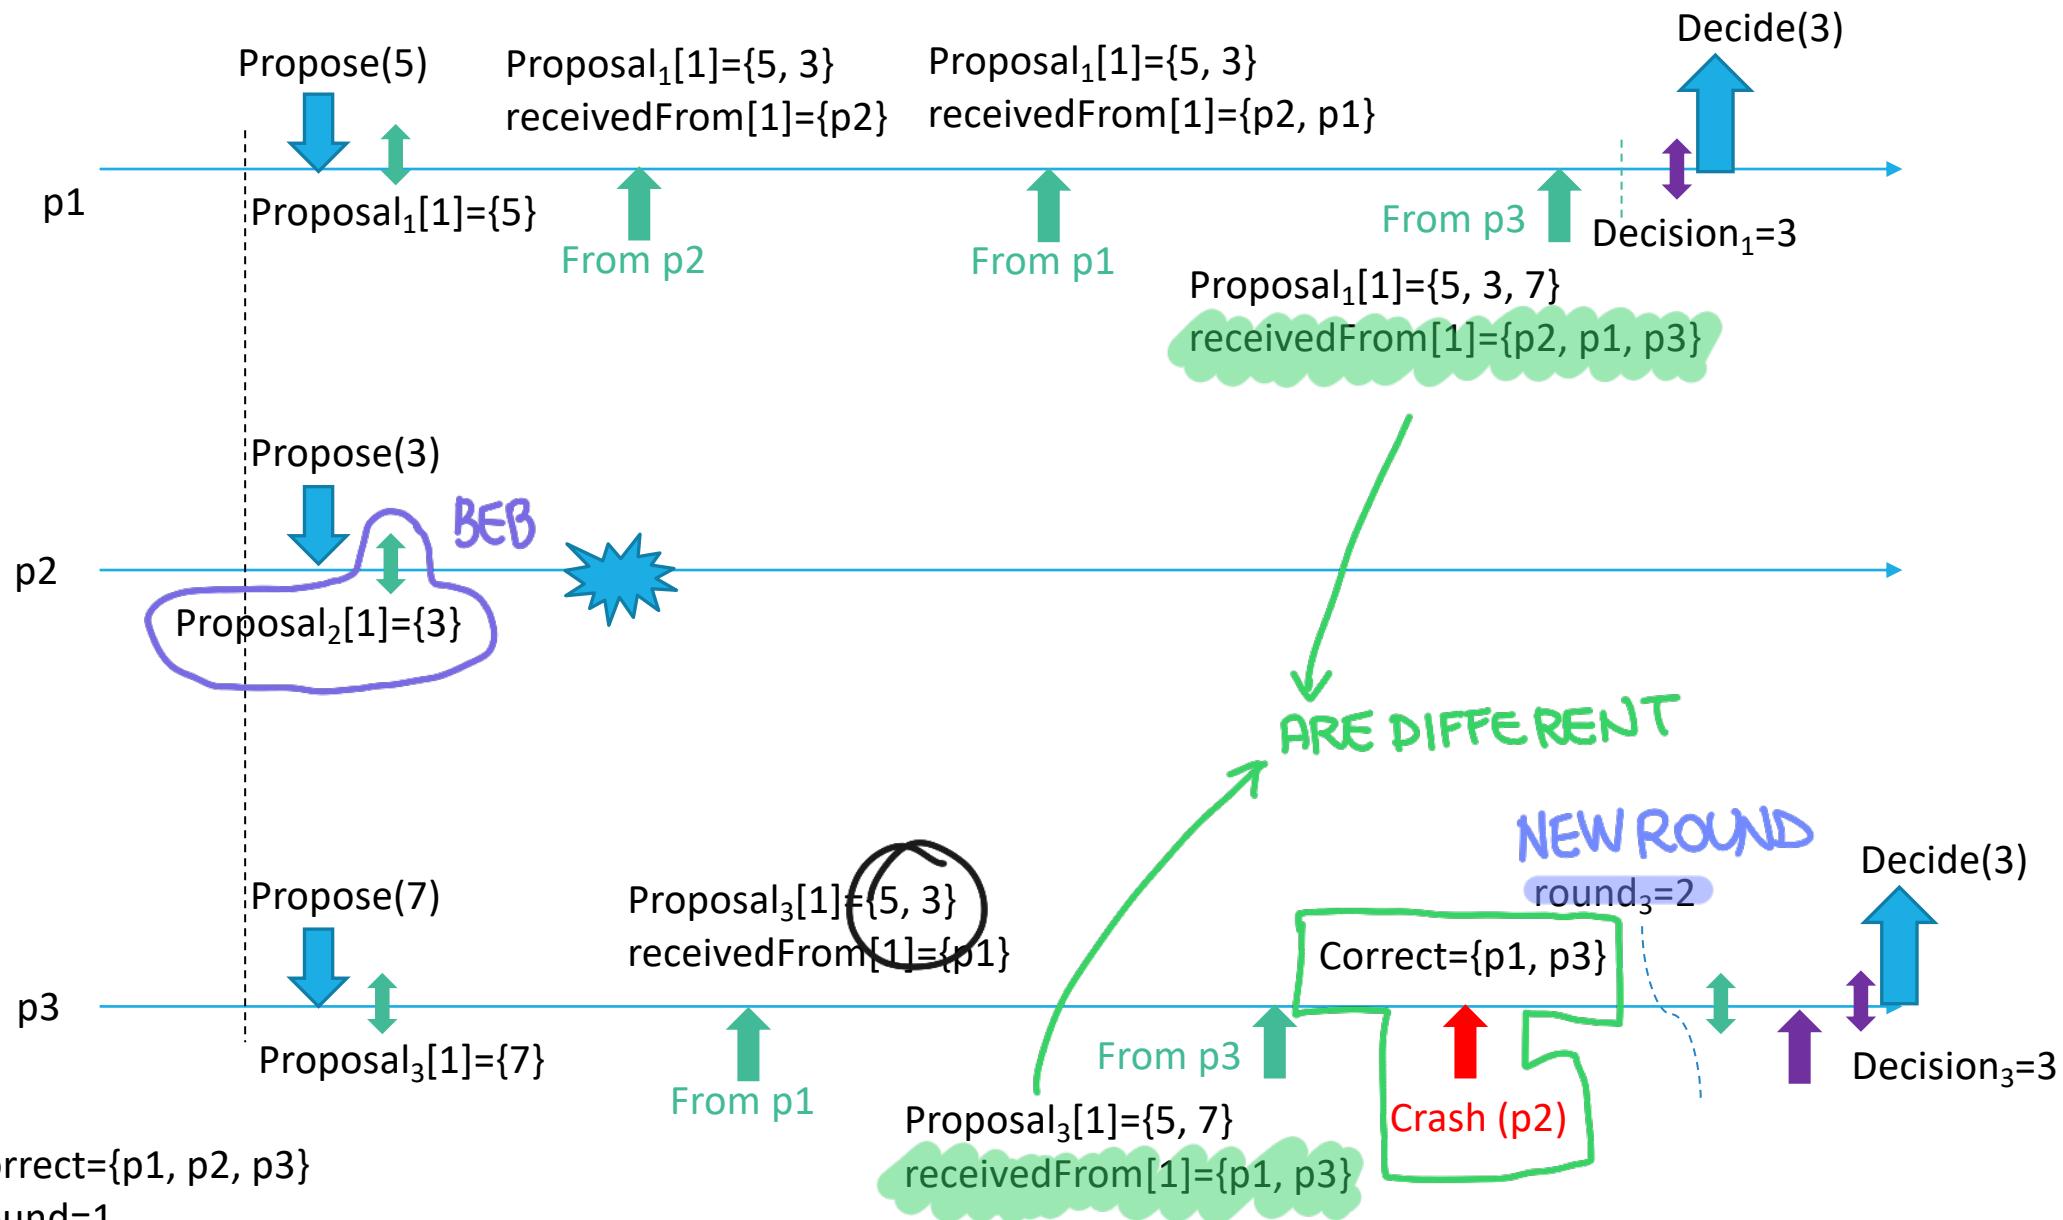
Correct= $\{p1, p2, p3\}$

Round=1

Decision= $\perp$

$\text{receivedFrom}[0]=\{p1, p2, p3\}$

$\text{Proposals}[0]=\{\}$



Correct={p1, p2, p3}

Round=1

Decision= $\perp$

receivedFrom[0]={p1, p2, p3}

Proposals[0]={}

# Correctness and Performance

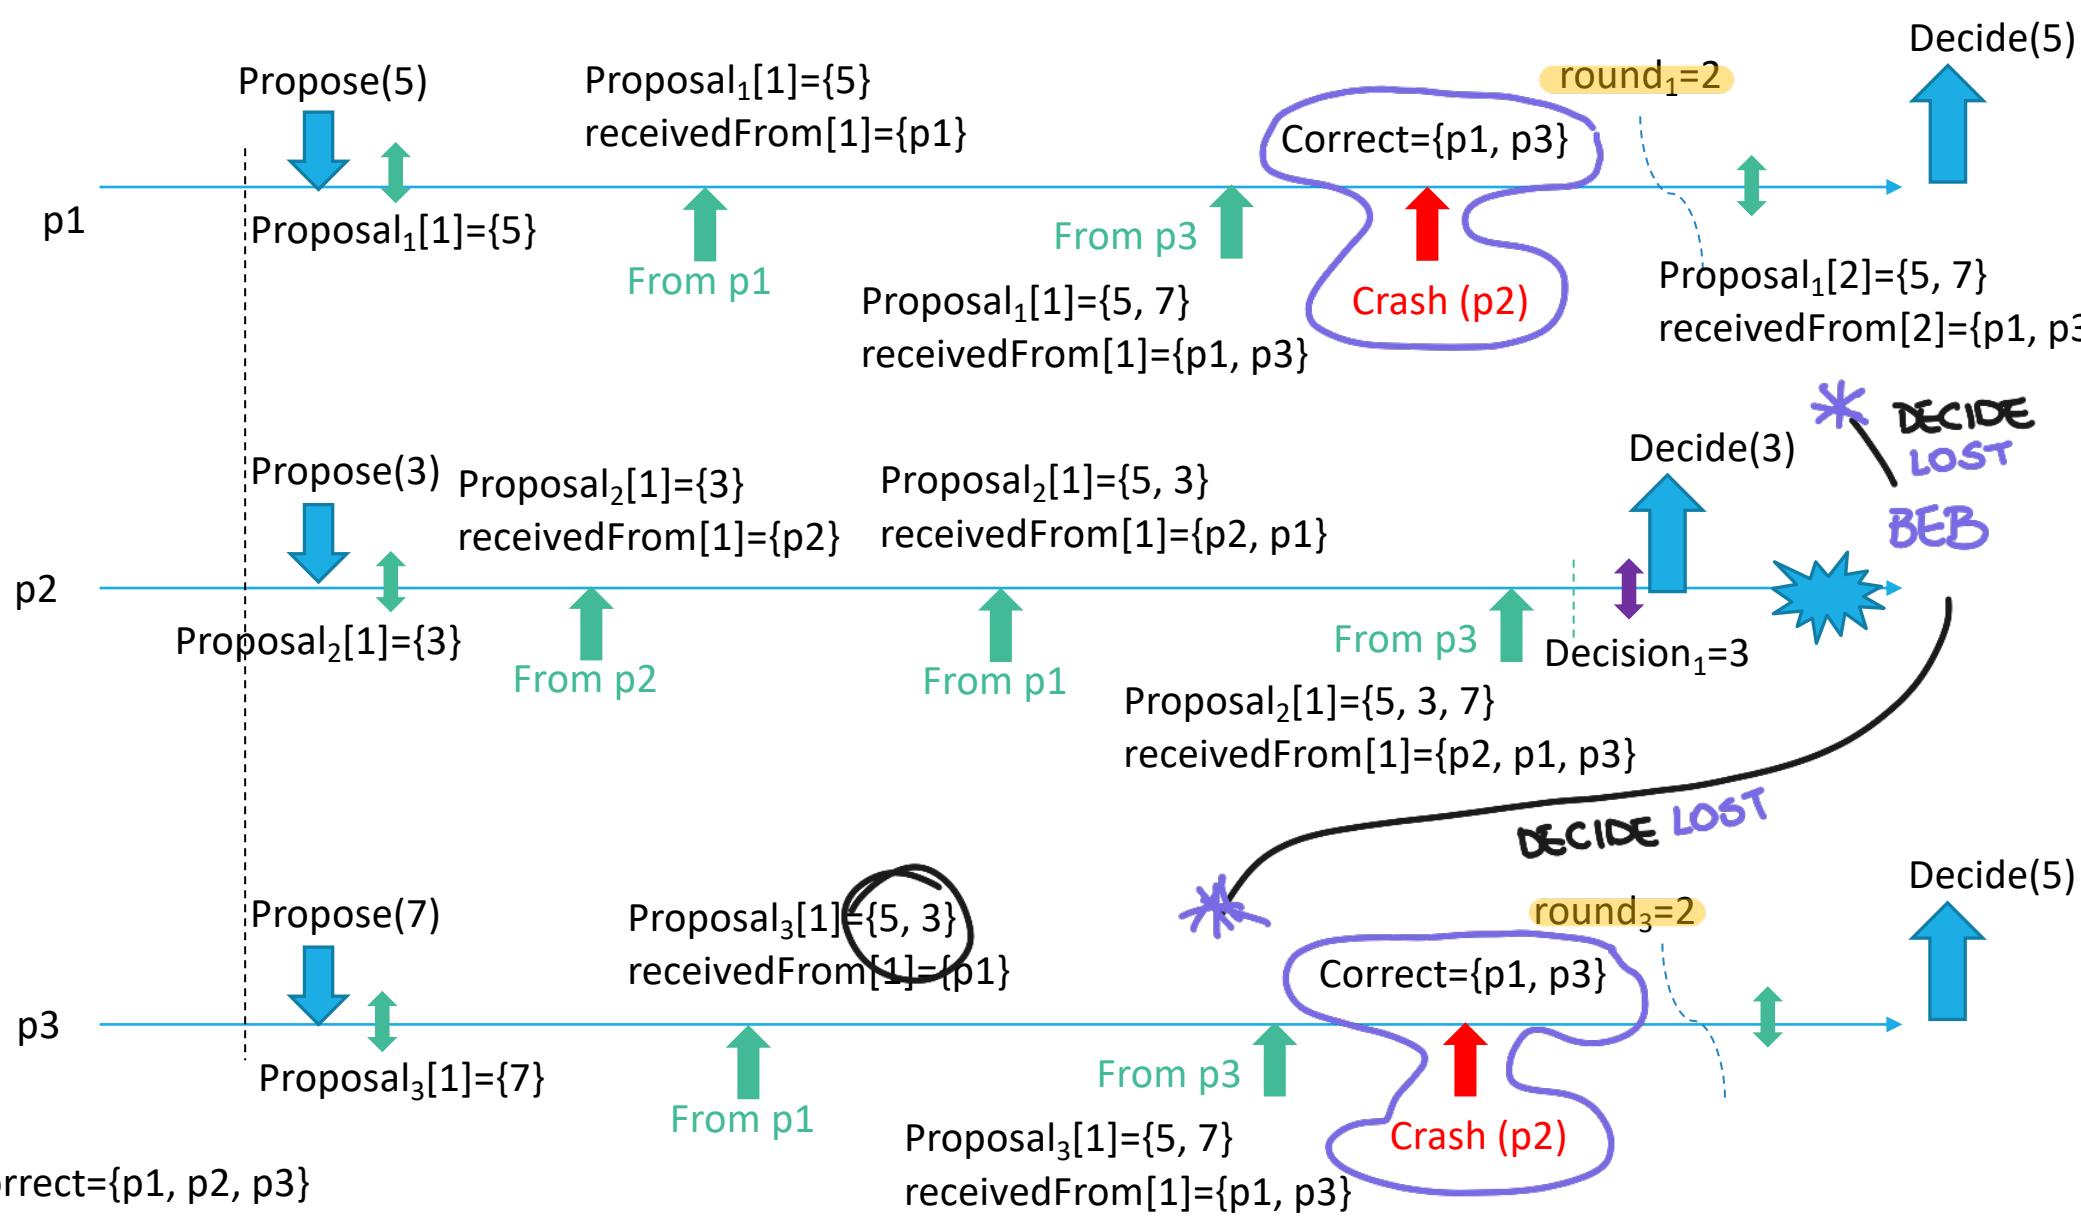
---

## Correctness

- Validity and integrity follow from the properties on the communication channels
- Termination. At most after  $N$  rounds processes decide
- Agreement. The same deterministic function is applied to the same values by correct processes

## Performance

- Best Case (No failures). One communication round ( $2N^2$  messages)
- Worst case ( $n-1$  failures).  $N^2$  messages exchanged for each communication step and at most  $N$  rounds =>  $N^3$  messages



Correct={p1, p2, p3}  
 Round=1  
 Decision= $\perp$   
 receivedFrom[0]={p1, p2, p3}  
 Proposals[0]={}

If you decide and you fail  $\rightarrow$  your decision influence the result

# Uniform Consensus Specification

---

## Module 5.2: Interface and properties of uniform consensus

---

### Module:

Name: UniformConsensus, instance  $uc$ .

### Events:

**Request:**  $\langle uc, \text{Propose} \mid v \rangle$ : Proposes value  $v$  for consensus.

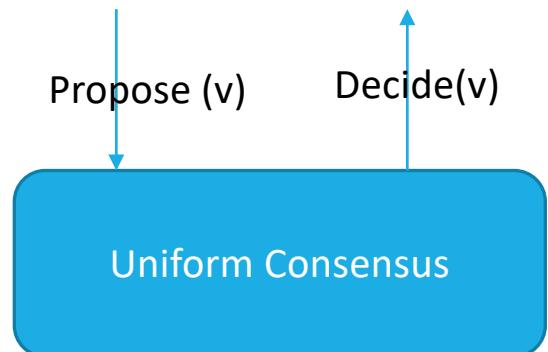
**Indication:**  $\langle uc, \text{Decide} \mid v \rangle$ : Outputs a decided value  $v$  of consensus.

### Properties:

**UC1–UC3:** Same as properties C1–C3 in (regular) consensus (Module 5.1).

**UC4:** *Uniform agreement:* No two processes decide differently.

---



# Uniform Consensus and Flooding Consensus algorithm

---

Does the flooding consensus algorithm (described in the previous slides) satisfy the Uniform Consensus specification?

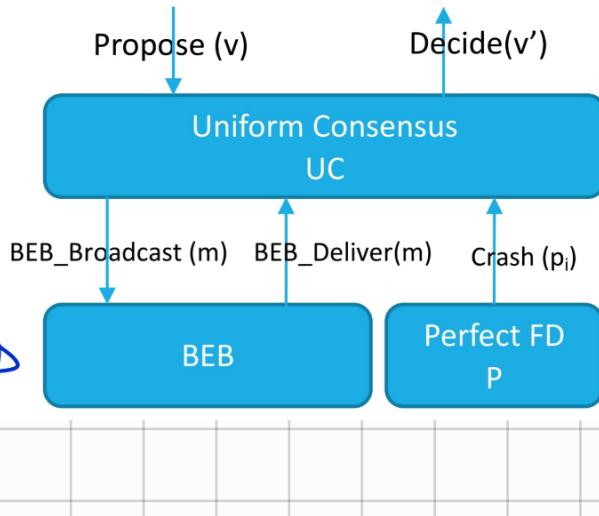
# 10: UNIFORM CONSENSUS

**C1: Termination:** Every correct process eventually decides some value.

**C2: Validity:** If a process decides  $v$ , then  $v$  was proposed by some process.

**C3: Integrity:** No process decides twice.

**UC4: Uniform agreement:** No two processes decide differently.



**Algorithm 5.3: Flooding Uniform Consensus**

**Implements:**

UniformConsensus, **instance uc**.

**Uses:**

BestEffortBroadcast, **instance beb**;  
PerfectFailureDetector, **instance P**.

**upon event**  $\langle uc, \text{Init} \rangle$  **do**

```

correct :=  $\Pi$ ;
round := 1;
decision :=  $\perp$ ;
proposalset :=  $\emptyset$ ; do not depend on the round
receivedfrom :=  $\emptyset$ ;

```

No more related to the round

I TAKE THE DECISION ONLY  
WHEN I HAVE ALL THE  
DECISIONS, AFTER m ROUND

**upon event**  $\langle P, \text{Crash} | p \rangle$  **do**

```
correct := correct \ {p};
```

**upon event**  $\langle uc, \text{Propose} | v \rangle$  **do**

```

proposalset := proposalset  $\cup$  {v};
trigger  $\langle beb, \text{Broadcast} | [\text{PROPOSAL}, 1, proposalset] \rangle$ ;

```

**upon event**  $\langle beb, \text{Deliver} | p, [PROPOSAL, r, ps] \rangle$  **such that**  $r = \text{round}$  **do**

```

receivedfrom := receivedfrom  $\cup$  {p};
proposalset := proposalset  $\cup$  ps;

```

**upon**  $\text{correct} \subseteq \text{receivedfrom} \wedge \text{decision} = \perp$  **do** = TRUE

*if round = N then*  
decision := min(proposalset);  
trigger  $\langle uc, \text{Decide} | \text{decision} \rangle$ ;

**else**

```

round := round + 1;
receivedfrom :=  $\emptyset$ ;
trigger  $\langle beb, \text{Broadcast} | [\text{PROPOSAL}, round, proposalset] \rangle$ ;

```

Decision only at the end

Cleaned at the beginning  
of each round

# Uniform Consensus Implementation in Synchronous Systems

## Algorithm 5.3: Flooding Uniform Consensus

**Implements:**

UniformConsensus, **instance** *uc*.

**Uses:**

BestEffortBroadcast, **instance** *beb*;  
PerfectFailureDetector, **instance**  $\mathcal{P}$ .

**upon event**  $\langle uc, Init \rangle$  **do**

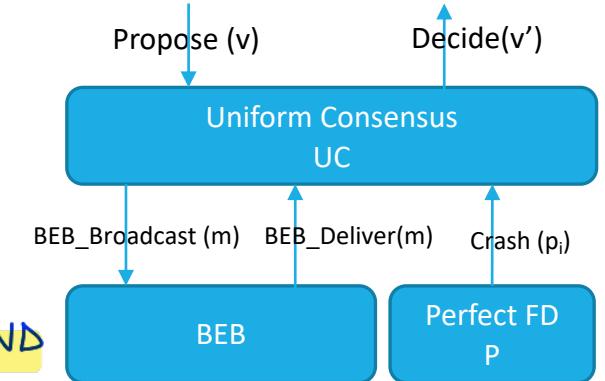
```

correct :=  $\Pi$ ;
round := 1;
decision :=  $\perp$ ;
proposalset :=  $\emptyset$ ; do not depend on the round
receivedfrom :=  $\emptyset$ ;
```

**upon event**  $\langle \mathcal{P}, Crash | p \rangle$  **do**  
 $correct := correct \setminus \{p\}$ ;

**upon event**  $\langle uc, Propose | v \rangle$  **do**  
 $proposalset := proposalset \cup \{v\}$ ;  
**trigger**  $\langle beb, Broadcast | [PROPOSAL, 1, proposalset] \rangle$ ;

I TAKE THE DECISION ONLY  
WHEN I HAVE ALL THE  
DECISIONS, AFTER *n* ROUND



**upon event**  $\langle beb, Deliver | p, [PROPOSAL, r, ps] \rangle$  **such that**  $r = round$  **do**

```

receivedfrom := receivedfrom  $\cup \{p\}$ ;  

proposalset := proposalset  $\cup ps$ ;
```

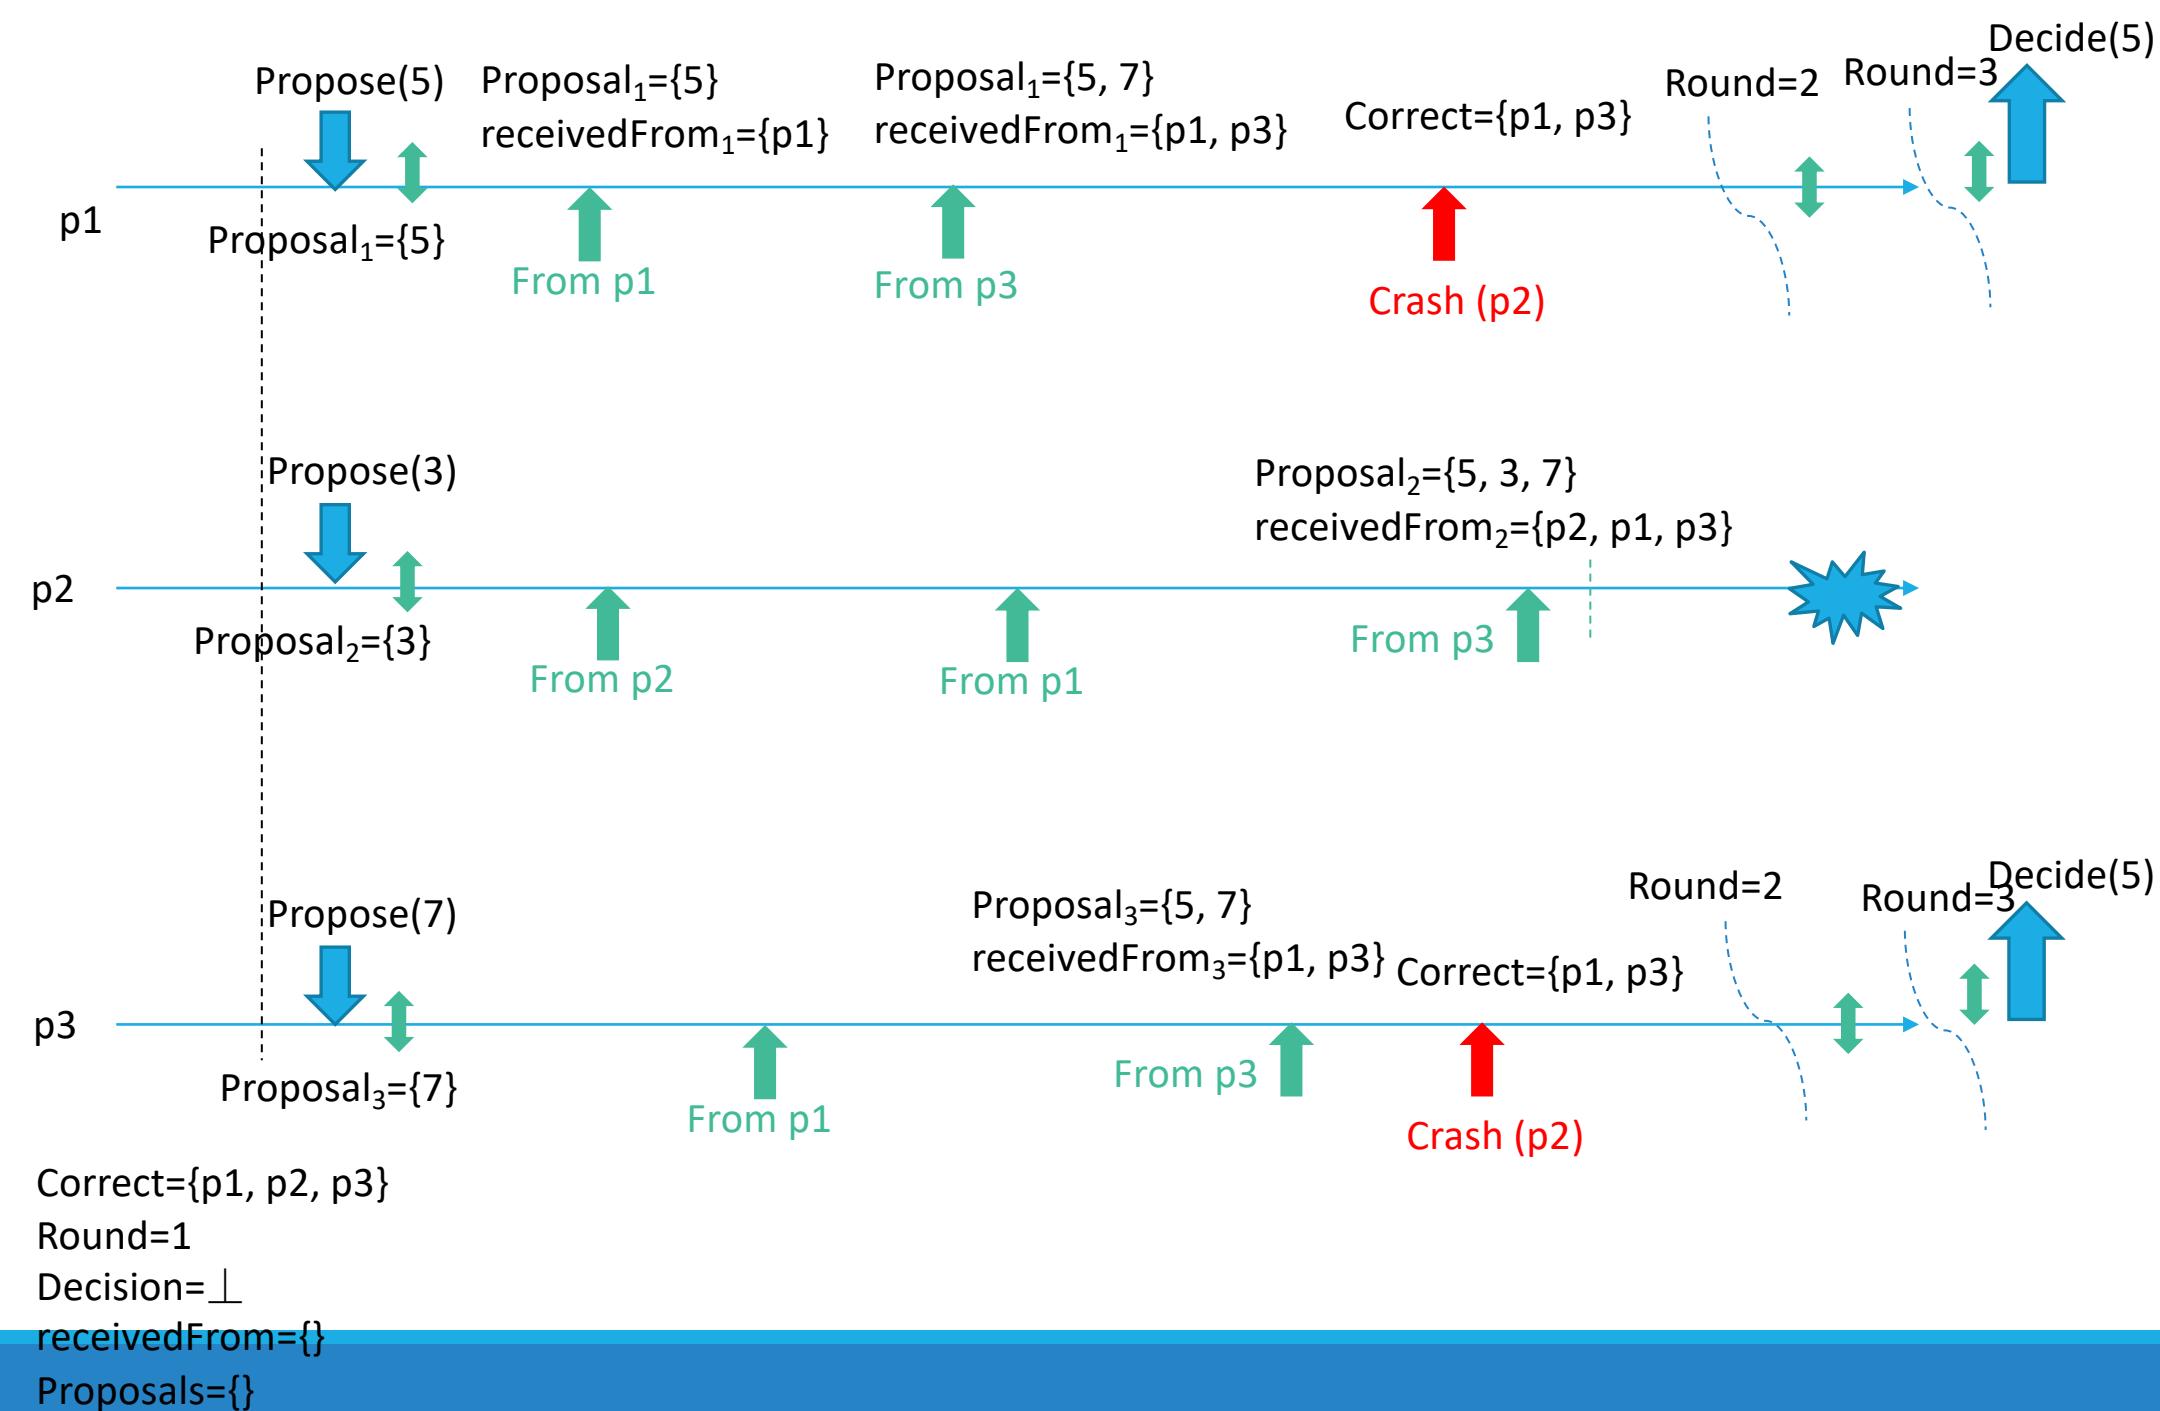
**upon**  $correct \subseteq receivedfrom \wedge decision = \perp$  **do** = TRUE

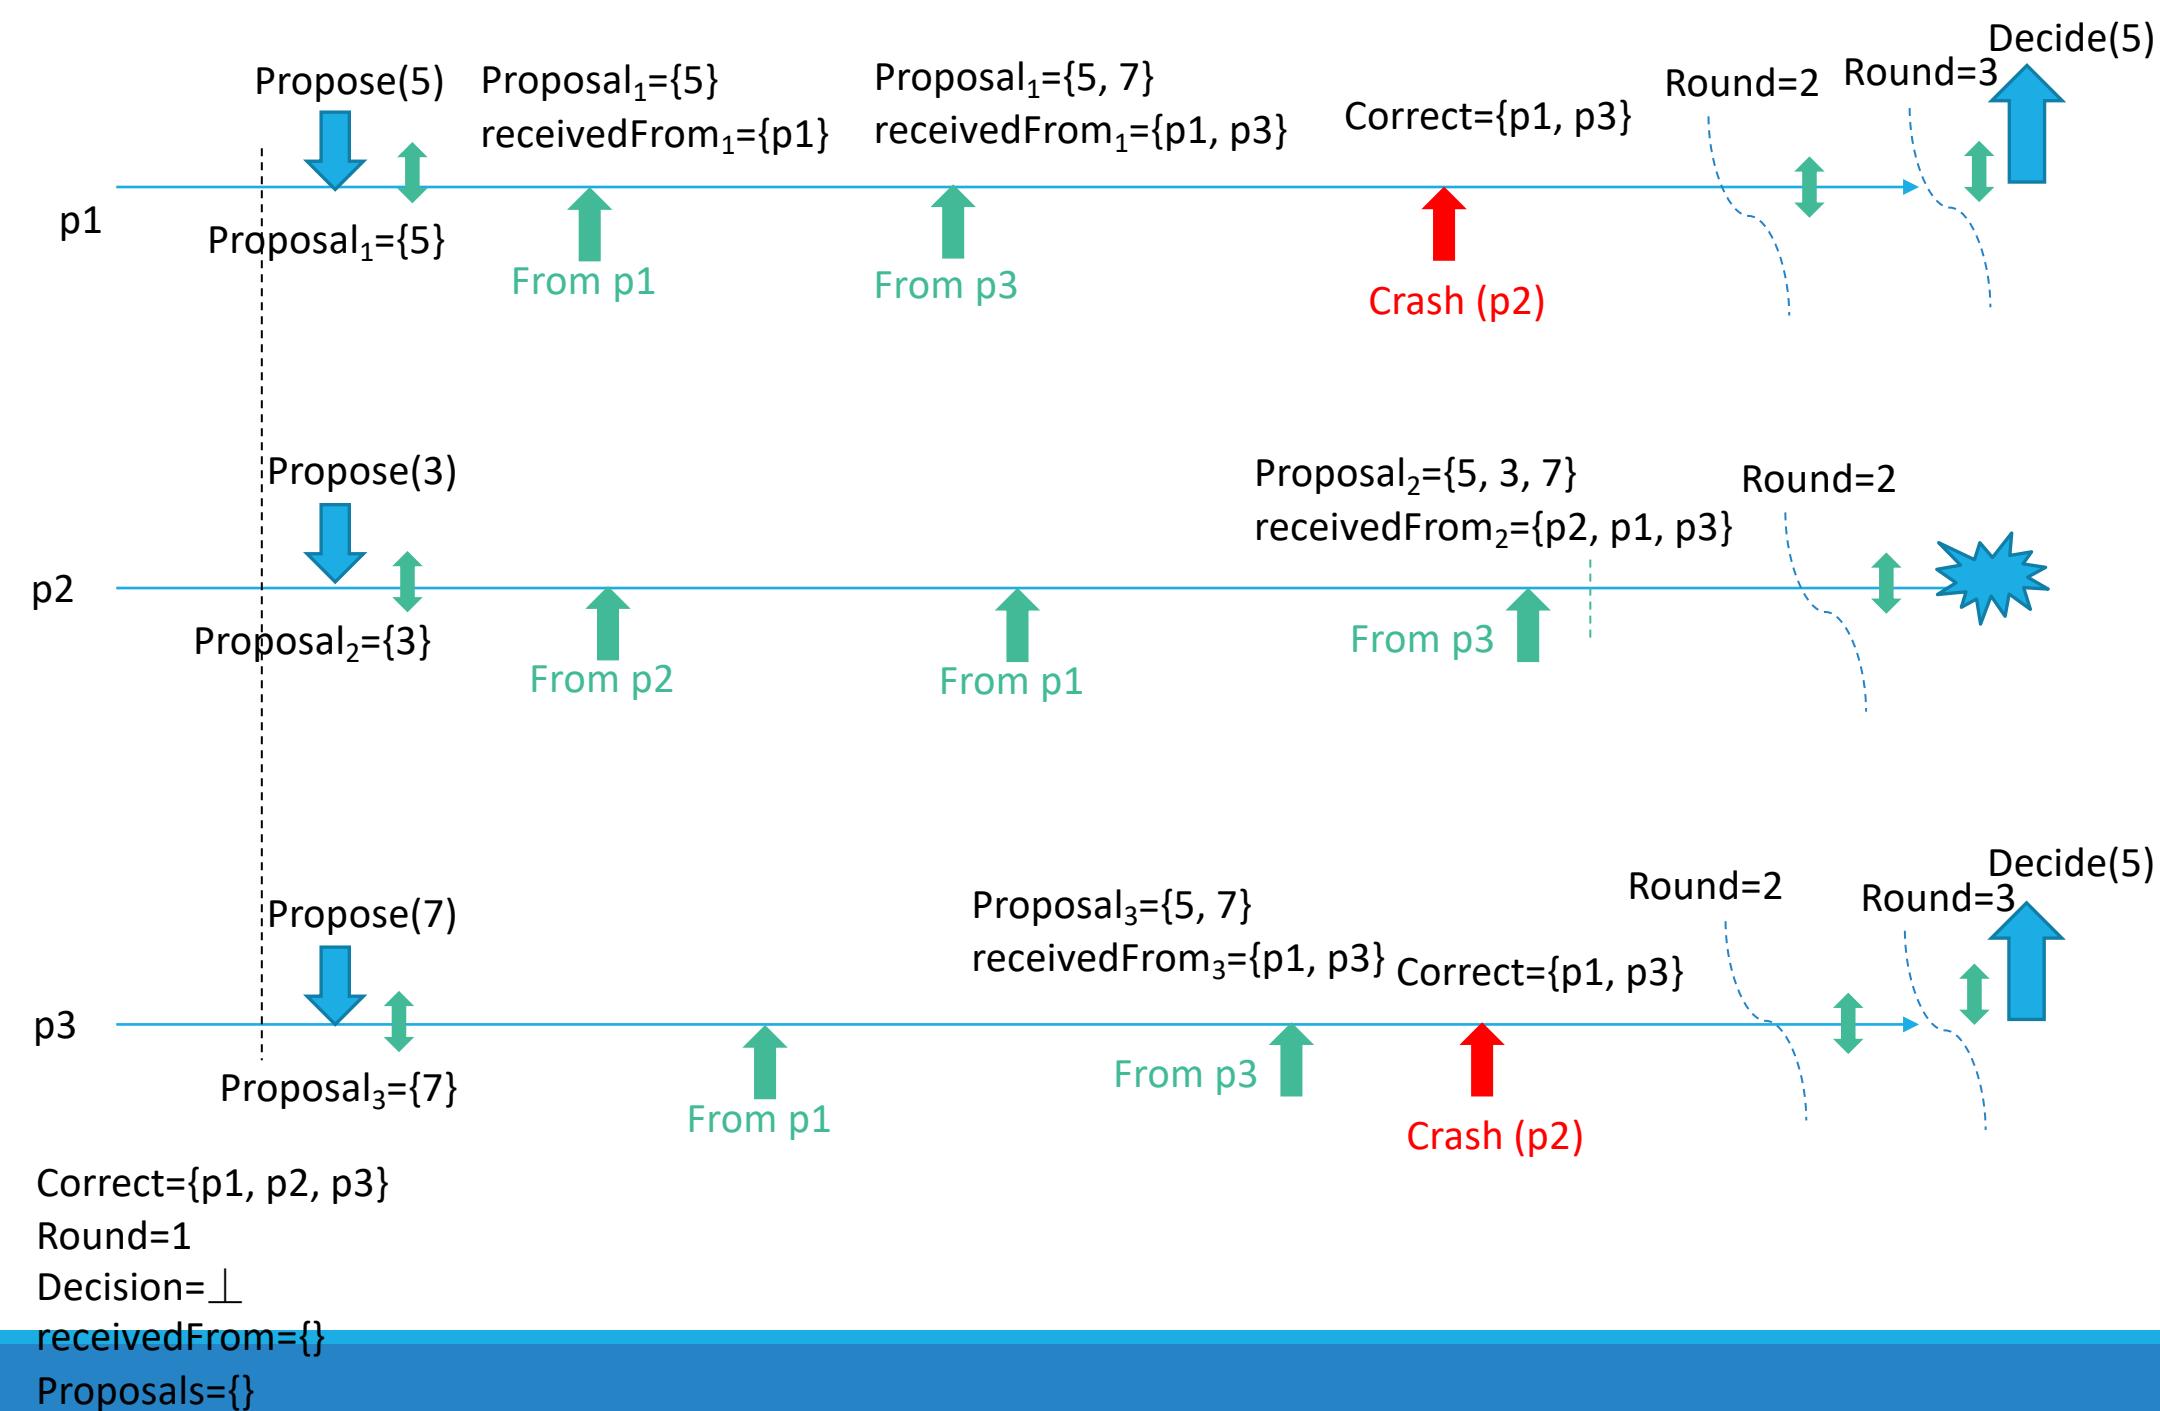
```

if round = N then
    decision := min(proposalset);
    trigger  $\langle uc, Decide | decision \rangle$ ;
else
    round := round + 1;
    receivedfrom :=  $\emptyset$ ;
    trigger  $\langle beb, Broadcast | [PROPOSAL, round, proposalset] \rangle$ ;
```

Decision only at the end

Cleaned at the beginning  
of each round





# Correctness and Performance

---

## Correctness

- The *validity* and *integrity* properties follow from the algorithm and from the properties of best-effort broadcast
- The *termination* property is ensured as all correct processes reach round N and decide in that round
  - the *strong completeness* property of the failure detector implies that no correct process waits indefinitely for a message from a process that has crashed, as the crashed process is eventually removed from *correct*.
- The *uniform agreement* holds because all processes that reach round N have the same set of values in their variable *proposalset*.

## Performance

- N communication steps and O(N<sup>3</sup>) messages for all correct processes to decide

# References

---

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 5, Sections 5.1.1, 5.1.2, 5.2.1, 5.2.2