

31|10|23

Dependable Distributed Systems
Master of Science in Engineering in
Computer Science

AA 2023/2024

LECTURE 13: TOTAL ORDER BROADCAST

System model

Static set of processes $\Pi = \{p_1 \dots p_n\}$

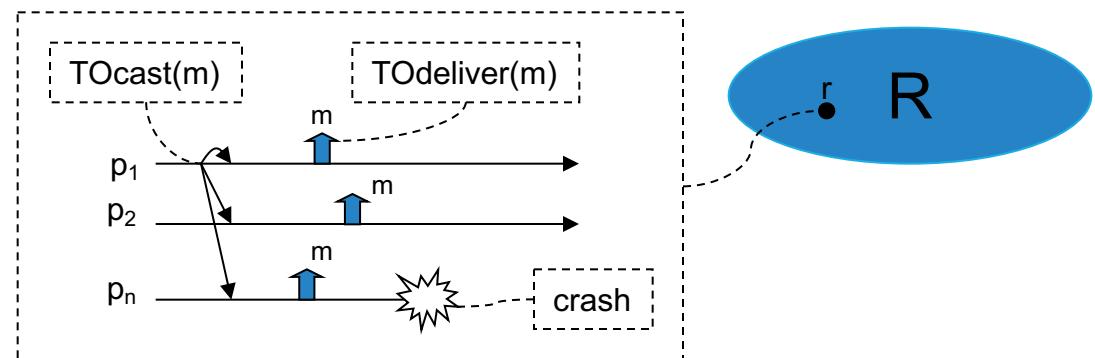
Message passing over perfect channels

- i.e., message exchanged between correct processes are reliably delivered

Asynchronous

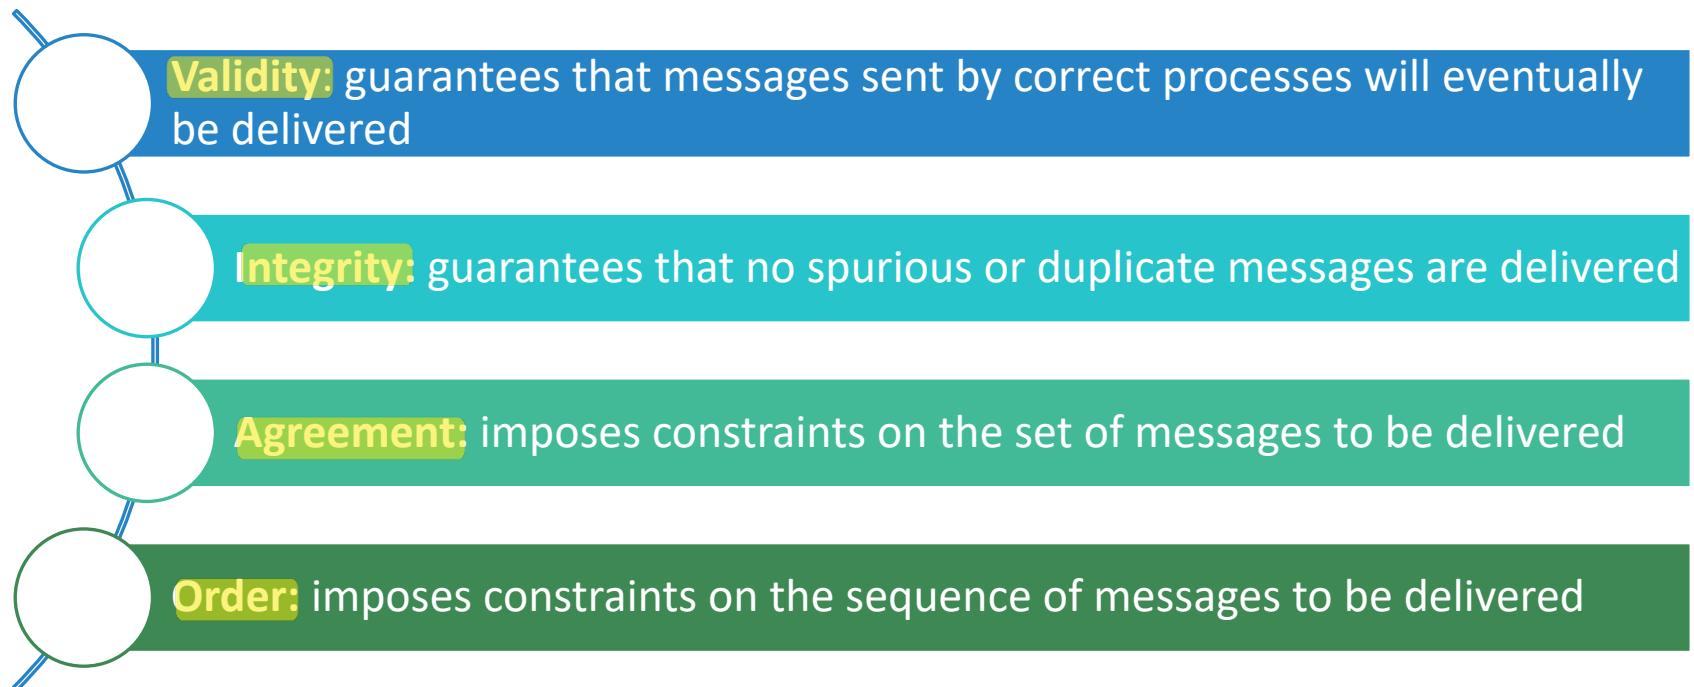
Crash fault model for processes

We characterize the system in terms of its possible runs R



TO specifications

Total order specifications are usually composed by four properties



improved version of RB



TO specifications

Total Order Broadcast = $\text{TO}(V,I,A,O)$

- V = Validity
- I = Integrity
- A = Agreement
- O = Order

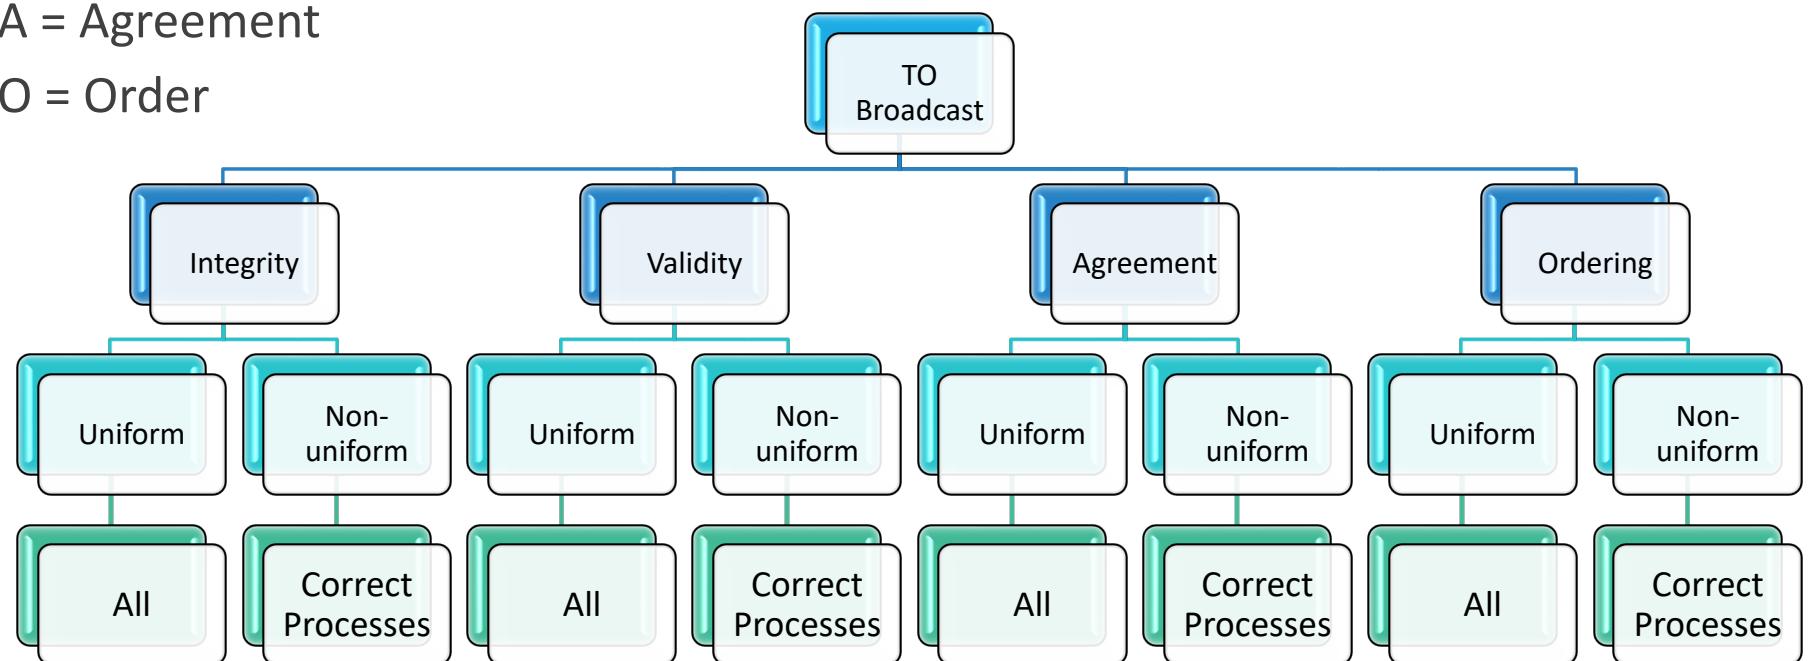
Distinct specifications arise from distinct formulations of each property

- uniform vs non-uniform
- A uniform property imposes restrictions on the behavior of (at least) correct processes

TO specifications

Total Order Broadcast = TO(V,I,A,O)

- V = Validity
- I = Integrity
- A = Agreement
- O = Order



TO Specifications

only crash, not recovery

Crash failure + Perfect channels \Rightarrow PP2PL

m can arrive or not (if fail)

NO UNIFORM VALIDITY

- **NUV:** if a correct process TOCAST a message m then some *correct* process will eventually deliver m

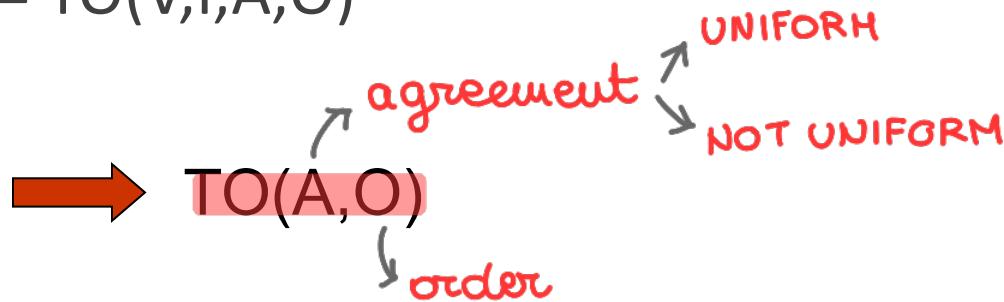
UNIFORM INTEGRITY

- **UI:** For any message m, every process p delivers m at most once and only if m was previously TOCAST by some (correct or not) process

TO specifications

Total Order Broadcast = $\text{TO}(V, I, A, O)$

- V = ~~Value~~
NUV
- I = ~~Integrity~~
UI
- ~~A~~ = Agreement
- ~~O~~ = Order



Distinct specifications arise from distinct formulations of each property

- uniform vs non-uniform
- A uniform property imposes restrictions on the behavior of (at least) correct processes

The Agreement property

UNIFORM AGREEMENT (UA)

If a process (correct or not) Todelivers a message m, then all correct processes will eventually Todeliver m

→ THE SET OF M DELIVERED
BY THE FAULTY IS A SUBSET
OF THE M DELIVERED BY THE CORRECT

CONSTRAINS THE SET OF DELIVERED MESSAGES

Correct processes always deliver the same set of messages M

Each faulty process p delivers a set M_p

$$\text{UA: } M_p \subseteq M$$

$$\text{NUA: } M_p \text{ can be s.t. } M_p - M \neq \emptyset$$

NON-UNIFORM AGREEMENT (NUA)

If a correct process Todelivers a message m, then all correct processes will eventually Todeliver m

- Agreement: UA

THE SET OF M DELIVERED BY THE FAULTY IS A SUBSET
OF THE M DELIVERED BY THE CORRECT

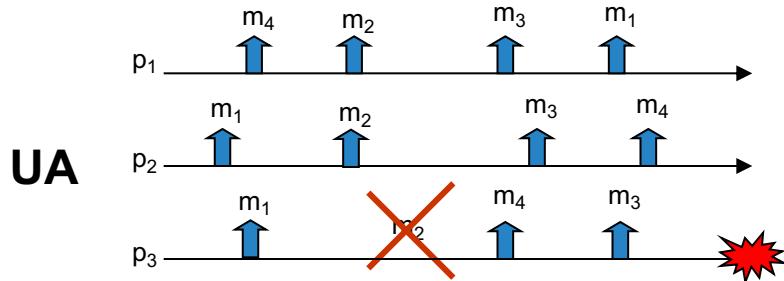
- Agreement: NUA

M_p can be s.t. $M_p - M \neq 0$

The Agreement property

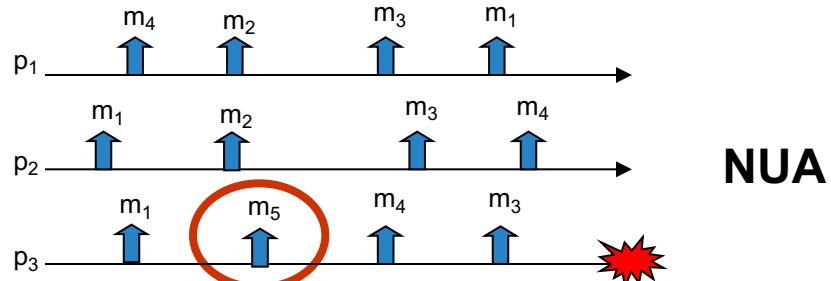
UNIFORM AGREEMENT (UA)

If a process (correct or not) Todelivers a message m , then all correct processes will eventually Todeliver m



NON-UNIFORM AGREEMENT (NUA)

If a correct process Todelivers a message m , then all correct processes will eventually Todeliver m



The Ordering Property

STRONG UNIFORM TOTAL ORDER
(SUTO)

If some process TODelivers some message m before message m' , then a process $\text{TODelivers} m'$ only after it has $\text{TODelivered} m$.



- same order
- same prefix of the set of delivered messages
- after an omission, disjoint sets of delivered messages

WEAK UNIFORM TOTAL ORDER
(WUTO)

If process p and process q both TODdeliver messages m and m' , then $p \text{ TODelivers } m \text{ before } m'$ if and only if $q \text{ TODelivers } m \text{ before } m'$.



- no restrictions on the set of delivered messages

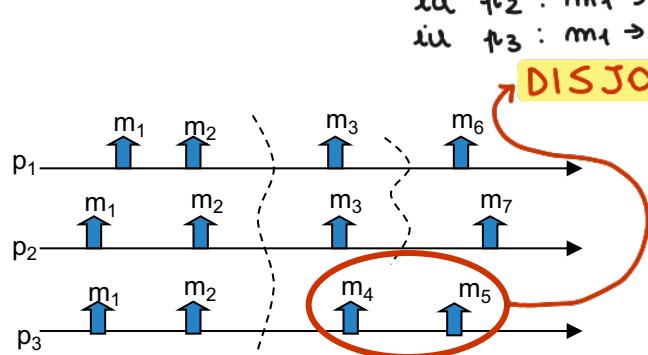
The Order Property

STRONG UNIFORM TOTAL ORDER (SUTO)

If some process TODelivers some message m before message m' , then a process $\text{TODelivers} m'$ only after it has $\text{TODelivered} m$.

↓
no agreement

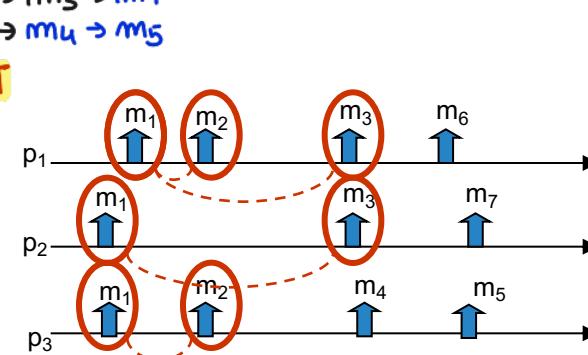
SUTO



WEAK UNIFORM TOTAL ORDER (WUTO)

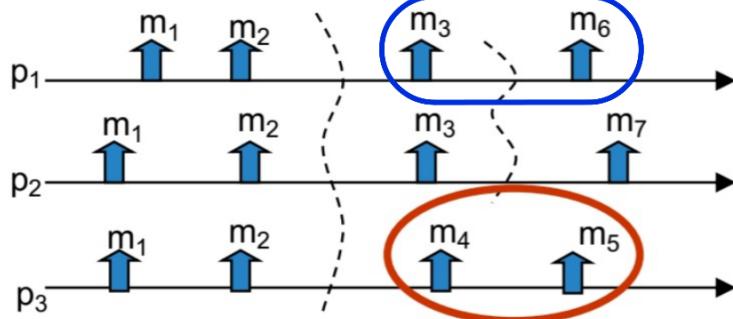
If process p and process q both TODeliver messages m and m' , then $p \text{ TODelivers } m \text{ before } m'$ if and only if $q \text{ TODelivers } m \text{ before } m'$.

WUTO



$$\begin{aligned} \text{i.u. } & p_1 : m_1 \rightarrow m_2 \rightarrow m_3 \rightarrow m_6 \\ \text{i.u. } & p_2 : m_1 \rightarrow m_2 \rightarrow m_3 \rightarrow m_7 \\ \text{i.u. } & p_3 : m_1 \rightarrow m_2 \rightarrow m_4 \rightarrow m_5 \end{aligned}$$

SUTO



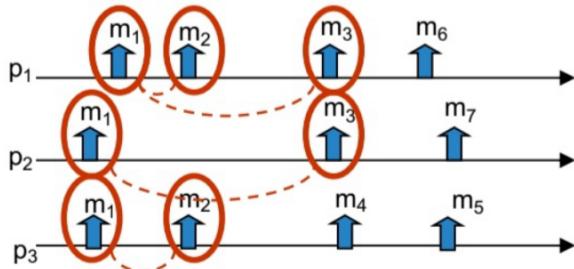
N.B.

p_1 : we have $m_1 \rightarrow m_2$

every process have to respect this

p_1 : we have m_3 ; the order imply that m_6 could be delivered only after the delivery of m_3

the other correct processes either deliver m_3 and then m_6 , or doesn't deliver m_6 because $m_3 \rightarrow m_6$



WUTO

$m_1 \rightarrow m_2 \rightarrow m_3 \rightarrow m_6$

$m_1 \rightarrow m_3$

$m_1 \rightarrow m_2$

Take two processes that deliver the same pair of m , I just want that the sequence is the same

$m_1 \rightarrow m_2$

but I can deliver $m_1 \rightarrow m_3$ also if I have not deliver m_2

The Order Property

SUTO and WUTO are uniform but they both have a non-uniform counterpart

STRONG **NON-UNIFORM TOTAL ORDER** (SNUTO)

If some correct process TODelivers some message m before message m', then a correct process TODelivers m' **only after** it has TODelivered m.



the FAULTY can do everything he wants

↓
no ordering respected

WEAK **NON-UNIFORM TOTAL ORDER** (WNUTO)

If correct processes p and q both TODeliver messages m and m', then p TODelivers m before m' if and only if q TODelivers m before m'.



the FAULTY can do everything he wants

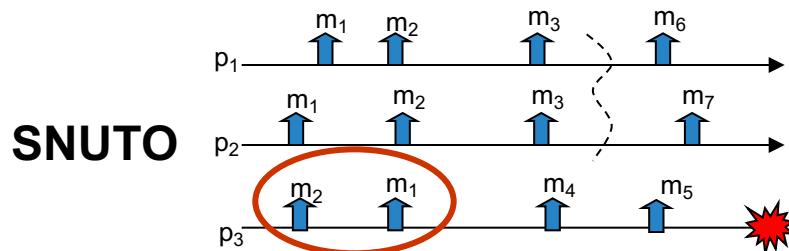


CORRECT p. may have the same prefix

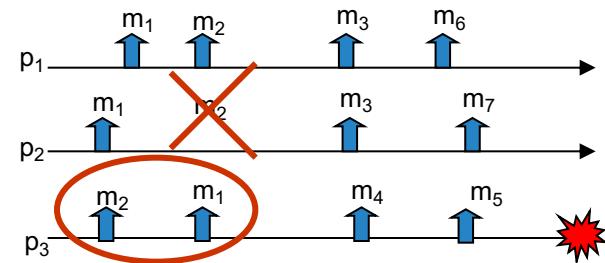
The Order property (2)

SUTO \Rightarrow WUTO

SNUTO \Rightarrow WNUTO



WNUTO

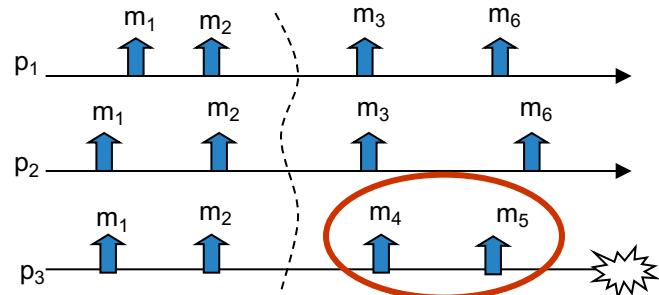


TO specifications

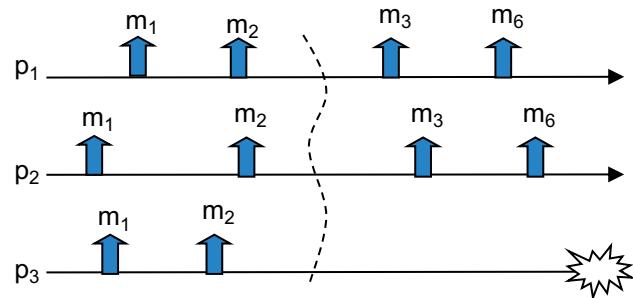
UA: forces the **faulty** p. to deliver a subset of the correct

SUTO: impose that the **correct** p. deliver the same sequence

TO(NUA,SUTO)



TO(UA,SUTO)
(Strongest total order)



TO(UA,SUTO)

- The strongest TO spec.

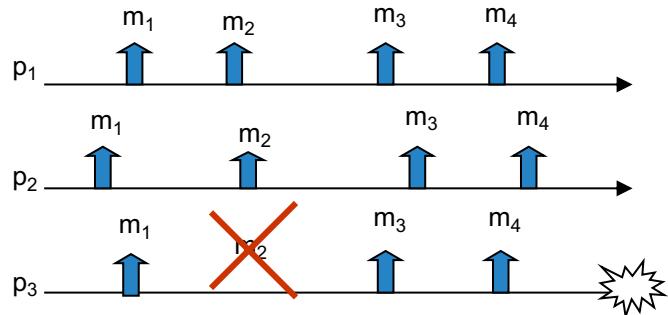
NUA: **faulty** are delivering something that the **correct** one not delivering

TO(NUA,SUTO)

SUTO: impose that the **correct** p. deliver the same sequence, same set

TO specifications (2)

TO(UA,WUTO)

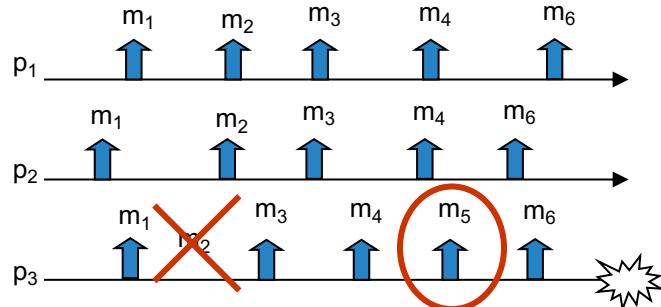


TO(UA,SUTO)
(Strongest total order)

TO(UA,WUTO)

TO(NUA,SUTO)

TO(NUA,WUTO)



TO(NUA,WUTO)

TO(UA, WUTO)

after the omission the set cannot
↓
diverge

UA: forces the faulty n. to deliver a subset of the correct

↓
I can tolerate just 1 MISSING EVENT

WUTO: correct deliver the same set in the same sequence

TO(NUA, WUTO)

(1)

NUA: faulty n. deliver same prefix, than an omission,
than I have messages deliver by the others

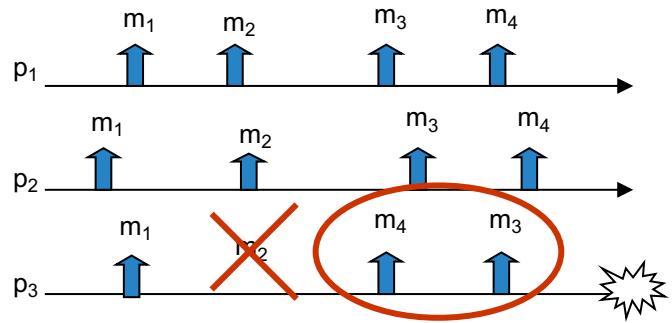
↓

no subset of the correct, no strong ordering

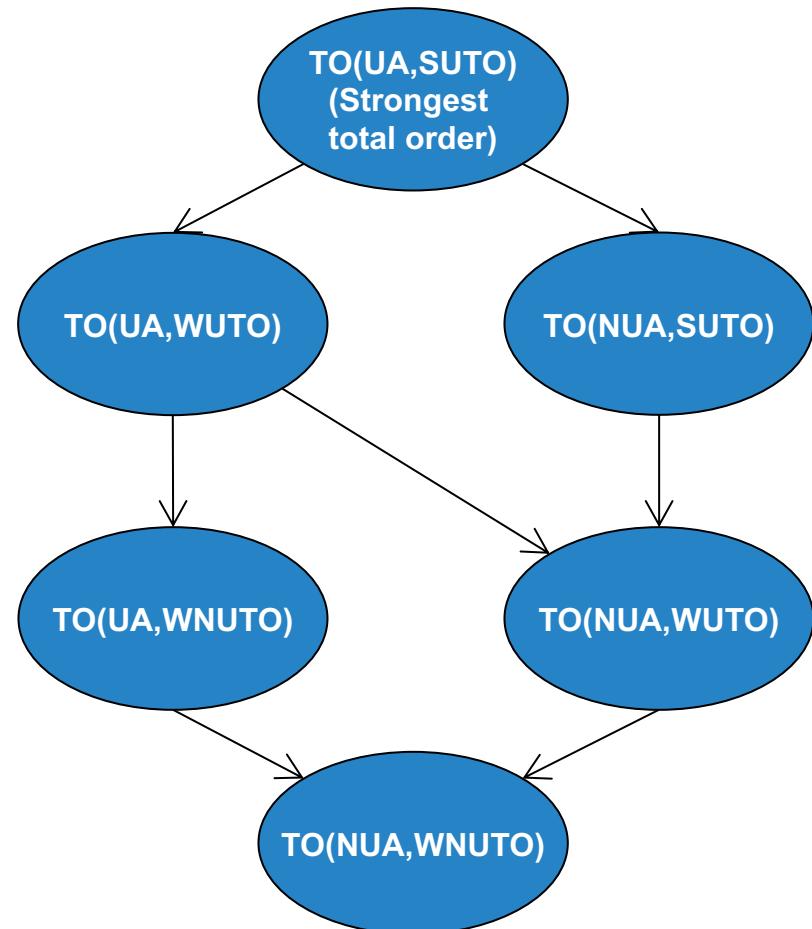
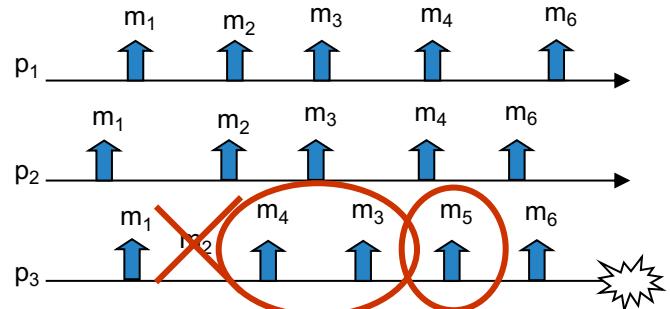
WUTO: correct deliver the same set in the same sequence

TO specifications (3)

TO(UA,WNUTO)



TO(NUA,WNUTO)



TO(UA, WNUTO)

UA: forces the faulty p. to deliver a subset of the correct

↓
I can tolerate just 1 MISSING EVENT

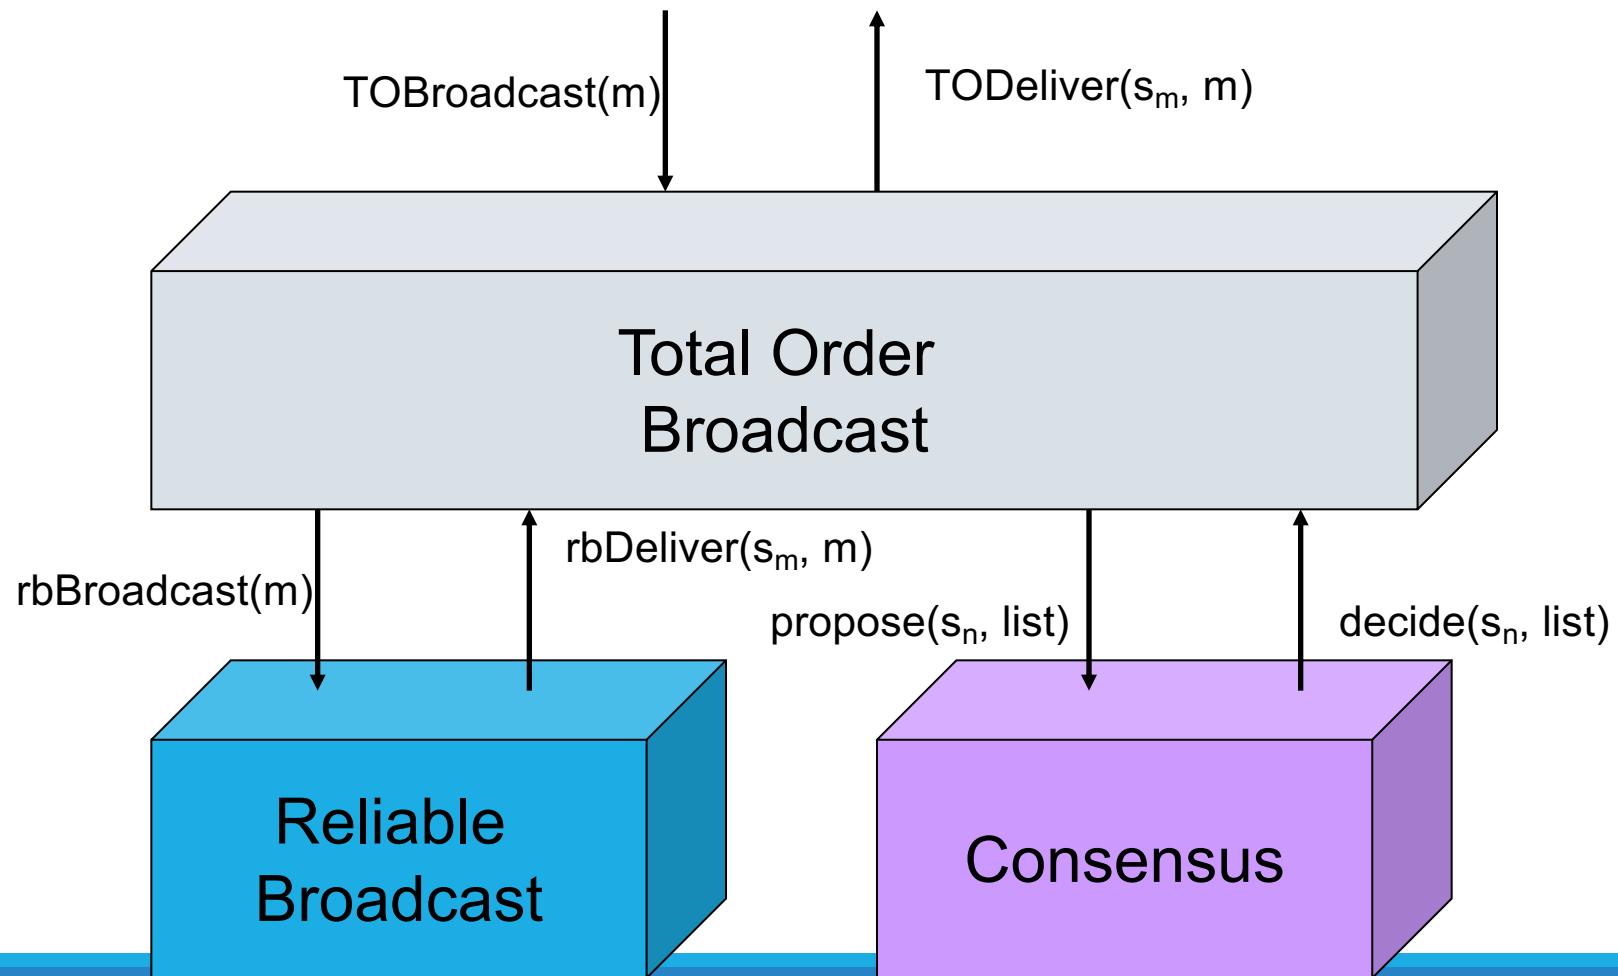
WNUTO: ordering is not strong, faulty can deliver in whenever order it wants

TO(NUA, WNUTO)

NUA: faulty p. can deliver something that the correct one not delivered, the order cannot be respected

WNUTO: ordering is not strong, faulty can deliver in whenever order it wants

Total Order Implementation



Total Order Algorithm

Algorithm 6.1: Consensus-Based Total-Order Broadcast

Implements:

TotalOrderBroadcast, **instance** *tob*.

Uses:

ReliableBroadcast, **instance** *rb*;
Consensus (multiple instances).

```
upon event < tob, Init > do
  unordered :=  $\emptyset$ ;
  delivered :=  $\emptyset$ ;
  round := 1;
  wait := FALSE;

upon event < tob, Broadcast | m > do
  trigger < rb, Broadcast | m >;

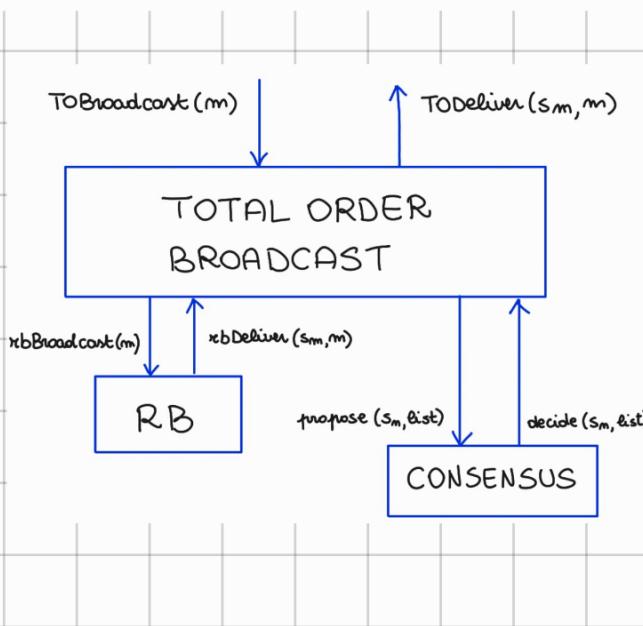
upon event < rb, Deliver | p, m > do
  if m  $\notin$  delivered then
    unordered := unordered  $\cup$  {(p, m)};

upon unordered  $\neq \emptyset$   $\wedge$  wait = FALSE do
  wait := TRUE;
  Initialize a new instance c.round of consensus;
  trigger < c.round, Propose | unordered >;

upon event < c.r, Decide | decided > such that r = round do
  forall  $(s, m) \in \text{sort}(\text{decided})$  do
    trigger < tob, Deliver | s, m >; // by the order in the resulting sorted list
  delivered := delivered  $\cup$  decided;
  unordered := unordered  $\setminus$  decided;
  round := round + 1;
  wait := FALSE;
```

13: TOTAL

- **Validity:** guarantees that messages sent by correct processes will eventually be delivered
- **Integrity:** guarantees that no spurious or duplicate messages are delivered
- **Agreement:** imposes constraints on the set of messages to be delivered
- **Order:** imposes constraints on the sequence of messages to be delivered



Algorithm 6.1: Consensus-Based Total-Order Broadcast

Implements:

TotalOrderBroadcast, instance *tob*.

Uses:

ReliableBroadcast, instance *rb*; Consensus (multiple instances).

Keep track of the m in the system that are not yet delivered; need to be ordered

you are not going anywhere

upon event $\langle \text{tob}, \text{Init} \rangle$ **do**

for repetition

unorderd := \emptyset ;

delivered := \emptyset ; *→ keep track of what is already delivered*

round := 1; *→ for counting how many consensus I am running*

wait := FALSE; *→ serialize the execution of the consensus and to avoid more than one instances of consensus*

upon event $\langle \text{tob}, \text{Broadcast} | m \rangle$ **do**

trigger $\langle \text{rb}, \text{Broadcast} | m \rangle$;

upon event $\langle \text{rb}, \text{Deliver} | p, m \rangle$ **do**

if $m \notin \text{delivered}$ then

unorderd := unorderd $\cup \{(p, m)\}$;

upon event $\langle \text{tob}, \text{Decide} | decided \rangle$ **such that** $r = \text{round}$ **do**

forall $(s, m) \in \text{sort}(\text{decided})$ **do** *is a list and need to be sorted*

trigger $\langle \text{tob}, \text{Deliver} | s, m \rangle$;

delivered := delivered $\cup \text{decided}$;

unorderd := unorderd $\setminus \text{decided}$;

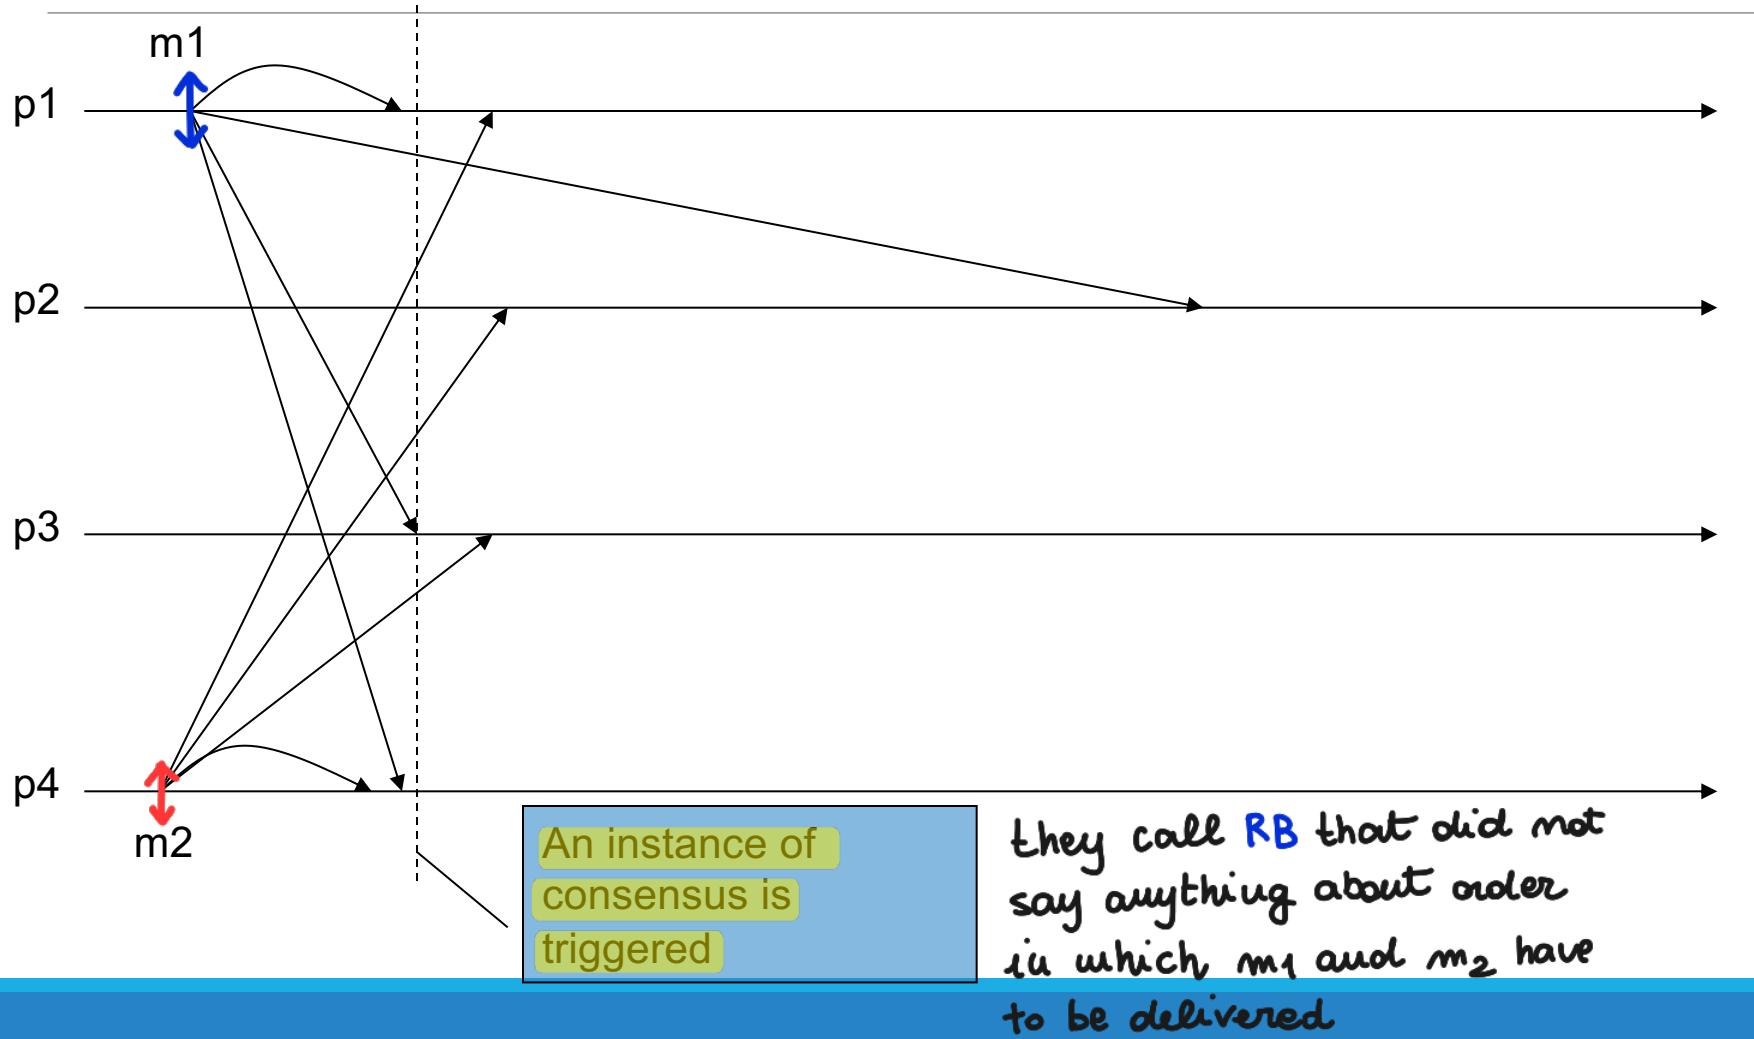
round := round + 1;

wait := FALSE;

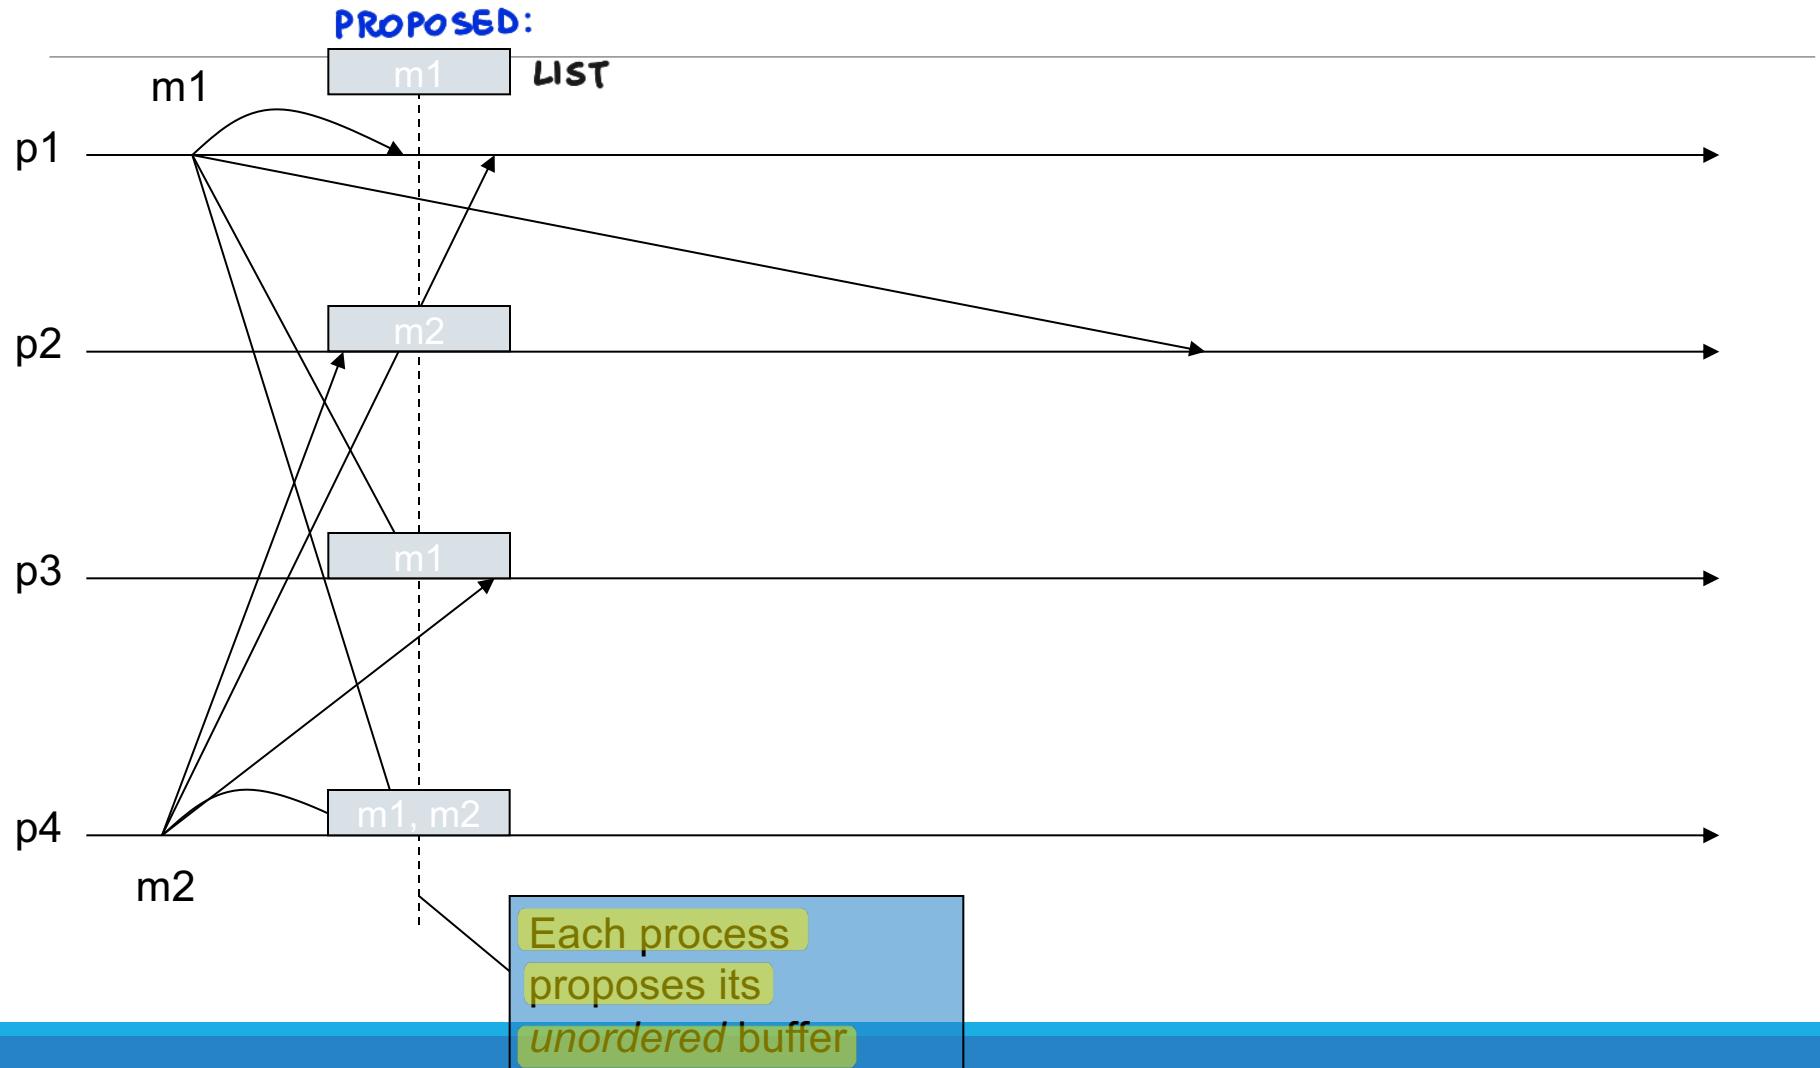
Example

4 processes

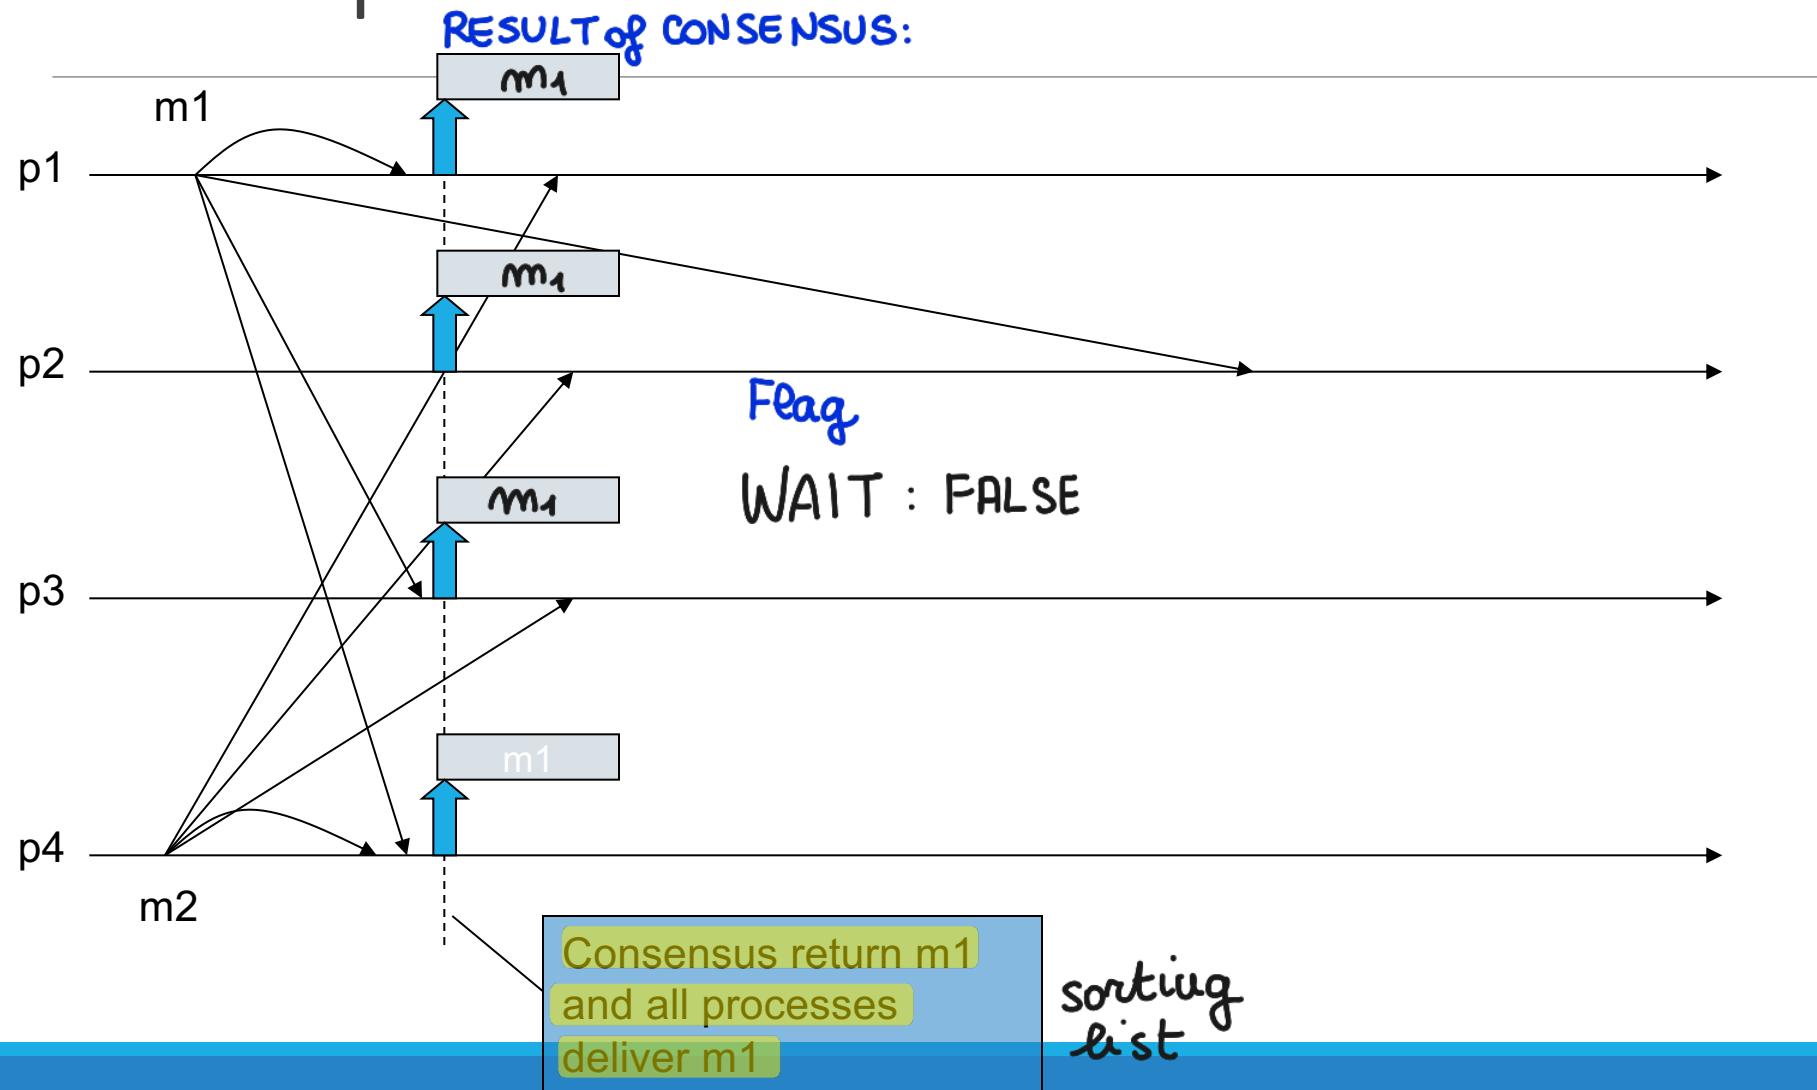
t^1
 t^2 } broadcast concurrently



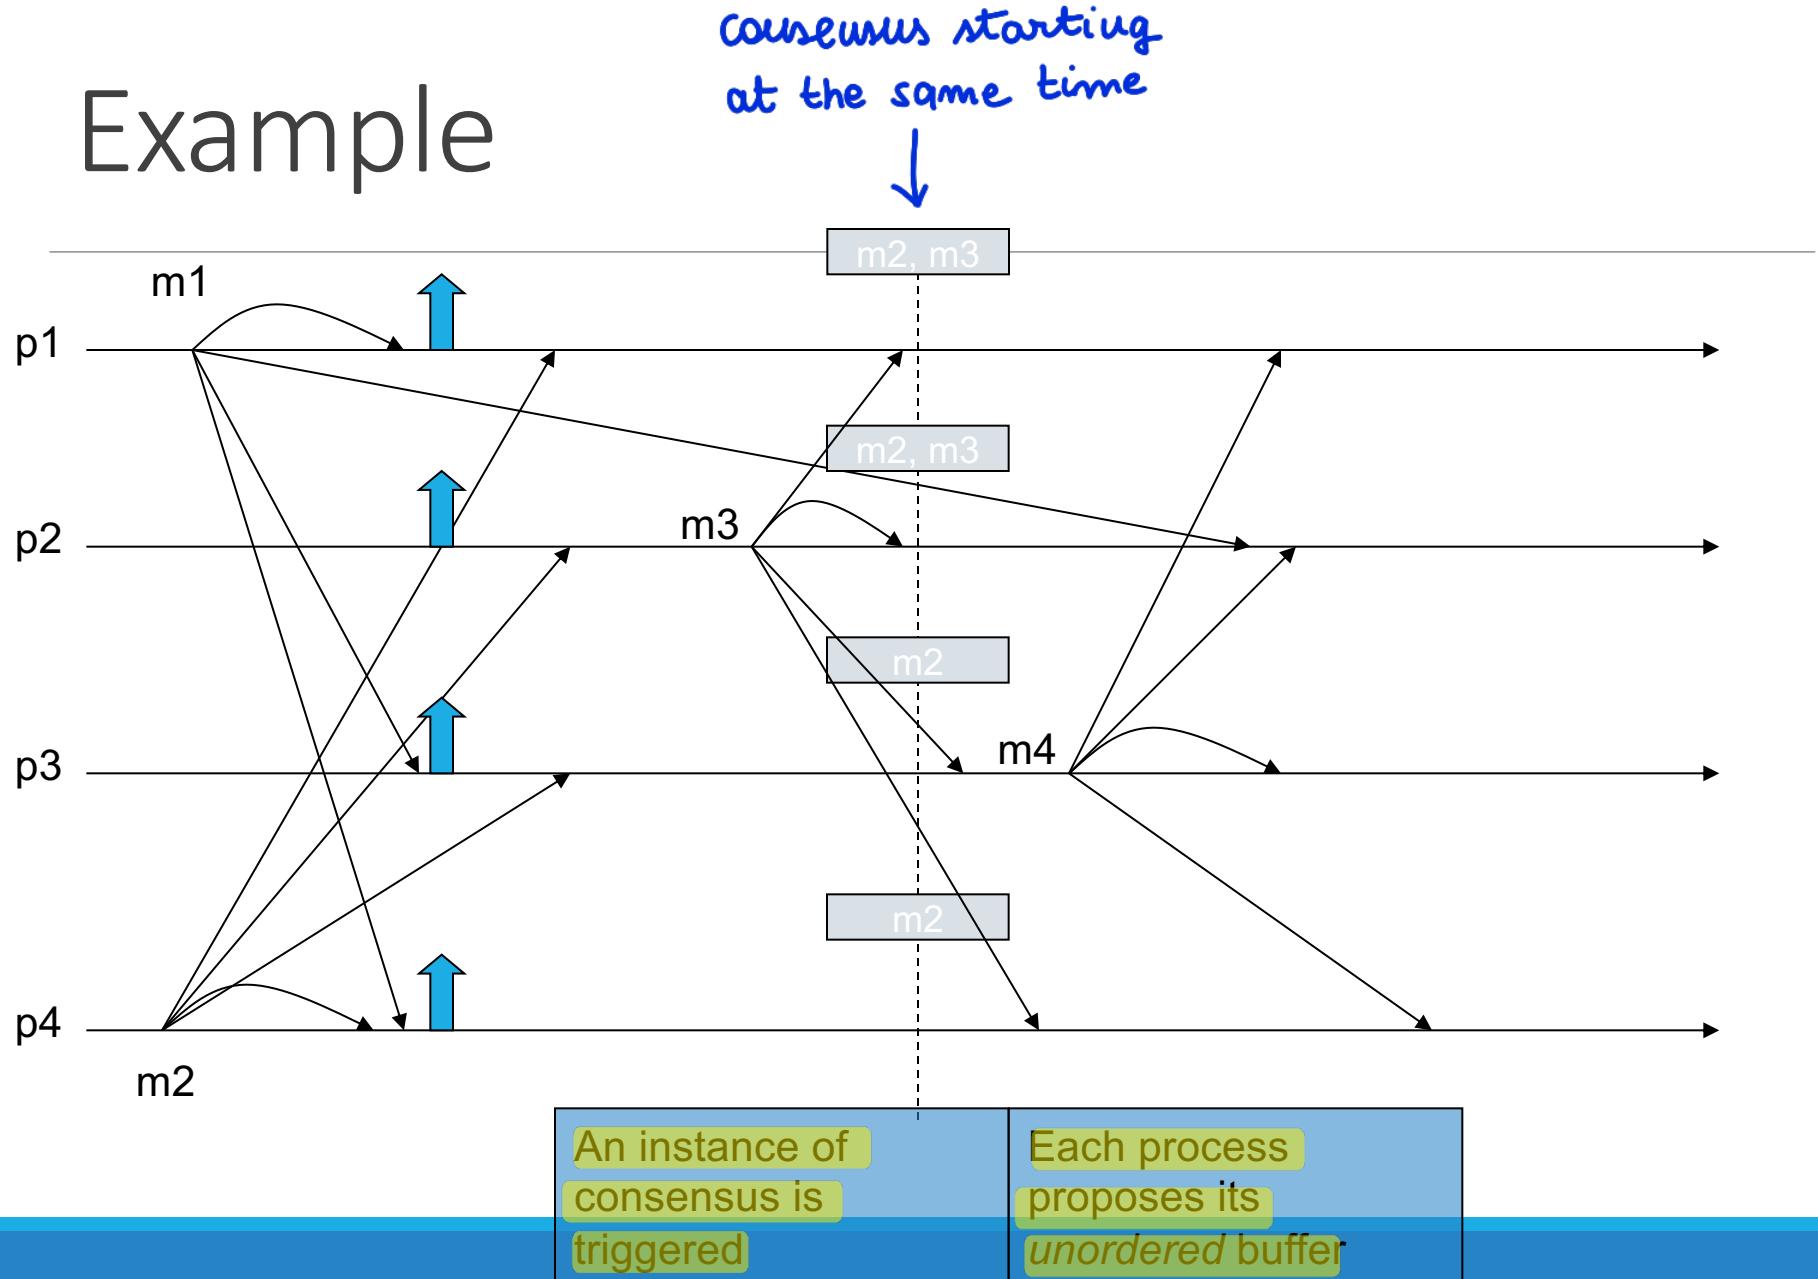
Example



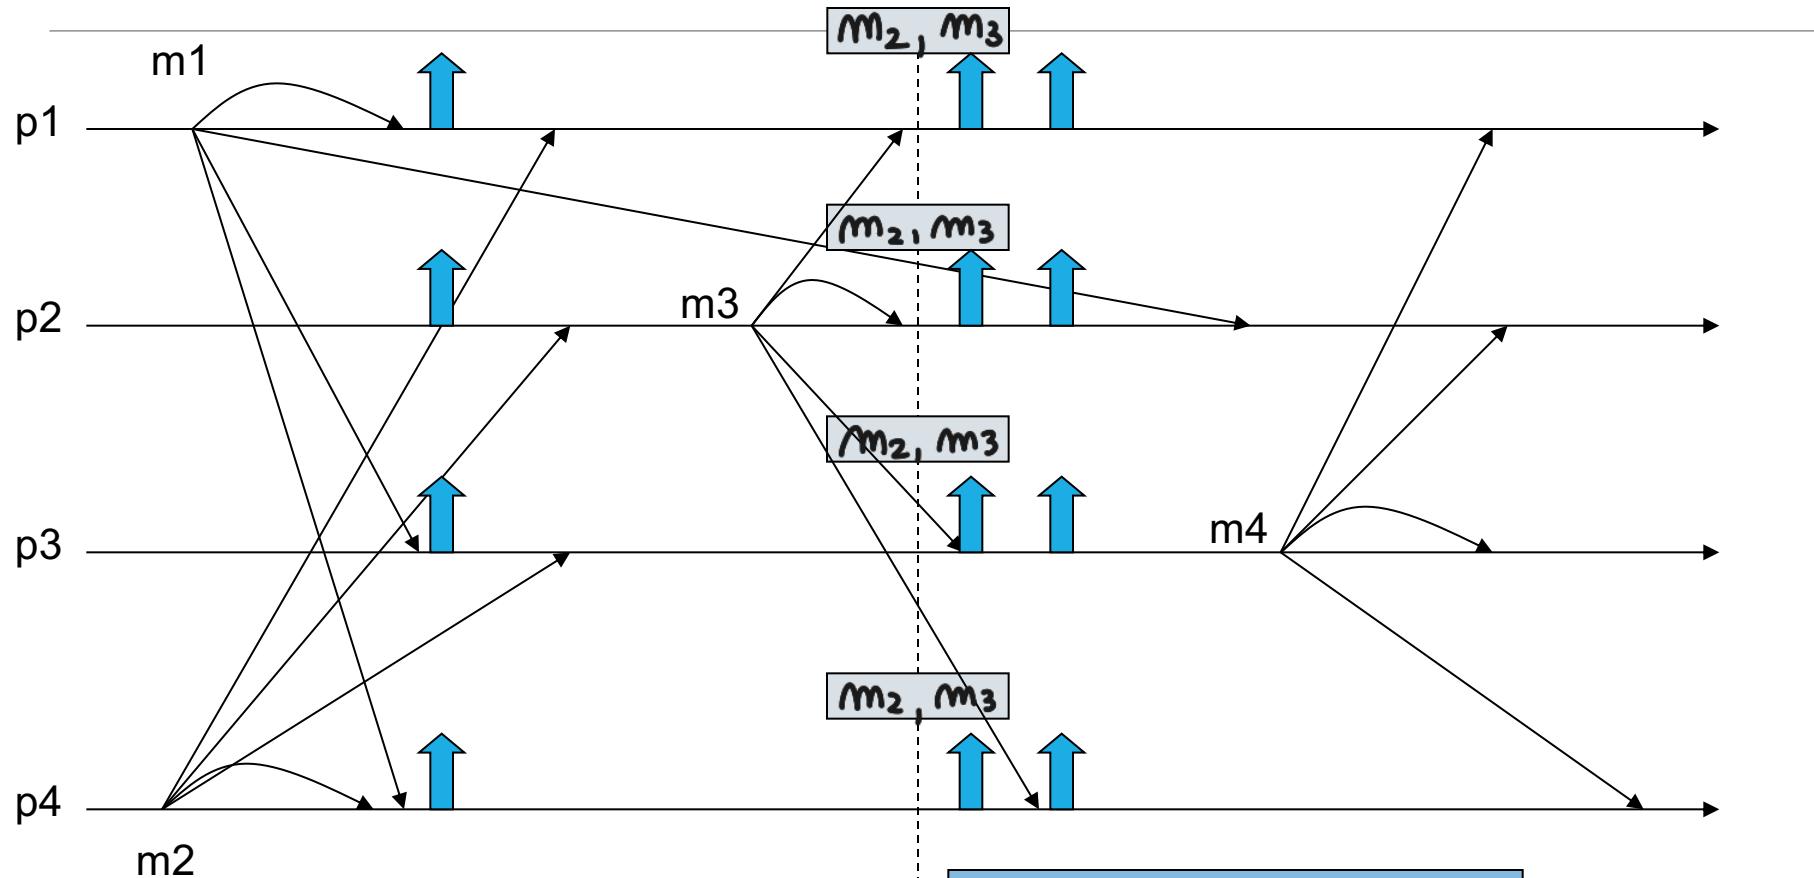
Example



Example



Example



Consensus return m_2, m_3
and all processes deliver
 m_2 and then m_3

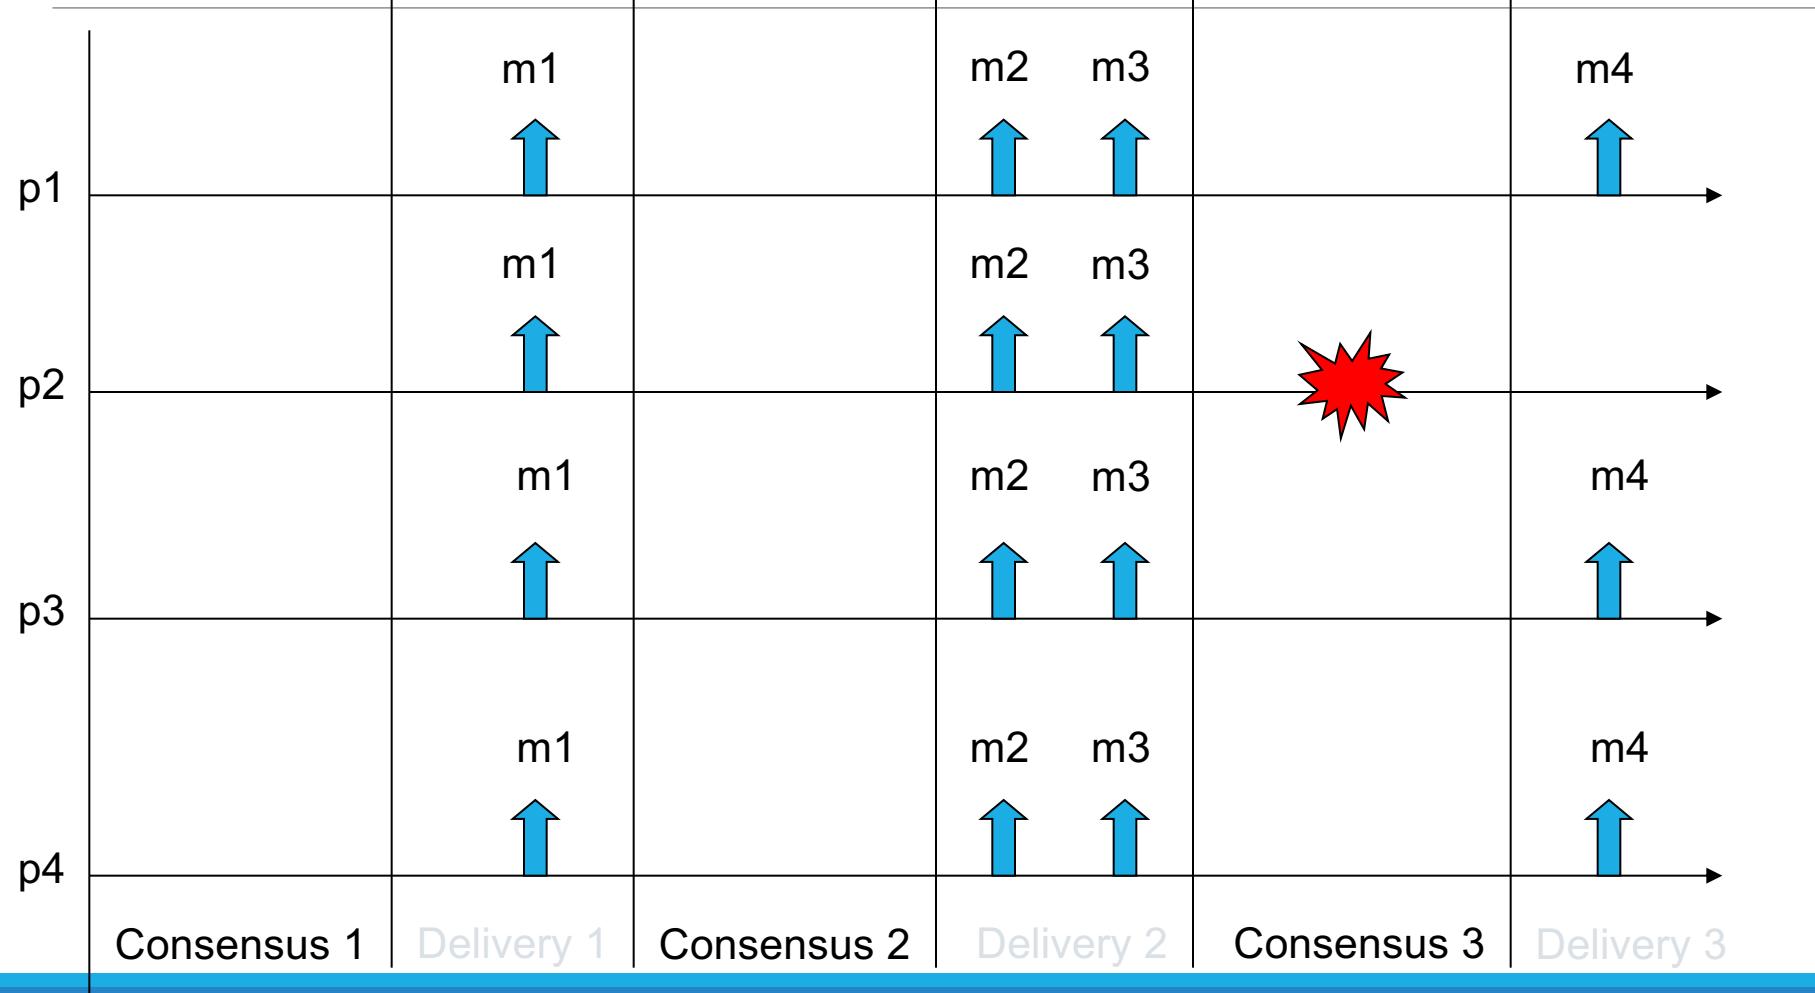
Exercise

Which TO specification is satisfied by this algorithm?

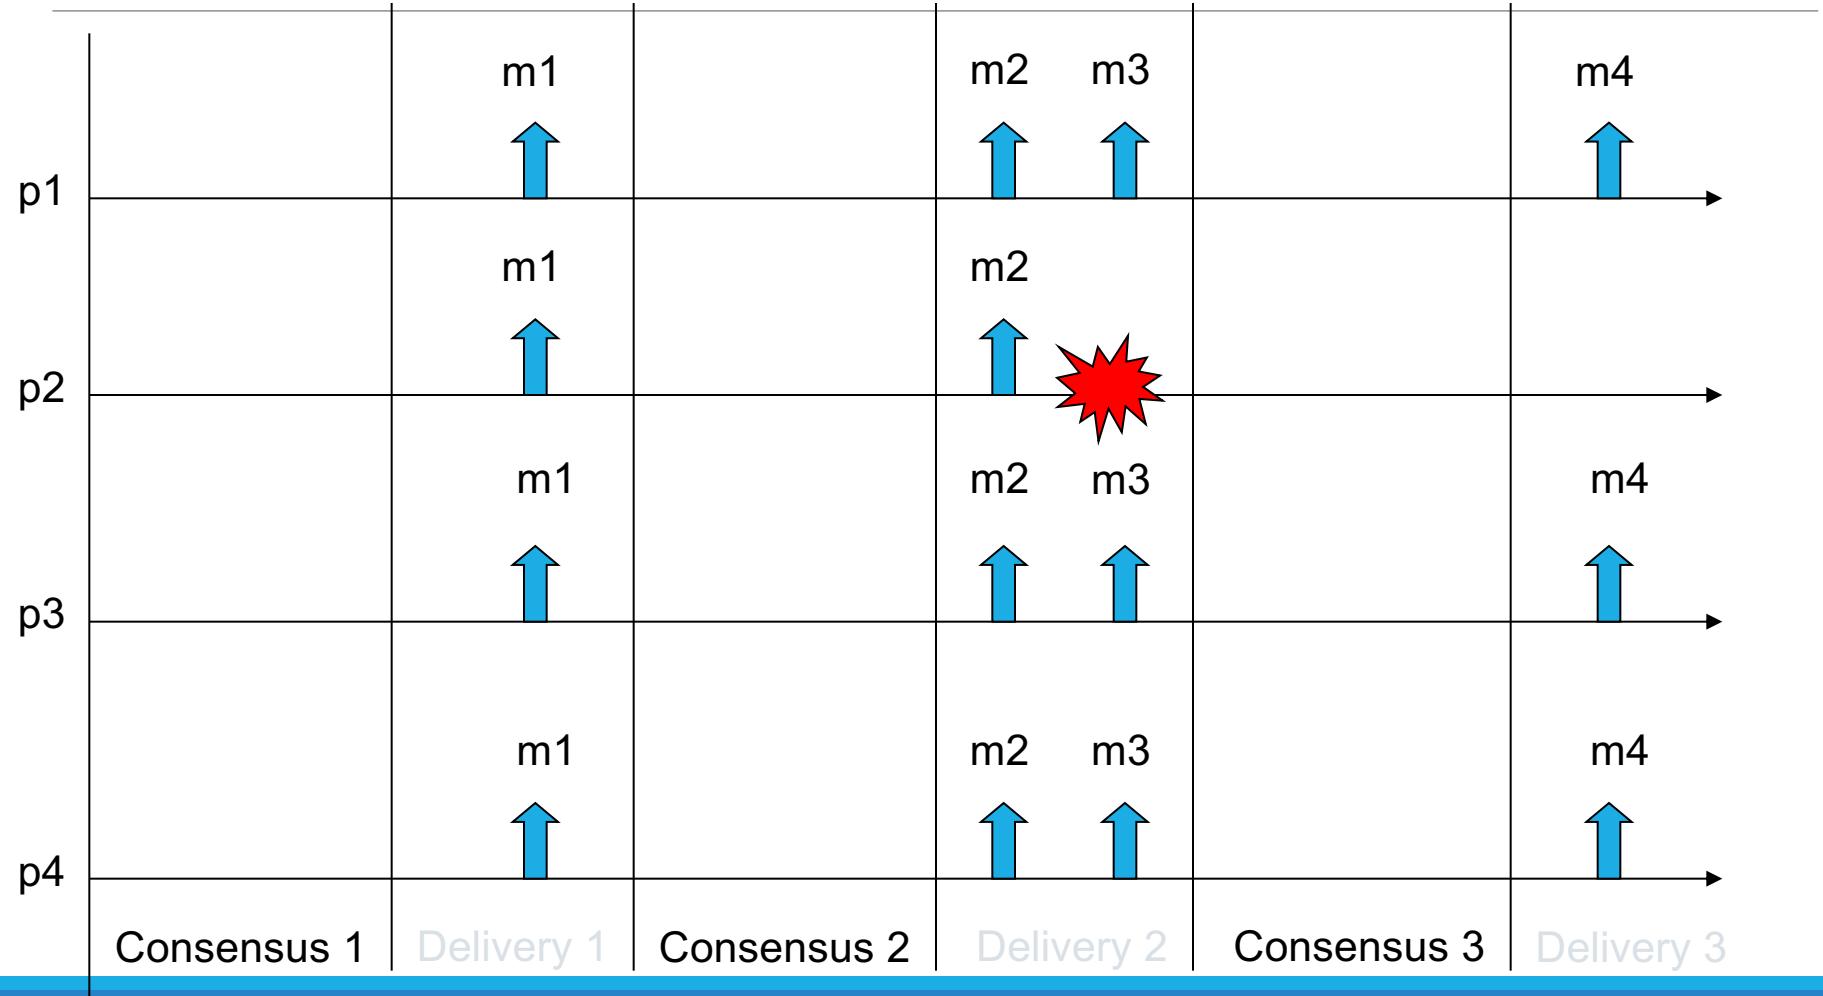
It depends from the assumptions about Reliable Broadcast and Consensus

Consensus \\ Reliable Broadcast	Uniform	Non Uniform
Uniform		
Non Uniform		

Example 1 (UC and URB)



Example 2 (UC and URB)



Uniform Consensus (UC) and Uniform Reliable Broadcast (URB)

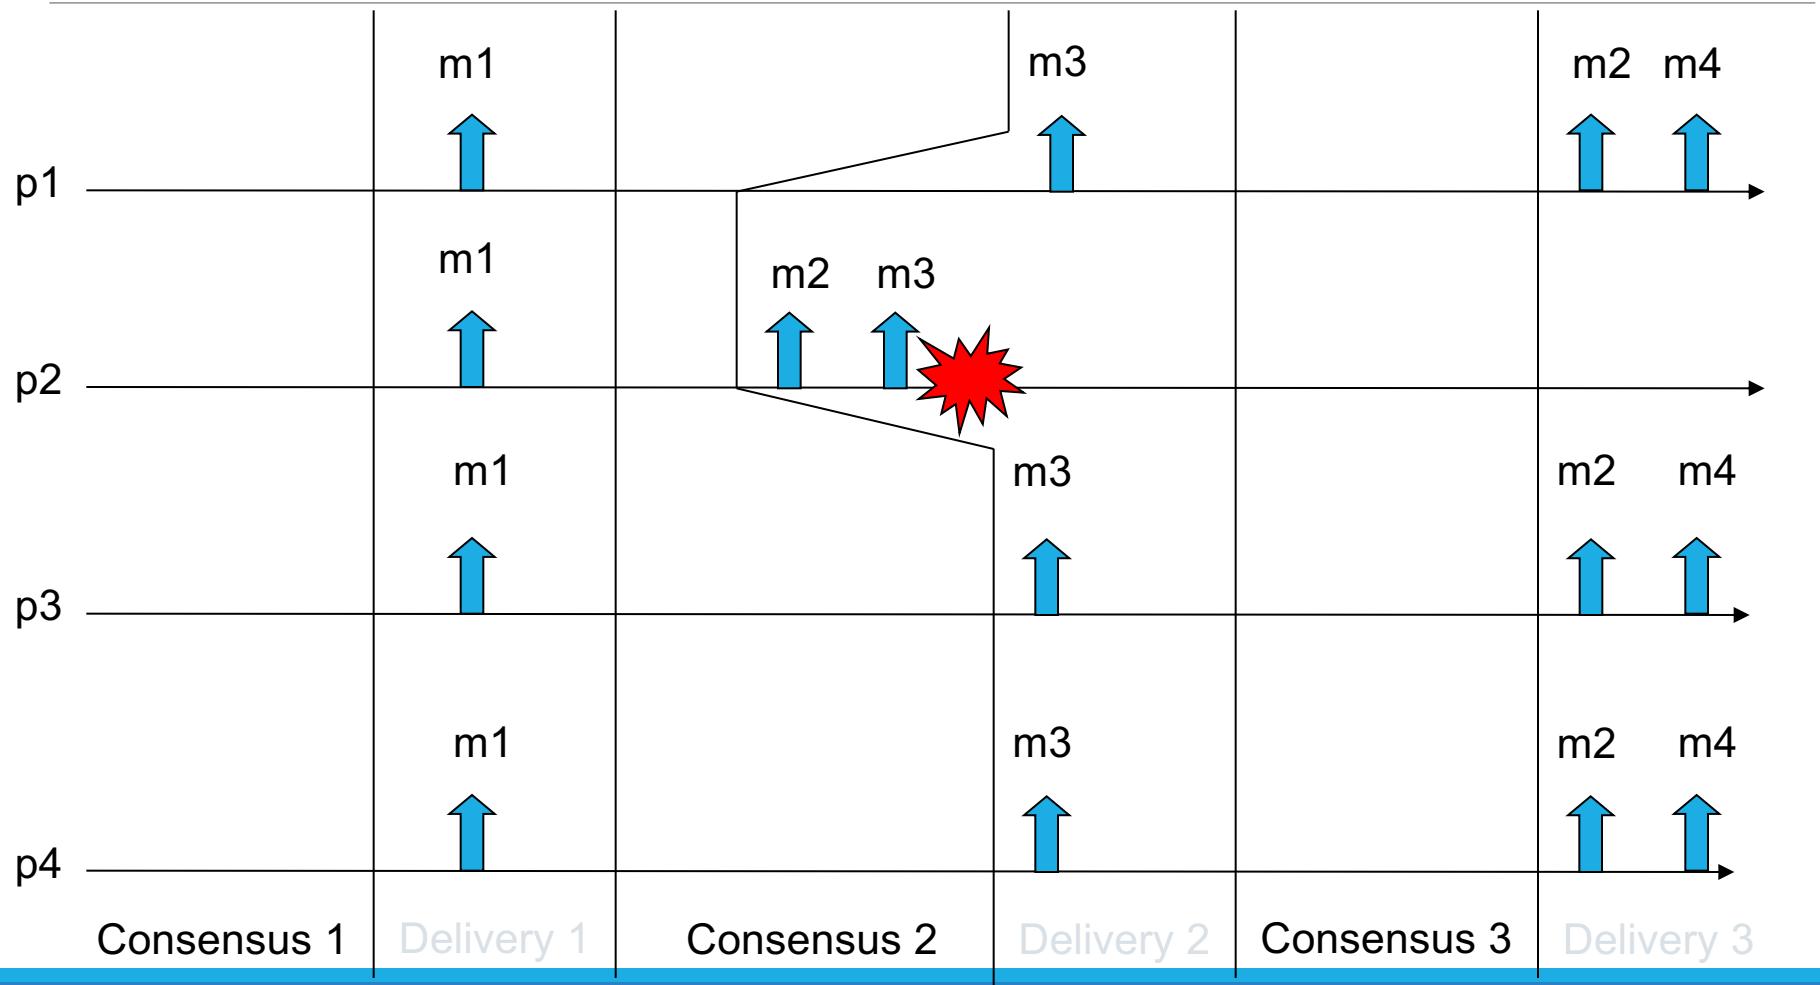
Assuming both Consensus and Reliable Broadcast uniform we have

TO (UA, SUTO)

Proof.

- Due to URB all the processes (even the faults) deliver the same set of messages
- The unordered buffer contains the same set of messages for each process
 - All the processes will deliver the same set of messages (UA)
- Due to UC, all processes (even the faults) decide for the same list of messages
- Messages are sorted by a deterministic rule
 - All processes will deliver the messages in the same order

Example (NUC and URB)



Non Uniform Consensus (NUC) and Uniform Reliable Broadcast (URB)

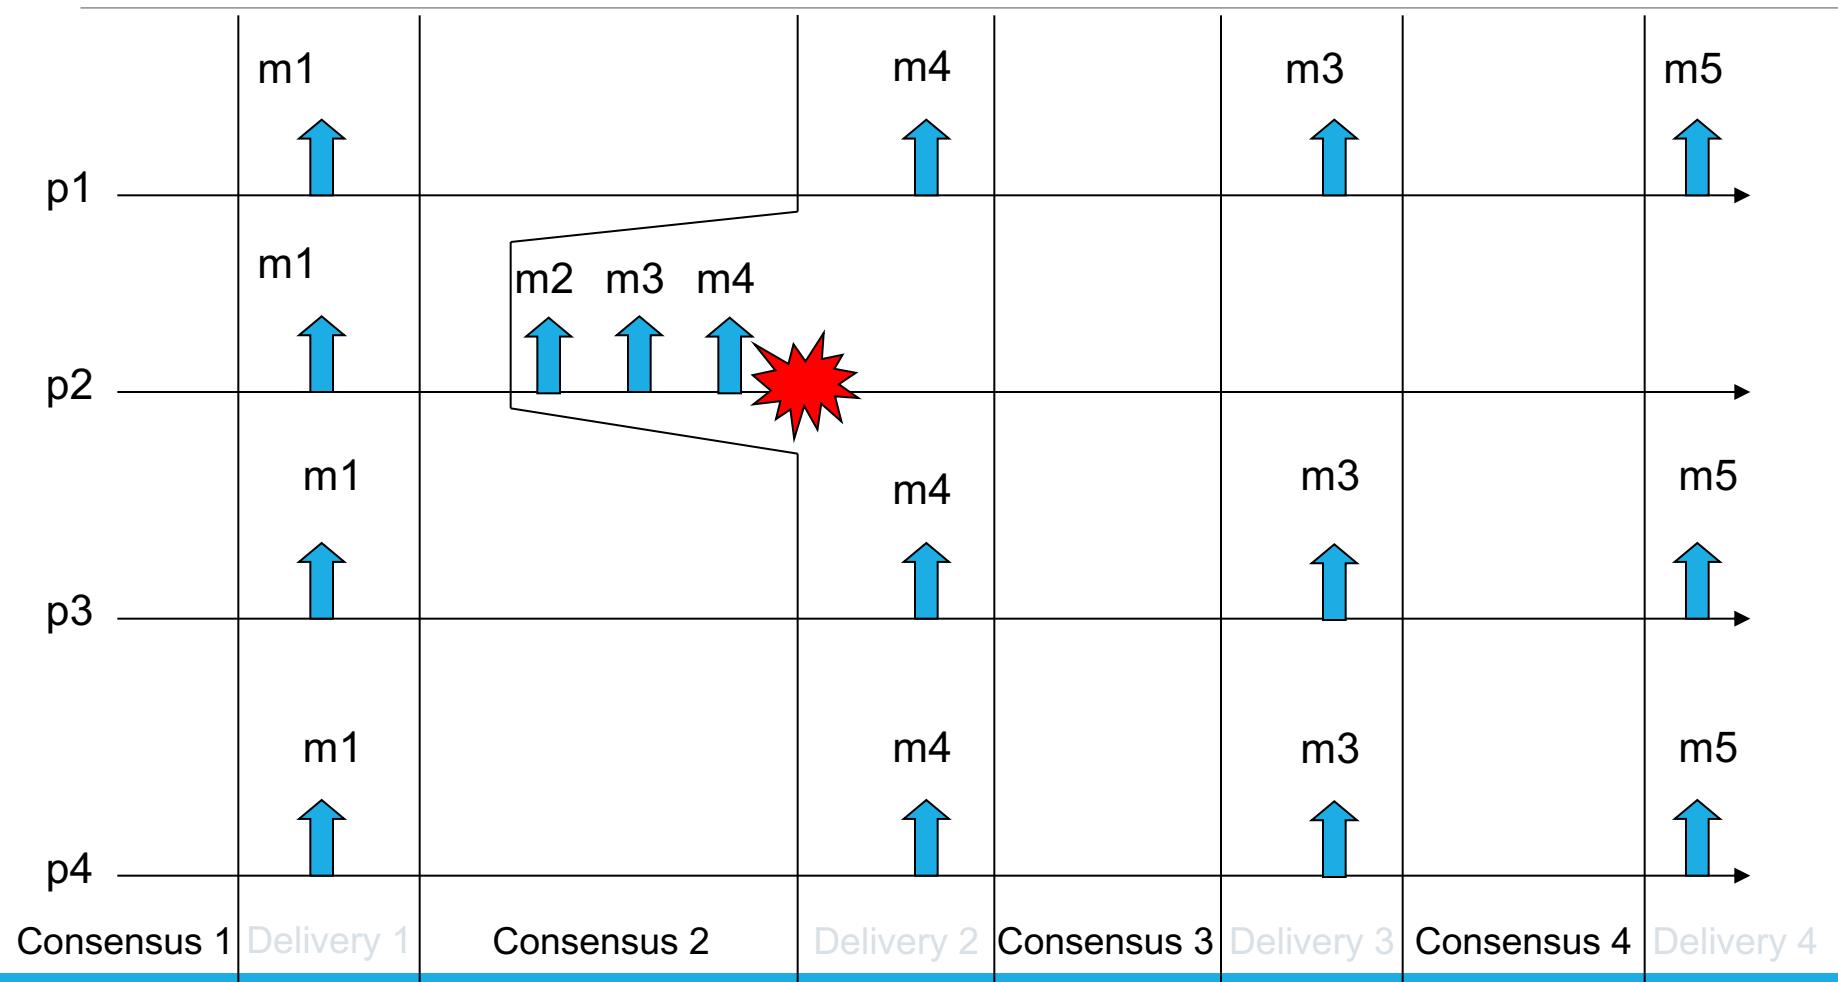
Assuming both Consensus and Reliable Broadcast uniform we have

TO (UA, WNUTO)

Proof.

- Due to URB all the processes (even the faults) deliver the same set of messages
- The unordered buffer contains the same set of messages for each process
 - All the processes will deliver the same set of messages (UA)
- Due to NUC, all correct processes decide for the same list of messages
- Faulty processes can decide differently
 - All correct processes will deliver the messages in the same order
 - Faulty processes will deliver, just before a crash, a different sequence of messages

Example (NUC and NURB)



Non Uniform Consensus (NUC) and Non Uniform Reliable Broadcast (NURB)

Assuming both Consensus and Reliable Broadcast uniform we have

TO (NUA, WNUTO)

Proof.

- Due to NURB correct processes deliver the same set of messages
- Faulty processes can deliver other messages
 - Only correct processes will deliver the same set of messages (NUA)
- Due to NUC, all correct processes decide for the same list of messages
- Faulty processes can decide differently
 - All correct processes will deliver the messages in the same order
 - Faulty processes will deliver, just before a crash, a different sequence of messages

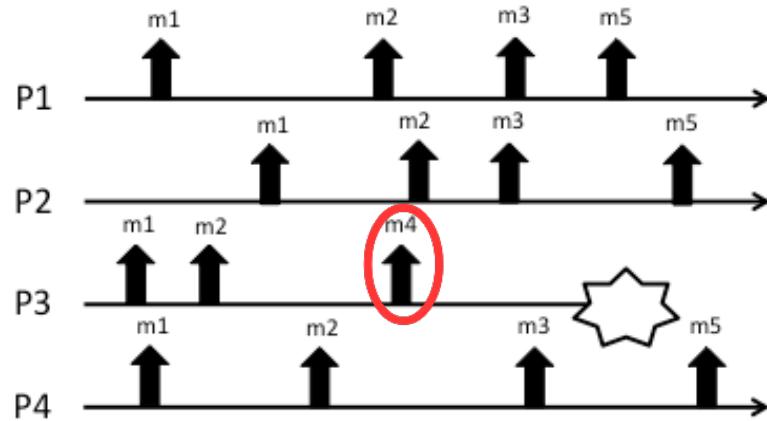
Consensus	Uniform	Non Uniform
Reliable Broadcast		
Uniform	UA SUTO	UA WNUTO
Non Uniform		NUA WNUTO

Exercice

Which specification is satisfied assuming UC and NURB?

Exercice

Consider the run depicted in the figure:



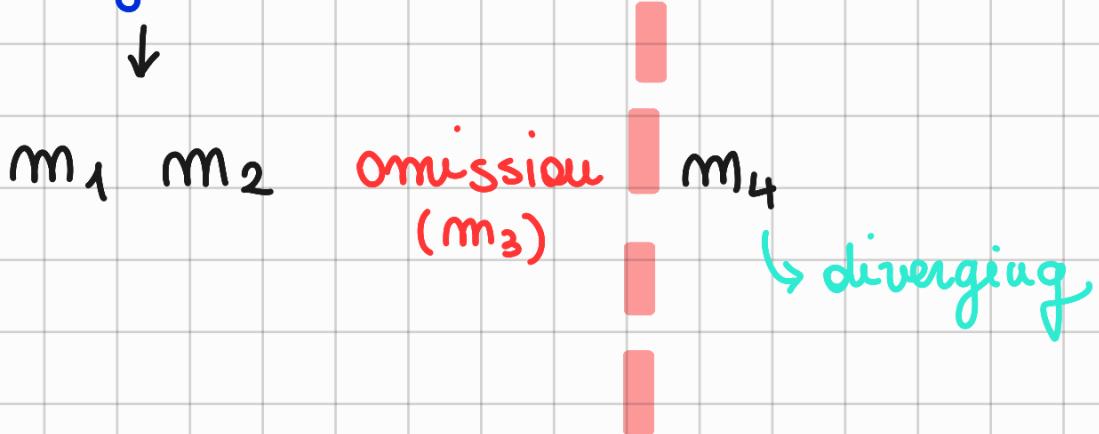
1. Which type of total ordering is satisfied by the run? Specify both the agreement and the ordering properties.
2. Modify the run in order to satisfy TO(UA, WUTO) but not TO (UA SUTO)
3. Modify the run in order to satisfy TO(NUA, WNUTO) but not TO(NUA, WUTO)

1) Correct deliver the same set: $m_1 m_2 m_3 m_5$

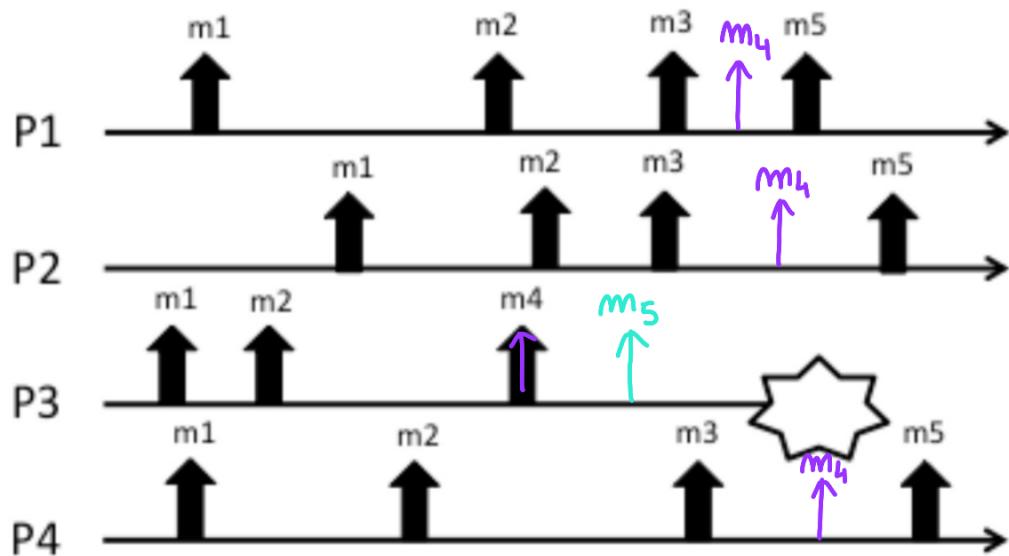
Faulty: deliver $m_1 m_2$ m_4 not delivered by correct

Agreement: NON-UNIFORM

Ordering: SUTO



2) TO(UA, WUTO) but not TO(UA, SUTO)



References

C. Cachin, R. Guerraoui and L. Rodrigues. Introduction to Reliable and Secure Distributed Programming, Springer, 2011

- Chapter 6 – Section 6.1

Stefano Cimmino, Carlo Marchetti, Roberto Baldoni "*A Guided Tour on Total Order Specifications*" WORDS Fall 2003: 187-194