

Dependable Distributed Systems

Master of Science in Engineering in Computer Science

AA 2023/2024

LECTURE 22A - OVERVIEW ON CAPACITY PLANNING

Schema

Recap dependability

Dependability is the ability of a system to deliver a service that can justifiably be trusted,

it is the ability to avoid service failures that are more frequent and more severe than is acceptable

A **service failure** is an event that occurs when the delivered service deviates from correct service

A **correct service** is delivered when the service implements its functional specifications in terms of

- **functionality**
- **performance**

Why do performance affect correct service?

SLA

It defines how a service should operate within agreed-upon boundaries
It is a formal agreement (contract) between a provider and a consumer of a service

SLAs determine what a user of an application can expect

It is composed by one or more Service Level Objectives (SLO) defined over Service Level Indicator (SLI) generally in terms of response time, throughput, system availability, reliability, etc.

- focus on metrics that users can understand
- set easy-to-measure goals
- strong relationship between IT costs and SLA

SLA composed by SERVICE LEVEL OBJECTIVES
(SLO)
defined over SERVICE LEVEL INDICATOR
(SLI)

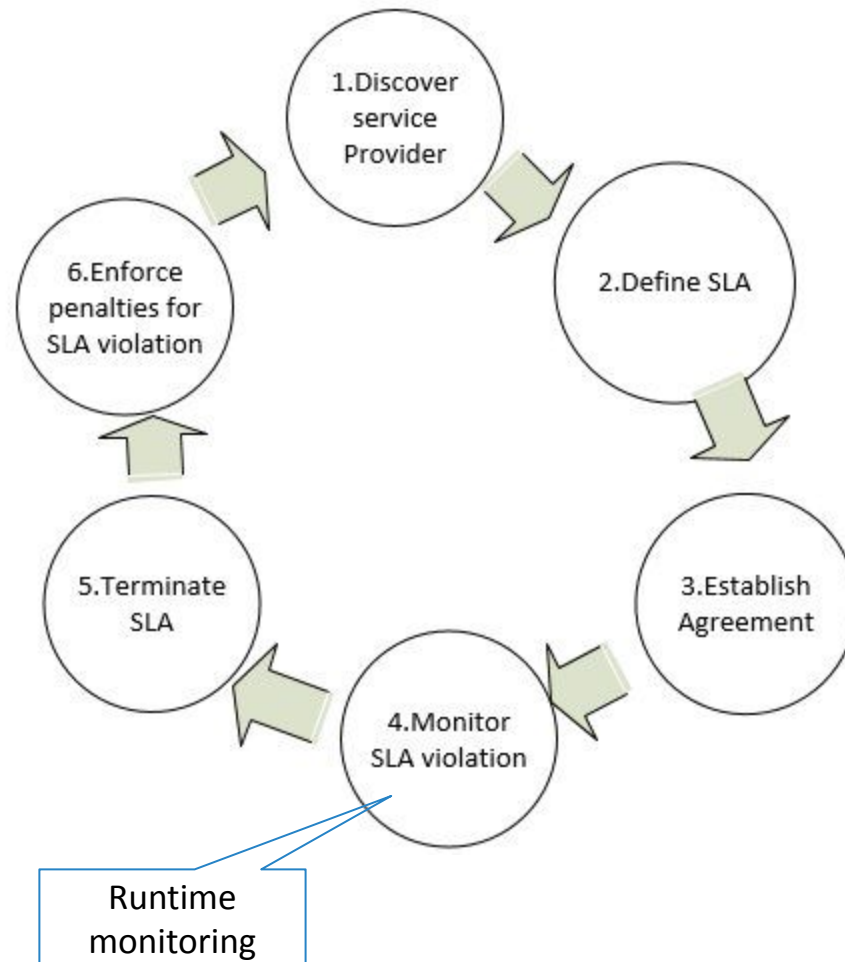
```
graph TD; SLI[SERVICE LEVEL INDICATOR (SLI)] --> RT[response time]; SLI --> TH[throughput]; SLI --> Dots[...];
```

response time

throughput

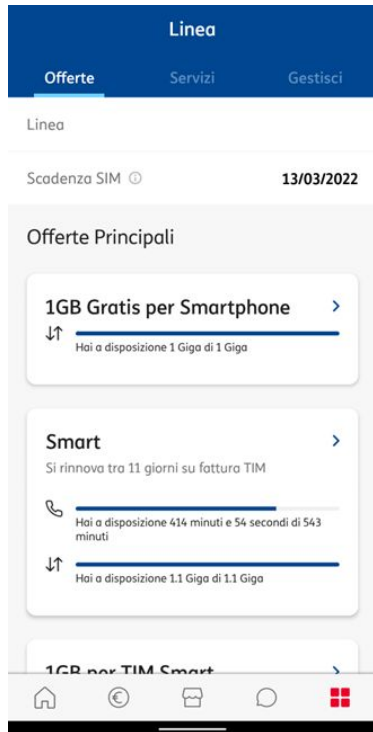
...

SLA life cycle



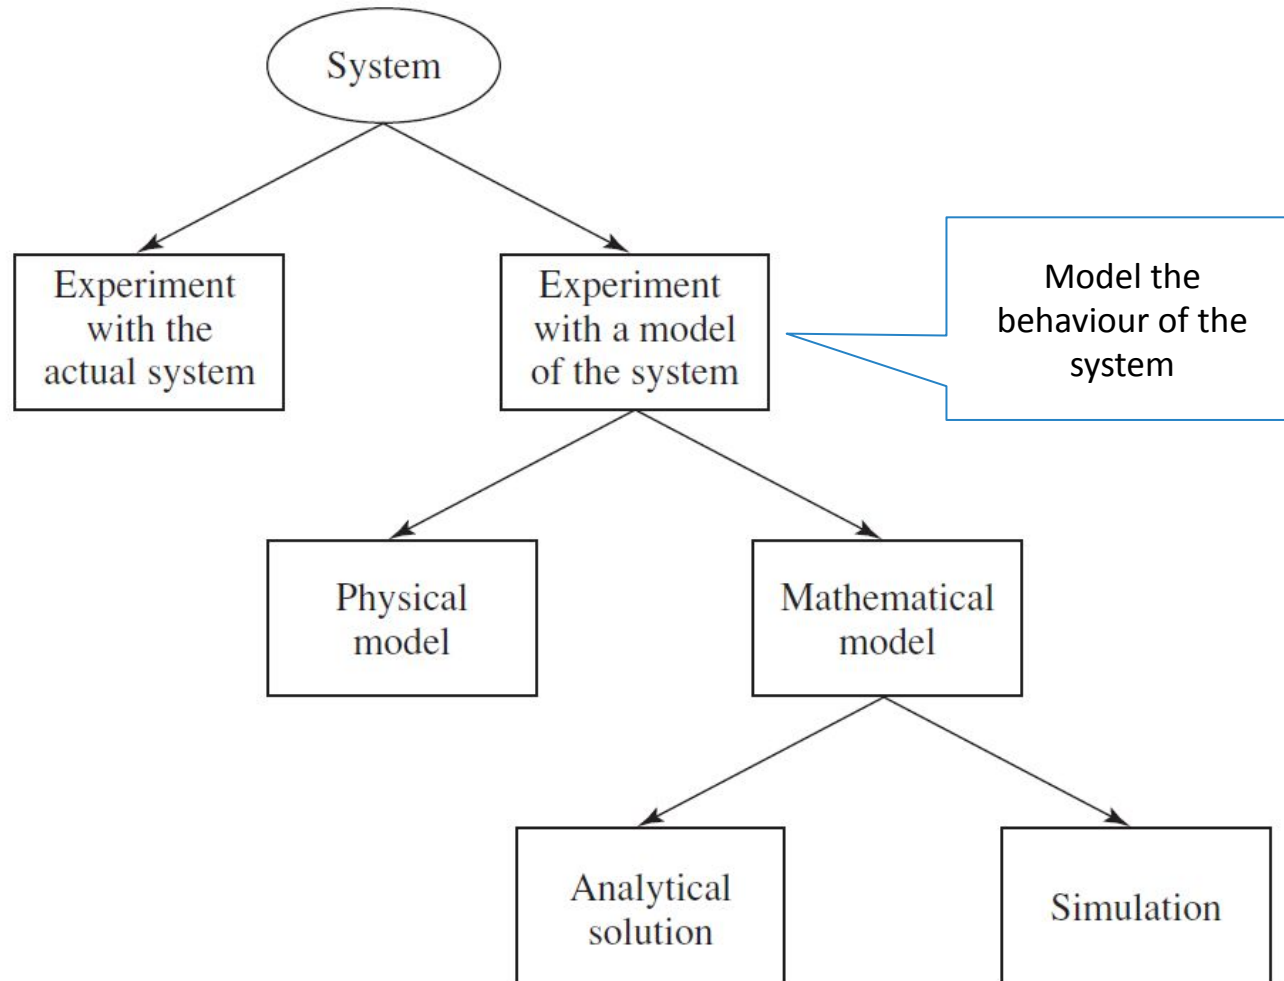
Why do performance affect correct service?

Users expectation

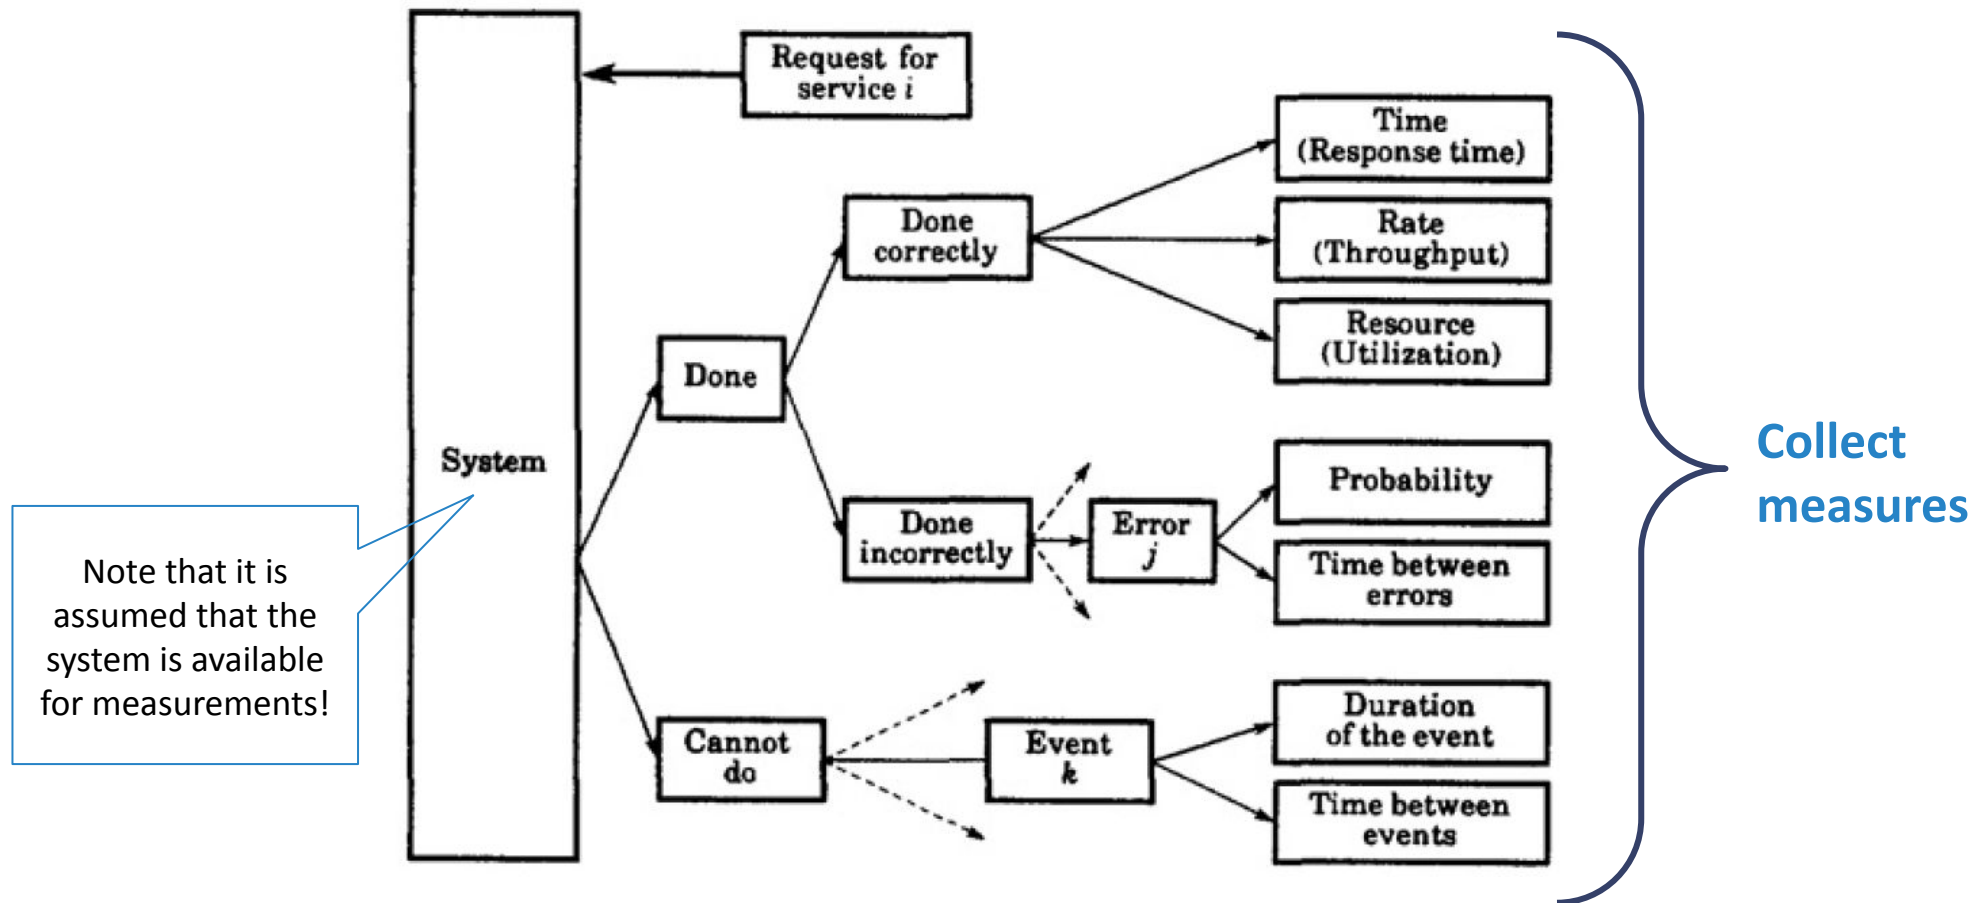


Users expectation **varies depending on what type of application** they are using and even what portion of the application they are interacting with

Ways to evaluate a system



Very Basics for Dependability Evaluation



How to achieve SLOs? Capacity Planning



OBJECTIVES

IT capacity planning consists in **estimating** the storage, hardware, software and connection infrastructure **resources required over some future period** of time to **correctly support service provisioning**

Alternatively

IT capacity planning is the process of **predicting when the service levels will be violated as a function of the workload evolution**, as well as the **determination of the most cost-effective way of delaying system saturation**.

_> **Adequate capacity**


Properly handle peaks and average behaviour

Why Do Capacity Planning?


Naive solution, **over-provisioning** : resource provisioning by taking into account ^{picco} peak loads

_> under-utilization

_> higher costs than required



over-provisioning




optimum

Why Do Capacity Planning?

On the other hand, in case of under-provisioning

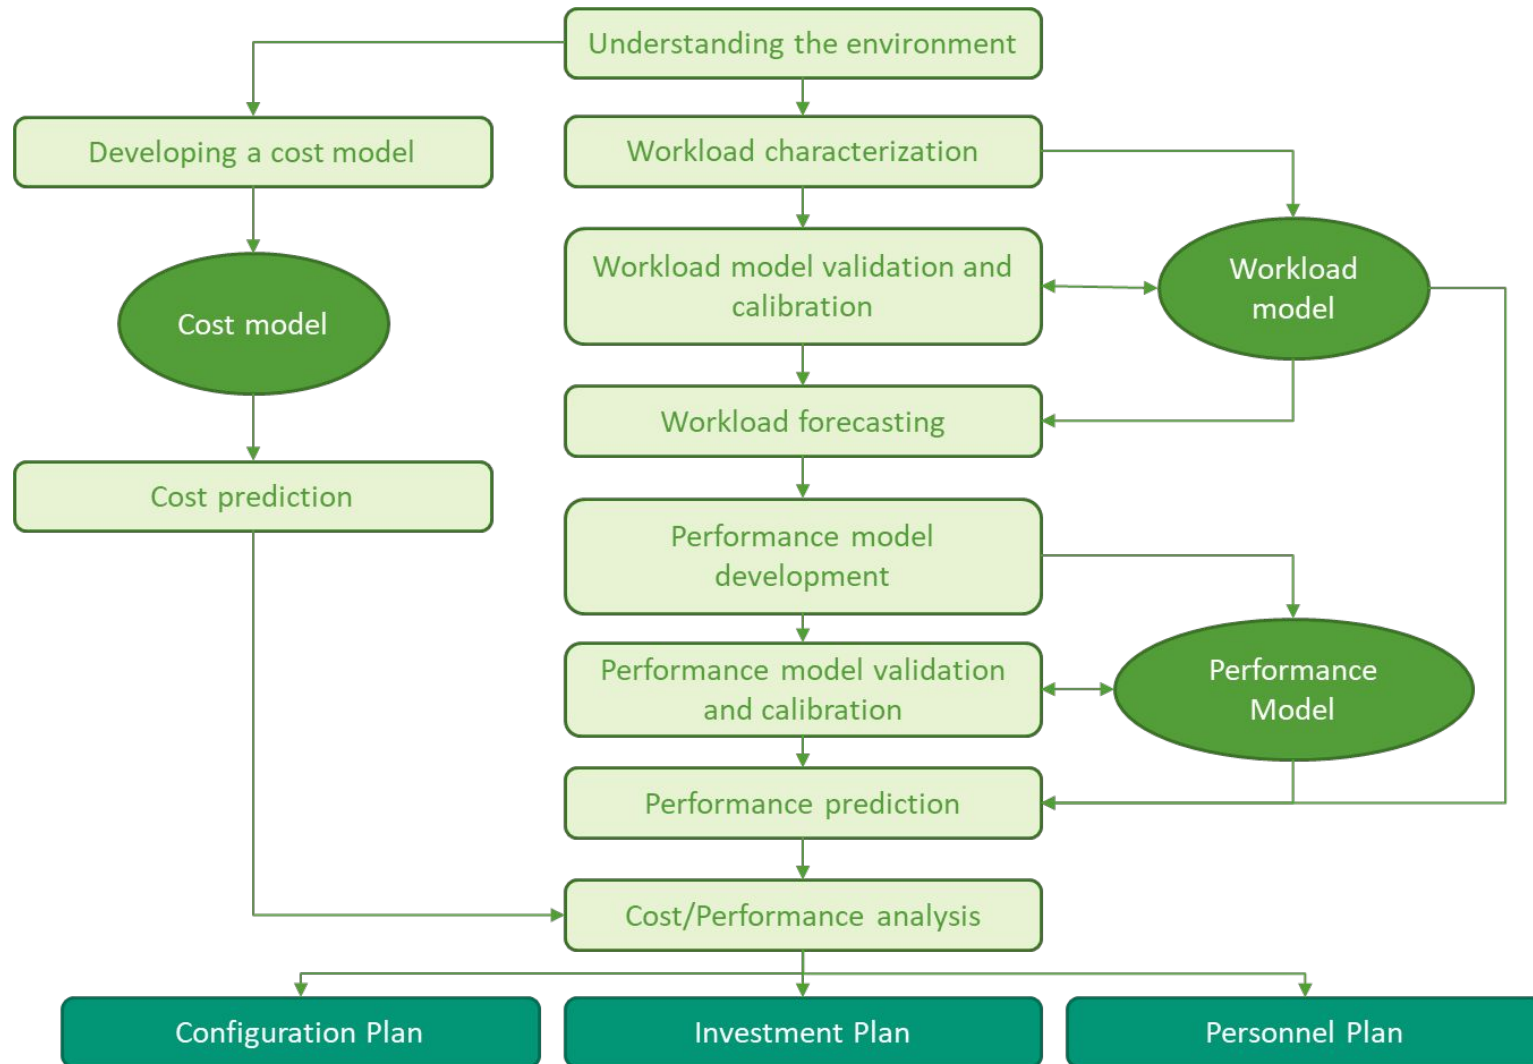
- **Overload**
- **Poor performance**
- **Losses** (users and SLA violations)



Why Do Capacity Planning?

- Avoid financial losses
- Ensure customer satisfaction
- Preserve company's external image
- It cheaper compared to deploying and testing
- Often you are in the design phase of a system
- **Capacity planning problem cannot be solved instantaneously**

A methodology for Capacity Planning



Understanding the environment

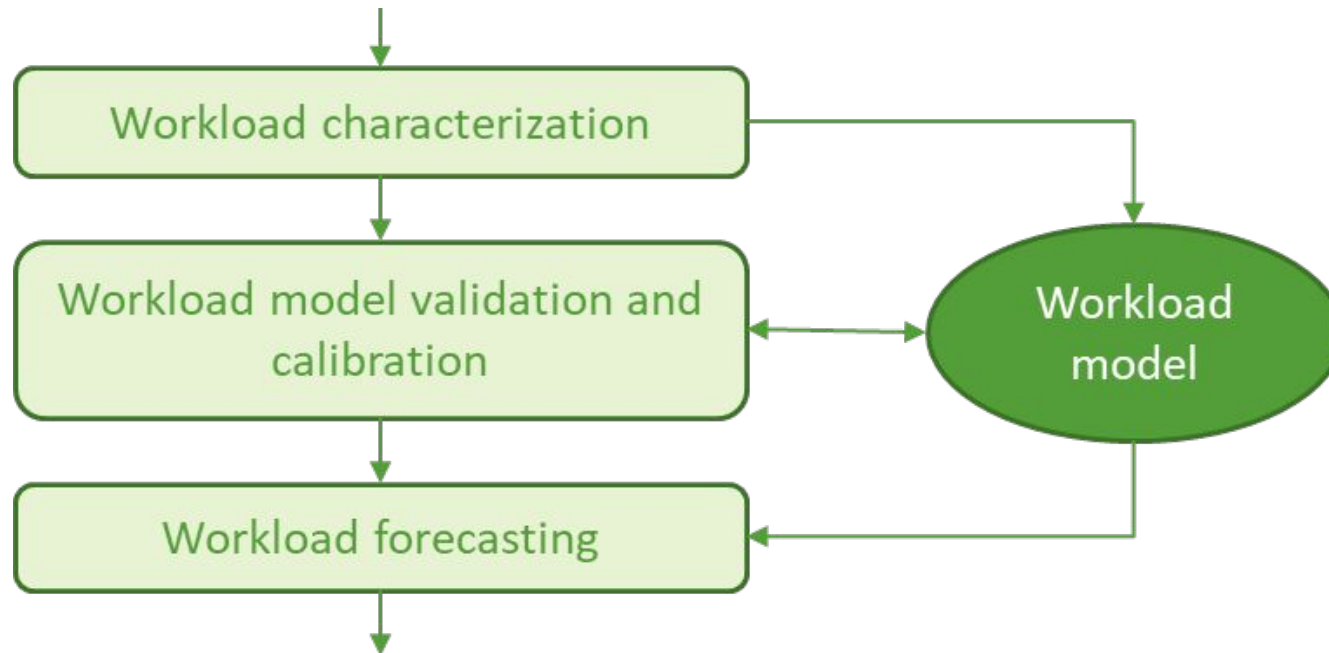
The goal is to learn what kind of

- **hardware** (clients and servers)
- **software** (OS, middleware, applications)
- **network connectivity and protocols**
- **SLA**
- ... (whatever may have an impact on the considered performance metrics)

are present in the environment

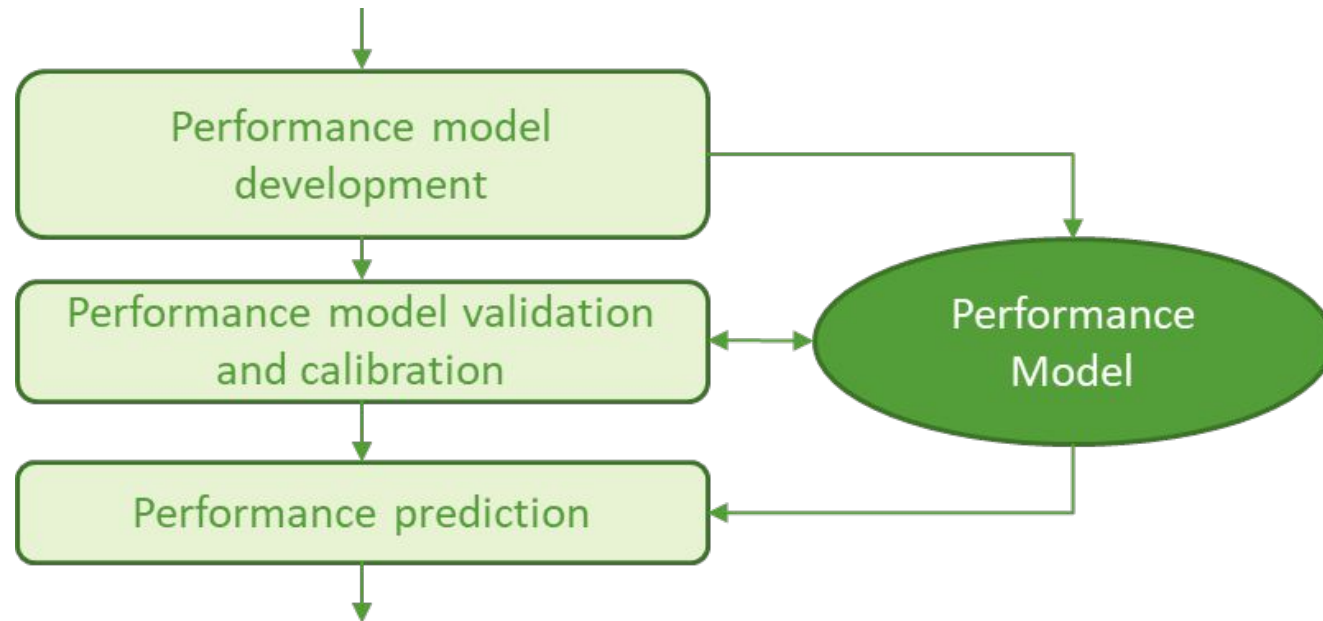
Workload model

The workload of a system is the set of all inputs that the system receives from its environment in a given period of time.



Performance model

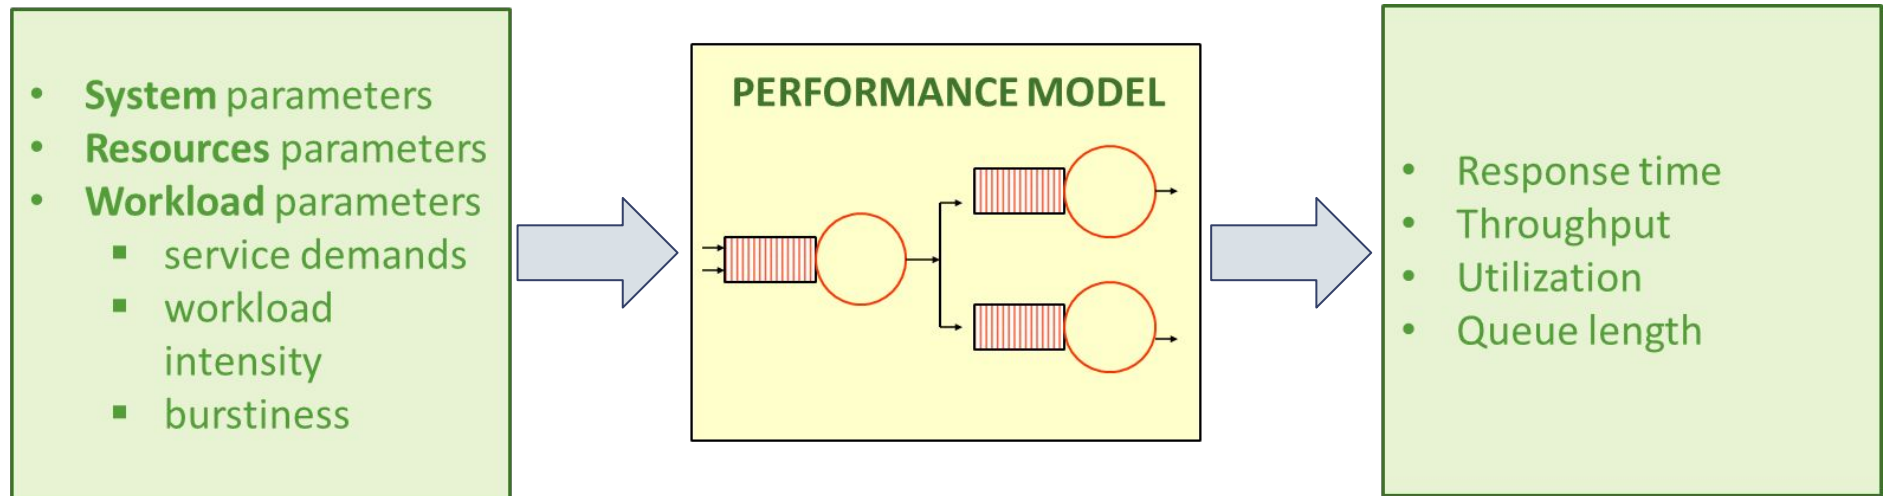
Predicts performance metrics based on system description and workload parameters



Estimates performance measures of a computer system for a given set of parameters

Outputs: response times, throughputs, system resources utilizations, queue lengths, etc.

Estimating performance measures



Parameters affecting performance metrics

System parameters examples:

- load-balancing disciplines
- network protocols
- max. num of connections supported
- ...

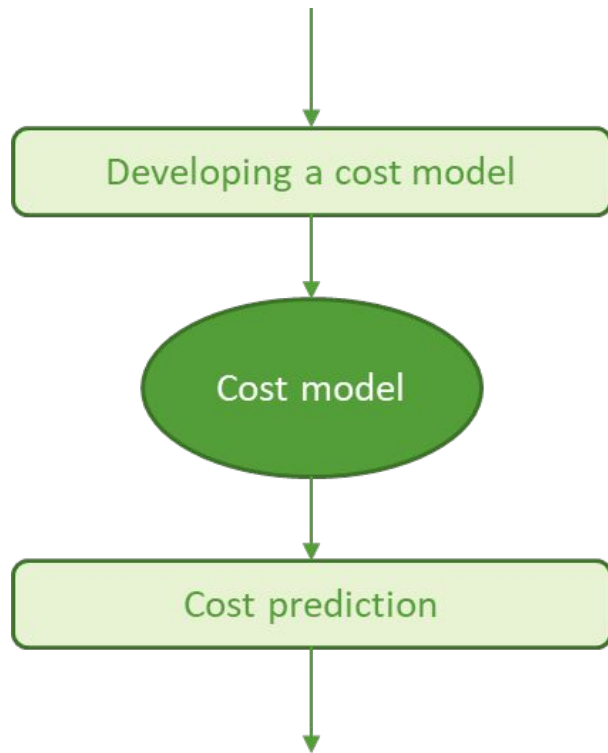
Resource parameters examples:

- disk latency, transfer rate
- network bandwidth;
- CPU speed
- ...

Workload parameters examples

- WL intensity parameters:
 - num. of requests
 - num. of clients running an application
 - Burstiness
 - ...
- WL service demand parameters:
 - CPU time per request
 - Disk usage per request
 - ...

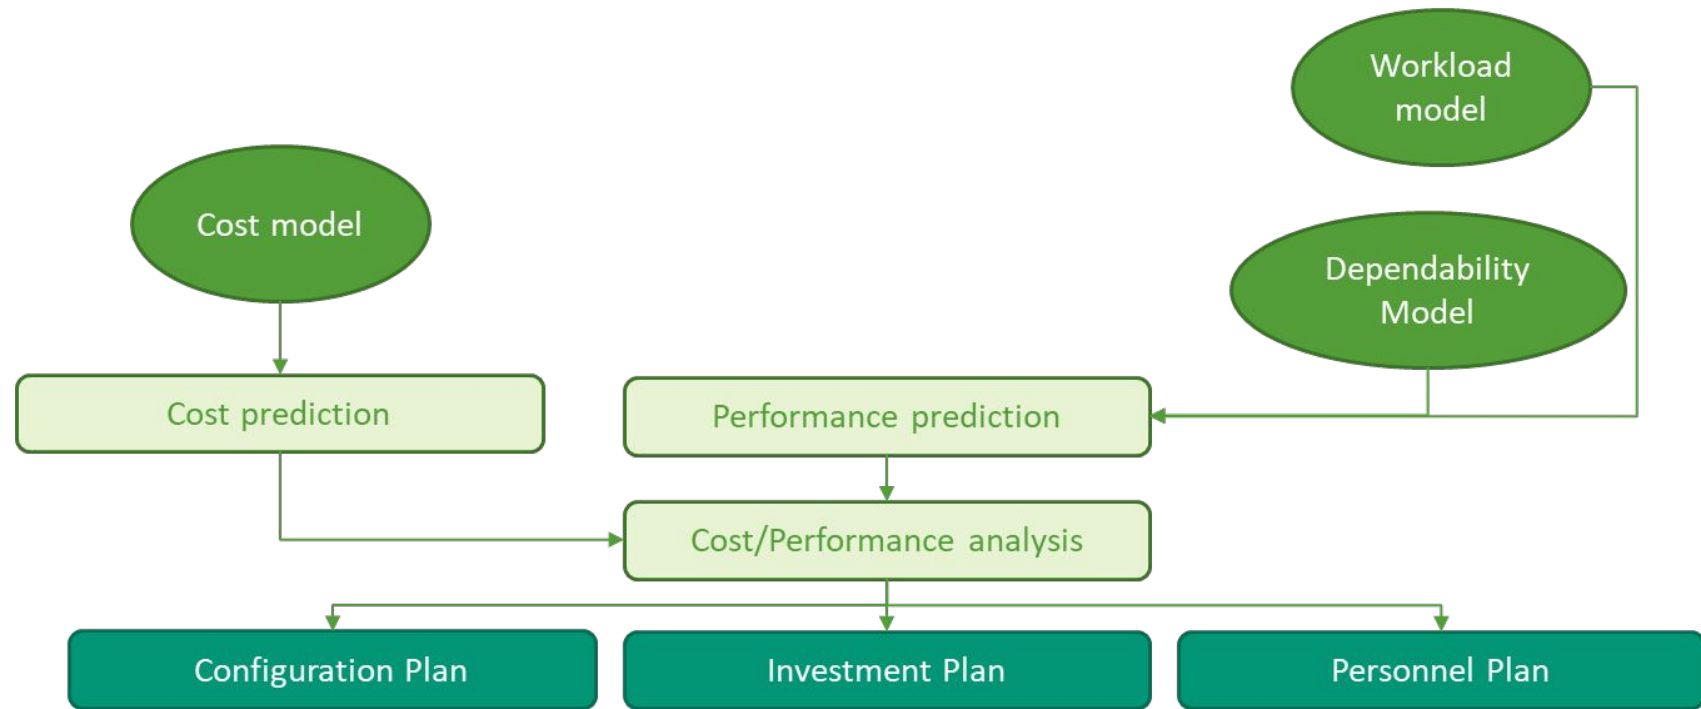
Cost model



Categories:

- **Hardware cost**: machines, disks, routers, etc.
- **Software cost**: operating systems, middleware, etc.
- **Telecommunication cost**
- ...


Cost/performance analysis



- Assess possible **scenarios**
- For each scenario, **predicts performance metrics and costs**
- Comparing scenarios, **get configuration, investment and personnel plans**
- Assess **payback**: ROI (return of investment), company's image, etc.

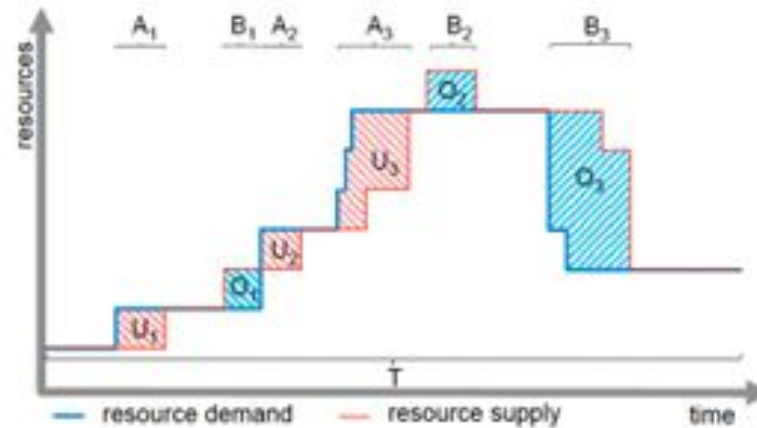
Elasticity

Elasticity is the **degree** to which a system is **able to adapt to workload changes** by provisioning and de-provisioning resources **in an autonomic manner**, such that at each point in time the available resources match the current demand as closely as possible.



Elasticity Metrics

- Consider over-provisioning and under-provisioning



- Various elasticity metrics, two examples:
 - Accuracy:** sum of areas of over-provisioning (O) and under-provisioning (U) for the duration of the measurement period T
 - Timing:** total amount of time spent in the over-provisioning (B) and under-provisioning (A) over the measurement period T

Scalability


Scalability is the property of a system to handle a growing amount of work

- Two directions for size scalability
 - **Vertical (scale-up)**: more powerful resources
 - **Horizontal (scale-out)**: more resources with same capacity



Scalability


Scalability is the property of a system to handle a growing amount of work



Scalability: Vertical VS Horizontal


- Less complex
- Upgrade limitations
- Single point of failure

Vertical



vs.


Horizontal



- Increased complexity
- No limit to the number of processes
- Increased resilience and fault tolerance
- Horizontal scale \Rightarrow Increase in performance metrics

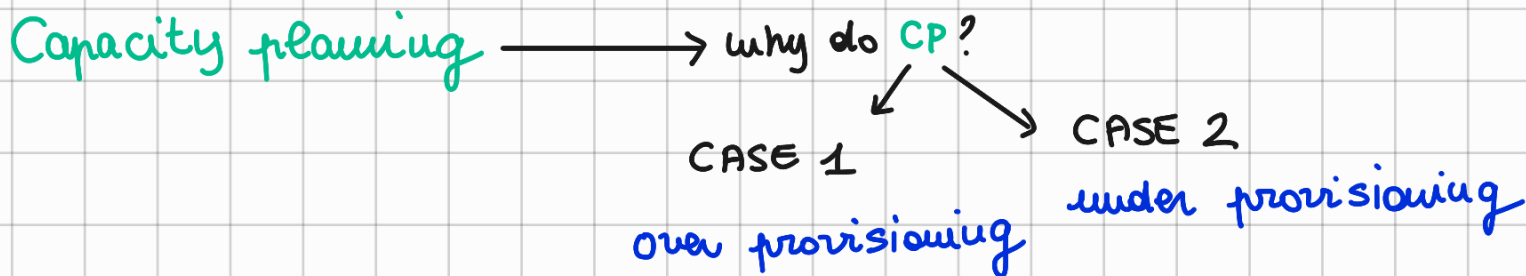
Reference

- D. A. Menascé, V. A. F. Almeida: *Capacity Planning for Web Services: metrics, models and methods*. Chapter 5 (Available in the library inside Dipartimento di Ingegneria informatica, automatica e gestionale Antonio Ruberti)

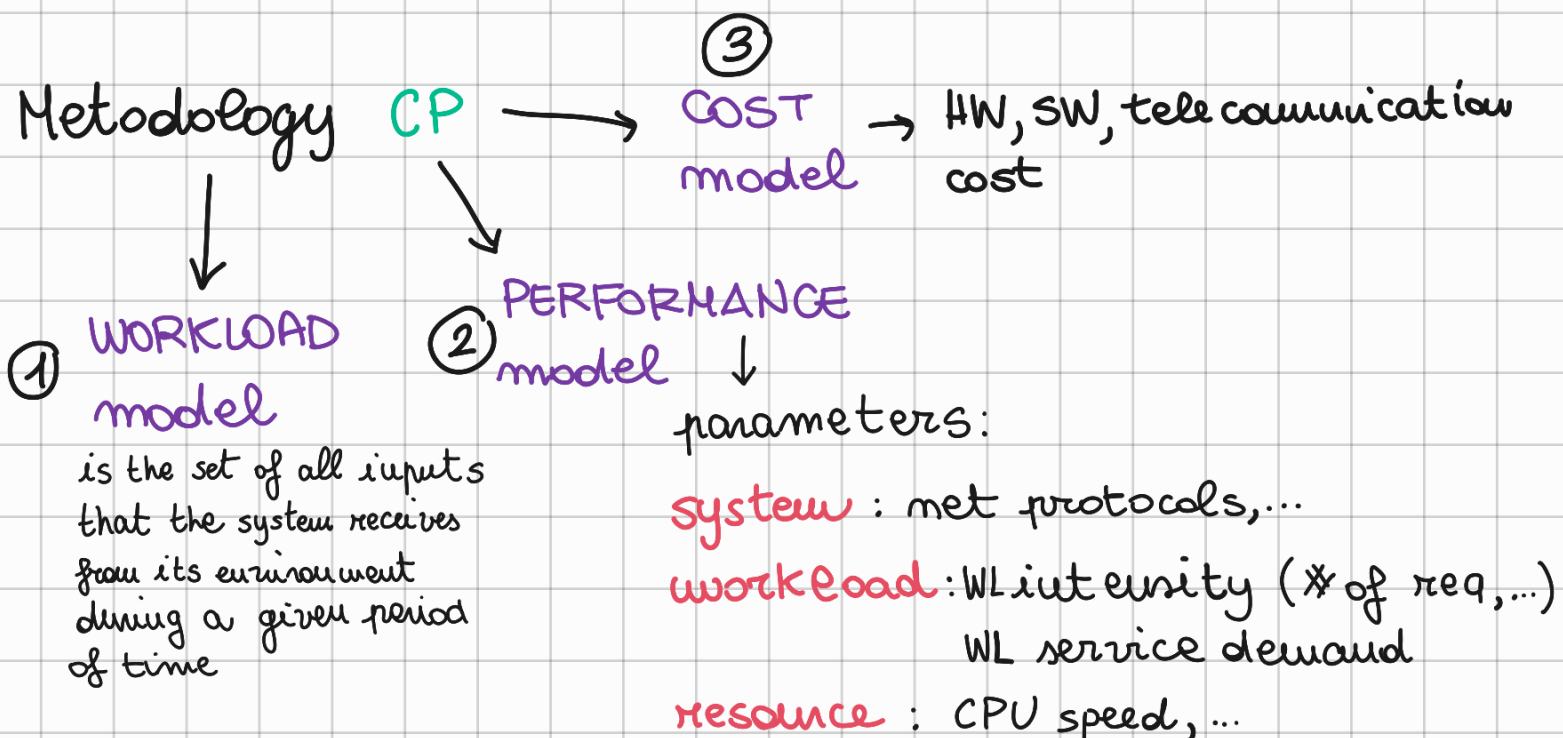


SLA composed by SERVICE LEVEL OBJECTIVES (SLO)
defined over SERVICE LEVEL INDICATOR (SLI) → throughput
↓
response time

How to achieve SLOs?
↓



GOAL: what kind of hardware, software, SLA, network, ... are presented in environment.



↓ when we have all of this
cost / perf. analysis
assess various scenarios

In DS → elasticity
metrics : accuracy
 timing
 ↓
scalability
size of scal : vertical SCALE-UP
 horizontal SCALE-OUT