

Distributed Systems
Master of Science in Engineering in
Computer Science

AA 2018/2019

BROADCAST IN PRESENCE OF BYZANTINE PROCESSES
ON DYNAMIC DISTRIBUTED SYSTEMS

Distributed system := ABSTRACTION, modeling a set of spatially separate entities, each of these with a certain computational power, that are able to communicate and to coordinate among themselves for reaching a common goal.

Dynamic = continuously changing

What does it change continuously?

- ▶ The set of processes composing the systems (**Churns**), i.e. processes may leave or may enter
- ▶ **The communication network**

Vehicular Area Network, Sensor Network, P2P Network, Social Network etc.

Dynamic Communication Network main issues

- ▶ Processes can directly exchange messages with a subset of all processes
- ▶ This subset changes over the time
- ▶ Processes may be isolated for a while

Dynamic Network Model

How does the network change?

A **model** that characterize the evolution is required

Many models proposed in the literature, one of the most general is the **TVG** model

Dynamic network model - TVG [CFQS12]

Time Varying Graph (TVG)

$$\mathcal{G} := (V, E, \rho, \zeta)$$

- ▶ V is the set of nodes;
- ▶ $E \subseteq V \times V$ is the set of edges (i.e., communication channels);
- ▶ $\rho : E \times \mathbb{N} \rightarrow \{0, 1\}$ is the *presence* function;
- ▶ $\zeta : E \times \mathbb{N} \rightarrow \mathbb{N}$ is the *latency* function;

\mathcal{G} can be alternatively described as a **sequence of static graphs (snapshots)** $\mathcal{S}_{\mathcal{G}} = G_0, G_1, \dots, G_T : G_i := (V, E_i)$, $E_i := \{e \in E \mid \rho(e, t_i) = 1\}$.

$G := (V, E)$ is called *underlying graph*

The problems that can be solved depend on the network evolution



$G1$



$G2$



G (underlying)


Problem: any process in G can broadcast (when all processes are correct)

In $\mathcal{G} = \{G1, G2\}$ ($\zeta(e, t) = 0 \forall t, e$) processes 2 and 3 can achieve broadcast, process 1 can't


Journey (or Dynamic Path)

A sequence of distinct nodes (p_1, \dots, p_n) is a *Journey* (or *Dynamic Path*) from p_1 to p_n if there exists a sequence of dates (t_1, \dots, t_n) such that, $\forall i \in \{1, \dots, n-1\}$ we have:


- ▶ $e_i = (p_i, p_{i+1}) \in E$, i.e. there exists an edge connecting p_i to p_{i+1} .
- ▶ $\forall t \in [t_i, t_i + \zeta(e_i, t_i)]$, $\rho(e_i, t) = 1$, i.e. p_i can send a message to p_{i+1} at date t_i .
- ▶ $\zeta(e_i, t_i) \leq t_{i+1} - t_i$, i.e. the aforementioned message is received by date t_{i+1} .




G_1




G_2



G_3



G_4



G (underlying)

with $\zeta(e, t) = 1 \ \forall t, e$

$(0, 2, 4)$ is a journey

$(0, 1, 4)$ is not a journey

$(0, 3, 4)$ is not a journey

Classes of TVG have been defined to identify the minimal conditions that dynamic network has to satisfy in order to solve specific distributed systems problems

- ▶ Class 1. Temporal Source:
 $\exists u \in V : \forall v \in V, u \rightsquigarrow v$. (broadcast feasible from at least one node)
- ▶ Class 2. Temporal Sink:
 $\exists u \in V : \forall v \in V, v \rightsquigarrow u$. (compute a function whose input is spread over all the nodes)
- ▶ Class 3. Connectivity over time:
 $\forall v, u \in V, v \rightsquigarrow u$. (every node can reach all other once)

- Class 5. Recurrent connectivity:

$\forall v, u \in V, \forall t \in \mathcal{T}, \exists \mathcal{J} \in \mathcal{J}_{(u,v)}^* : \text{departure}(\mathcal{J}) > t.$
(routing can always be achieved over time)

- Class 6. Recurrence of edges:

$\forall e \in E, \forall t \in \mathcal{T}, \exists t' > t : \rho(e, t') = 1$ and G is connected.

- Class 7. Time-bounded recurrence of edges:


$\forall e \in E, \forall t \in \mathcal{T}, \exists t' \in [t, t + \Delta), \rho(e, t') = 1$ and G is connected.

- Class 8. Periodicity of edges:

$\forall e \in E, \forall t \in \mathcal{T}, \forall k \in \mathbb{N}, \rho(e, t) = \rho(e, t + kp)$ some $p \in \mathbb{T}$ and G is connected.

- Class 10. T -interval connectivity:

$\forall i \in \mathbb{N}, \forall T \in \mathbb{N}, \exists G' \subseteq G : V_{G'} = V_G, G'$ is connected and $\forall j \in [i, i + T - 1), G' \subseteq G_j$.



Reliable Delivery in Dynamic Networks

Issue: can we assume Perfect Point-to-Point Links? In particular, is the *Reliable Delivery* achievable?

A channel may be available for limited time and may fail in delivering a message due to the transmission latency

How implement a Reliable Channel?


- ▶ a process knows exactly ρ and ζ of its channel;
- ▶ the channel is up enough to received an *ack* message;
- ▶ infinite message retransmission.

NOTE: the Journey definition guarantees *Reliable Delivery* from p to q


Broadcast Latency in 1-interval connected networks [KLO10]

Broadcast Latency on **static** Multi-Hop Networks
 $O(n)$ (line topology)


Broadcast Latency on **1-interval connected networks**
 $O(n)$ ($\forall t$ at least one not-informed node is connected to an informed node)



G_1



G_2



G_3

Byzantine Reliable Broadcast Specification

Module:

Name: Byzantine Reliable Broadcast, **instance** *brb*, with source *s*.

Events:


Request: $\langle brb, Broadcast | m \rangle$: Broadcasts a message *m* to all processes. Executed only by process *s*.


Indication: $\langle brb, Deliver | p, m \rangle$: Delivers a message *m* broadcast by *p*.

Properties:

RB1: Safety: If some correct process delivers a message *m* with source *p* and process *p* is correct, then *m* was previously broadcast by *p*.

RB2: Liveness: If a correct process *p* broadcasts a message *m*, then every correct process eventually delivers *m*.






Failure Assumptions

Byzantine Failures


Globally Bounded



$f=1$

up to f faulty
processes arbitrarily
spread over the
system

Locally Bounded



$f=1$

up to f faulty
processes in the
neighborhood of every
process

*Specific Spatial
Distribution,
Probabilistic
Distribution, etc.*


Dynamic Network
Globally Bounded Failure Model
[MTD15]

System Model


- ▶ n processes (each one with an unique identifier);
- ▶ **Dynamic Communication Network - TVG**
- ▶ processes can be correct or Byzantine faulty;
- ▶ up to f processes can be faulty (*globally bounded failures*);
- ▶ **processes have no global knowledge (except the value of f)**;
- ▶ Authenticated Channels

Graph Theory - Vertex Connectivity

Static Graphs: *Menger Theorem* : **Vertex Cut = Disjoint Paths**




Disjoint Paths



Min Cut

Graph Theory - Vertex Connectivity on Dynamic Networks

Dynamic Graphs: **dynamic Vertex Cut** \geq **dynamic Disjoint Paths** (between two endpoints)




2 disjoint dynamic paths
between p and q

Min-Cut between p and
 q is 3.

Maurer et al. Algorithm


Extends **Dolev's algorithm on Dynamic Networks**

Idea: leverage the authenticated channels to **collect the ID's of the processes traversed** by a message




Message format

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example




$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example



$$f = 1$$

Dolev (Static DS) Algorithm Graphical Example



$$f = 1$$

Maurer et al. Algorithm

Same propagation algorithm

The **verification algorithm** checks for a **dynamic min-cut** with size $f + 1$.

Every message is retransmitted every time a process detects a network change in its neighborhood.


Correctness - Safety

Same as Dolev:

All the messages generated by a Byzantine process are labeled with its ID

\implies Byzantine processes are not able to generate Traversed_Processes with minimum cut lower than f .

\implies Maurer et al. algorithm enforces *safety*



$[0, 1, b], [p, 3, b, 5]$

Correctness - Liveness (Communication Between two Endpoints)

The Byzantine Reliable Communication (single source, single destination) from process p to q , considering at most f Byzantine processes, is achievable if and only if the **Dynamic Minimum Cut between p and q is at least $2f+1$** (i.e. the minimum number of nodes to remove from the network in such a way that no dynamic path exist between p and q)

Maurer et al. algorithm Analysis

Same as Dolev:

Message Complexity: **Exponential in the number of processes** (considering only correct processes)

Delivery Complexity: Solve an **NP-Complete problem**
(Min-Cut \Rightarrow Minimum Hitting Set)

\Rightarrow Not practically employable

Is it possible to do better? open research

Additional Issue

Static distributed systems

strict broadcast condition : $2f + 1$ -connected network

The **vertex connectivity** of a graph can be **polynomially** verified through a max-flow algorithm

Dynamic distributed systems

strict communication condition : dynamic min-cut at least $2f + 1$


The computation of the **dynamic min-cut** is a **NP-Complete** problem.

Dynamic Network
Locally Bounded Failure Model
[BFT18]

System Model

- ▶ n processes (each one with an unique identifier);
- ▶ **Dynamic Communication Network - TVG**
- ▶ processes can be correct or Byzantine faulty;
- ▶ up to f processes can be faulty in the neighborhood of every process (*locally bounded failures*);
- ▶ **processes have no global knowledge (except the value of f)**;
- ▶ Authenticated Channels


CPA Algorithm (Static DS)



$$f = 1$$

- ▶ **the source broadcasts the message;**
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.


CPA Algorithm (Static DS)



$f = 1$

- ▶ the source broadcasts the message;
- ▶ **a neighbor of the source directly accepts and relays the message;**
- ▶ a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.


CPA Algorithm (Static DS)



$f = 1$

- ▶ the source broadcasts the message;
- ▶ **a neighbor of the source directly accepts and relays the message;**
- ▶ a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.


CPA Algorithm (Static DS)



$f = 1$

- ▶ the source broadcasts the message;
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ **a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.**


CPA Algorithm (Static DS)



$f = 1$

- ▶ the source broadcasts the message;
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ **a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.**


CPA Algorithm (Static DS)



$$f = 1$$

- ▶ the source broadcasts the message;
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ **a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.**


CPA Algorithm (Static DS)



$$f = 1$$

- ▶ the source broadcasts the message;
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ **a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.**

CPA Algorithm (Static DS)



$f = 1$

- ▶ the source broadcasts the message;
- ▶ a neighbor of the source directly accepts and relays the message;
- ▶ **a process that receives the same message from $f + 1$ distinct neighbors accepts and relays the message.**


Does the CPA algorithm work on Dynamic Distributed Systems?

The **safety** property **is still guaranteed** for the same reasons of static systems:

- (i) Every process relays a message only if it has been delivered
- (i) At most f faulty process are present in the neighborhood

\implies **CPA** algorithm **enforces safety** also on Dynamic Distributed Systems


The **liveness** property on static DS requires the existence of a specific partition (MKLO) to be guaranteed



$$k = 3$$

► The source is placed in L_0 ;


The **liveness** property on static DS requires the existence of a specific partition (MKLO) to be guaranteed



$$k = 3$$

- ▶ The source is placed in L_0 ;
- ▶ **The neighbors of the source are placed in level L_1 ;**


The **liveness** property on static DS requires the existence of a specific partition (MKLO) to be guaranteed



$$k = 3$$

- ▶ The source is placed in L_0 ;
- ▶ The neighbors of the source are placed in level L_1 ;
- ▶ **Any other node is places in the first level such that it has at least k neighbors in the previous levels.**

The **liveness** property on static DS requires the existence of a specific partition (MKLO) to be guaranteed



$$k = 3$$

- ▶ The source is placed in L_0 ;
- ▶ The neighbors of the source are placed in level L_1 ;
- ▶ **Any other node is placed in the first level such that it has at least k neighbors in the previous levels.**


Correctness Necessary and Sufficient condition (Static DS)

Necessary condition: MKLO with $k = f+1$

Sufficient condition: MKLO with $k = 2f+1$

Strict condition: MKLO with $k = f+1$ removing any possible placement of the Byzantine processes (NP-Complete Problem)

On Dynamic DS a MKLO is not enough



- 1) An edge may disappear while transmitting a message
- 2) The order of appearance of edges matters

1) An edge may disappear while transmitting a message

A process has to face the unreliability of the channels:

- ▶ it knows exactly ρ and ζ of its channel;
- ▶ infinite message retransmission (if an ack is received it stops).

We are assuming no global knowledge, thus processes have to infinitively retransmit the delivered messages.

The CPA algorithm can be ported on Dynamic Distributed Systems without further changes

Liveness Conditions

Given the TVG, we identify through the predicate $\text{RCD}(p_i, p_j, t')$ the **edge appearances that allow to reliably deliver** a message transmitted over a channel

$$\text{RCD}(p_i, p_j, t') = \begin{cases} \text{true} & \text{if } \rho(< p_i, p_j >, \tau) = 1, \forall \tau \in [t', t' + \zeta(e_{ij}, t')]. \\ \text{false} & \text{otherwise.} \end{cases}$$

Liveness Conditions

Liveness condition \Rightarrow RCD + MKLO


= Temporal Minimum K-level Ordering (TMKLO)

$$\mathcal{A}_k(p_j, t) = \begin{cases} 1 & \text{if } p_j = p_s \text{ with } t \geq t_{br} & \text{(AK1)} \\ 1 & \text{if } \exists t' \geq t_{br} : \text{RCD}(p_s, p_j, t') = \text{true with } t \geq t' + \zeta(e_{s,j}, t') & \text{(AK2)} \\ 1 & \text{if } \exists p_1, \dots, p_k : \forall i \in [1, k], \mathcal{A}_k(p_i, t_i) = 1 \text{ and} \\ & \quad \exists t'_i \geq t_i : \text{RCD}(p_j, p_i, t'_i) = \text{true with } t \geq t'_i + \zeta(e_{i,j}, t'_i) & \text{(AK3)} \\ 0 & \text{otherwise} \end{cases}$$


TMKLO := $p \in L_{t_i}$ iff $t_i = \min t \in \mathbb{N}$ such that $\mathcal{A}_k(p, t_i) = 1$

TMKLO example


constant latency $\delta = 2$




G0




G1



G2



G3



G4

$$TMKLO_{k=2} : \quad L_0 = \{0\}, \quad L_1 = \{1\}, \quad L_2 = \{3\}, \\ L_4 = \{2, 4\}$$

Correctness - Liveness

Necessary condition: TMKLO with $k = f + 1$

Sufficient condition: TMKLO with $k = 2f + 1$

How verify the Liveness Conditions

Liveness Conditions \implies TMKLO computation

TMKLO computation \implies full TVG knowledge

$$\mathcal{G} := (V, E, \rho, \zeta)$$


Complexity TMKLO computation: $= O(|V| + |T||E|)$

Dynamic Networks are rarely completely defined

Exceptions are transportation networks, satellite network, etc.
("constrained networks")


Most of the times just global features are available:
connectivity, number of nodes, possible edges that may exists
etc.

Liveness Conditions with weaker knowledge



A TMKLO is necessary to enforce CPA liveness
⇒ further assumptions on the TVG are required

Liveness Conditions with weaker knowledge



C1. Temporal Source

C6. Recurrence of edges

C7. Time-bounded recurrence of edges

What happens moving in C6 or C7?

CPA Liveness Condition in C6 and C7

C6: every edge reappears infinitively often

C7. Time-bounded recurrence of edges

Assuming:

- ▶ on every edge the *RCD* predicate is true infinitively often
- ▶ the underlying graph G is known

$$\text{MKLO}(G) \implies \exists \text{TMKLO}(\mathcal{G})$$

Broadcast Latency

How much it takes a broadcast to complete?

Broadcast Latency (BL) := the length of the period between the broadcast start and the last delivery of a correct process.

C1, full knowledge about \mathcal{G} , broadcast start at t_{br} ,

$\exists TMKLO_{2f+1} : P_{2f+1} = \{L_{t_0}, L_{t_1} \dots L_{t_x}\},$

$t_{max}^{2f+1} = t_x$ (the time associated to the last level of P_{2f+1}).

$t_{max}^{f+1} = t_x$ (the time associated to the last level of P_{f+1}).

$$t_{max}^{f+1} - t_{br} \leq BL \leq t_{max}^{2f+1} - t_{br}$$

Broadcast Latency in C6

C6. every edge reapers infinitively often

assuming that: (i) on every edge the *RCD* predicate is true infinitively often

(ii) the underlying graph G is known

$$\text{MKLO}(G) \implies \exists \text{TMKLO}(\mathcal{G})$$

but it is unknown, because only the knowledge about G is available

\implies **no bounds for BL in C6**

Broadcast Latency in C7

C7. Time-bounded recurrence of edges,
reappearance bound Δ known,

assuming that:

- (i) on every appearance makes the *RCD* predicate true,
- (ii) the underlying graph G is known,
- (iii) $\delta_{\max} = \max(\zeta(e, t))$.

$\exists MKLO_{2f+1} : M_{2f+1} = \{L_0, L_1 \dots L_x\},$

$S_{2f+1} = |M_{2f+1}|$ (the number of levels in M_{2f+1})

$$\mathbf{BL} \leq \mathbf{S}_{2f+1}(\delta_{\max} + \mathbf{\Delta})$$

Final Remarks on Dynamic Distributed Systems

- ▶ The complete characterization about the evolution of the systems is seldom available (all the information about a TVG for example)
- ▶ Instead, most of the times global information are available (extracted by a network analysis)
- ▶ Solving problems becomes more challenging due to the dynamicity of the system
- ▶ More attention dedicated the approximate/weaker solution, otherwise too unrealistic assumption have to be considered.

References I

- [BFT18] Silvia Bonomi, Giovanni Farina, and Sébastien Tixeul.
Reliable broadcast in dynamic networks with locally bounded byzantine failures.
In Taisuke Izumi and Petr Kuznetsov, editors, *Stabilization, Safety, and Security of Distributed Systems - 20th International Symposium, SSS 2018, Tokyo, Japan, November 4-7, 2018, Proceedings*, volume 11201 of *Lecture Notes in Computer Science*, pages 170–185. Springer, 2018.
URL: <https://hal.archives-ouvertes.fr/hal-01712277>.
- [CFQS12] Arnaud Casteigts, Paola Flocchini, Walter Quattrociocchi, and Nicola Santoro.
Time-varying graphs and dynamic networks.
IJPEDS, 27(5):387–408, 2012.
URL: <https://doi.org/10.1080/17445760.2012.668546>.
- [KLO10] Fabian Kuhn, Nancy A. Lynch, and Rotem Oshman.
Distributed computation in dynamic networks.
In Leonard J. Schulman, editor, *Proceedings of the 42nd ACM Symposium on Theory of Computing, STOC 2010, Cambridge, Massachusetts, USA, 5-8 June 2010*, pages 513–522. ACM, 2010.
URL: <https://doi.org/10.1145/1806689.1806760>.

References II

- [MTD15] Alexandre Maurer, Sébastien Tixeul, and Xavier Défago.
Communicating reliably in multihop dynamic networks despite byzantine failures.
In *34th IEEE Symposium on Reliable Distributed Systems, SRDS 2015, Montreal, QC, Canada, September 28 - October 1, 2015*, pages 238–245. IEEE Computer Society, 2015.
URL: <https://doi.org/10.1109/SRDS.2015.10>.